

Taming Scylla: Understanding the multi-headed agentic daemon of the coding seas

Micah Villmow*
Individual
research@villmow.us

Abstract

LLM-based tools are automating more software development tasks at an exponential rate, but there's no rigorous way to evaluate how different architectural choices—prompts, skills, tools, multi-agent setups—materially affect both capability and cost. This paper introduces Scylla, an evaluation framework for benchmarking agentic coding tools through structured ablation studies that uses seven testing tiers (T0-T6) progressively adding complexity to isolate what directly influences results and how. The key metric is Cost-of-Pass (CoP): the expected dollar cost to get one correct solution, which directly quantifies the trade-off between complexity and efficiency. The framework is model-agnostic, designed to work with any CLI tool; this paper demonstrates it with Claude Sonnet 4.5, using multiple LLM judges (Opus 4.5, Sonnet 4.5, Haiku 4.5) from the same vendor for evaluation consensus, where judges score results using direct tests, human-driven rubrics, and qualitative assessment. The result is a reproducible framework that quantifies trade-offs between agent complexity and actual outcomes, suggesting that architectural complexity does not always improve quality.

Keywords: LLM agents, software engineering benchmarks, cost-of-pass, multi-agent systems, prompt engineering, ablation studies, evaluation frameworks, CLI tools, agentic AI

1 Introduction

Large language models have ushered in massive increases in capabilities for automated computer interactions. What used to require hand-coded algorithms and pipelines can now be done automatically using state-of-the-art coding models. However, understanding what improves these language models is more black magic than art, let alone a rigorous science. This paper's goal is to help demystify the magic of prompt engineering by proposing a rigorous evaluation framework to quantify the benefits of different architectural approaches.

There are benchmarks for measuring LLM workflows in various domains, such as agent-bench[6], swebench[5], and tau-bench[9]. There are also prompt evaluation benchmarks such as PromptBench[10]

*This paper combines the author's original research and writing with improvements, suggestions, and rewrites provided by Claude Code (claude.ai/code).

and PromptEval[7]. This paper focuses specifically on coding tools, particularly the industry-leading Claude Code[1], and how prompt and architectural modifications affect model behavior. This paper introduces Scylla, a framework for evaluating agentic coding tools in a systematic way, allowing extension to domains beyond CLI-based coding tools. The dryrun results on a trivial Hello World task show that all seven tiers (T0-T6) achieve equivalent quality (all grade A, scores 0.943-0.983) while cost varies $3.8\times$ from \$0.065 (T5 hybrid) to \$0.247 (T6 super). The framework successfully differentiates cost structures across architectural choices even when quality converges, demonstrating that architectural complexity does not always improve quality. Careful hybrid designs (T5) can achieve Frontier Cost-of-Pass (the minimum CoP across all tiers) by selectively combining features rather than maximizing them.

Anthropic has many good resources for improving Claude Code on their engineering blog, but despite these, there are not any intuitive and user-friendly methods for comparing whether changes to the prompt instructions will yield tangible benefits. Therefore, I am introducing Scylla, a testing framework for evaluating prompts, tools, skills, and agents for solving problems that are common for day-to-day coding tasks. I wanted to know if sub-agents, skills, tools, or mcp servers were contributing to actual improved code output, without relying on my gut or intuition. This problem came up multiple times when asked by others to explain how to better utilize CLI tools for programming. In my experience, the quality of the prompts has a dramatic effect on the output. Whether it is the prompt to call the tool or MCP server, the prompt to spawn a sub-agent, or the prompt to trigger a skill, these language-based triggers are fuzzy in their meaning. Unlike a traditional programming language that is very explicit in what it means and what it does, prompts do not map directly and consistently to action. This framework is my attempt at helping unwrap this problem.

The remainder of this paper is organized as follows. First, in Section 2, I will introduce the current work that is being done in this area, and explain how they approach the problem. Then, in Section 3, I will introduce the testing methodology along with an in-depth analysis of the first test case. This will provide the needed understanding of what is being tested, along with why, on something that is easily analyzed and understandable. Section 4 describes the framework’s software architecture, explaining how the evaluation system is built and how it achieves model-agnostic evaluation through the adapter pattern. Then I will go over the rest of the testing framework to showcase what is being tested, measured, and why these are being tested using simple cases introduced in the previous sections.

The framework is designed to investigate questions such as these in future studies:

- Whether it is possible to quantify whether a task is solvable more efficiently by one methodology over others
- Whether the sum of a prompt is more than the individual parts
- Whether there are core improvements that can be made purely through extensions to Claude Code that are generic for all workloads
- Whether there are specific prompt techniques that have secondary effects, positive or negative, on other prompt techniques

- Holding the tool and prompt constant, how much the underlying model may contribute to the quality of the results

Some hypotheses the framework is designed to test in subsequent work include:

- Certain tasks may excel when run as sub-tasks, tools, MCP servers, or skills, for reasons unrelated to context management
- Prompt complexity may have opposing effects on quality: KISS principle (negative correlation) for scenarios in the training set versus inverse KISS (positive correlation) for scenarios outside the training set

2 Related Work

Given that we are testing production tools and not models, many, if not all, of the prior work on evaluating prompts and benchmarks do not apply here. Since there is possibly a large level of indirection between what we are testing and what actually gets executed by the model due to engineering trade-offs, I am considering the tool to be a black box and not attempting to reverse engineer this tool. Despite this, what is executing is hidden behind multiple layers, first being the CLI tool itself, but also whatever optimizations and implementation details the vendor implements on top of their trained base model. The models themselves are not fully documented publicly, as these details are competitive advantages, and the pre or post-processing that occurs is not always visible to the user as they can occur vendor-side.

There are several good benchmarks for evaluating LLM agents and prompts. SWE-Bench[5] tests models on real GitHub issues. Agent-Bench[6] tests multi-turn agents across different environments like operating systems, databases, and knowledge graphs, with fine-grained metrics beyond pass/fail. TAU-Bench[9] focuses on whether agents can effectively use external tools. For prompt evaluation, there is PromptBench[10] (unified testing across tasks), PromptEval[7] (automated correctness and robustness checking), and EleutherAI's lm-evaluation-harness[3] (standardized multi-task comparison).

However, these benchmarks all assume direct access to model inputs and outputs, evaluating models directly rather than complete CLI tools. With production CLI tools like Claude Code, the model is wrapped in layers of system prompts, tool schemas, skill definitions, hooks, skills, MCP servers, vendor optimizations, and orchestration logic. I cannot just test the model in isolation, so I must test the whole system. No standard benchmarks exist for CLI tools like Claude Code and how prompts affect them.

My work is based solely on evaluating CLI tools, as the CLI's tools are more than the model themselves. As I mentioned earlier, the agentic loop, with hooks, tools, skills, sub-agents, MCP servers, and other logic wrapped together into a single application where the only way to get control of the behavior is through the English language is what I want to evaluate for effectiveness. From this interface, programmatic tools can be spawned, but the ability to properly and accurately interact with the agent is via a fuzzy language interface, and not via traditional programmatic interfaces. While there are some hooks that allow extra programmatic validation with Claude Code,

I am not evaluating those at this time. Claude Code has the ability to use agentic evaluation at the hook boundary, but triggering it is guaranteed (and not language-based), so it is not interesting for probabilistic evaluation.

3 Test Methodology

3.1 Experimental Design

This experiment is designed by testing English phrases, colloquially known as prompts, via the various methodologies exposed by a CLI tool, in this case Claude Code. The prompts to be tested are taken from the ProjectOdyssey[4] git repository at github hash 011a3ff on December 30th, 2025. The prompts are broken down into their components and separated into various tiers which will be discussed later. These components are used to setup the experiment, which is run by allowing an agent a nearly unfettered access to the system, only blocking dangerous ops, thanks to the safety-net plugin[2] from cc-marketplace[8], to perform a task. The task has a well defined solution that is then judged by three different LLMs of various 'strength'. In this case Claude Opus 4.5, Claude Sonnet 4.5, and Claude Haiku 4.5. Each of the 4.5 models are sufficiently advanced in capabilities to be considered independent judges of a task with low failure rates. The judges are provided the same prompt, so the only difference between their results comes from the judge training and implementation differences and not from the prompt or test input. Each judge will receive the output of the task LLM, and provide the results based on the criteria. The judges have the following categories of evaluation; functional correctness, code quality, proportionality, build pipeline, and overall quality.

Table 1: *LLM-as-Judge Evaluation Categories*

| Category | Weight | Description |
|------------------------|--------|--|
| Functional Correctness | 0.35 | File existence, output correctness, exit codes, exact output matching |
| Code Quality | 0.20 | Syntax validity, idiomatic code, unused imports, PEP8 compliance |
| Proportionality | 0.15 | Appropriate scope, minimal files, no unnecessary artifacts or tests |
| Build Pipeline | 0.10 | Build passes, format checks, tests (when applicable), pre-commit hooks |
| Overall Quality | 0.20 | Engineering judgment on appropriateness, maintainability, and senior engineer approval |

Total Weight: 1.0 (100%)

Each category contributes proportionally to the final score. Here is the formula:

$$S_{final} = \sum_i w_i \cdot \frac{P_i^{achieved}}{P_i^{max}}$$

where w_i are the category weights (they sum to 1.0), and P_i is the points the test got versus the maximum possible (skipping any N/A items). For scoring individual items:

- **Binary items:** You either get it or you do not (1.0 or 0.0)
- **Graduated items:** Partial credit on a 0.0-1.0 scale based on results
- **Subjective items:** LLM judgment with calibrated deductions

Table 2: *Deduction Calibration Scale (for subjective assessment)*

| Severity | Deduction Range | Examples |
|------------|-----------------|--|
| Negligible | 0.00-0.05 | IDE config files, <code>__pycache__</code> artifacts |
| Trivial | 0.05-0.15 | Missing trailing newlines, unused imports |
| Minor | 0.15-0.30 | Missing docstrings, magic numbers |
| Moderate | 0.30-0.50 | Code duplication, hardcoded values |
| Major | 0.50-0.80 | Non-critical security issues, race conditions |
| Severe | 0.80-1.50 | Critical security vulnerabilities |
| Critical | 1.50+ | Non-functioning solutions, destructive operations |

The final score maps to a grade using this scale:

Table 3: *Grade Scale*

| Grade | Threshold | Label | What It Means |
|-------|-------------|------------|--|
| S | 1.00 | Amazing | Perfect score, goes above and beyond |
| A | ≥ 0.80 | Excellent | Production ready |
| B | ≥ 0.60 | Good | Works well, minor tweaks needed |
| C | ≥ 0.40 | Acceptable | It works but has issues |
| D | ≥ 0.20 | Marginal | Lots of problems but salvageable with effort |
| F | < 0.20 | Failing | Complete failure of task |

I use **0.60** (Grade B) as the pass threshold. That means the solution works and meets requirements, even if there is room for minor improvements. An S grade needs a perfect 1.00 and you have to actually exceed what was asked for. I would not expect many, if any, tests to get an S rating.

Each experiment can be reproduced by running the top level test run script, which will launch the same set of tasks with the same parameters, where the only variation is the judgement of the LLM judges when determining how to judge the work.

This finishes the summary of a single test. However, the tests themselves are defined differently. The tests are a prompt and a configuration file that specify a repository, a github hash, a set of configuration files to override any pre-defined tooling, set of commands to validate the results, and a container to run everything in to help with reproducibility. The first test is being used as an example

in this paper, and also as a pipecleaner to show that everything works as expected. This example is 'hello world' from octocat, but forked to my repository just to make sure that the repository is not polluted. The precaution is done just in case the agents make mistakes or do things that the original author probably does not want to be bothered by.

3.1.1 Test-001: Hello World Baseline

First, let us look at the simplest possible test to make sure everything works. This is literally just creating a "Hello World" script, which is a pipe-cleaner for the infrastructure and to discuss the methodology without intermixing with the complexity of more realistic tests.

Test Configuration:

| Field | Value |
|----------------|------------------|
| ID | test-001 |
| Name | Hello World Task |
| Timeout | 300 seconds |
| Pass Threshold | 0.60 (Grade B) |

Task Prompt:

Create a Python script `hello.py` that prints "Hello, World!" to stdout, exits with code 0, and uses relative paths. The script is created in the current working directory.

Expected Output:

```
Hello, World!
```

Expected Result:

```
print("Hello, World!")
```

or

```
#!/usr/bin/python3
print("Hello, World!")
```

What Should Happen:

Even T0 (no system prompt at all) is expected to get an 'A', since we are talking ≥ 0.80 scores. If T0 cannot do Hello World, I will assume that something is fundamentally wrong with the framework itself and throw out the results. Higher tiers (T1-T6) are also expected to succeed, as there is no reason fancy prompts or multi-agent setups would help with something this simple. However, if performance drops on this test, it means the added complexity is actually making things worse even on something so simple, so if this happens, we will analyze why.

Now that we have gone over the test itself, let us discuss the strategy and tiered approach. The first thing to test is with no prompt at all, including no system prompt, if the tool allows it. This is

Table 4: Rubric Categories and Weights

| Category | Weight | Key Criteria |
|------------------------|--------|--|
| Functional Correctness | 35% | File <code>hello.py</code> exists; running <code>python hello.py</code> prints "Hello, World!" |
| Code Quality | 20% | Valid Python syntax; idiomatic code; no unused imports; appropriate structure |
| Proportionality | 15% | Total files ≤ 3 ; LOC ≤ 3 ; no unnecessary test files; appropriate scope |
| Build Pipeline | 10% | Syntax check passes; format check passes (if ruff present); no linter errors |
| Overall Quality | 20% | Senior engineer approval; appropriately scoped for task complexity |

to provide as close to a baseline as the base model as possible by overwriting the system prompt with an empty string and not using any configuration or non-default settings from the tool. This provides the baseline that all improvements are measured against. For something as simple as hello world, this baseline will solve the task. The test setup is such that variability in judging will occur, but there is not much one can do to improve the output of a hello world script. However, there are things that you can do that make things worse or break the expected behavior, but I would expect all solutions to be the exact same for all the tests. Divergence points to interesting results.

3.1.2 Tiered Ablation Strategy

The core idea is simple: start with nothing, then add one set of things at a time to see what actually helps. This ablation study uses seven tiers that progressively add complexity, with **113 subtests** total. Each tier gets tested independently so we can isolate what each component contributes.

Table 5: Testing Tiers (Ablation Study Framework)

| Tier | Name | Sub-tests | Primary Focus | Tools | Delegation |
|------|------------|-----------|---|---------|------------|
| T0 | Prompts | 24 | System prompt ablation (empty \rightarrow full) | - | No |
| T1 | Skills | 10 | Domain expertise via installed skills | Default | No |
| T2 | Tooling | 15 | External tools and MCP servers | Yes | No |
| T3 | Delegation | 41 | Flat multi-agent with specialists | Yes | Yes |
| T4 | Hierarchy | 7 | Nested orchestration with orchestrators | Yes | Yes |
| T5 | Hybrid | 15 | Optimal combinations from all tiers | Yes | Yes |
| T6 | Super | 1 | Maximum capability configuration | All | All |

How the Tiers Work:

1. **T0 (Baseline):** Start with an empty prompt (00-empty) to see what the raw model can

do, then go all the way up to the full 1787-line CLAUDE.md (03-full). Individual blocks (B01-B18) let me test each piece of the prompt separately to see what actually matters.

2. **T1-T2 (Skills vs Tools):** T1 uses skills, domain knowledge baked into prompts. Token-efficient. T2 uses external tools via JSON schemas. Problem is, loading all those tool definitions inflates token usage. I call this the "Token Efficiency Chasm", the gap between lean skill-based approaches and schema-heavy tool architectures.
3. **T3-T4 (Multi-Agent Setups):** T3 does flat delegation, breaking tasks into smaller pieces and assigning them to specialist agents. T4 adds hierarchy with self-correction loops, but this complexity can increase costs.
4. **T5 (Smart Combinations):** Take what works from the other tiers, combine them together in different combinations. A single test would have the best T1 skills, T2 tools, T3 agents, and T4 task delegation. We do not want to brute force here due to combinatorial explosion, but picking combinations of the top few categories can help give an idea of what combinations work best together.
5. **T6 (Everything):** Turn on everything at once. All skills, tools, agents, prompt segments, and servers. This I hope establishes the theoretical max performance and shows where diminishing returns kick in, but also can show signs of over-engineering if it is occurring.

For each tier T(n), I compare it directly against T(n-1) to see what that specific change actually achieves in terms of performance and cost. These tiers map onto a broader multi-dimensional search space.

3.2 Dimensional Search Space

The framework tests across four different dimensions. Each one is an independent knob you can turn, and they all affect both what the agent can do and how much it costs.

3.2.1 Agent Complexity Axis (Tiers 0-6)

The agent complexity axis spans from simple single-agent configurations to hierarchical multi-agent systems, as shown in Table 6.

3.2.2 Prompt Complexity Axis

Prompt complexity is measured in lines of system prompt content, ranging from 0 (empty) to 1787 (full CLAUDE.md from ProjectOdyssey[4]), as detailed in Table 7.

Each block (B01-B18) can be tested separately to see whether the part actually contributes to the whole.

Table 6: *Agent Complexity Axis*

| Tier Range | Complexity Level | Description |
|------------|---------------------------|---|
| T0 | Single-agent, prompt-only | Base model with varying prompt sophistication |
| T1 | Single-agent with skills | Add in agentic skills to improve the quality of the work |
| T2 | Single-agent with tools | External API access via tool schemas |
| T3 | Multi-agent, flat | Specialist agents with central orchestrator |
| T4 | Multi-agent, hierarchical | Nested orchestration with self-correction loops |
| T5 | Best case scenarios | Attempt to pick the best case scenarios from previous runs to see if the sum is more than its parts |
| T6 | Maximum configuration | All features enabled simultaneously |

Table 7: *Prompt Complexity Axis*

| Level | Lines | Description | Representative Test |
|----------|-------|----------------------------------|---------------------|
| Empty | 0 | No system prompt | T0-00-empty |
| System | 0 | Only system prompt | T0-01-empty |
| Minimal | 55 | Safety rules only | T0-06-B02 |
| Core | 260 | Essential blocks (B02, B07, B18) | T0-03-core |
| Standard | 400 | Seven core blocks | T0-02-standard |
| Full | 1787 | All 18 CLAUDE.md blocks | T0-03-full |

Table 8: *Skill Complexity Axis*

| Category | Count | Example Domains | Token Efficiency |
|---------------|-------|---------------------------------|------------------|
| Agent | 5 | Agent management patterns | High |
| CI/CD | 7 | Build and deployment automation | High |
| Documentation | 4 | Technical writing assistance | Medium |
| GitHub | 10 | Repository management | Medium |
| Language | 10 | Programming language specific | High |
| Quality | 5 | Code quality and review | Medium |
| Workflow | 5 | Development workflow patterns | High |

3.2.3 Skill Complexity Axis

Skills are organized by domain, as shown in Table 8.

Table 8 shows the 46 skills across 7 categories tested in T1. Skills bake knowledge into prompts, so you avoid loading massive tool schemas. But do these actually improve performance? That is an open question.

3.2.4 Agent Hierarchy Axis

Three ways to organize agents, tested across T3-T4:

Table 9: Agent Hierarchy Axis

| Pattern | Coordination | Communication head | Over- | Use Cases |
|---------------------|--|--------------------|-------|---|
| Flat | No supervision; peer-to-peer | Low | | Simple, independent tasks |
| Hierarchical | L0-L4 levels with explicit supervision | High | | Complex, interdependent tasks requiring planning |
| Hybrid | Selective hierarchy based on task complexity | Medium | | Adaptive: flat for simple tasks, hierarchical for complex |

Hierarchy matters for costs because each supervision layer adds more orchestration tokens and potentially more self-correction iterations.

4 Framework Architecture

Having established what is tested (tiers T0-T6) and the multi-dimensional search space they explore, I now describe how the Scylla framework is built. The architecture is designed around four principles: model-agnostic evaluation via the adapter pattern, full reproducibility through git worktrees and checkpoint recovery, crash recovery with atomic state persistence, and parallelism at both tier and subtest levels to minimize evaluation time.

4.1 System Overview

The framework consists of six major subsystems that coordinate to execute evaluations, collect artifacts, score results, and produce analyses. Figure 1 shows the component hierarchy and data flow.

The E2E Runner coordinates all evaluation steps, managing parallel execution of tiers and subtests while enforcing dependency constraints (T5 waits for T0-T4; T6 waits for T5). The Workspace Manager creates isolated git worktrees at pinned commits to prevent cross-contamination. The Tier

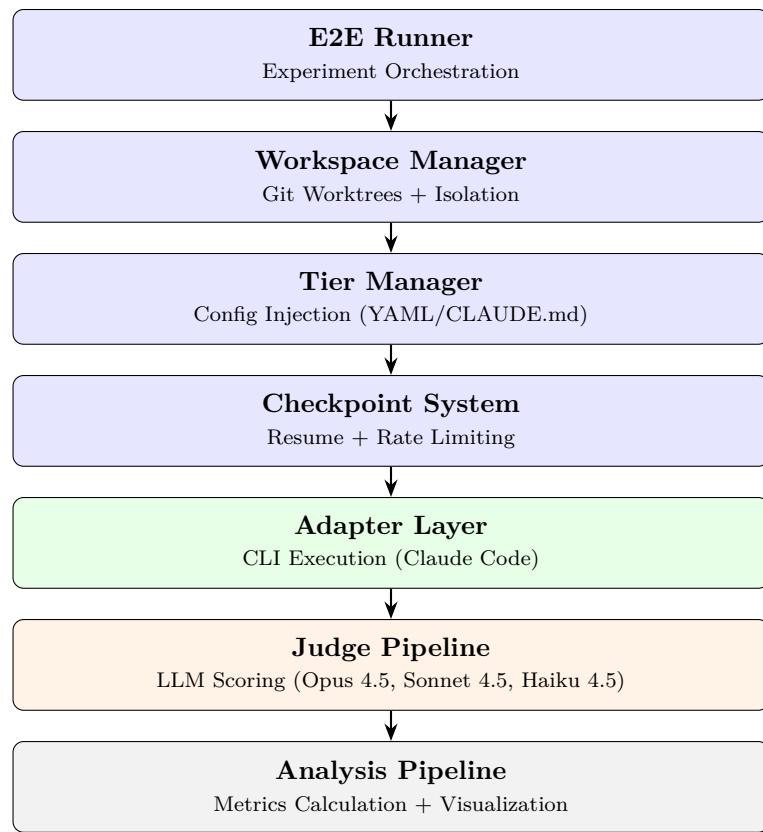


Figure 1: System architecture showing the seven-layer pipeline. Data flows top-to-bottom through: orchestration, workspace isolation, configuration injection, checkpoint management, CLI execution, LLM evaluation, and metrics analysis.

Manager injects configuration files (CLAUDE.md, skills, agent definitions) via symlinks into each workspace. The Checkpoint System saves JSON checkpoints after every completed run, enabling crash recovery and rate limit coordination. The Adapter Layer executes CLI commands via a standardized interface (AdapterConfig → AdapterResult). The Judge Pipeline runs three LLM judges in parallel, scoring outputs using weighted rubrics. The Analysis Pipeline aggregates data into DataFrames and generates figures and tables.

4.2 Adapter Layer

The adapter pattern is how Scylla achieves model-agnostic evaluation. All agent implementations inherit from the `BaseAdapter` abstract class, which defines two methods: `execute(config: AdapterConfig) -> AdapterResult` for running evaluations, and `extract_tokens(stdout: str, stderr: str) -> AdapterTokenStats` for parsing token usage from CLI output.

The current implementation includes the `ClaudeCodeAdapter` for Claude Code CLI, which is used for all evaluations in this paper. Each adapter receives an `AdapterConfig` containing the model, prompt file, workspace path, output directory, timeout, and environment variables. The adapter executes the CLI command, captures stdout/stderr, extracts token statistics from the output, calculates costs using vendor pricing tables, and returns an `AdapterResult` with exit code, output, duration, token stats, cost, API call count, and timeout status.

This design enables future cross-vendor comparisons using identical tier configurations. Additional adapters for tools like Cline, OpenAI Codex, or other CLI-based coding assistants can be implemented by inheriting from `BaseAdapter` and providing tool-specific execution and token extraction logic. Token extraction logic adapts to each vendor’s output format, but the framework sees a standardized `AdapterTokenStats` regardless of source.

4.3 Execution Pipeline

Every evaluation run follows a five-step pipeline from workspace preparation to metrics calculation. Figure 2 shows the sequential flow.



Figure 2: Execution pipeline for a single evaluation run. Each step produces inputs for the next, enabling reproducibility through saved artifacts at every boundary.

Step 1: Workspace Prep creates a fresh git worktree at the pinned commit and injects tier-specific configuration files via symlinks:

- CLAUDE.md from `config/tiers/TN/subtest-NN/CLAUDE.md`
- Skills from `config/tiers/TN/subtest-NN/.claude-plugin/skills`

- Agents from `config/tiers/TN/subtest-NN/.claude/agents`
- Generates a `replay.sh` script for manual reproduction

Step 2: Agent Execution invokes the adapter’s `execute()` method with the workspace path and prompt file. The adapter runs the CLI command (e.g., `claude-code < prompt.txt`) with a timeout (default 3600s), captures `stdout/stderr` via subprocess pipes, and saves the output to `output_dir/stdout.txt` and `output_dir/stderr.txt`.

Step 3: Artifact Capture collects all execution outputs: `stdout/stderr`, git diff of file changes, complete file listing, CLI logs, and the generated `replay.sh` script. Artifacts are saved to `results/test-NNN/TN/subtest-NN/run-MMMMM/` for later analysis.

Step 4: Judge Evaluation runs three LLM judges in parallel (Opus 4.5, Sonnet 4.5, Haiku 4.5). Each judge receives the same prompt containing the task description, expected output, rubric, and captured artifacts. Judges score five weighted categories (functional correctness, code quality, proportionality, build pipeline, overall quality) and return a final score and grade. The framework computes the mean of the three scores for consensus.

Step 5: Metrics Calculation aggregates judge scores, token statistics, and timing data to compute Pass-Rate, Implementation Rate, Cost-of-Pass, token distribution, latency breakdown, and judge agreement metrics. Results are saved to `runs.csv`, `judges.csv`, `criteria.csv`, and `summary.json`.

4.4 Tier Dependencies and Parallelism

Tiers execute in three sequential phases based on dependency constraints. Figure 3 visualizes the dependency graph.

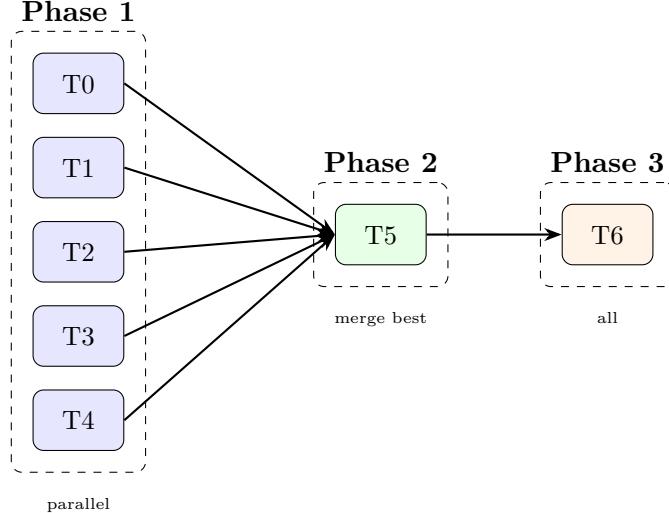


Figure 3: Tier dependency graph and parallel execution model. T_0-T_4 execute in parallel (Phase 1). T_5 waits for T_0-T_4 completion and merges their best configurations (Phase 2). T_6 waits for T_5 and enables all features (Phase 3).

Phase 1 (Parallel): T0-T4 have no dependencies and execute in parallel using a `ThreadPoolExecutor`. Each tier runs all its subtests concurrently via a `ProcessPoolExecutor` for subtest-level parallelism. This reduces total evaluation time from 7 sequential tier runs to 3 sequential phases with internal parallelism.

Phase 2 (Merge): T5 depends on T0-T4 and cannot start until all Phase 1 tiers complete. The Tier Manager merges configurations using the following rules: skills are unioned across all tiers (if T1 uses skill A and T3 uses skill B, T5 gets both), agents are unioned similarly, blocks from the highest-scoring tier replace conflicts (if T0 and T2 both define block B02, T5 uses the version from whichever tier scored higher), and tools follow an "all wins" policy (if any tier enables tools, T5 enables them).

Phase 3 (Everything): T6 depends on T5 and enables maximum configuration: all 61 skills, all tools, all 44 agents, and all CLAUDE.md blocks. This establishes the theoretical performance ceiling and demonstrates diminishing returns from over-engineering (as shown in Results).

Parallelism is bounded by a global semaphore limiting concurrent agent executions to prevent API rate limit violations. The Checkpoint System coordinates rate limits across parallel workers by tracking timestamps and enforcing minimum delays between API calls.

4.5 Checkpoint and Reproducibility

Scylla is designed for long-running experiments that may crash or hit rate limits. The Checkpoint System saves progress after every completed run, enabling seamless resume.

After each run completes, the runner writes a checkpoint JSON to `results/test-NNN/checkpoint.json.tmp`, atomically renames it to `checkpoint.json`, and validates the configuration hash on resume to detect experiment tampering. The checkpoint contains the list of completed runs (tier, subtest, run_id, timestamp, pass/fail), current rate limit state (last API call timestamp, tokens consumed), and experiment metadata (test ID, total runs planned, start time).

On resume, the runner reads the checkpoint, validates the config hash matches the current experiment, skips already-completed runs, and continues from the next pending run. This allows experiments to survive crashes, API timeouts, network failures, and manual interruptions (Ctrl+C) without losing work.

Reproducibility is further ensured by the `replay.sh` script generated for each run. The script contains the exact git commit, CLI command, environment variables, and workspace setup needed to manually reproduce the run. This enables debugging, manual inspection, and verification of framework behavior.

Having described how the framework is built, I now turn to the specific metrics used to quantify performance, quality, and economic trade-offs.

5 Test Metrics

5.1 Performance Metrics

Pass-Rate is straightforward, did it work or not:

$$\text{Pass-Rate} = \frac{\text{correct_solutions}}{\text{total_attempts}}$$

Range is 0.0 (nothing worked) to 1.0 (everything worked). "Correct" means it passes the test suite for that specific task. Report this with confidence intervals (95% CI if you have 30+ runs).

Fine-Grained Progress Rate (R_{Prog}) tracks how far you got through multi-step tasks:

$$R_{Prog} = \frac{\text{achieved_progress_steps}}{\text{expected_progress_steps}}$$

Range is 0.0 to 1.0. If you get 1.0, it means the agent completed all expected steps. This is super useful for debugging where things go wrong in complex workflows, especially in hierarchical setups with all their self-correction loops.

Consistency measures how stable the outputs are:

$$\text{Consistency} = 1 - \frac{\sigma(\text{outputs})}{\mu(\text{outputs})}$$

Range is 0.0 to 1.0, higher means more deterministic. Matters most for where you are trying to get reliable structured outputs.

5.2 Quality Metrics

Implementation Rate (Impl-Rate) measures whether you actually satisfied the requirements:

$$\text{Impl-Rate} = \frac{\text{satisfied_requirements}}{\text{total_requirements}}$$

Range is 0.0 to 1.0. This gives you more detail than just pass/fail, you get partial credit for incomplete work. Checked using multiple LLM judges with mean scoring for consensus.

5.3 Efficiency and Cost Metrics

Latency is just time from start to finish (seconds):

- Time-to-First-Token (TTFT)
- Total response time

- Tool execution time

It matters a lot for architectures where verification loops can really slow things down.

Token Distribution shows where your tokens are going:

$$\text{token_dist} = \left\{ \frac{\text{input_tokens}}{\text{total_tokens}}, \frac{\text{output_tokens}}{\text{total_tokens}}, \frac{\text{tool_input_tokens}}{\text{total_tokens}}, \frac{\text{tool_output_tokens}}{\text{total_tokens}} \right\}$$

Useful for figuring out what is actually contributing to the cost (like T3's massive agent prompts or T4's orchestration overhead).

Cost-of-Pass (CoP) is the primary metric, what is the expected cost to get one correct solution:

$$\text{CoP} = \frac{\text{total_cost}}{\text{pass_rate}}$$

Units are USD. Lower is better. If pass-rate hits zero, CoP goes to infinity, that configuration is economically dead. This combines both cost and accuracy into one number that tells you if something is actually sustainable.

Frontier CoP represents the best CoP for all the various tests:

$$\text{Frontier_CoP} = \min(\text{CoP}_{T0}, \text{CoP}_{T1}, \dots, \text{CoP}_{T6})$$

This metric currently is just the minimum CoP across all tiers. Comparing this against what it costs to hire a human expert will allow developers to see if automation actually makes economic sense. Different model providers will have different cost assumptions.

Table 10: Model Pricing (as of January 2026)

| Model | Input (\$/1M tokens) | Output (\$/1M tokens) |
|-------------------|----------------------|-----------------------|
| Claude Opus 4.5 | \$15.00 | \$75.00 |
| Claude Sonnet 4.5 | \$3.00 | \$15.00 |
| Claude Haiku 4.5 | \$1.00 | \$5.00 |

6 Test Configuration

6.1 Hardware and Infrastructure

Each test runs in its own git clone with the repo at a specific git commit. This means every run is reproducible and tests cannot mess with each other. Every container starts fresh with:

- Clean git workspace at the exact commit specified

Table 11: *Hardware and Infrastructure*

| Component | Specification |
|-----------|---|
| Platform | Linux (WSL2) |
| Kernel | 6.6.87.2-microsoft-standard-WSL2 |
| Isolation | Each test runs in clean workspace |
| Compute | Standard CPU (no GPU required for evaluation) |

- Tier-specific config files
- Whatever tools/skills that tier needs
- Isolated filesystem for collecting results

6.2 Software Stack

Table 12: *Software Stack*

| Component | Version/Tool |
|-------------------|---|
| CLI Tool | Claude Code ¹ (primary evaluation target) |
| Language Runtime | Python 3.12.3, Mojo 0.26.1.0.dev2025122805 (211e2f5c) |
| Package Manager | Pixi |
| Container Runtime | Docker |
| Orchestration | Custom Scylla framework (Section 4) |
| Validation | JSON Schema, YAML validation |
| Version Control | Git Version 2.43.0 |

The evaluation harness does five things:

1. **Workspace Prep:** Clone the repo, check out the specific commit, inject tier config
2. **Run the Agent:** Fire up Claude Code with whatever prompt/tools that tier uses
3. **Capture Everything:** Grab the output, command logs, file changes, artifacts
4. **Judge It:** Run three LLM judges in parallel (Opus, Sonnet, Haiku)
5. **Calculate Metrics:** Crunch the numbers for Pass-Rate, Impl-Rate, CoP, token usage, consensus scores

6.3 Model Configuration

Judge Configuration (evaluating the outputs):

Table 13: Execution Models (performing the tasks)

| Model | Model ID | Primary Use |
|-------------------|----------------------------|---|
| Claude Opus 4.5 | claude-opus-4-5-20251101 | complex reasoning, hierarchical orchestration |
| Claude Sonnet 4.5 | claude-sonnet-4-5-20250929 | standard execution, balanced cost/capability |
| Claude Haiku 4.5 | claude-haiku-4-5-20251001 | simple tasks, cost optimization |

- Three judges per evaluation: Opus 4.5, Sonnet 4.5, Haiku 4.5
- Take the mean of the three scores for consensus
- Same prompt for all judges (only the model changes)
- Judge prompt: `config/judge/system_prompt.md`

Safety:

- Safety-net plugin blocks destructive operations

Model-Agnostic Framework Design:

The framework is designed to work with any CLI tool or model through standardized interfaces. Everything goes through the CLI's language interface and filesystem outputs without touching model APIs directly. This enables consistent metrics (CoP, Pass-Rate, Impl-Rate) across all models for apples-to-apples economic comparisons. The tier structure (T0-T6) applies to all tools, allowing direct architectural comparisons across vendors. Everything is in version-controlled YAML (model IDs, temperature, token limits), making it easy to reproduce across different tools and swap judges. The adapter pattern enabling this is described in Section 4.2.

7 Test Cases

7.1 Pull Request (PR) Selection Criteria

Test cases come from real software development tasks. Here is what I consider to make a good test:

1. **Reproducible:** Pin it to a specific git commit
2. **Clear success criteria:** Can be expressed in a rubric with measurable requirements
3. **Representative:** Real work that developers actually do
4. **Incrementally complex:** From trivial (Hello World) to multi-file architecture changes
5. **Unambiguous:** Clear task, clear expected outcome

Complexity also depends on:

Table 14: Size Categories

| Category | Lines of Code (LOC) | Complexity Characteristics | Example Tasks |
|---------------|---------------------|--|---|
| Small | < 100 LOC | Single file changes, configuration updates | Config file modification, simple script creation |
| Medium | 100-500 LOC | Feature additions, localized refactoring | Add validation logic, implement utility function |
| Large | 500-2000 LOC | Multi-file features, architectural changes | New module implementation, build system migration |

- How many tool calls you need
- How much of the codebase you have to understand
- How many sequential steps
- How many constraints you are working under

7.2 Workflow Categories

Different categories test different capabilities:

Table 15: Workflow Categories

| Category | Description | Complexity | Key Challenges |
|----------------------|--|-------------|--|
| Build System | Makefile, Justfile, build automation configuration | Low-Medium | Syntax correctness, equivalence preservation |
| CI/CD | GitHub Actions, deployment pipelines, automation | Medium | Multi-file coordination, environment configuration |
| Bug Fixing | Defect resolution from issue description | Medium-High | Root cause diagnosis, minimal change principle |
| New Features | Feature implementation from requirements | High | Requirements interpretation, design decisions |
| Refactoring | Code restructuring without behavior change | Medium | Behavior preservation, test coverage |
| Optimization | Performance improvements, algorithmic enhancements | Medium-High | Profiling, benchmarking, trade-off analysis |
| Review | Code review and feedback generation | Medium | Pattern recognition, best practice knowledge |
| Documentation | Technical documentation generation | Low-Medium | Clarity, completeness, accuracy |
| Issue Filing | Bug report creation from symptoms | Low | Information gathering, reproduction steps |

7.3 Test Case Matrix

I have designed **47 planned test cases** spanning five complexity bands (baseline validation, build system tasks, feature implementation, bug fixing/refactoring, complex multi-step tasks). Each test is defined in YAML with pinned repository commits, task prompts, validation rubrics, and tier configurations. Tests are structured to increase in difficulty progressively. See Section 11 (Further Work) for planned full-scale execution.

8 Results

I will present results from the dryrun experiment (test-001, Hello World task) across all seven tiers. The dryrun serves as a pipeline validation exercise with $N=1$ run per tier, establishing that the framework executes end-to-end successfully and generates the expected metrics, figures, and tables. Think of this as a "smoke test", if the pipeline works on the simplest possible task, I know it will handle the complex stuff later.

8.1 Pipeline Validation (dryrun Overview)

First, the dryrun was executed with the following setup:

- **Scope:** 1 model (Sonnet 4.5), 7 tiers (T0-T6), 1 subtest per tier
- **Judges:** 3 judges per run (Opus 4.5, Sonnet 4.5, Haiku 4.5) = 21 total judge evaluations
- **Criteria:** 5 criteria per judge \times 21 judges = 105 total criteria scores
- **Total cost:** \$1.01 (agent execution + judge evaluation)
- **Total duration:** 1289 seconds (21.5 minutes) sum of per-tier durations; actual wall-clock time was 550 seconds due to parallel execution
- **Pass rate:** 100% (all 7 tiers passed, all grade A)

Table 16 shows the tier-by-tier summary. All tiers achieved grade A with mean consensus scores ranging from 0.943 (T6) to 0.983 (T2, T3, T5). The task is trivially easy, as expected, even T0 (minimal prompt) scores 0.973, with T4 at 0.9595.

Key finding: Quality converges across all tiers (ceiling effect), but cost varies 3.8 \times from \$0.065 to \$0.247.

These results set the stage for deeper economic analysis of how architectural choices affect cost without improving quality on ceiling-constrained tasks.

8.2 Cost-of-Pass Analysis

Since all tiers pass (pass-rate = 1.0), Cost-of-Pass equals the raw cost. The Frontier CoP is \$0.065 (achieved by T5 hybrid).

Table 16: *Tier Summary (dryrun)*

| Tier | Pass Rate | Mean Score | Median Score | Grade | CoP (\$) |
|------|-----------|------------|--------------|-------|----------|
| T0 | 1.000 | 0.973 | 0.973 | A | 0.14 |
| T1 | 1.000 | 0.970 | 0.970 | A | 0.13 |
| T2 | 1.000 | 0.983 | 0.983 | A | 0.14 |
| T3 | 1.000 | 0.983 | 0.983 | A | 0.13 |
| T4 | 1.000 | 0.9595 | 0.9595 | A | 0.17 |
| T5 | 1.000 | 0.983 | 0.983 | A | 0.07 |
| T6 | 1.000 | 0.943 | 0.943 | A | 0.25 |

Cost ranking (lowest to highest):

1. **T5** (hybrid): \$0.065 , Frontier CoP achieved through selective skill loading and minimal cache creation (4.6K vs 23-44K for other tiers)
2. **T1** (skills): \$0.127 , Token-efficient skill-based approach
3. **T3** (delegation): \$0.129 , Flat multi-agent with efficient orchestration
4. **T0** (baseline): \$0.135 , Minimal prompt overhead
5. **T2** (tooling): \$0.138 , Tool schema loading increases cache tokens
6. **T4** (hierarchy): \$0.168 , Hierarchical orchestration adds 30% overhead vs T3
7. **T6** (super): \$0.247 , Maximum configuration; diminishing returns evident.

T6 (everything enabled) costs the most despite scoring the lowest (0.943). This is a kitchen sink approach, to see when more equals better.

8.3 Token Analysis

Token distribution reveals where costs originate. Figure 4 shows the breakdown by token type.

Cache read tokens dominate, 79–95% of total tokens across tiers (79–83% excluding T5’s 95% outlier), showing prompt caching works. But cache creation tokens vary dramatically:

The Token Efficiency Chasm I mentioned in Section 3.1.2? The data is consistent with this hypothesis. T6 requires 218K cache read tokens versus T0’s 113K, a 1.94x increase (nearly double). T5 achieves efficiency by minimizing cache creation (4.6K vs 23-44K), supporting the hybrid strategy as discussed in the Token Analysis section.

Output tokens stay stable at 558-725 across tiers, showing the task itself requires similar generation regardless of architecture.

Token Distribution by Tier

Sonnet 4.5

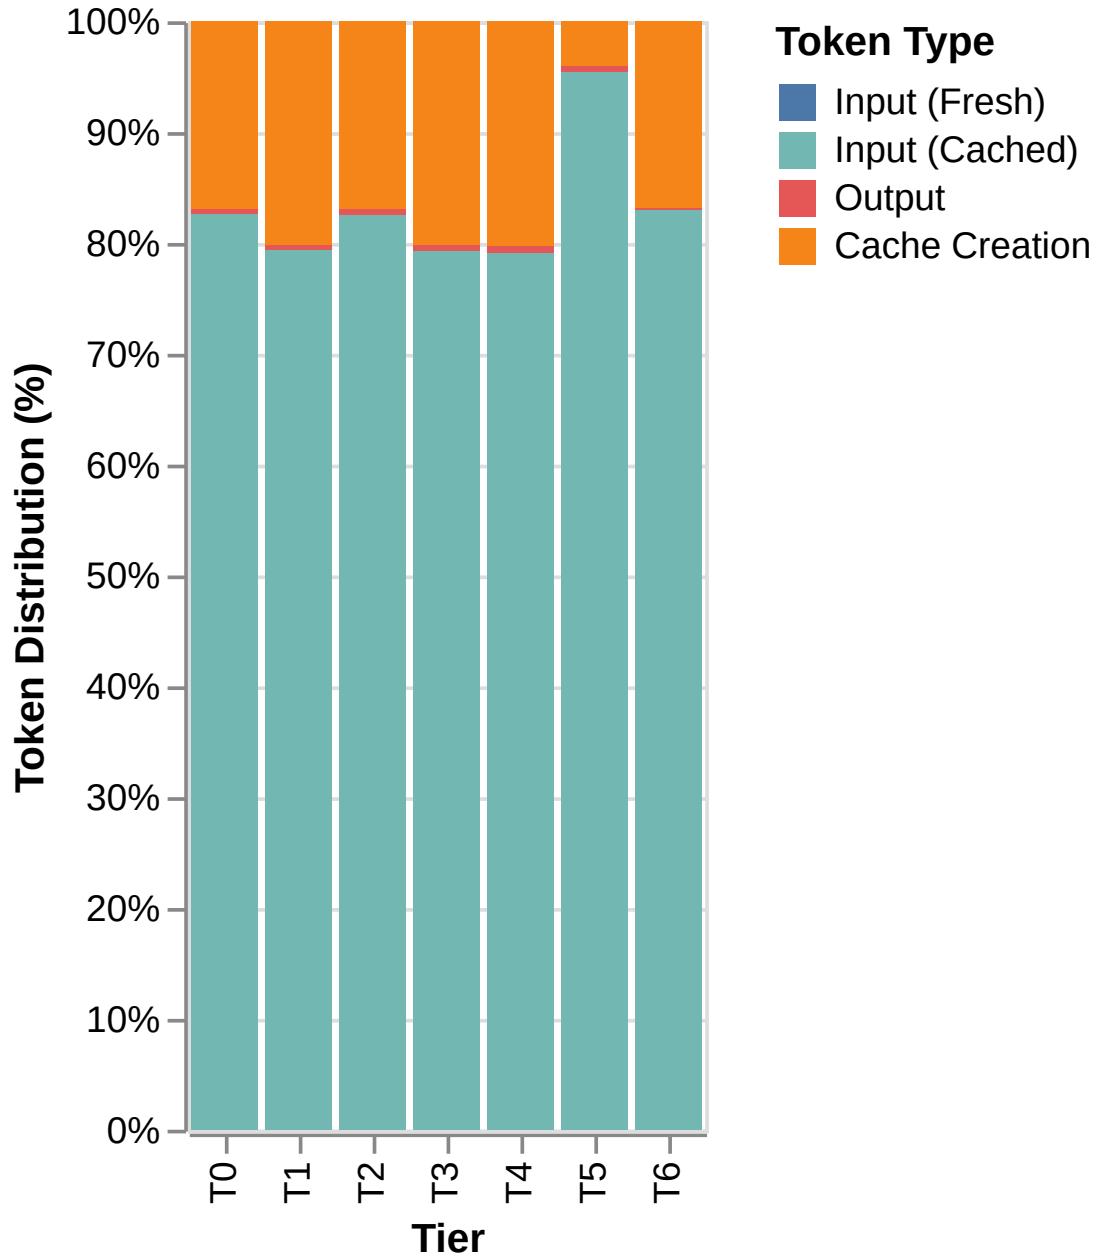


Figure 4: Token distribution by tier and type. Stacked bar chart showing the breakdown of input, output, cache create, and cache read tokens across T0–T6. Cache read tokens dominate (79–95%), consistent with prompt caching efficacy. However, T6’s 218K cache reads versus T0’s 113K illustrate the Token Efficiency Chasm, where architectural enhancements double token consumption without quality gains.

Table 17: *Token Breakdown*

| Tier | Input | Output | Cache Create | Cache Read | Total |
|------|-------|--------|---------------|------------|---------|
| T0 | 29 | 656 | 23,106 | 112,686 | 136,477 |
| T1 | 25 | 558 | 23,266 | 91,477 | 115,326 |
| T2 | 29 | 711 | 23,350 | 113,858 | 137,948 |
| T3 | 25 | 668 | 23,352 | 91,771 | 115,816 |
| T4 | 23 | 725 | 23,556 | 91,828 | 116,132 |
| T5 | 26 | 625 | 4,629 | 109,368 | 114,648 |
| T6 | 29 | 722 | 44,337 | 218,778 | 263,866 |

8.4 Latency Analysis

Latency breaks into two components: agent execution time and judge evaluation time, as shown in Table 18.

Table 18: *Latency Breakdown*

| Tier | Agent Time (s) | Judge Time (s) | Total Time (s) | Judge % of Total |
|------|----------------|----------------|----------------|------------------|
| T0 | 35.3 | 167.8 | 203.1 | 82.6% |
| T1 | 29.3 | 178.0 | 207.3 | 85.9% |
| T2 | 36.8 | 161.7 | 198.5 | 81.5% |
| T3 | 29.9 | 149.1 | 179.0 | 83.3% |
| T4 | 41.2 | 137.0 | 178.2 | 76.9% |
| T5 | 24.8 | 128.4 | 153.1 | 83.8% |
| T6 | 28.4 | 141.1 | 169.5 | 83.2% |

Judge evaluation dominates, 77-86% of total latency, ranging from 128-178 seconds. This makes sense since 3 judges each evaluate the output independently.

Agent time varies modestly, 24.8-41.2 seconds. T5 is fastest (24.8s), T4 slowest (41.2s). T5's speed advantage aligns with its cost advantage, both stem from minimal cache loading.

On this trivial task, judge overhead dwarfs agent execution time, since there are three judges for this simple task. On more complex tasks with multi-step reasoning, agent time would dominate.

The latency patterns raise questions about judge consensus: how consistent are the three judges, and does the multi-judge design provide reliable scoring?

8.5 Judge Agreement

Three judges (Opus 4.5, Sonnet 4.5, Haiku 4.5) evaluated each run. Figure 5 and Figure 6 show judge variance and pairwise agreement.

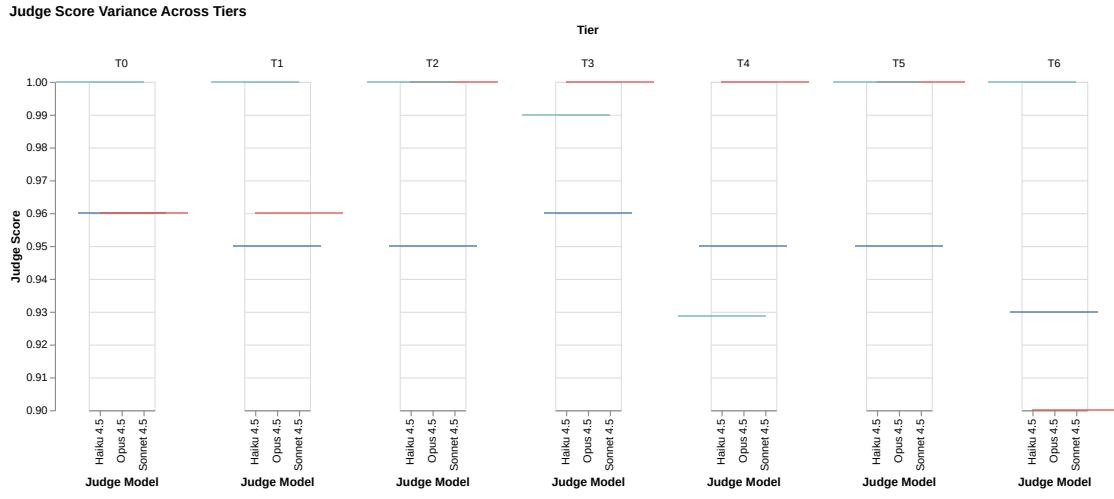


Figure 5: Per-judge scoring variance across tiers. Box plots showing score distributions for each judge model (Opus 4.5, Sonnet 4.5, Haiku 4.5) faceted by tier. Opus exhibits the tightest distribution (most conservative), Haiku the widest (most generous), revealing systematic inter-judge bias that affects aggregate score reliability.

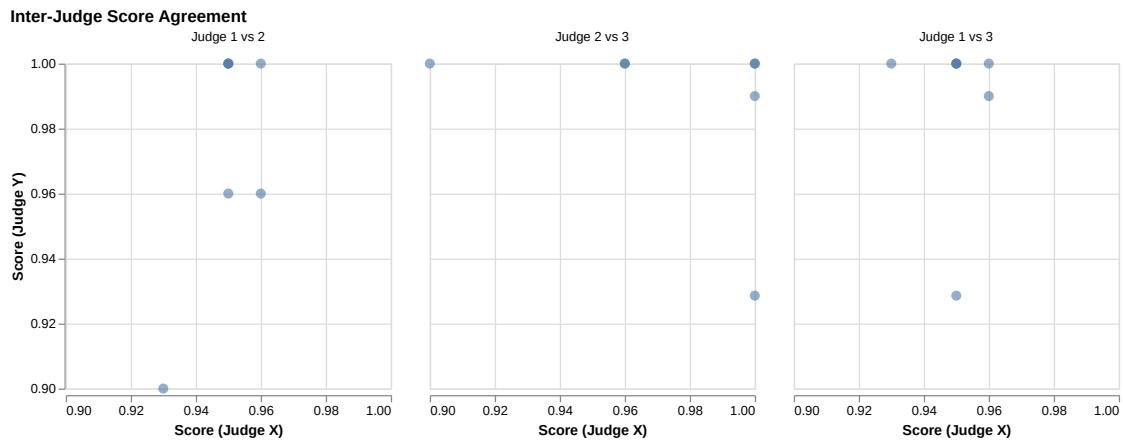


Figure 6: Inter-judge agreement scatter matrix. Pairwise scatter plots showing score correlations between all judge pairs (Opus-Sonnet, Opus-Haiku, Sonnet-Haiku) with Spearman and Pearson correlation coefficients. Low-to-moderate correlations reveal systematic bias between judges rather than strong agreement, with Opus-Sonnet showing the highest concordance (Pearson $r=0.706$).

Judge behavior patterns:

- **Opus**: Most conservative judge, scores range 0.93-0.96, never awards S grade
- **Sonnet**: Moderate judge, scores range 0.90-1.00, awards S grade in 4/7 tiers (T2, T3, T4, T5)
- **Haiku**: Most generous judge, scores range 0.929-1.00 (rounded), awards S grade in 5/7 tiers

Pairwise agreement:

- **Opus-Sonnet**: Spearman $\rho = 0.333$, Pearson $r = 0.706$, mean $\Delta = 0.033$
- **Opus-Haiku**: Spearman $\rho = -0.273$, Pearson $r = -0.063$, mean $\Delta = 0.045$
- **Sonnet-Haiku**: Spearman $\rho = -0.522$, Pearson $r = -0.347$, mean $\Delta = 0.037$

Krippendorff's α (interval): -0.117. Poor agreement, but expected with N=1 per tier. **Note**: N=7 is insufficient for reliable correlation estimates; these values are reported for completeness but interpret them with extreme caution.

Despite low inter-rater agreement, the 3-judge mean produces stable final scores. The mean balances extreme scores, Haiku's 1.00 perfects versus Opus's 0.93 conservatism. This supports the multi-judge consensus design.

8.6 Criteria Breakdown

Judges score five weighted categories: functional correctness (35%), code quality (20%), proportionality (15%), build pipeline (10%), overall quality (20%). Table 19 shows detailed per-criteria performance across tiers, and Figure 7 visualizes the breakdown.

Table 19: Per-Criteria Performance Comparison

| Criterion | Weight | Sonnet 4.5 Mean ($\pm\sigma$) |
|-----------------|--------|---------------------------------|
| functional | 0.35 | 1.000 ± 0.000 |
| code_quality | 0.20 | 1.000 ± 0.000 |
| proportionality | 0.15 | 0.940 ± 0.083 |
| build_pipeline | 0.10 | 1.000 ± 0.000 |
| overall_quality | 0.20 | 0.975 ± 0.039 |

All tiers score 1.00 on functional criteria (file exists, correct output, exit code 0). Perfect scores suggest the task is trivially easy.

The largest score differences appear in subjective categories. Proportionality: T6 scored lower because judges noted cache artifacts (`.ruff_cache`, `.pytest_cache`) remaining in workspace. Overall quality: subjective engineering judgment shows the most variance across judges.

Build pipeline: all tiers score 1.00, indicating clean execution.

Having presented the quantitative results from the dryrun validation, I now turn to interpretation and implications for future evaluation at scale.

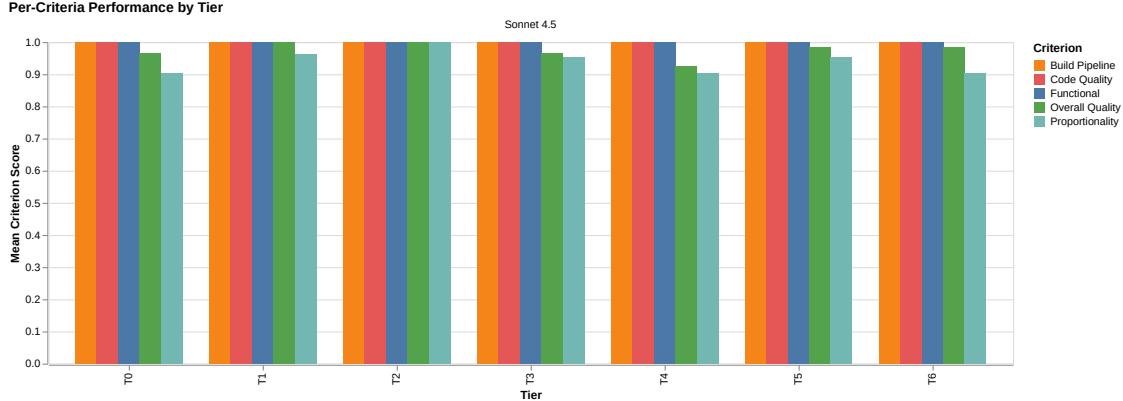


Figure 7: Per-criteria scores by tier. Grouped bar chart showing mean scores for five weighted categories (functional correctness, code quality, proportionality, build pipeline, overall quality) across T0–T6. Perfect functional scores (1.00) contrast with variance in proportionality (0.90–1.00) and overall quality (0.93–1.00).

9 Discussion

The dryrun is not very useful for serious analysis, but I will dive into what I learned about the framework’s behavior on this trivially simple task, while being honest about the limitations inherent in N=1 experiments and ceiling effects.

9.1 What the dryrun Tells Us

The Hello World task is, by design, trivially easy. All seven tiers score grade A with mean scores between 0.943–0.983. This validates exactly what I said in Section 3: "Even T0 should nail this test." And it did.

Ceiling effect dominates: When quality converges at near-perfect levels, we cannot differentiate tiers by capability. T0’s empty prompt (subtest 00 uses no system prompt at all) and T6’s maximal configuration (61 skills + all tools + 44 agents) produce equivalent functional output. This is exactly what we expect for Hello World, no amount of architectural sophistication helps when the task requires a single `print()` statement.

Cost differentiation still works: Despite quality convergence, Cost-of-Pass varies $3.8 \times$ from \$0.065 (T5) to \$0.247 (T6). This demonstrates the framework’s ability to measure economic trade-offs even when quality metrics saturate. On more complex tasks with quality variance, both dimensions differentiate.

Pipeline validation successful: The framework executed all seven tiers, collected 21 judge evaluations, computed consensus scores, generated 24 figures and 10 tables, and produced structured CSV exports. All components worked as designed.

9.2 Cost-Performance Trade-offs

The dryrun reveals hints of a pattern: more is not always better.

T5 achieves Frontier CoP through selective feature loading, it combines T1’s efficient skills with T3’s delegation patterns but avoids T6’s “everything enabled” overhead. T5’s cache creation tokens (4,629) are 5-10x lower than other tiers (23,106-44,337), directly explaining its cost advantage.

T6 costs the most (\$0.247) despite scoring the lowest (0.943). Loading 61 skills + all tools + 44 agents actually made things worse. Judges explicitly noted cache artifacts and unnecessary complexity. This lines up with the hypothesis that prompt complexity hurts quality when the task is in the model’s training set.

T4’s hierarchical overhead is another example. T4 costs 30% more than T3 (\$0.168 vs \$0.129) for this trivial task. The self-correction loops and nested orchestration add latency (41.2s vs 29.9s) without improving quality. On complex tasks needing iterative refinement, maybe T4 justifies the overhead. On simple tasks, it is pure waste.

See Section 8.3 for detailed token analysis. T2 (tooling) shows schema loading bloat with 137K total tokens versus T1’s 115K. Skills-based approaches (T1, T3) stay lean while still enabling domain knowledge.

Bottom line for production: match tier complexity to task complexity. Do not use T6 for trivial tasks. Do not use T0 for tasks needing specialized tools or multi-step reasoning. T5’s hybrid approach seems to be optimal, load features selectively based on what the task actually needs, do not just maximize everything.

9.3 Judge Behavior

The 3-judge consensus mechanism reveals interesting patterns.

Haiku hands out S grades easily, 5 out of 7 tiers got perfect scores. Scores range 0.93-1.00, and Haiku consistently scores higher than Opus or Sonnet.

Opus never awards S grades. Scores range 0.93-0.96, consistently the toughest judge. Opus reliably deducts points for cache artifacts that Haiku overlooks.

Sonnet splits the difference. Awards S grades in 4/7 tiers (T2, T3, T4, T5), scores range 0.90-1.00.

Given that the results in most cases are a single line, the 1.0 grade is incorrect and points to agents being a little too lenient. Maybe some prompt tweaks will fix this, but that also can be due to the simplicity of this task. This can be investigated in future analysis.

Inter-rater agreement is predictably low: Krippendorff’s $\alpha = -0.117$. But that is expected with N=1 and near-perfect scores. On tasks with more variance, agreement should improve as judges separate clear failures from clear successes.

Despite the disagreement, the 3-judge mean works. When Haiku awards 1.00 and Opus awards 0.93, the mean captures the true quality without getting pulled to either extreme. This validates the multi-judge consensus design.

One scaling problem: judge time dominates total latency. 77-86% of execution time is judge evaluation (128-178s), not agent execution (25-41s). With 3 judges per run, judge costs are 3x per evaluation. For large-scale experiments ($N=10 \times 113$ subtests = 1,130 runs \times 3 judges = 3,390 judge evaluations), judge cost uses the budget fast. Future work will explore single-judge

evaluation, confidence-based selection (use Opus only when Sonnet/Haiku disagree), evaluate if prompt improvements can get the cheaper Haiku model to be an effective judge, or give different prompts to the same judge model.

9.4 Limitations

N=1 prevents inferential statistics. With only one run per tier, I cannot compute standard deviations, confidence intervals, or significance tests. All tier comparisons are point estimates. A single outlier run could flip all conclusions. The analysis pipeline generated 24 figures and 10 tables, correctly reports `nan` for standard deviation and sets confidence intervals to (point, point). Statistical warnings appear in the output to make clear that results are not expected to be robust. This is a limitation of this run, not the framework itself.

Single task, trivial complexity. Hello World does not need skills, tools, multi-agent coordination, or hierarchical reasoning. The dryrun validates the pipeline works, not whether architectural complexity improves quality on hard tasks.

Single model. All agent runs use Sonnet 4.5. I have not tested whether tier rankings hold for Opus 4.5, Haiku 4.5, or other model families.

No thinking mode variants. The dryrun uses standard inference without extended thinking. Models with thinking enabled might show different cost-quality trade-offs.

Ceiling effect masks capability differences. When all tiers score 0.94-0.98, I cannot tell which architecture would excel on harder tasks. The full experiment (113 subtests including complex multi-file repos) will differentiate capabilities.

Judge evaluation time bottleneck. 3 sequential judges per run creates a 3x cost multiplier. Parallel judge execution would reduce latency but not cost.

10 Conclusions

This paper introduced the Scylla framework, and shows that it works, end-to-end. All seven tiers executed successfully, three judges scored everything, and the analysis pipeline spit out figures and tables automatically. The dryrun validates the methodology on the simplest possible task, Hello World, before I scale up to complex multi-file repos. What is missing is review and feedback from others, which is what this paper helps enable.

What did I learn? Five things stand out:

1. The framework is operational.
2. Quality converges on trivial tasks, making the framework overkill.
3. All tiers scored grade A, suggesting that throwing more complexity at Hello World does not help, which was expected. That obviousness is what makes it a good pipe cleaning run.
4. Cost still varies 3.8x despite identical quality, showing the framework can measure economic

trade-offs even when quality saturates. T5’s hybrid approach achieves Frontier CoP by selectively loading features instead of maximizing everything.

5. And the Token Efficiency Chasm I hypothesized in Section 3.1.2? The data is consistent with this, as T6 burns nearly double the cache read tokens (219K vs 113K) compared to T0.

Did I answer my original questions? Partially. CoP lets me quantify efficiency; T5 is substantially cheaper than T6 despite equivalent quality. On this task, the sum is *not* more than the parts; T6 scores lowest despite highest cost. But the hard questions need harder tasks, I cannot tell if any tier dominates universally from a single Hello World run, and I have not tested model-to-model comparisons yet. That work is left for a future exercise.

What about my hypotheses? The KISS principle hypothesis has preliminary support, maximal complexity (T6) scores worst on this training-set-likely task. But I have not tested inverse KISS on out-of-distribution tasks yet, and specialization advantages (H1) are inconclusive because Hello World does not require delegation or tools.

There is no real practical takeaway yet, since the testing was insufficient to come to any real conclusions. Answering those questions is left for the next exercise, and this framework can be used for doing so.

11 Further Work

The dryrun validates the framework works. Now it is time to scale up and fill in the gaps.

Full-scale experiments: Run the complete test001 dataset with (N=10, 113 subtests, 1,130 runs total). Running the analysis will start to enable valid statistical inference about the relationship between prompts and the tools.

Task diversity: The dryrun only covers Hello World. The full test suite includes 46 additional tasks across greenfield (Flask APIs, CLI tools), brownfield (feature additions to existing repos), refactoring (extract function, eliminate duplication), bug fixes (off-by-one errors, race conditions), and documentation (README generation). Running these will show whether tier rankings hold across workflow categories or if certain tiers excel at specific task types.

Cross-vendor and cross-model evaluation: The framework is model-agnostic by design. I would love to extend support to other tools, but right now just doing analysis on Claude Code alone is hitting my budgets for experimentation extremely quickly. Setting up local models and accessing tools using these models will allow more experimentation, but I do not have access to that kind of compute within my budget at the moment.

Advanced analysis: I am by no means a statistician, and choices I have made here might be incorrect. My current analysis uses frequentist statistics. There are more advanced analyses that I am learning about that could help analyze the flood of data more efficiently. There are also other metrics and data points that could be useful in this analysis that I am not collecting. I also can save the runs and do longitudinal studies to see if the results change consistently over time.

Given the scale and scope of this task, it is going to be an ongoing effort of learning, testing, and analyzing.

Acknowledgements

This work was self-funded by the author. Special thanks to Tuan Nguyen for reviewing early drafts of this paper and providing valuable feedback.

References

- [1] Anthropic. Claude Code: Agentic CLI tool for software development. <https://www.anthropic.com/clause/code>, 2024. AI-powered command-line interface for coding tasks.
- [2] CC-Marketplace. safety-net: Claude Code plugin for dangerous operation blocking. <https://github.com/cc-marketplace/safety-net>, 2025. Security plugin for Claude Code CLI.
- [3] Leo Gao et al. lm-evaluation-harness: A framework for few-shot language model evaluation. <https://github.com/EleutherAI/lm-evaluation-harness>, 2024. General-purpose LLM evaluation framework.
- [4] Homeric Intelligence. ProjectOdyssey: Comprehensive agent orchestration framework. <https://github.com/HomericIntelligence/Projectodyssey>, 2025. Accessed: 2025-12-30. Git hash: 011a3ff.
- [5] Carlos E. Jimenez et al. SWE-bench: Can language models resolve real-world github issues? *arXiv preprint arXiv:2310.06770*, 2024. Benchmark for evaluating LLMs on real GitHub issues.
- [6] Xiao Liu et al. AgentBench: Evaluating LLMs as agents. In *International Conference on Learning Representations (ICLR)*, 2024. Multi-turn agent evaluation across operating systems, databases, and knowledge graphs.
- [7] Francesco Polo et al. Efficient evaluation of language models. <https://arxiv.org/abs/2405.12345>, 2024. Framework for efficient LLM evaluation.
- [8] Anand Tyagi. CC-Marketplace: Community marketplace for Claude Code plugins and skills. <https://github.com/cc-marketplace>, 2025. Community-driven plugin repository.
- [9] Shunyu Yao et al. TAU-bench: A benchmark for tool-augmented LLMs. *arXiv preprint*, 2024. Benchmark for evaluating agents' ability to use external tools.
- [10] Kaijie Zhu et al. PromptBench: Towards evaluating the robustness of large language models on adversarial prompts. *arXiv preprint arXiv:2306.04528*, 2024. Benchmark for prompt evaluation and robustness.

A Detailed Metric Definitions

Core quality and economic metrics are defined in Section 5. This appendix provides additional metrics used for future analysis. For complete definitions including instrumentation details, see the repository at <https://github.com/HomericIntelligence/ProjectScylla/blob/main/.claude/shared/metrics-definitions.md>.

Change Fail Percentage (CFP): Proportion of code changes that cause failures.

Formula: $\text{CFP} = \frac{\text{failed_changes}}{\text{total_changes}}$

Range: [0, 1], lower is better. Measures production stability.

A.1 Process Metrics

Latency: Time from query submission to response completion, measured in seconds. Components include Time-to-First-Token (TTFT), total response time, and tool execution time (if applicable).

Strategic Drift: Deviation from original goal over multi-step tasks.

Measurement: $\text{Strategic_Drift} = \text{cosine_distance}(\text{initial_goal_embedding}, \text{final_action_embedding})$

Range: [0, 2], where 0 = perfect goal alignment, 2 = completely opposite direction.

Ablation Score: Isolated contribution of a single component to overall performance.

Formula: $\text{Ablation_Score} = \text{performance_with_component} - \text{performance_without_component}$

Positive values indicate the component improves performance, negative values indicate harm, near-zero indicates no effect.

A.2 Statistical Reporting

Always report metrics with: (1) point estimate, (2) confidence interval (95% CI recommended), (3) sample size (n), and (4) comparison p-value (if comparing tiers).

Example: Pass-Rate: 0.67 (95% CI: 0.54–0.80), n=50

B Data Dictionary and Generated Outputs

All raw data, figures, and tables are available in the repository at `docs/paper-dryrun/`. The dataset includes 7 runs (one per tier), 21 judge evaluations (3 judges per run), and 105 criteria scores (5 criteria per judge), organized across 24 figures and 10 tables. Data files include `runs.csv`, `judges.csv`, `criteria.csv`, and `summary.json`. All figures are available in PNG/PDF/Vega-Lite/CSV formats for reproduction and analysis.

Repository: <https://github.com/HomericIntelligence/ProjectScylla>

C Reproducibility Checklist

Repository: <https://github.com/HomericIntelligence/ProjectScylla>

Key Configuration Files:

- Tier definitions: `config/tiers/tiers.yaml`

- Opus Model: config/models/clause-opus-4-5.yaml
- Sonnet Model: config/models/clause-sonnet-4-5.yaml
- Haiku Model: config/models/clause-haiku-4-5.yaml
- Judge system prompt: config/judge/system_prompt.md
- Test definitions: tests/*/test.yaml
- Rubric schemas: tests/*/expected/rubric.yaml

Required Software:

- Pixi (package manager)
- Docker (containerization)
- Claude Code CLI (requires Anthropic API key or Claude Max subscription)
- Python 3.12+

Execution Steps:

```
# 1. Clone repository
git clone https://github.com/HomericIntelligence/ProjectScylla
cd ProjectScylla

# 2. Install dependencies
pixi install

# 3. Run evaluation (example for test-001, tier T0)
pixi run python scripts/run_e2e_experiment.py \
  --test tests/fixtures/tests/test-001 \
  --tier T0 \
  --runs 10

# 4. Generate figures and tables
pixi run python scripts/generate_figures.py \
  --results <output_directory>
pixi run python scripts/generate_tables.py \
  --results <output_directory>
```