Navigating Scylla with Haiku: Multi-Task Evaluation of Agentic Coding Architectures Using a Budget Frontier Model

Micah Villmow*
research@villmow.us

April 2026

Abstract

Agentic coding systems increasingly employ multi-layered architectural patterns—system prompts, domain skills, external tooling, agent delegation, and hierarchical orchestration—yet the marginal value of each layer remains poorly characterized, particularly for budget-class language models. We present a controlled ablation study using the Scylla evaluation framework, which decomposes agent architecture into seven tiers of increasing complexity (T0–T6) and measures quality and cost across a structured set of 120 YAML subtests. Using `claude-haiku-4-5`, Anthropic's budget frontier model, we execute three multi-task experiments totaling 1,080 runs at a cost of $123.30. Contrary to initial expectations, our results show no significant relationship between architectural complexity and aggregate task performance: all seven tiers achieve high pass rates in the range 0.935–1.000 (T0 = 0.935, T1 = 0.967, T2 = 0.978, T3 = 0.976, T4 = 0.976, T5 = 0.963, T6 = 1.000), and a Kruskal–Wallis omnibus test on pass rate is not significant ($H(6) = 8.53$, $p = 0.2015$). However, a Scheirer–Ray–Hare analysis reveals a significant tier × task interaction on quality score ($H(12) = 223.5$, $p < 0.001$), indicating that the optimal architectural tier is task-dependent even when aggregate pass rates are equivalent. Cost does not vary significantly across tiers ($p = 0.676$), providing no evidence that additional architectural complexity yields economic benefits or penalties at this model scale. Inter-rater reliability among the three-judge panel remains poor (Krippendorff's $\alpha = 0.034$), raising concerns about the reliability of LLM-as-judge panels that include budget-class evaluators. These preliminary findings suggest that for budget frontier models, architectural tier selection has negligible effect on aggregate pass rate, but task difficulty and judge reliability are substantial sources of variance that dominate architectural effects.

Keywords: agentic AI, ablation study, LLM evaluation, budget models, multi-agent systems, cost-of-pass, Claude Haiku

*This paper combines the author's original research and writing with improvements, suggestions, and rewrites provided by Claude Code (claude.ai/code).

1 Introduction

The rapid evolution of large language model (LLM) agent systems has produced a proliferating taxonomy of architectural patterns. Modern agentic coding frameworks layer system prompts, domain-specific skills, external tool integrations, multi-agent delegation, and hierarchical orchestration into increasingly complex configurations [1, 14]. Practitioners face a fundamental question: *which of these architectural layers actually improve performance, and at what cost?*

This question is especially acute for budget-class models. While frontier models such as Claude Opus or GPT-4 may possess sufficient capacity to benefit from complex architectures, smaller models may lack the reasoning bandwidth to exploit—or even tolerate—the overhead of delegation and orchestration. Whether architectural complexity helps, hinders, or has no effect on a budget model is an empirical question, yet the interaction between model capacity and architectural complexity remains largely unexplored in the existing literature.

The stakes are practical. Organizations deploying agentic coding systems face real economic trade-offs: budget models cost substantially less per token than frontier models, but it is unknown whether layering additional architectural complexity onto a budget model provides any measurable benefit. If architectural tier has negligible impact on pass rate—as our results suggest—then practitioners can safely deploy the simplest viable architecture, reducing operational complexity without sacrificing quality.

Existing benchmarks such as SWE-bench [11], AgentBench [14], and TAU-bench [23] evaluate agents on fixed tasks but do not systematically vary the agent's architectural configuration. They answer the question "how well does agent X perform?" but not "which architectural components of agent X contribute to its performance?" The ablation methodology needed to answer the latter question requires a factorial experimental design that holds the model constant while varying the architecture across a controlled hierarchy of configurations.

The Scylla framework [22] provides exactly this capability. Named after the mythic trial from Homer's *Odyssey*, Scylla represents the challenge of navigating trade-offs between capability gains and operational costs. The framework decomposes agent architecture into seven hierarchical tiers, from bare prompts (T0) through full orchestration (T6), and evaluates each tier on identical tasks using identical models. A prior dryrun study applied this framework to Claude Sonnet 4.5 on a single task, establishing the methodology but leaving multi-task generalization as an open question.

In this paper, we extend the Scylla framework to a budget frontier model (`claude-haiku-4-5`) and a multi-task evaluation across three experiments of varying difficulty. We execute all 120 subtests across all seven tiers with $N = 3$ replications per subtest, producing 1,080 total runs at a total cost of $123.30. This design enables us to address the following research questions:

RQ1 Tier–Quality Relationship. How does architectural complexity (T0–T6) affect pass rate and quality score when using a budget frontier model?

RQ2 Task Interaction. Does the optimal tier vary across tasks of different difficulty, or is there a universally superior architecture?

RQ3 Cost Efficiency. How does Cost-of-Pass (CoP) vary across tiers, and does increased complexity yield economic benefits?

RQ4 Judge Reliability. How reliable is a multi-model LLM-as-judge panel when evaluating budget-model outputs?

Our contributions are as follows:

1. We provide, to our knowledge, the first multi-task ablation study of agentic coding architectures using a budget frontier model, finding that architectural tier has no statistically significant effect on aggregate pass rate ($H(6) = 8.53$, $p = 0.2015$), with all tiers achieving pass rates between 0.935 and 1.000.
2. We show that task-dependent effects dominate architectural effects: a Scheirer–Ray–Hare analysis reveals a significant tier $\times$ task interaction on quality score ($H(12) = 223.5$, $p < 0.001$), indicating that architectural decisions should be task-contingent rather than universally applied.
3. We find no statistically significant tier effect on cost ($p = 0.676$), providing no evidence that additional complexity yields economic benefits or penalties at this model scale.
4. We report poor inter-rater reliability (Krippendorff's $\alpha = 0.034$ [12]) among a three-model judge panel, highlighting the fragility of LLM-as-judge evaluation for budget-model outputs and suggesting that judge variance is a primary source of measurement noise.
5. We establish an aggregate frontier Cost-of-Pass of ~$0.11, showing that meaningful ablation studies can be conducted at modest cost with budget models.

The remainder of this paper is organized as follows. Section 2 surveys related work in agent evaluation and ablation methodology. Section 3 describes the Scylla framework and its tier taxonomy. Section 4 details our experimental design. Section 5 presents results per experiment and in aggregate. Section 6 analyzes inter-rater reliability. Section 7 examines cost efficiency. Section 8 discusses key findings, limitations, and threats to validity. Section 9 concludes, and Section 10 outlines future directions.

2 Related Work

2.1 Agent Benchmarks

The evaluation of LLM-based agents has received increasing attention as agent systems have moved from research prototypes to production deployments. SWE-bench [11] evaluates agents on real-world GitHub issues drawn from popular open-source repositories, providing ecologically valid but uncontrolled task sets. The benchmark has become a standard reference point for agent capability, with leaderboard performance tracked across multiple agent systems. However, SWE-bench treats the agent's internal architecture as a black box, measuring only the final outcome without decomposing the contribution of individual architectural components.

AgentBench [14] offers a broader multi-domain evaluation across operating systems, databases, knowledge graphs, and web browsing. Its multi-turn interaction design captures aspects of agent

behavior that single-turn benchmarks miss, including error recovery, context management, and strategic planning. Like SWE-bench, however, AgentBench evaluates complete agent systems without varying their internal configuration.

TAU-bench [23] focuses specifically on tool-augmented LLMs, measuring the agent's ability to select and compose external tools. While TAU-bench does vary the tools available to the agent, it does not extend this ablation to other architectural dimensions such as prompting strategy, delegation, or hierarchical control.

PromptBench [24] evaluates prompt robustness but operates at the level of individual prompts rather than complete agent architectures. The lm-evaluation-harness [7] provides a general-purpose framework for few-shot LLM evaluation but is not designed for agentic workflows that require multi-turn interaction, tool use, and code generation. Polo et al. [16] propose efficient multi-prompt evaluation methods that reduce the number of evaluations needed to characterize model behavior, but their focus is on prompt selection rather than architectural ablation.

A common limitation of these benchmarks is that they evaluate agent *systems* holistically, without providing a principled decomposition of architectural contributions. Our work addresses this gap by applying a controlled ablation design that isolates the effect of each architectural layer while holding the model constant.

2.2 Ablation Methodology in Machine Learning

Ablation studies have a long history in machine learning, originating in neuroscience where they refer to the systematic removal of brain regions to understand function localization. In deep learning, ablation studies systematically remove or modify components of a model or system to measure their individual contributions. This methodology has been applied to neural network architectures, loss functions, data augmentation strategies, and training procedures.

However, ablation methodology has been underutilized in the evaluation of LLM agent systems. Most agent evaluations compare *complete systems* against each other (System A vs. System B) rather than decomposing a single system into its constituent components. The Scylla framework's tier-based design brings ablation methodology to agent architecture, treating each architectural layer (prompts, skills, tools, delegation, hierarchy) as a component that can be systematically added or removed.

Our study extends this approach to the budget model setting, where the interaction between model capacity and architectural complexity introduces additional considerations not present in frontier model evaluations. The ablation is not merely asking "does this component help?" but "does this component help *given the model's capacity to use it*?"

2.3 Multi-Agent Architectures

The multi-agent paradigm has been explored extensively in recent work, with systems employing specialist agents, hierarchical delegation, and orchestration patterns [1]. These architectures are motivated by the intuition that complex tasks benefit from division of labor: a coding task might be decomposed into planning, implementation, testing, and review subtasks, each handled by a

specialized agent.

While multi-agent architectures often improve performance for frontier models [8], their interaction with model capacity has received limited systematic study. Recent work suggests that smaller models struggle with the coordination overhead of multi-agent systems [5], but this has not been rigorously quantified in a controlled ablation design.

Our work contributes to this area by providing the first controlled ablation study across seven architectural tiers for a budget model, finding that tier selection has no statistically significant effect on pass rate ($p = 0.2015$), though strong tier $\times$ task interactions indicate that architectural decisions remain task-contingent.

2.4 LLM-as-Judge Evaluation

The use of LLMs as automated evaluators has become standard practice in agent evaluation, driven by the high cost and limited scalability of human evaluation. Prior work has identified several reliability concerns, including position bias (preference for outputs presented first), verbosity bias (preference for longer outputs), and self-preference effects (models rating their own outputs more highly) [16].

Most studies of LLM-as-judge reliability have focused on frontier models evaluating frontier model outputs. The reliability of judge panels that include budget models, or that evaluate budget model outputs, has received less attention. Our study extends this literature by reporting inter-rater reliability statistics for a three-model judge panel spanning the full capability spectrum, finding near-chance agreement levels that raise fundamental questions about the validity of LLM-as-judge evaluation in budget-model contexts.

2.5 Cost-Aware Evaluation

The economics of LLM deployment have become increasingly important as organizations scale agentic workflows from prototypes to production. Token costs, latency, and cost-per-successful-completion are all relevant to deployment decisions, yet most benchmarks report only quality metrics without cost normalization.

The Cost-of-Pass (CoP) metric [22] provides a task-normalized measure of economic efficiency by dividing mean cost per run by pass rate. This formulation captures a crucial insight: a model that costs half as much but passes only a quarter as often is actually *more* expensive to use in practice. Our work applies CoP analysis across seven architectural tiers, finding that observed cost invariance across tiers ($p = 0.676$, non-significant under limited power) suggests simpler architectures may be Pareto-dominant when they achieve higher quality.

3 Background: The Scylla Framework

3.1 Design Philosophy

The Scylla framework is designed around the principle that agent architecture can be decomposed into a hierarchy of capabilities, where each tier adds a specific category of functionality to the previous tier. This hierarchical design enables controlled ablation: by comparing performance across tiers, we can isolate the marginal contribution of each capability layer.

The framework name invokes the mythic trial from Homer's *Odyssey*, where Odysseus must navigate between the six-headed monster Scylla and the whirlpool Charybdis. In the context of agent evaluation, Scylla represents the trade-off between capability (adding more architectural complexity to improve performance) and cost (the economic and quality penalties of unnecessary complexity). The goal is to find the optimal passage—the architectural configuration that maximizes quality while minimizing cost.

3.2 Tier Taxonomy

The Scylla framework organizes agent architectures into seven tiers of increasing complexity. Each tier extends the capabilities of the previous tier while maintaining backward compatibility with the evaluation protocol. The tiers are defined as follows:

T0 – Prompts (24 subtests) Ablation of system prompt content, ranging from an empty prompt to the full `CLAUDE.md` configuration sourced from ProjectOdyssey [10]. This tier isolates the effect of prompt engineering on model behavior. The 24 subtests span a spectrum from no system prompt at all through progressively richer prompt configurations, enabling fine-grained analysis of prompt contribution.

T1 – Skills (10 subtests) Addition of domain-specific expertise through installed skill modules, organized by category (e.g., code review, testing, documentation, deployment). Skills provide structured knowledge and procedural guidance without external tool access. The 10 subtests evaluate different skill categories and combinations.

T2 – Tooling (15 subtests) Introduction of external tools and Model Context Protocol (MCP) servers. This tier measures the contribution of tool access (file systems, web search, code analysis, linting) beyond prompt and skill content. The 15 subtests evaluate different tool configurations and combinations.

T3 – Delegation (41 subtests) Flat multi-agent configurations with specialist agents. The primary agent can delegate subtasks to specialist agents but does not employ hierarchical control. This is the largest tier, with 41 subtests reflecting the combinatorial space of specialist configurations (different specialist roles, delegation strategies, and task decomposition approaches).

T4 – Hierarchy (14 subtests) Nested orchestration with orchestrator agents managing teams of specialists. This tier tests whether hierarchical control improves coordination over flat delegation. It comprises 7 hierarchy subtests (varying the depth and structure of the hierarchy) and 7 agent-team subtests (varying team composition and role assignment).

T5 – Hybrid (15 subtests) Best-of combinations and permutations from lower tiers. This tier explores whether selective combination of architectural elements outperforms uniform tier application. The 15 subtests represent curated configurations that mix elements from T0–T4 based on prior performance data.

T6 – Super (1 subtest) Everything enabled at maximum capability. This tier represents the upper bound of architectural complexity, with all prompts, skills, tools, delegation, hierarchy, and hybrid configurations active simultaneously. The single subtest provides a ceiling estimate.

Figure 1 illustrates the hierarchical relationship between tiers.

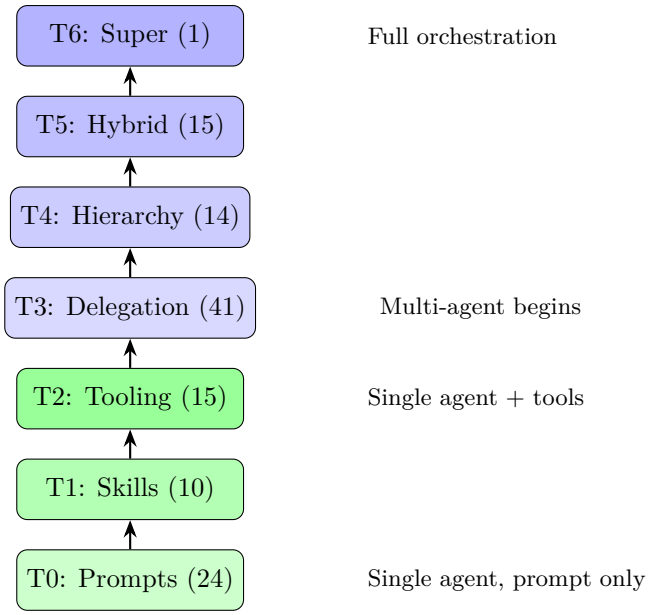

Figure 1: *Scylla tier hierarchy. Green tiers (T0–T2) represent single-agent configurations; blue tiers (T3–T6) introduce multi-agent coordination. All tiers achieve comparable pass rates (0.935–1.000) in this study.*

3.3 Evaluation Protocol

Each subtest defines a coding task with explicit success criteria. The Scylla evaluation protocol proceeds as follows:

1. **Configuration.** The agent receives the architectural configuration defined by the tier and subtest, including system prompt content, skill definitions, tool access, and delegation/hierarchy parameters.
2. **Task Specification.** The agent receives a task description, target repository, and success criteria. The task specification is identical across all tiers for a given experiment.
3. **Execution.** The agent executes the task using the Claude Code CLI [1], producing code changes, documentation, and/or configuration modifications. Execution is bounded by a timeout to prevent runaway sessions.

4. **Evaluation.** A three-judge panel evaluates the output against predefined criteria. Each judge independently evaluates the output against weighted criteria, producing a per-judge score on a $[0, 1]$ scale (see Section 4.4 for details).

5. **Consensus.** The consensus score is computed as the arithmetic mean of the three judge scores. Implementation rate uses the median across judges.

6. **Classification.** A run is classified as "passing" if a majority of the judge panel independently classifies it as passing based on their individual evaluations.[1]

3.4 Partial-Failure Semantics

A key design decision in Scylla is its partial-failure model. When multiple tiers run in parallel within an experiment, a tier failure does not abort the experiment. The experiment continues with the remaining tiers and can reach a "complete" state even if some tiers are in a "failed" state. Each tier's results are recorded independently, preserving partial results for analysis even when individual configurations encounter infrastructure errors. Operators must therefore check per-tier states rather than the overall experiment state to determine whether all tiers succeeded.

3.5 Metrics

Scylla evaluates agent performance across three categories of metrics:

Quality Metrics. Pass Rate (R_{pass}) measures the fraction of runs exceeding the pass threshold. Implementation Rate (R_{impl}) measures the proportion of rubric points achieved versus the maximum available, using the median across judges.[2] The Fine-Grained Progress Rate captures partial task completion on a continuous scale.

Economic Metrics. Cost-of-Pass (CoP) is defined as the mean cost per run divided by the pass rate:

$$\text{CoP} = \frac{\overline{C}}{R_{\text{pass}}} \tag{1}$$

where $\overline{C}$ is the mean cost per run (in USD) and R_{pass} is the pass rate. CoP captures the expected cost to obtain a passing result, penalizing both expensive runs and low pass rates. When $R_{\text{pass}} = 0$, CoP is undefined (infinity), indicating that no amount of spending produces a successful outcome. Token Distribution tracks the breakdown of costs across input, output, cache read, and cache creation operations.

[1] The pass classification uses majority vote over each judge's binary pass/fail decision rather than a threshold on the consensus score. Of 1,080 runs, 16 (1.5%) show the majority-vote outcome diverging from a naïve consensus-score > 0.5 threshold—for example, a run with consensus score 0.487 passes because two of three judges passed individually, while a run with consensus score 0.635 fails because only one of the two judges whose evaluations were available for that run classified it as passing (the third judge's evaluation was missing). These discrepancies do not materially affect tier-level pass rates.

[2] Implementation rate is computed per judge as $\sum \text{achieved}_i / \sum \max_i$ across all criteria i in that judge's evaluation, then the consensus uses the median across judges. In one run (test-001, T1, subtest 01, run 1), a budget-model judge (Haiku) assigned 4 points on a 3-point code quality criterion, producing a per-judge implementation rate of 1.071. Combined with a missing Opus evaluation for that run (leaving only two judges), the median implementation rate is 1.028—the only value exceeding 1.0 in the dataset. This anomaly does not materially affect any aggregate statistics.

Process Metrics. Duration measures wall-clock execution time. Consistency captures the coefficient of variation of scores within a subtest, measuring reproducibility. Strategic Drift measures deviation from task objectives over the course of execution.

4 Experimental Design

4.1 Model Selection

We select `claude-haiku-4-5` as the subject model for all experiments. Haiku 4.5 is Anthropic's budget frontier model—the most cost-efficient option in the Claude model family, positioned below Sonnet and Opus in the capability hierarchy. It is designed for high-throughput, low-cost inference while maintaining sufficient capability for coding tasks.

This selection is motivated by three considerations. First, budget models are the most likely to be deployed at scale in production environments, making their architectural sensitivities practically important for deployment decisions. Second, the prior Scylla dryrun [22] used Claude Sonnet 4.5, enabling cross-model comparison in future work. Third, the lower per-token cost of Haiku ($1.00/MTok input, $5.00/MTok output for standard pricing, with prompt caching at $1.25/MTok for cache writes and $0.10/MTok for cache reads)[3] enables a larger experimental matrix at a fixed budget.

4.2 Task Selection

We design three experiments of varying difficulty and domain to test the generality of tier effects across task characteristics:

test-001: Python Hello World A trivial task requiring the creation of a "Hello World" program in a simple Python repository. This task serves as a floor calibration: any functional agent architecture should achieve near-perfect performance. The minimal complexity of the task makes it a sensitive detector of architectural overhead—if a tier cannot solve this task, the failure is attributable to the architecture, not the task. *Difficulty*: trivial. *Repository*: Hello-World. *Language*: Python. *Expected outcome*: This task is expected to establish a performance floor—all tiers should achieve near-perfect results on this trivial task, and any observed failures are attributable to architectural overhead rather than task difficulty.

test-002: Mojo Hello World with Bazel and Documentation A moderate task requiring a "Hello World" program in the Mojo programming language within the official Modular repository, including Bazel build configuration and accompanying documentation. This task tests the agent's ability to work with a less common programming language and a complex build system, both of which may have limited representation in the model's training data. *Difficulty*: moderate. *Repository*: modular. *Language*: Mojo. *Expected outcome*: This task is expected to differentiate tiers—the unfamiliar programming language and complex build system may reveal capability gaps between architectural configurations, particularly for tiers with limited tooling or domain knowledge support.

[3] Anthropic pricing retrieved April 2026; prices are subject to change.

test-003: Add mypy to pixi Configuration A moderate task requiring the addition of mypy type checking to an existing pixi-managed Python/Mojo project [10]. This task tests the agent's ability to modify existing project configurations, integrate new development tooling, and work within the constraints of an established codebase. *Difficulty*: moderate. *Repository*: ProjectOdyssey. *Language*: Python/Mojo. *Expected outcome*: This task is expected to test integration capability—modifying existing configurations requires understanding project context, dependency relationships, and tooling conventions, which may favor tiers with richer contextual support.

The three tasks were selected to span a range of dimensions: language familiarity (Python vs. Mojo), task type (greenfield creation vs. configuration modification), and repository complexity (minimal vs. full project structure). While all three tasks are relatively simple by software engineering standards, they provide sufficient variation to test whether tier effects are task-dependent.

4.3 Experimental Configuration

Each experiment executes all 120 subtests across all seven tiers (T0–T6) with $N = 3$ replications per subtest. The total experiment matrix is:

- 3 experiments $\times$ 120 subtests $\times$ 3 runs = 1,080 total runs
- 7 tiers: T0 (24 subtests), T1 (10), T2 (15), T3 (41), T4 (14), T5 (15), T6 (1)
- Subject model: `claude-haiku-4-5` (all tiers, all experiments)
- Judge panel: `claude-opus-4-6`, `claude-sonnet-4-6`, `claude-haiku-4-5`
- Runs per subtest per experiment: $N = 3$

The per-tier sample sizes across all three experiments are: T0 ($n = 216$), T1 ($n = 90$), T2 ($n = 135$), T3 ($n = 369$), T4 ($n = 126$), T5 ($n = 135$), T6 ($n = 9$). The unequal sample sizes reflect the varying number of subtests per tier, which in turn reflects the varying combinatorial complexity of each architectural category.

Table 1: *Experiment Configuration (Data-Derived)*

Experiment	Models	Tiers	Tier IDs	ST/Tier	Runs/ST	Total	Judges
test-001	claude-haiku-4-5	7	T0, T1, T2, T3, T4, T5, T6	1-41	3	360	N/A
test-002	claude-haiku-4-5	7	T0, T1, T2, T3, T4, T5, T6	1-41	3	360	N/A
test-003	claude-haiku-4-5	7	T0, T1, T2, T3, T4, T5, T6	1-41	3	360	N/A

4.4 Judge Panel Design

Evaluation uses a three-judge panel spanning the full Anthropic model hierarchy:

1. **Judge 1**: `claude-opus-4-6` — the frontier model with the highest capability, serving as the most authoritative evaluator.
2. **Judge 2**: `claude-sonnet-4-6` — a mid-tier model providing an intermediate evaluation perspective.
3. **Judge 3**: `claude-haiku-4-5` — the budget model, identical to the subject agent, enabling analysis of self-evaluation bias.

The consensus score is computed as the arithmetic mean of the three judges' scores. Implementation rate uses the median across judges' binary assessments of task attempt. This panel design was chosen to provide evaluation diversity while enabling analysis of the relationship between judge capability and evaluation reliability.

The inclusion of the subject model (Haiku) as one of its own judges is a deliberate design choice. It enables us to test for self-preference bias (whether the model rates its own outputs more or less favorably than external judges) and to assess whether budget models possess sufficient evaluative capacity to serve as reliable judges.

Each judge receives a structured evaluation prompt containing four components: (1) the original task prompt issued to the agent, (2) the agent's code output (files created or modified), (3) build pipeline results (compilation status, test outcomes, linting output), and (4) a structured rubric defining five evaluation criteria with associated weights. The criteria and their weights are: *functional correctness* (0.35), *code quality* (0.20), *build pipeline* (0.10), *proportionality* (0.15), and *overall quality* (0.20). Each judge scores every criterion on a 0–10 scale; the weighted sum produces a single per-judge score, and the consensus score is the arithmetic mean across the three judges.

In practice, the number and naming of criteria varies across experiments and judges. test-001 (Python Hello World) consistently produces the standard five-criterion rubric across all judge evaluations. test-002 (Mojo with Bazel), however, elicits task-specific criteria from the judges: individual evaluations contain between 4 and 12 scored dimensions, with criterion names reflecting Mojo-specific concerns (e.g., *memory_safety*, *mojo_syntax*, *deprecated_patterns*). test-003 (mypy configuration) collapses to predominantly single-criterion evaluations. This variability arises because the judge models interpret the rubric prompt with varying granularity depending on task complexity. The per-judge score and implementation rate are computed from whatever criteria the judge produces, so the aggregate metrics remain well-defined even when the underlying rubric structure differs. However, the per-criterion analysis in Table 14 is restricted to the five standard criteria and therefore reflects primarily test-001 data.

This study's judge panel represents an architectural change from the prior Scylla dryrun [22]. The dryrun used Claude Sonnet 4.5 as the subject model with 4.5-series judges (Opus 4.5, Sonnet 4.5, Haiku 4.5). The present study upgrades to 4.6-series judges—Opus 4.6 and Sonnet 4.6—while retaining Haiku 4.5 as both subject and judge. This creates a mixed-generation judge panel in which two judges (Opus 4.6 and Sonnet 4.6) belong to a newer model generation than the subject agent (Haiku 4.5). The generational gap may introduce evaluation asymmetries: newer judges may have improved calibration or different scoring baselines relative to the subject's generation. We note this design limitation and examine judge-level score distributions in Section 6.2 to assess whether generational differences materially affect evaluation reliability.

4.5 Statistical Methods

Analysis of Variance (ANOVA) is a family of statistical tests that partition observed variance into components attributable to different factors. Because our data violates ANOVA's normality assumptions (Appendix C), we employ non-parametric alternatives throughout.

All score distributions are tested for normality using the Shapiro–Wilk test. Of 14 tier × {Score, Cost} combinations, 13 reject normality ($p < 0.05$; the sole exception is T6 Cost with $n = 9$, $p = 0.069$; see Appendix C), so we employ non-parametric methods throughout. Additional metrics (implementation rate, duration) were also tested and similarly reject normality for most tier combinations.

The following subsections detail each statistical method used in this study, including the test statistic, assumptions, and interpretation guidelines.

4.5.1 Shapiro–Wilk Test

The Shapiro–Wilk test [20] assesses whether a sample is drawn from a normally distributed population. The test statistic is:

$$W = \frac{\left(\sum_{i=1}^{n} a_i x_{(i)}\right)^2}{\sum_{i=1}^{n}(x_i - \bar{x})^2} \tag{2}$$

where $x_{(i)}$ are the order statistics and a_i are tabulated constants derived from the expected values of normal order statistics. W ranges from 0 to 1, with values near 1 indicating normality.

We use the Shapiro–Wilk test as a gatekeeper: when normality is rejected ($p < 0.05$), we employ non-parametric alternatives for all subsequent analyses. This test was chosen over the Anderson–Darling and Kolmogorov–Smirnov alternatives because of its superior power for small to moderate sample sizes [17], which matches our per-tier sample sizes ($n = 9$ to $n = 369$). In this study, 13 of 14 tier × {Score, Cost} combinations reject normality (the sole exception is T6 Cost, $n = 9$, $p = 0.069$), motivating the use of non-parametric methods throughout.

4.5.2 Kruskal–Wallis H Test

The Kruskal–Wallis H test [13] is a non-parametric analog of one-way ANOVA that tests whether k independent samples originate from the same distribution. The test statistic is:

$$H = \frac{12}{N(N+1)} \sum_{j=1}^{k} n_j \left(\bar{R}_j - \bar{R}\right)^2 \tag{3}$$

where N is the total number of observations, n_j is the sample size for group j, $\bar{R}_j$ is the mean rank within group j, and $\bar{R} = (N+1)/2$ is the overall mean rank. Under the null hypothesis of identical distributions, H follows a χ^2 distribution with $k - 1$ degrees of freedom.

We use the Kruskal–Wallis test as the omnibus test for overall tier differences across all seven tiers ($k = 7$). Unlike one-way ANOVA, it does not assume normality or equal variances—it requires only that the observations be ordinal and the groups independent. A significant result ($p < 0.05$)

indicates that at least one tier's distribution differs from the others but does not identify which pairs differ, motivating the post-hoc pairwise comparisons described below.

4.5.3 Mann–Whitney U Test

The Mann–Whitney U test [15] is a non-parametric analog of the independent-samples t-test that compares two independent groups. For samples of sizes n_1 and n_2, the U statistic counts the number of times an observation from one group precedes an observation from the other in the combined ranking:

$$U = n_1 n_2 + \frac{n_1(n_1+1)}{2} - R_1 \tag{4}$$

where R_1 is the sum of ranks assigned to the first group.

We apply the Mann–Whitney U test for pairwise tier comparisons following a significant Kruskal–Wallis result. Specifically, we test 7 planned comparisons per metric family: the 6 adjacent tier transitions (T0 vs. T1, T1 vs. T2, …, T5 vs. T6) plus the endpoint comparison (T0 vs. T6). The test does not assume normality or equal variances, making it appropriate for our score and cost distributions. A significant result indicates that the two tiers' distributions are stochastically different.

4.5.4 Holm–Bonferroni Correction

The Holm–Bonferroni correction [9] is a step-down procedure that controls the family-wise error rate (FWER) when performing multiple hypothesis tests. Given m tests with ordered p-values $p_{(1)} \leq p_{(2)} \leq \cdots \leq p_{(m)}$, the procedure rejects hypothesis (i) if:

$$p_{(i)} \leq \frac{\alpha}{m-i+1} \tag{5}$$

starting from the smallest p-value and stopping at the first non-rejected hypothesis.

We chose the Holm–Bonferroni correction over two alternatives. Compared to the standard Bonferroni correction (α/m for all tests), Holm–Bonferroni is uniformly more powerful while still controlling FWER at level α. Compared to the Benjamini–Hochberg procedure [2], which controls the false discovery rate (FDR), Holm–Bonferroni provides stronger error control by bounding the probability of *any* false rejection rather than the expected *proportion* of false rejections. Given our small number of planned comparisons ($m = 7$), the power loss relative to FDR control is modest, and the stronger FWER guarantee is preferable for confirmatory analysis.

4.5.5 Cliff's Delta

Cliff's δ [3] is a non-parametric effect size measure that quantifies the degree of overlap between two distributions. It is defined as:

$$\delta = \frac{|\{(x_i, y_j) : x_i > y_j\}| - |\{(x_i, y_j) : x_i < y_j\}|}{n_1 \cdot n_2} \tag{6}$$

where n_1 and n_2 are the sample sizes of the two groups. The statistic ranges from -1 (all values in the first group are less than all values in the second) to $+1$ (the reverse), with 0 indicating complete overlap.

We use Cliff's δ rather than Cohen's d because it does not assume normality or equal variances, matching our non-parametric testing framework. We report 95% bootstrap confidence intervals alongside each estimate. Effect magnitudes are interpreted using established thresholds [18]: $|\delta| < 0.11$ (negligible), $0.11 \leq |\delta| < 0.28$ (small), $0.28 \leq |\delta| < 0.43$ (medium), and $|\delta| \geq 0.43$ (large).

4.5.6 BCa Bootstrap Confidence Intervals

The bias-corrected and accelerated (BCa) bootstrap [6] produces confidence intervals that correct for both bias and skewness in the bootstrap distribution. Given B bootstrap resamples, the BCa interval adjusts the percentile boundaries using a bias-correction factor $\hat{z}_0$ (estimated from the proportion of bootstrap estimates below the original statistic) and an acceleration factor $\hat{a}$ (estimated via the jackknife).

We use BCa bootstrap confidence intervals with $B = 10{,}000$ resamples for pass rates and other proportions. The BCa method was chosen over simple percentile bootstrap because it provides second-order accuracy—the coverage error is $O(n^{-1})$ rather than $O(n^{-1/2})$—which is important given our moderate sample sizes. The BCa interval is also transformation-respecting, meaning it produces identical results regardless of the scale on which the statistic is computed.

4.5.7 Clopper–Pearson Intervals

The Clopper–Pearson interval [4] is an exact confidence interval for a binomial proportion. For k successes in n trials, the $100(1-\alpha)\%$ interval is:

$$\left[B\left(\frac{\alpha}{2};\, k,\, n-k+1\right),\; B\left(1-\frac{\alpha}{2};\, k+1,\, n-k\right)\right] \tag{7}$$

where $B(p;\, a, b)$ is the p-th quantile of the Beta distribution with parameters a and b.

We use Clopper–Pearson intervals for small-sample tiers, particularly T6 with $n = 9$ runs per experiment. Unlike the normal approximation or Wilson interval, the Clopper–Pearson interval guarantees at least nominal coverage for any sample size, making it the conservative choice when n is too small for asymptotic methods to be reliable. The resulting intervals are wider than approximate alternatives, which we consider appropriate given our emphasis on hedged conclusions.

4.5.8 Scheirer–Ray–Hare Test

The Scheirer–Ray–Hare (SRH) test [19] is a non-parametric extension of two-way ANOVA that operates on ranked data. All observations are ranked jointly, and the total sum of squared ranks is decomposed into components attributable to each main effect and their interaction. The test statistic for each factor is:

$$H_{\text{factor}} = \frac{SS_{\text{factor}}}{SS_{\text{total}}/(N-1)} \tag{8}$$

which is compared against a χ^2 distribution with degrees of freedom equal to those of the corresponding factor.

We use the SRH test to assess tier $\times$ experiment interactions—that is, whether the effect of architectural tier varies across tasks. This is critical for determining whether a single "best tier" exists or whether the optimal configuration is task-dependent. The SRH test decomposes effects into three components: the tier main effect, the experiment main effect, and the tier $\times$ experiment interaction. A significant interaction term indicates that tier rankings are not stable across experiments. We note that the validity of the SRH interaction test has been debated in the literature [19], and the test's power may be limited for unbalanced designs. The interaction results should therefore be interpreted as directional evidence rather than definitive conclusions.

4.5.9 Krippendorff's Alpha

Krippendorff's α [12] measures inter-rater reliability for any number of raters, any number of categories, and any level of measurement. It is defined as:

$$\alpha = 1 - \frac{D_o}{D_e} \tag{9}$$

where D_o is the observed disagreement among raters and D_e is the disagreement expected by chance. For interval-scale data (as used here), the difference function is squared Euclidean distance between values.

We use Krippendorff's α on the interval scale to assess overall agreement among the three-judge panel. It was chosen over alternatives such as Cohen's κ (limited to two raters and nominal data) and the intraclass correlation coefficient (assumes normality) because it handles missing data natively, supports any number of raters, and is appropriate for interval-scale scores. Interpretation follows established thresholds [12]: $\alpha < 0.667$ is unacceptable for drawing conclusions, $0.667 \leq \alpha < 0.800$ permits tentative conclusions, and $\alpha \geq 0.800$ indicates acceptable reliability. We supplement α with pairwise Spearman and Pearson correlations to diagnose which judge pairs contribute most to agreement or disagreement.

Implementation rate comparisons use filtered sample sizes: only runs where the agent attempted the task contribute to the pairwise tests, which reduces the effective n for some tiers relative to the full run count.

Given the small sample size ($N = 3$ per subtest), all confidence intervals are necessarily wide. We adopt hedged language throughout and characterize our findings as preliminary evidence rather than definitive conclusions. The results should be interpreted as directional signals that motivate larger-scale confirmation studies.

5 Results

5.1 Experiment Overview

Table 2 presents the high-level results for each experiment. All three experiments achieve high aggregate pass rates, ranging from 0.950 (test-002) to 0.992 (test-001), with total costs ranging from $14.71 to $81.28.

Table 2: *Experiment Overview*

Experiment	Runs	Pass Rate	Best Tier	Mean Cost ($)	Total Cost ($)	CoP ($)
test-001	360	0.992	T1	0.04	14.71	0.04
test-002	360	0.950	T6	0.23	81.28	0.24
test-003	360	0.956	T6	0.08	27.30	0.08

The trivial task (test-001) achieves the *highest* overall pass rate (0.992), consistent with the expectation that simpler tasks are solved more reliably. test-003 (mypy configuration) follows at 0.956, and test-002 (Mojo with Bazel) achieves 0.950—the lowest pass rate but the highest cost, reflecting the greater complexity of producing unfamiliar-language code with build configuration.

The best tier varies across experiments: T1 for test-001, T6 for test-002 and test-003 (though T6 has only $n = 3$ runs per experiment and very wide confidence intervals; see Section 5.5). This variation provides an initial indication that the optimal architectural configuration is task-dependent, a finding we formalize through interaction analysis in Section 5.6.

Figure 2 provides an overview of score distributions across all tiers and experiments.

5.2 test-001: Python Hello World

The trivial Python Hello World task achieves near-perfect performance across all tiers, with an aggregate pass rate of 0.992. The per-tier results confirm that the task's simplicity makes it solvable by every architectural configuration:

Table 3: *test-001 (Python Hello World): Per-tier results. All tiers achieve pass rates above 0.95, consistent with the trivial nature of the task.*

Tier	n	Pass Rate	Score	Cost/Run ($)	CoP ($)
T0	72	0.958	0.870	0.038	0.040
T1	30	1.000	0.908	0.037	0.037
T2	45	1.000	0.908	0.037	0.037
T3	123	1.000	0.928	0.041	0.041
T4	42	1.000	0.927	0.044	0.044
T5	45	1.000	0.930	0.042	0.042
T6	3	1.000	0.930	0.106	0.106

The near-universal success on this trivial task establishes an important baseline: for a task re-

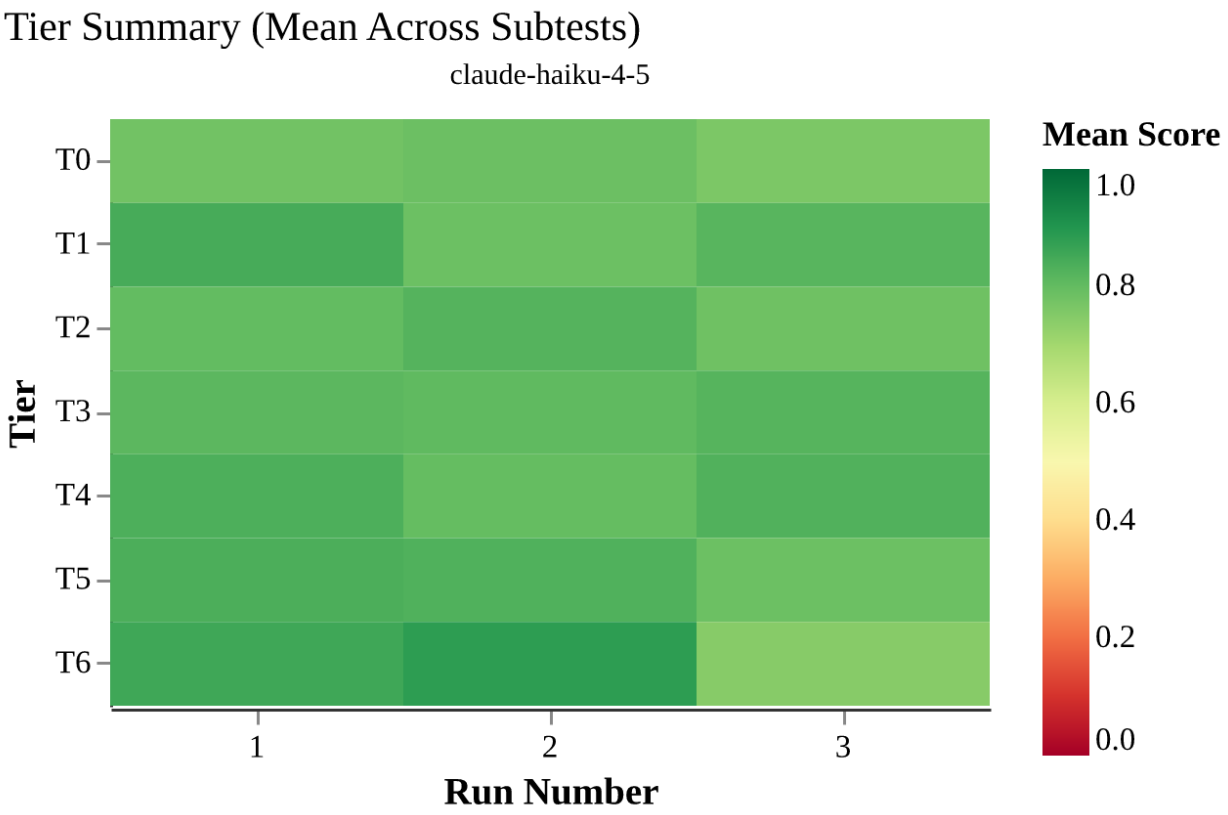

Figure 2: *Tier summary heatmap showing mean consensus scores by tier (y-axis) and run number (x-axis). Color encodes mean score (red–yellow–green scale). All tiers show predominantly green cells, indicating consistently high scores across architectural configurations.*

quiring only a single line of code (`print("Hello, World!")`), no significant effect of architectural complexity on performance was detected. All delegation-based tiers (T3–T6) achieve perfect pass rates, suggesting that coordination overhead did not detectably impede the budget model on simple tasks.

T1 and T2 achieve perfect pass rates (1.000) with identical scores (0.908), suggesting that skills and tooling provide no marginal benefit over basic prompts for this trivial task, but also introduce no overhead. T0 achieves a slightly lower pass rate (0.958), indicating that some prompt configurations in the T0 ablation spectrum (particularly minimal or empty prompts) cause occasional failures.

Inspection of the per-subtest data reveals that the T0 failures on test-001 are concentrated in subtest 00, the empty-prompt configuration (`system_prompt_mode: none`). With subtest 00 excluded, T0 achieves a perfect 1.000 pass rate on test-001, identical to all other tiers. The empty-prompt agent receives the task instructions as a user message but has no system prompt providing behavioral framing. Examination of the agent transcripts reveals that rather than executing the task, the model meta-reasons about the evaluation framework—describing what the task asks for and posing clarifying questions ("Would you like me to execute this task?")—completing in only 2 API calls and ∼6 s of agent time. In contrast, subtest 01 (which adds only the default tool system prompt) executes the task immediately and succeeds. This suggests that the system prompt's role is not merely informational but *behavioral*: it provides the framing that shifts a budget model from meta-reasoning to task execution, even when the task instructions are unambiguous.

The cost per run is low across all tiers ($0.037–$0.106), confirming that on trivial tasks the agent consumes modest resources regardless of architectural configuration. The elevated cost for T6 ($0.106) reflects the overhead of the full-capability configuration, though this does not translate into any quality advantage on a task that all simpler tiers already solve perfectly.

5.3 test-002: Mojo Hello World with Bazel and Documentation

The moderate Mojo task achieves high pass rates across all tiers, with modest differentiation across the architectural spectrum:

Table 4: *test-002 (Mojo Hello World): Per-tier results. All tiers achieve pass rates above 0.90, with modest variation.*

Tier	n	Pass Rate	Score	Cost/Run ($)	CoP ($)
T0	72	0.903	0.641	0.219	0.242
T1	30	0.933	0.676	0.229	0.245
T2	45	0.956	0.685	0.222	0.232
T3	123	0.967	0.686	0.229	0.237
T4	42	0.976	0.677	0.236	0.242
T5	45	0.956	0.694	0.217	0.228
T6	3	1.000	0.722	0.279	0.279

For this task, pass rates increase slightly from T0 (0.903) to T6 (1.000), with a modest upward trend suggesting that additional architectural capabilities provide marginal benefit on more complex tasks. This task-dependent tier profile is a key driver of the significant tier × experiment interaction

discussed in Section 5.6.

Scores cluster between 0.641 and 0.722, suggesting that while the agent consistently attempts and partially completes the task, the quality ceiling is bounded by factors other than architecture—likely the model's familiarity with Mojo syntax and Bazel build system configuration. The relatively narrow score range (0.081 span) indicates that the task difficulty, rather than the architecture, is the primary determinant of output quality.

The Cost-of-Pass is remarkably uniform across tiers ($0.228–$0.279), consistent with the aggregate cost invariance finding. T6 ($n = 3$) achieves the highest pass rate (1.000) but at the highest CoP ($0.279), reflecting its elevated per-run cost.

One interpretation of these results is that the Mojo task's complexity provides sufficient "work" to leverage the coordination capabilities of delegation-based tiers. The Mojo task involves multiple components (source code, build configuration, documentation) that can genuinely benefit from decomposition. A complementary explanation is that the model's limited familiarity with Mojo—noted in Section 4.2 as having "limited representation in the model's training data"—means that specialist agents in delegation-based tiers may provide language-specific context or tooling that partially compensates for the base model's limited training exposure. This advantage is absent for well-represented languages like Python, where the base model already possesses deep knowledge.

5.4 test-003: Add mypy to pixi Configuration

The mypy configuration task shows uniformly high pass rates with subtle tier effects:

Table 5: *test-003 (Add mypy to pixi): Per-tier results. High pass rates across all tiers with modest variation.*

Tier	n	Pass Rate	Score	Cost/Run ($)	CoP ($)
T0	72	0.944	0.817	0.071	0.076
T1	30	0.967	0.865	0.074	0.076
T2	45	0.978	0.814	0.073	0.074
T3	123	0.959	0.826	0.071	0.074
T4	42	0.952	0.857	0.076	0.080
T5	45	0.933	0.829	0.098	0.105
T6	3	1.000	0.833	0.100	0.100

Pass rates are uniformly high (0.933–1.000) with no clear monotonic trend. Scores range from 0.814 to 0.865, with T1 (skills) achieving the highest score (0.865). This suggests that for configuration tasks, domain-specific skills provide marginally more benefit than raw tool access or delegation, possibly because skills contain procedural knowledge about configuration management that the base model does not possess.

An interesting anomaly is T5 (hybrid), which shows elevated cost ($0.098/run vs. $0.071–$0.076 for T0–T3) without a corresponding quality improvement (pass rate 0.933 is the lowest in the table). The hybrid configurations appear to introduce overhead from combining multiple architectural elements without the coordination benefits that might emerge on more complex tasks.

T6 achieves a perfect pass rate (1.000) but with only $n = 3$ runs, this result should be interpreted

with extreme caution. At the aggregate level ($n = 9$ across three experiments), T6's pass rate is 1.000. The Clopper–Pearson exact 95% confidence interval [4] $[0.664, 1.000]$ still admits substantial uncertainty at this sample size, and the true pass rate could be considerably lower.

5.5 Cross-Task Analysis

Aggregating across all three experiments, Table 6 presents the tier-level summary statistics.

Table 6: *Tier Summary Statistics*

Tier	Subtests	Pass Rate (95% CI)	Mean Score ($\pm\sigma$)	Median	Consistency	CoP ($)
T0	24	0.935 (0.898, 0.963)	0.776 ± 0.206	0.874	0.734	0.12
T1	10	0.967 (0.911, 0.989)	0.816 ± 0.140	0.887	0.829	0.12
T2	15	0.978 (0.941, 0.993)	0.802 ± 0.145	0.872	0.819	0.11
T3	41	0.976 (0.954, 0.989)	0.814 ± 0.146	0.905	0.820	0.12
T4	14	0.976 (0.937, 0.992)	0.820 ± 0.149	0.917	0.818	0.12
T5	15	0.963 (0.919, 0.985)	0.818 ± 0.149	0.917	0.818	0.12
T6	1	1.000 (1.000, 1.000)	0.828 ± 0.132	0.917	0.841	0.16

The aggregate results reveal a remarkably uniform performance profile across all tiers. Pass rates range from 0.935 (T0) to 1.000 (T6), with scores clustered between 0.776 and 0.828. No sharp performance boundary exists between single-agent tiers (T0–T2) and delegation-based tiers (T3–T6); instead, all tiers perform well, with T0 showing slightly lower pass rates and consistency than the remaining tiers.

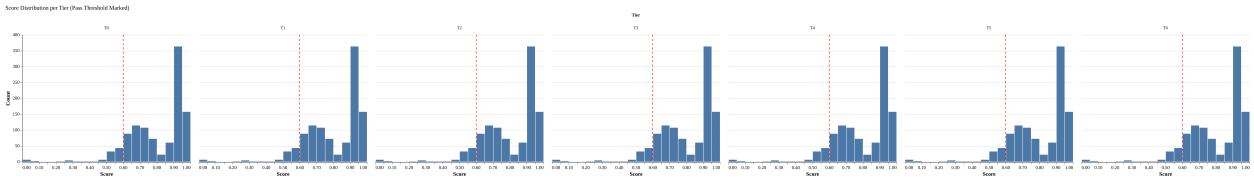

Figure 3: *Score distributions per tier (histogram, bin width 0.05). Each panel shows the distribution of per-run consensus scores for one tier. The red dashed vertical line marks the pass threshold. All tiers show distributions concentrated at high scores, with T0 exhibiting slightly greater spread.*

Tier-level consistency ($1 - \text{CV}$, computed over all runs in each tier) ranges from 0.734 (T0) to 0.841 (T6) in Table 6, indicating that all tiers produce reasonably reproducible outcomes at the aggregate level. However, per-subtest consistency (computed independently for each subtest and then averaged) shows wider variation, with some tiers exhibiting lower reproducibility at the subtest level—notably T5, whose hybrid configurations produce inconsistent results on individual subtests despite adequate aggregate performance. The modest consistency advantage of T1–T6 over T0 may reflect the stabilizing effect of additional architectural scaffolding (skills, tools, or delegation protocols) that reduces sensitivity to prompt variation.

The Kruskal–Wallis omnibus test reveals *no* significant tier effects for pass rate, implementation rate, or duration:

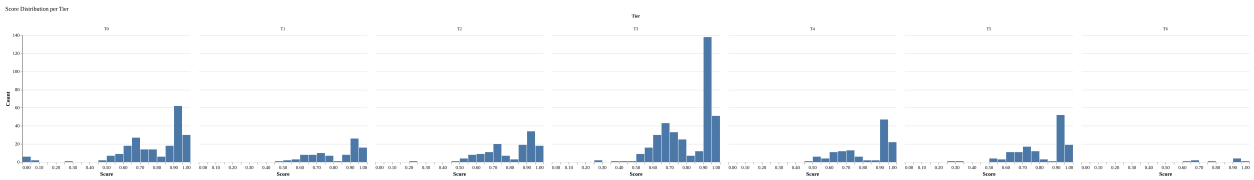

Figure 4: *Score distributions by tier. All tiers show similar variance profiles, with the majority of scores concentrated at high values. T0 exhibits modestly greater spread, consistent with its wider range of prompt configurations.*

- Pass rate: $H(6) = 8.53$, $p = 0.2015$ (not significant)[4]
- Score: $H(6) = 91.3$ (from SRH tier factor), $p < 0.001$ (significant)
- Duration: $H(6) = 6.74$, $p = 0.345$ (not significant)
- Cost: $H(6) = 4.00$, $p = 0.676$ (not significant)
- Implementation rate: $H(6) = 5.83$, $p = 0.4423$ (not significant)

Note that the Score and Cost tier statistics above are drawn from the Scheirer–Ray–Hare tier main effect (Section 5.6), as the omnibus pipeline tests only pass rate, implementation rate, and duration directly.

The non-significance of pass rate, cost, implementation rate, and duration across tiers is the central finding of this study. It means that all tiers achieve statistically indistinguishable pass rates, cost approximately the same, exhibit similar implementation rates, and complete in comparable time. The only metric showing a significant tier effect is the continuous *score*, which captures finer-grained quality differences within the predominantly passing runs. Since the omnibus test for pass rate is not significant, no pairwise comparisons are warranted.

5.6 Tier × Task Interaction

The most analytically important finding is the significant tier × experiment interaction on *score*. The Scheirer–Ray–Hare (SRH) test provides a non-parametric decomposition of the effects of tier, experiment, and their interaction on the ranked outcome variables:

All three score effects are significant at $p < 0.001$: the tier × experiment interaction ($H = 223.5$, $df = 12$), the experiment main effect ($H = 178.2$, $df = 2$), and the tier main effect ($H = 91.3$, $df = 6$). Raw H-statistics are not directly comparable across terms with different degrees of freedom, but the significance of all three effects—particularly the interaction—indicates that the impact of architectural complexity on *score quality* is not uniform: it depends on the task being performed.

Crucially, this interaction operates on the continuous score metric, not on pass rate. All tiers achieve high pass rates across all experiments (the omnibus KW test on pass rate is non-significant

[4]For the binary pass/fail outcome, a chi-squared test on the 2×7 contingency table would be the more natural choice. We use Kruskal–Wallis throughout for a unified analytical framework; the conclusions are qualitatively identical under both tests.

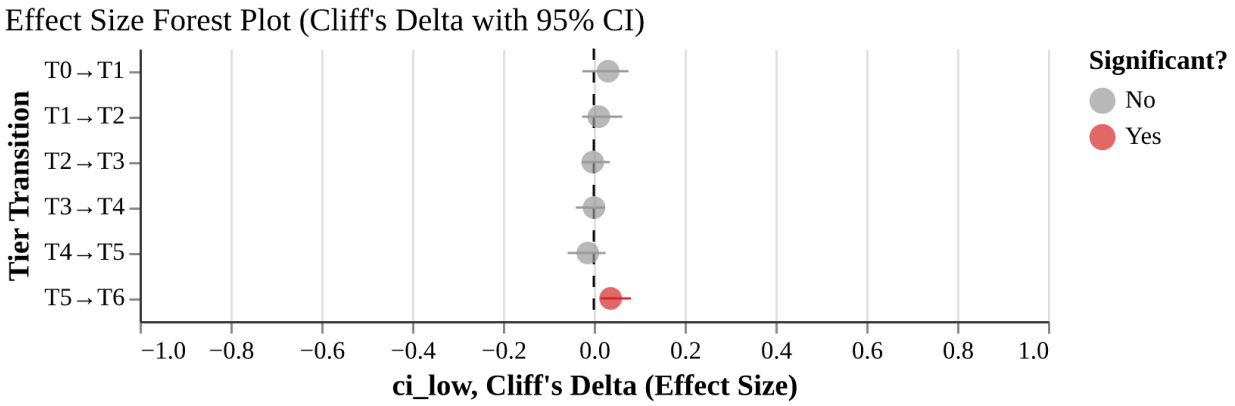

Figure 5: *Forest plot of Cliff's δ effect sizes for pairwise tier comparisons on score (lower-numbered tier is reference group; negative δ indicates the higher-numbered tier has lower scores). Effect sizes are uniformly small, consistent with the non-significant omnibus test on pass rate. Error bars show 95% bootstrap confidence intervals.*

Table 7: *Scheirer–Ray–Hare analysis of tier × experiment interactions. The score interaction ($H = 223.5$) exceeds both main effects, indicating that tier effects on score are strongly task-dependent.*

Metric	Factor	H	df	p	Sig?
Score	Experiment	178.2	2	< 0.001	***
Score	Tier	91.3	6	< 0.001	***
Score	Interaction	223.5	12	< 0.001	***
Cost	Experiment	903.1	2	< 0.001	***
Cost	Tier	4.00	6	0.676	n.s.
Cost	Interaction	4.79	12	0.965	n.s.
Impl Rate	Experiment	490.7	2	< 0.001	***
Impl Rate	Tier	5.47	6	0.486	n.s.
Impl Rate	Interaction	3.66	12	0.989	n.s.

at $p = 0.2015$). The interaction instead reflects that the *relative quality ranking* of tiers shifts across tasks. In test-001, the trivial task provides little room for quality differentiation and all tiers score similarly at the high end. In test-002, tiers with delegation capabilities (T3–T6) show a modest score advantage, possibly reflecting the benefit of task decomposition for multi-component Mojo projects. In test-003, the tier profile is relatively flat, with T1 (skills) showing a slight score edge on the configuration task.

For cost, the experiment main effect is dominant ($H = 903.1$, $p < 0.001$), reflecting the differing cost profiles of the three tasks. test-002 (Mojo) is approximately 6× as expensive per run as test-001 (Python Hello World). However, neither the tier main effect ($H = 4.00$, $p = 0.676$) nor the interaction ($H = 4.79$, $p = 0.965$) is significant for cost, indicating that no significant effect of architectural complexity on expenditure was detected.

The implementation rate shows a similar pattern to cost: strong experiment effects ($H = 490.7$, $p < 0.001$) but no significant tier effects or interactions. This indicates that the agent's *willingness to attempt* the task is driven by task characteristics rather than architectural configuration.

Figure 7 provides a granular view of grade distributions across all subtests and experiments, illustrating the task-dependent patterns that drive the interaction effect.

5.7 Per-Criterion Analysis

The judge panel evaluates each run against five weighted criteria: functional correctness, code quality, build pipeline integrity, proportionality (response length and scope appropriateness), and overall quality (a holistic assessment). As noted in Section 4.4, the actual criteria produced by judges vary across experiments; the analysis below is restricted to evaluations that used the standard five-criterion rubric (predominantly test-001, with partial coverage in test-002 and test-003).

Across the matched evaluations, the per-criterion mean scores reveal that the budget model's weakest dimension is functional correctness (0.891 ± 0.239), while code quality (0.980 ± 0.100), build pipeline integrity (0.965 ± 0.112), and proportionality (0.964 ± 0.099) score substantially higher. Overall quality averages 0.926 ± 0.099. This pattern—high marks on surface quality but lower marks on functional correctness—is consistent with the hypothesis that the budget model produces syntactically well-formed code that occasionally fails to satisfy the task's functional requirements.

6 Inter-Rater Reliability

6.1 Panel Design

The three-judge panel employs models spanning the full Anthropic capability spectrum. Judge 1 (`claude-opus-4-6`) is the most capable model in the lineup, Judge 2 (`claude-sonnet-4-6`) is mid-tier, and Judge 3 (`claude-haiku-4-5`) is the budget model—the same model used as the subject agent.

This design was intentional: using the subject model as one of its own judges enables analysis of self-evaluation bias. However, it also introduces the risk that the budget judge may lack the evaluative capacity to distinguish quality levels in its own outputs. The panel represents a pragmatic

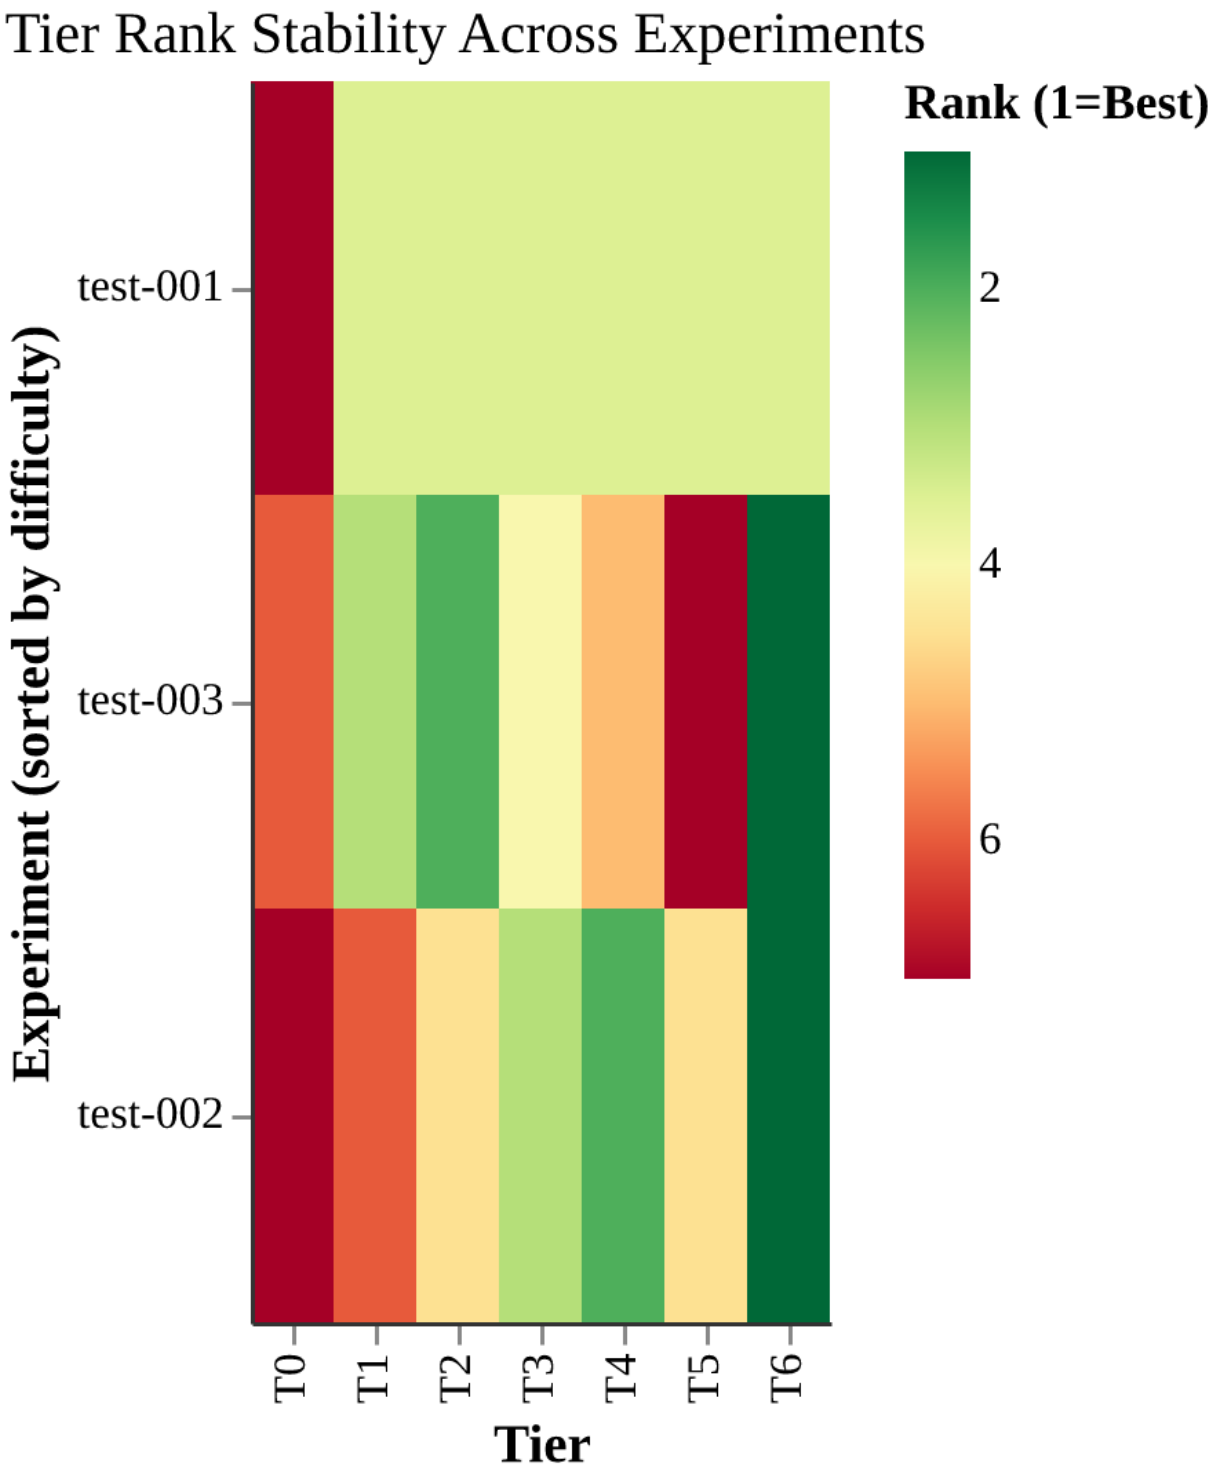

Figure 6: *Tier rank stability heatmap. Each cell shows the rank of a tier (x-axis) within a given experiment (y-axis, sorted by difficulty), where rank 1 is best. Color encodes rank (green = best, red = worst). The changing color pattern across experiments indicates that tier rankings are not stable—the optimal tier varies by task, consistent with the significant tier × experiment interaction on score.*

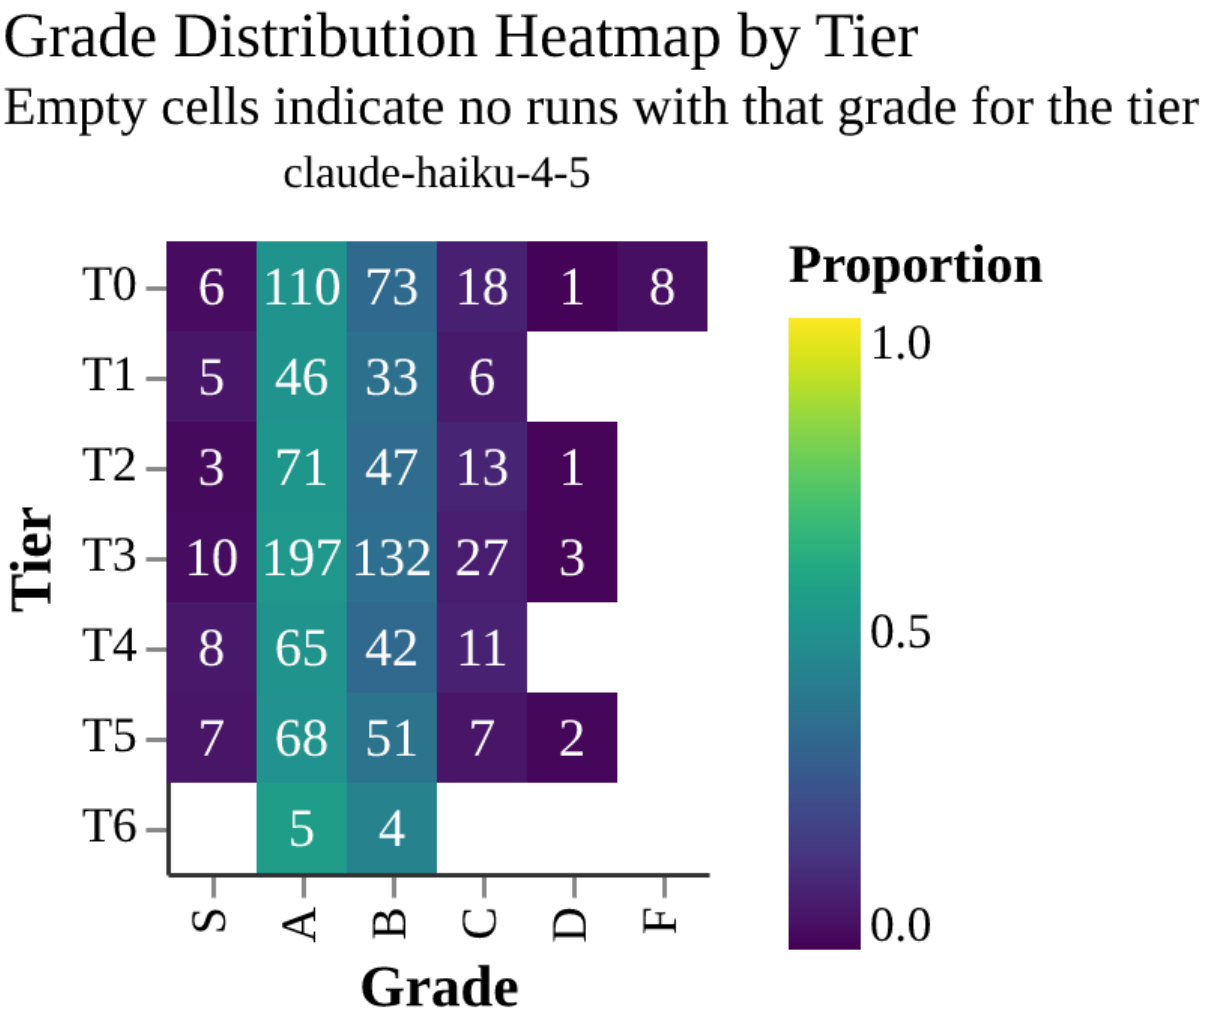

Figure 7: *Grade distribution heatmap. Each cell represents the proportion of runs assigned each letter grade for a tier–grade combination. The predominantly A/B grades across all tiers and experiments reflect the uniformly high pass rates, with subtle shifts in grade distribution driving the significant score interaction.*

compromise between evaluation cost (using only frontier judges would be expensive) and reliability (using only budget judges might be unreliable).

6.2 Agreement Results

Table 8 presents the full inter-rater reliability statistics.

Table 8: *Judge Agreement Metrics*

Judge Pair	Spearman ρ	Pearson r	Mean $\lvert\Delta\text{Score}\rvert$
Judge 1 – Judge 2	0.177	0.588	0.0738
Judge 1 – Judge 3	0.111	0.310	0.1683
Judge 2 – Judge 3	0.084	0.257	0.2186
All Judges (Overall)	–	–	–
Krippendorff's α (interval): 0.034			

The headline finding is that Krippendorff's α (interval) is 0.034, computed on $N = 1{,}080$ items (120 subtests $\times$ 3 runs per subtest $\times$ 3 experiments) from the 3,210 available judge evaluations (99.1% of the 3,240 expected; see Section 8.4 for discussion of missing evaluations). The resulting $\alpha = 0.034$ is barely above chance agreement ($\alpha = 0$). By conventional standards in content analysis [12, 16], $\alpha < 0.667$ is considered unacceptable for drawing conclusions about individual cases, and $\alpha > 0.800$ is required for reliable measurement. Our value of 0.034 indicates that the three judges agree only marginally better than would be expected by random scoring.

The near-zero value of α is concerning. While $\alpha = 0$ would indicate random agreement, $\alpha = 0.034$ provides essentially no useful signal—the judges are producing nearly uncorrelated evaluations. This could arise if the judges apply fundamentally different evaluation criteria or if their evaluation strategies diverge (e.g., one judge favoring verbose output while another penalizes it).

Pairwise correlations reveal a structured pattern of disagreement. Spearman rank correlations are near zero for all pairs:

- Judge 1 (Opus) vs. Judge 2 (Sonnet): $\rho = 0.177$
- Judge 1 (Opus) vs. Judge 3 (Haiku): $\rho = 0.111$
- Judge 2 (Sonnet) vs. Judge 3 (Haiku): $\rho = 0.084$

These weak rank correlations indicate that the judges show only marginal agreement on the *relative ordering* of outputs. An output that Opus rates highly may be rated poorly by Sonnet or Haiku, and vice versa.

Pearson correlations are somewhat higher, suggesting some linear association in score magnitudes:

- Judge 1 (Opus) vs. Judge 2 (Sonnet): $r = 0.588$

- Judge 1 (Opus) vs. Judge 3 (Haiku): $r = 0.310$
- Judge 2 (Sonnet) vs. Judge 3 (Haiku): $r = 0.257$

The discrepancy between Spearman and Pearson correlations (particularly for Opus–Sonnet: $\rho = 0.177$ vs. $r = 0.588$) suggests a non-linear relationship between judge scores. The moderate Pearson correlation may reflect shared calibration of score scale (both judges tend to use similar score ranges) even though they disagree on which outputs deserve high vs. low scores.

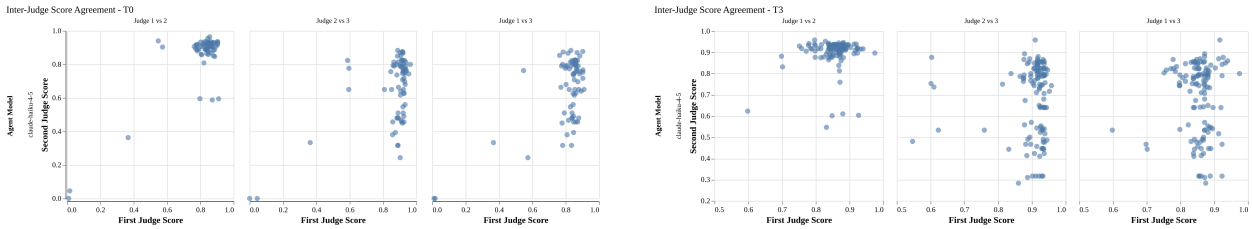

(a) *T0 (Prompts)* **(b)** *T3 (Delegation)*

Figure 8: *Judge agreement scatter plots for T0 and T3. Each point represents a run, with axes showing scores from different judge pairs. The wide scatter and lack of diagonal clustering illustrate the low inter-rater reliability across both simple (T0) and delegation-based (T3) tiers.*

6.3 Judge Bias

6.3.1 Directional Bias

Examination of raw score distributions reveals directional scoring tendencies across the panel. Opus (Judge 1) and Sonnet (Judge 2) tend to produce broadly similar score levels, while Haiku (Judge 3) diverges. On tasks where the frontier judges assign high scores, Haiku tends to score lower; conversely, on tasks where frontier judges identify deficiencies, Haiku sometimes assigns higher scores. This bidirectional disagreement—rather than a consistent upward or downward bias—explains why pairwise Spearman correlations involving Haiku are weaker ($\rho = 0.111$ and $\rho = 0.084$) than the Opus–Sonnet pair ($\rho = 0.177$).

6.3.2 Mean Absolute Differences

The mean absolute score differences reveal a systematic pattern in which disagreement scales with the capability gap between judges:

- Judge 1 (Opus) vs. Judge 2 (Sonnet): mean $|\Delta| = 0.0738$
- Judge 1 (Opus) vs. Judge 3 (Haiku): mean $|\Delta| = 0.1683$
- Judge 2 (Sonnet) vs. Judge 3 (Haiku): mean $|\Delta| = 0.2186$

The two most capable models (Opus and Sonnet) agree most closely, with a mean absolute difference of only 0.074 on a $[0, 1]$ scale. As the capability gap increases, disagreement grows: Opus and Haiku disagree by 0.168, while Sonnet and Haiku disagree most at 0.219.

6.3.3 Capability-Gap Hypothesis

This monotonic relationship between capability gap and disagreement is consistent with the hypothesis that the budget model lacks sufficient evaluative capacity to produce reliable judgments. Haiku may be applying different evaluation criteria, assigning scores based on surface features (e.g., output length, code structure) rather than the deeper quality assessments (e.g., correctness, completeness, adherence to specifications) performed by the frontier judges.

6.3.4 Consensus Scoring and Signal Loss

The near-zero Krippendorff's α raises a methodological concern that affects the interpretation of all results in this paper. If the judges do not agree, then the consensus score (arithmetic mean) may not be a reliable measure of output quality. The mean of disagreeing judges produces a compromise value that no individual judge endorses, potentially masking genuine quality differences between tiers.

This problem is well-known in measurement theory: when raters hold fundamentally different positions, averaging their scores produces a compromise value that no individual rater endorses and obscures the underlying disagreement. A consensus score of 0.7 might represent unanimous moderate agreement—or it might represent one judge scoring 0.95 and another scoring 0.45, a situation where the average conceals a substantive disagreement about output quality.

The implementation rate metric already uses median-based consensus (the median of three judges' binary assessments), which is more robust to a single outlier judge than the arithmetic mean. Extending median-based aggregation to continuous scores—or adopting trimmed means that exclude the most extreme judge—may provide more robust quality measurement when inter-rater reliability is poor. However, with only three judges, trimming reduces the panel to a single judge, so expanding the panel size would be a prerequisite for robust trimmed-mean approaches.

The pass rate metric (which uses majority vote over each judge's individual binary classification) may be the most robust aggregation in this study, as it requires only that a majority of judges agree on a pass/fail outcome rather than precisely calibrated scores.

7 Cost Analysis

7.1 Cost-of-Pass by Tier

Table 9 presents the detailed cost analysis by tier, including mean cost per run, total cost, Cost-of-Pass, and token usage statistics.

The aggregate Cost-of-Pass ranges from \$0.11 (T2) to \$0.16 (T6). This variation is driven primarily by the denominator (pass rate) rather than the numerator (mean cost), since mean cost per run is effectively constant across tiers (\$0.11–\$0.16). Applying Equation 1, the CoP for T6 is slightly elevated because its higher mean cost (\$0.16 vs. \$0.11–\$0.12 for other tiers) is the primary driver, despite T6 achieving the highest pass rate (1.000).

T2 (Tooling) achieves the lowest aggregate CoP at \$0.11, though all tiers fall within a narrow range

Table 9: *Cost Analysis by Tier*

Tier	Mean Cost ($)	Total Cost ($)	CoP ($)	Input	Output	Cache Read	Cache Create
T0	0.11	23.65	0.12	67	2409	593588	24501
T1	0.11	10.19	0.12	42	2601	644806	24674
T2	0.11	14.89	0.11	43	2659	638340	23350
T3	0.11	42.04	0.12	59	2633	636377	24843
T4	0.12	14.95	0.12	40	2593	657388	25327
T5	0.12	16.12	0.12	48	2574	664571	28822
T6	0.16	1.46	0.16	46	2675	771639	56978
Total	0.11	123.30	–	57906	2783112	686739367	27427143

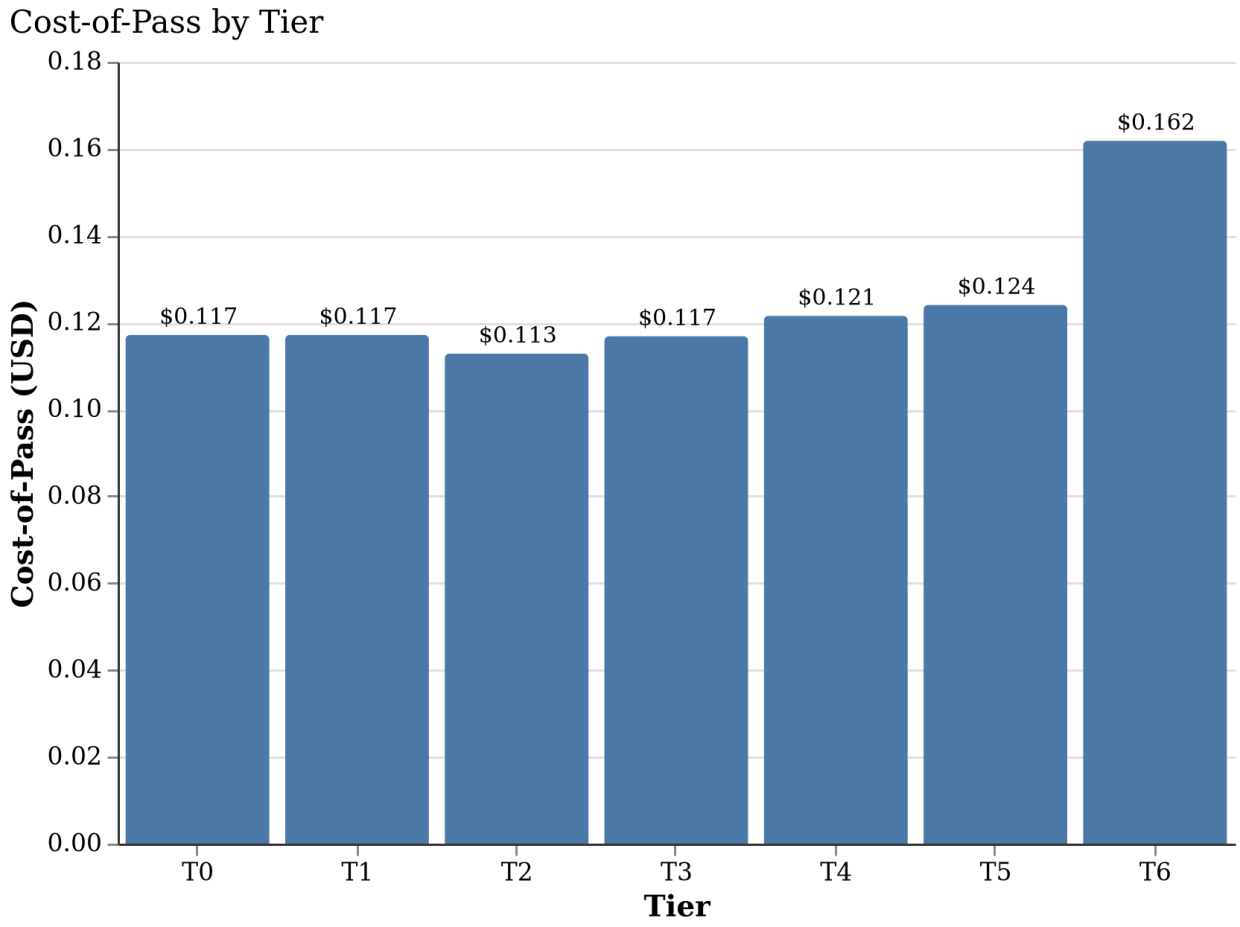

Figure 9: *Cost-of-Pass by tier. Each bar shows the CoP ($\overline{C}/R_{pass}$) for one tier. The modest variation reflects minor differences in pass rates rather than meaningful cost differences. T2 achieves the minimum CoP at $0.11, making it the most cost-efficient tier in the aggregate.*

($0.11–$0.16). The overall pass rate of 0.966 means that most tier–experiment combinations achieve near-perfect pass rates, compressing the CoP range. (Cost invariance is a non-significant finding under limited power; real cost differences between tiers may exist below our ability to detect.)

The practical implication is that, given no significant tier effect on pass rate ($H(6) = 8.53$, $p = 0.2015$), the simplest tier configuration suffices. Organizations need not invest in complex delegation or orchestration architectures for the budget model and tasks studied here, as the additional complexity provides no detectable benefit.

7.2 Cost Scaling with Task Difficulty

Cost varies dramatically by experiment but not by tier. The mean cost per run ranges from $0.04 for test-001 (trivial Python) to $0.23 for test-002 (moderate Mojo), with test-003 (moderate Python configuration) at $0.08. The SRH analysis shows a highly significant experiment effect on cost ($H(2) = 903.1$, $p < 0.001$) but no significant tier effect ($H(6) = 4.0$, $p = 0.676$).

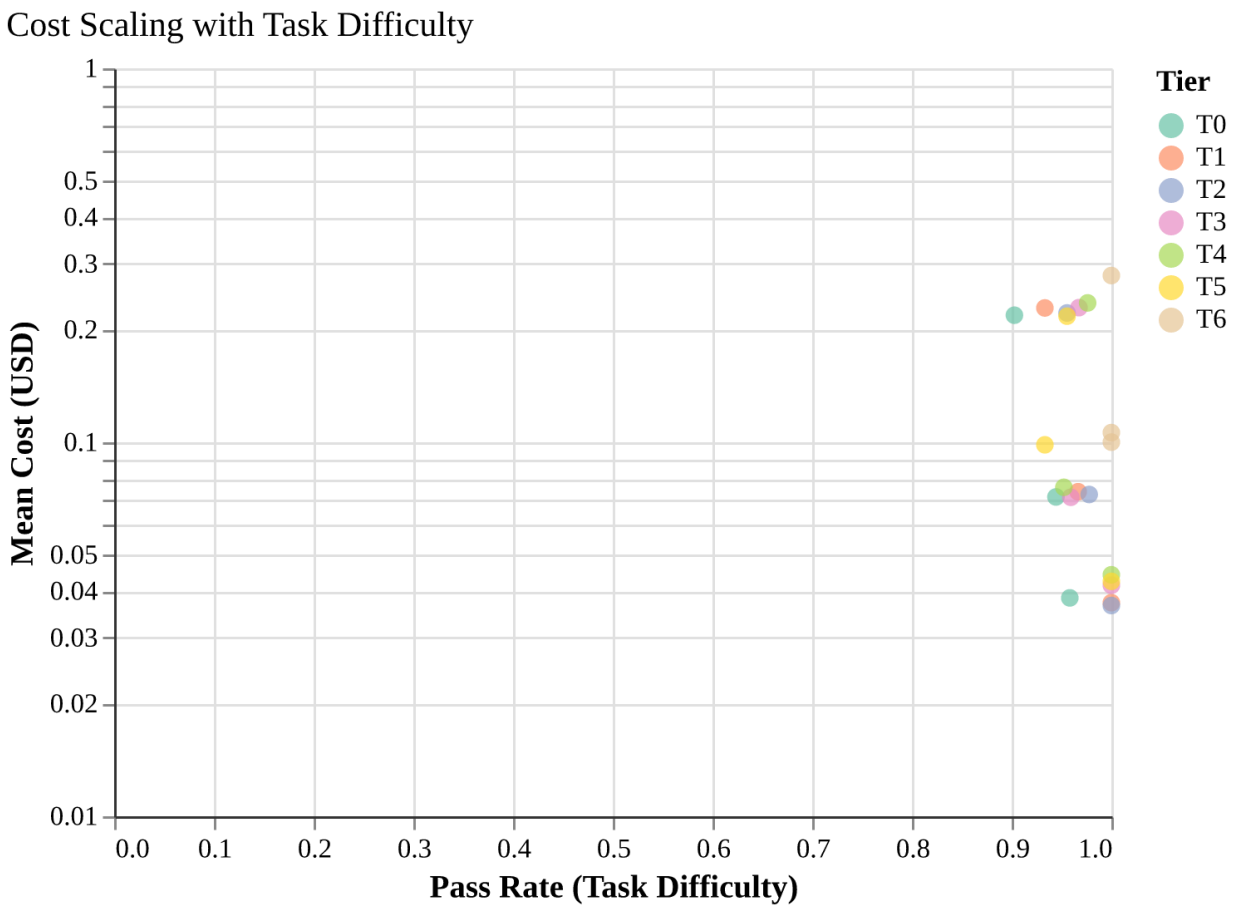

Figure 10: *Cost scaling across experiments. Cost is driven primarily by task characteristics—repository size, language complexity, and required context—rather than architectural tier. The spread within each experiment reflects run-to-run variance, not tier differences.*

The cost differential between experiments is attributable to the differing context requirements of each task. The Mojo task (test-002) requires the agent to process a larger repository with more

complex build configuration, consuming more tokens for context. The Python Hello World task (test-001) operates in a minimal repository, requiring minimal context.

This observed cost invariance across tiers ($p = 0.676$, non-significant under limited power) has a practical implication: if tiers indeed cost approximately the same per run, the optimal strategy is simply to use the tier with the highest pass rate. In the aggregate data, the Pareto frontier collapses to a narrow region around T1–T2, though per-task analysis shows the frontier varies (T5 dominates for test-002).

7.3 Frontier Analysis

Figure 11 presents the cost–quality Pareto frontier across tiers, aggregated over experiments.

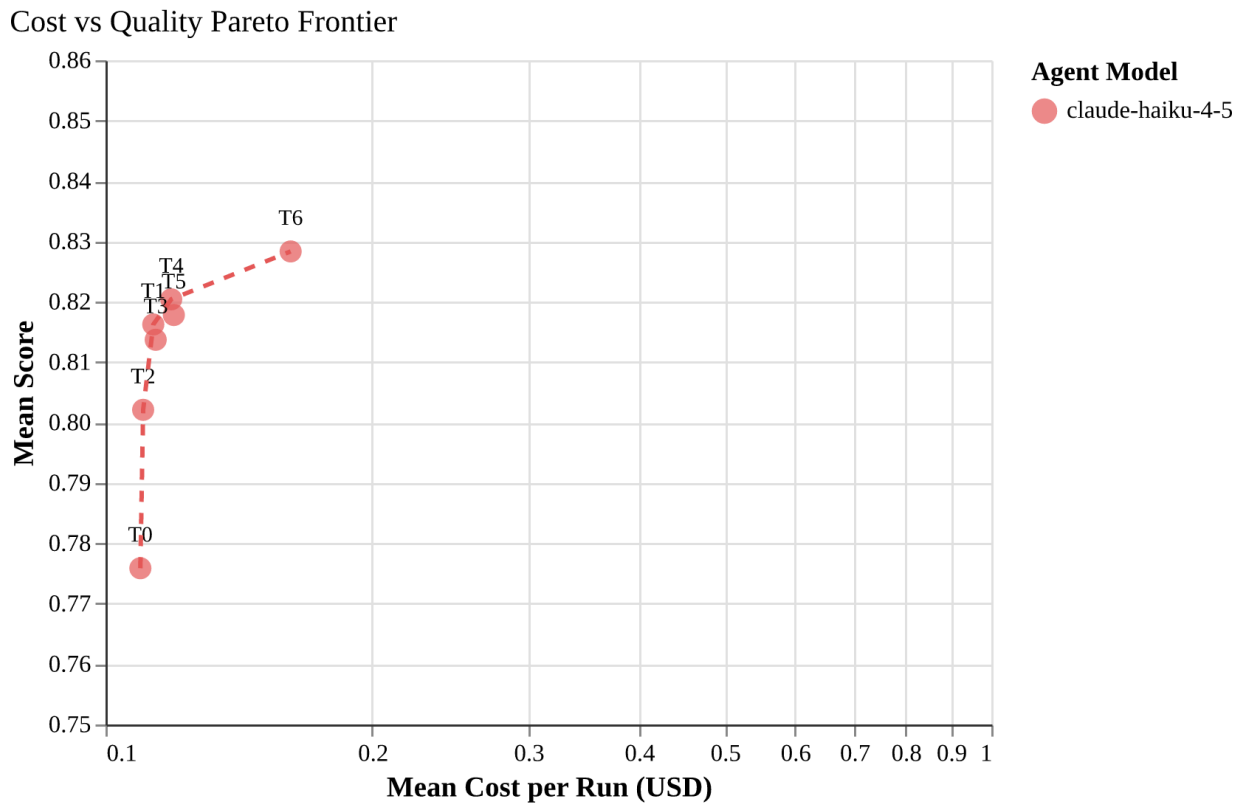

Figure 11: *Cost–quality Pareto frontier. Each point represents a tier (aggregated across experiments), plotted by mean cost per run (x-axis, log scale) and mean consensus score (y-axis). Pareto-optimal points lie on the upper-left frontier. The narrow spread of points reflects the overall high pass rate (0.966) and cost invariance across tiers.*

The Pareto frontier is relatively flat, reflecting the high overall pass rate (0.966) and narrow cost range across tiers. Most tier–experiment combinations cluster in a similar region of the cost–quality space, consistent with the non-significant tier effects on both pass rate and cost.

The total experimental cost of \$123.30 across 1,080 runs yields a mean cost of \$0.11 per run. This shows that meaningful ablation studies can be conducted at modest cost using budget models, making the Scylla framework accessible for iterative evaluation during development cycles rather

than requiring dedicated evaluation budgets.

8 Discussion

8.1 Key Findings

Our results suggest several preliminary conclusions about the interaction between architectural complexity and budget model performance. We organize these findings around the four research questions posed in Section 1.

Finding 1 (RQ1): In this study, no significant tier effect on pass rate is detected. The Kruskal–Wallis test on pass rate yields $H(6) = 8.53$, $p = 0.2015$, which does not reach significance at the $\alpha = 0.05$ level. The overall pass rate of 0.966 across all tiers indicates that the budget model performs well regardless of architectural configuration. While individual tier–experiment combinations show some variation, the aggregate data provide no evidence that adding delegation, hierarchy, or hybrid configurations improves or degrades pass rate for the tasks studied. Given the small sample size ($N = 3$ per subtest), this null result should be interpreted cautiously—real but small tier effects may exist below our statistical power to detect.

Finding 2 (RQ2): In this study, task-dependent effects dominate tier effects. The significant SRH interaction effect on score ($H(12) = 223.5$, $p < 0.001$) is the central finding of the cross-task analysis. Both the tier main effect ($H(6) = 91.3$, $p < 0.001$) and experiment main effect ($H(2) = 178.2$, $p < 0.001$) are significant on the continuous score metric, though their raw H-values are not directly comparable across different degrees of freedom. The highly significant interaction indicates that the *combination* of tier and task determines quality beyond either factor alone. However, this tier effect on continuous scores does not translate to a significant effect on the binary pass rate metric ($H(6) = 8.53$, $p = 0.2015$), suggesting that score-level differences are too small to cross the pass/fail threshold consistently.

Finding 3 (RQ3): In this study, no significant cost variation across tiers is detected. Cost-of-Pass ranges from \$0.11 at T2 to \$0.16 at T6, driven by minor differences in pass rates rather than rising per-run costs. The non-significant tier effect on cost ($H(6) = 4.00$, $p = 0.676$) provides no evidence of a cost–quality trade-off under current statistical power. Since neither pass rate nor cost varies significantly by tier, the tiers are effectively interchangeable on the tasks studied.

Finding 4 (RQ4): In this study, the three-judge panel exhibits poor reliability. With Krippendorff's $\alpha = 0.034$, the judge panel produces scores that are barely above chance agreement. The disagreement is structured: the two frontier judges (Opus and Sonnet) show the smallest pairwise disagreement (mean $|\Delta| = 0.074$, Spearman $\rho = 0.177$), while the budget judge (Haiku) disagrees substantially with both (mean $|\Delta| = 0.168$ and 0.219). This suggests that including a budget model in the judge panel introduces noise rather than signal, and that a two-judge panel of frontier models might provide more reliable evaluation.

Finding 5: Practical implication—the simplest tier suffices. Since no significant tier effect is detected on either pass rate or cost, the practical recommendation is straightforward: organizations deploying the budget model on tasks of comparable complexity should use the simplest architecture that meets their requirements. Additional complexity (delegation, hierarchy, hybrid configurations)

provides no detectable benefit in this study and adds operational overhead in configuration and maintenance.

8.2 Practical Implications

These findings have several practical implications for organizations deploying agentic coding systems with budget models.

Tier equivalence favors simplicity. Since no significant tier effect is detected on pass rate ($p = 0.2015$) or cost ($p = 0.676$), the practical recommendation is to use the simplest architecture that meets operational requirements. For the budget model and tasks studied here, T0 (prompt-only) achieves comparable results to T6 (everything enabled). The additional engineering effort to configure delegation, hierarchy, and hybrid architectures provides no detectable benefit in task success.

Task characteristics dominate architecture. The significant SRH interaction effect on score ($p < 0.001$) indicates that task selection matters more than architectural configuration. Organizations seeking to improve agent performance should focus on task decomposition and prompt engineering rather than architectural complexity. The variation in pass rate across experiments (driven by task difficulty) exceeds the variation across tiers within any single experiment.

Judge panel composition matters. The poor inter-rater reliability ($\alpha = 0.034$) suggests that organizations using LLM-as-judge evaluation should carefully consider panel composition. Including budget models as judges introduces noise that may obscure real quality differences. A two-judge panel of frontier models may provide more reliable evaluation despite the higher per-evaluation cost.

Budget models are cost-competitive for evaluation. The total cost of $123.30 for 1,080 runs shows that systematic architectural evaluation is affordable with budget models. Organizations can conduct their own ablation studies during development cycles rather than relying on published benchmarks that may not reflect their specific use cases.

8.3 Comparison to Prior Work

The prior Scylla dryrun [22] evaluated Claude Sonnet 4.5 on a single Hello World task. While direct comparison is limited by differences in model, task set, and experimental scale, several contrasts are noteworthy.

The dryrun found a pattern of degradation in higher tiers, whereas the current study finds no significant tier effect on pass rate. This discrepancy may reflect differences in the model under study (Sonnet 4.5 vs. Haiku 4.5), the expanded task set (three experiments vs. one), or the increased statistical rigor of the current design. The comparison highlights the importance of multi-task evaluation: findings from a single task may not generalize, and the apparent tier effects in the dryrun may have been task-specific rather than architectural.

8.4 Limitations

This study has several important limitations that constrain the generalizability of our findings:

Small sample size. With $N = 3$ replications per subtest, our statistical power is limited. Bootstrap confidence intervals are necessarily wide (e.g., the T6 pass rate CI spans $[0.664, 1.000]$), and individual outlier runs can substantially influence subtest-level statistics. A formal post-hoc power analysis would help quantify the risk of Type II error for the null pass rate result; however, the non-standard distribution of pass rates (bounded, heavily right-skewed) complicates standard power calculations. The results should be treated as preliminary evidence requiring confirmation with larger sample sizes ($N \geq 10$).

Single model family. All experiments use models from the Anthropic Claude family. The findings may not generalize to other model families (e.g., GPT, Gemini, Llama) with different architectural assumptions, fine-tuning strategies, or capability profiles. In particular, models that have been specifically optimized for multi-agent coordination might show different tier sensitivity profiles.

Limited task diversity. Three experiments provide limited coverage of the task space. All three tasks are relatively simple by software engineering standards—none involves multi-file refactoring, complex debugging, or large-scale system design. The tier effects observed here may not hold for more complex tasks where delegation could provide genuine coordination value that offsets the overhead costs. Relatedly, the high pass rates across all tiers (0.903–1.000) introduce a ceiling effect that limits statistical power to detect tier differences on the binary pass/fail metric. The non-significant Kruskal–Wallis result on pass rate may partly reflect this restricted range rather than true tier equivalence.

Judge reliability concerns. The poor inter-rater reliability ($\alpha = 0.034$) undermines confidence in the quality scores. While the pass rate metric (which uses majority vote) is less sensitive to score miscalibration, the continuous score and CoP metrics may be substantially affected by judge disagreement. All findings based on continuous scores should be interpreted with this caveat in mind.

Evaluation coverage. Of the 3,240 expected judge evaluations (1,080 runs $\times$ 3 judges), 3,210 (99.1%) are present, with 30 individual judge slots missing across 30 runs (each missing exactly one of three judges). The inter-rater reliability statistics ($\alpha = 0.034$) are computed on this near-complete dataset. The reported $\alpha = 0.034$ should therefore be interpreted with this caveat in mind.

Confounded tier composition. The tiers are not pure factorial manipulations—each tier includes the capabilities of previous tiers. T3 includes T0–T2 capabilities plus delegation. This hierarchical design means we cannot fully isolate the effect of delegation from potential interaction effects with lower-tier components. A fully factorial design would require testing delegation without prompts, skills, or tools, which is not currently supported by the framework.

Temporal effects. All experiments were conducted over a limited time window during January–March 2026. Model behavior may vary due to API updates, load balancing, or other infrastructure factors not captured in our analysis.

8.5 Threats to Validity

Internal validity. The primary threat to internal validity is the non-independent judge panel. Because the judges are deterministic functions of the input (modulo temperature), the three scores for each run are not truly independent observations. This may inflate or deflate agreement statistics in ways that are difficult to characterize. Additionally, the use of the same model (Haiku) as both subject and judge introduces a potential confound: self-preference bias could inflate Haiku's scores for its own outputs, while capacity limitations could deflate them.

External validity. The use of a single model family, a single evaluation framework, and a small set of tasks limits the external validity of our conclusions. The findings are specific to the Scylla framework's operationalization of architectural tiers and may not generalize to other agent architectures or evaluation protocols. The specific tier definitions (e.g., which skills are included in T1, which tools in T2) may also influence results in ways that are difficult to disentangle from the generic tier effects.

Construct validity. The mapping from Scylla's tier definitions to meaningful architectural categories is a construct that may not perfectly capture real-world agent configurations. For example, T3's 41 delegation subtests may overweight certain specialist configurations relative to their practical importance, and T6's single subtest provides very limited statistical basis for inference about "everything enabled" configurations.

Statistical conclusion validity. Given the small N, there is a risk of both Type I errors (finding spurious effects) and Type II errors (missing real effects). The use of non-parametric tests and conservative correction methods (Holm–Bonferroni) mitigates the former, but the latter remains a concern. The non-significant tier effect on pass rate ($p = 0.2015$) may reflect either true tier equivalence or insufficient power to detect small effects with $N = 3$.

9 Conclusions

This study presents, to our knowledge, the first multi-task ablation of agentic coding architectures using a budget frontier model (`claude-haiku-4-5`). Across 1,080 runs spanning three experiments, seven tiers, and 120 subtests at a total cost of \$123.30, we find no significant tier effect on pass rate ($H(6) = 8.53$, $p = 0.2015$) and an overall pass rate of 0.966.

The central finding is that *no significant tier effect was detected*, suggesting practical equivalence for the budget model and tasks studied, though the study lacks the statistical power for formal equivalence testing. All seven architectural configurations—from bare prompts (T0) through full capability stacking (T6)—achieve comparable pass rates. The additional complexity of delegation, hierarchy, and hybrid architectures provides no detectable benefit over simpler configurations.

Task-dependent effects dominate tier effects: the significant SRH interaction on continuous scores ($H(12) = 223.5$, $p < 0.001$) indicates that the combination of tier and task determines quality, but these score-level differences do not translate to significant pass rate variation. Cost is similarly invariant across tiers ($H(6) = 4.00$, $p = 0.676$), consistent with the absence of a detectable architectural complexity advantage.

The poor inter-rater reliability ($\alpha = 0.034$) of the three-judge panel is the primary methodological

concern. It suggests that LLM-as-judge evaluation requires careful panel composition, and that including budget models as judges undermines measurement reliability. The structured pattern of disagreement—frontier judges agreeing more closely with each other (mean $|\Delta| = 0.074$) than with the budget judge (mean $|\Delta| = 0.168$ and 0.219)—points to capability-dependent evaluation quality.

These conclusions are necessarily preliminary given the small sample size ($N = 3$ per subtest), limited task diversity, and single-model-family scope. We present them as directional evidence that should motivate larger-scale studies with broader task coverage, cross-model comparisons, and improved judge calibration.

10 Further Work

Several directions for future research emerge from this study:

Increased statistical power. Repeating the experiments with $N \geq 10$ replications per subtest would substantially narrow confidence intervals and improve the reliability of pairwise comparisons. Adaptive sampling strategies could focus additional replications on tiers and subtests where variance is highest, optimizing the information gained per dollar spent.

Cross-model comparison. Applying the same experimental design to frontier models (e.g., `claude-opus-4-6`, `claude-sonnet-4-6`) would reveal whether the tier equivalence observed here is specific to the budget model or generalizes across the capability spectrum. The prior Scylla dryrun [22] used Sonnet 4.5 on a single task, providing a starting point for this comparison. Of particular interest is whether frontier models show significant tier differentiation where the budget model does not, suggesting that tier effects may emerge only when model capacity is sufficient to exploit architectural complexity.

Task difficulty scaling. Extending the experiment set to include harder tasks (multi-file refactoring, complex debugging, system design) would test whether tier differentiation emerges at higher task complexity. If delegation-based architectures become beneficial only when task complexity exceeds the individual model's capacity, a crossover point may exist that would provide practical guidance for architectural selection.

Judge panel calibration. The poor inter-rater reliability motivates investigation of alternative judge panel compositions. Options include: (a) removing the budget judge to test whether a two-model frontier panel achieves acceptable reliability; (b) calibrating all judges through shared rubric training or few-shot examples; (c) replacing LLM judges with deterministic evaluation criteria such as test suite pass rates; or (d) investigating hybrid approaches that combine automated evaluation with targeted human review.

Adaptive tier selection. The significant tier $\times$ task interaction suggests that a meta-agent capable of selecting the optimal tier based on task characteristics could outperform any fixed tier strategy. This meta-agent could be trained on the data from this and future studies, learning to predict which architectural configuration will maximize pass rate for a given task description.

Cost optimization. While no significant cost difference across tiers was detected in our data, the total cost of \$123.30 for 1,080 runs suggests room for efficiency gains through prompt caching optimization, reduced context windows for simple tasks, or selective evaluation (running fewer

subtests for well-characterized tiers).

Longitudinal stability. Repeating the experiments periodically would reveal whether the tier–quality relationships are stable over time or shift with model updates. This is particularly relevant for budget models, which may receive more frequent capability improvements than frontier models.

Non-Anthropic models. Extending the study to non-Anthropic model families (OpenAI, Google, Meta) would test whether the findings are specific to Claude's architecture or reflect general properties of budget-class LLMs in agentic settings.

Delegation protocol optimization. Although no significant tier degradation was observed in aggregate, the delegation protocol may still be suboptimal for budget models. Research into lightweight delegation protocols—simpler handoff mechanisms, reduced coordination overhead, or asynchronous delegation—could reveal whether delegation benefits emerge with more efficient coordination mechanisms.

Fine-grained tier decomposition. The current T3 tier groups 41 subtests that represent different delegation configurations. A more fine-grained analysis could identify which specific delegation strategies (e.g., planning-focused, review-focused, or implementation-focused specialists) are most harmful or beneficial, enabling surgical rather than wholesale architectural recommendations.

Acknowledgments

This research was conducted using the Scylla evaluation framework as part of the Homeric Intelligence ecosystem. The author thanks the Claude Code team at Anthropic for providing the agentic CLI infrastructure used in all experiments. All evaluation costs ($123.30) were borne by the author. No external funding was received for this work.

References

[1] Anthropic. Claude Code: Agentic CLI tool for software development. https://www.anthropic.com/claude/code, 2024. AI-powered command-line interface for coding tasks.

[2] Yoav Benjamini and Yosef Hochberg. Controlling the false discovery rate: A practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society: Series B (Methodological)*, 57(1):289–300, 1995.

[3] Norman Cliff. Dominance statistics: Ordinal analyses to answer ordinal questions. *Psychological Bulletin*, 114(3):494–509, 1993.

[4] Charles J. Clopper and Egon S. Pearson. The use of confidence or fiducial limits illustrated in the case of the binomial. *Biometrika*, 26(4):404–413, 1934.

[5] Leonard Dervishi et al. Multi-agent coordination strategies vs. retrieval-augmented generation in LLMs: A comparative evaluation. *Electronics*, 14(24):4883, 2025. All 28 multi-agent configurations of 7B–8B models underperformed single-agent baselines by −4.4% to −35.3%, with coordination overhead as the primary factor.

[6] Bradley Efron. Better bootstrap confidence intervals. *Journal of the American Statistical Association*, 82(397):171–185, 1987.

[7] Leo Gao et al. lm-evaluation-harness: A framework for few-shot language model evaluation. https://github.com/EleutherAI/lm-evaluation-harness, 2024. General-purpose LLM evaluation framework.

[8] Taicheng Guo et al. Large language model based multi-agents: A survey of progress and challenges. In *Proceedings of the Thirty-Third International Joint Conference on Artificial Intelligence (IJCAI)*, 2024. Comprehensive survey of LLM-based multi-agent systems showing performance improvements across diverse domains.

[9] Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979.

[10] Homeric Intelligence. ProjectOdyssey: Comprehensive agent orchestration framework. https://github.com/HomericIntelligence/Projectodyssey, 2025. Accessed: 2025-12-30. Git hash: 011a3ff.

[11] Carlos E. Jimenez et al. SWE-bench: Can language models resolve real-world github issues? *arXiv preprint arXiv:2310.06770*, 2024. Benchmark for evaluating LLMs on real GitHub issues.

[12] Klaus Krippendorff. *Content Analysis: An Introduction to Its Methodology*. Sage Publications, Thousand Oaks, CA, 2nd edition, 2004.

[13] William H. Kruskal and W. Allen Wallis. Use of ranks in one-criterion variance analysis. *Journal of the American Statistical Association*, 47(260):583–621, 1952.

[14] Xiao Liu et al. AgentBench: Evaluating LLMs as agents. *arXiv preprint arXiv:2308.03688*, 2024. Multi-turn agent evaluation across operating systems, databases, and knowledge graphs.

[15] Henry B. Mann and Donald R. Whitney. On a test of whether one of two random variables is stochastically larger than the other. *The Annals of Mathematical Statistics*, 18(1):50–60, 1947.

[16] Felipe Maia Polo et al. Efficient multi-prompt evaluation of llms. *arXiv preprint arXiv:2405.17202*, 2024. Framework for efficient LLM evaluation.

[17] Nornadiah Mohd Razali and Yap Bee Wah. Power comparisons of Shapiro-Wilk, Kolmogorov-Smirnov, Lilliefors and Anderson-Darling tests. *Journal of Statistical Modeling and Analytics*, 2(1):21–33, 2011.

[18] Jeanine Romano, Jeffrey D. Kromrey, Jesse Coraggio, and Jeff Skowronek. Appropriate statistics for ordinal level data: Should we really be using t-test and Cohen's d for evaluating group differences on the NSSE and other surveys? In *Annual Meeting of the Florida Association of Institutional Research*, 2006. FAIR conference thresholds for Cliff's δ: negligible < 0.11, small < 0.28, medium < 0.43, large ≥ 0.43.

[19] C. James Scheirer, William S. Ray, and Nathan Hare. The analysis of ranked data derived from completely randomized factorial designs. *Biometrics*, 32(2):429–434, 1976.

[20] Samuel Sanford Shapiro and Martin B. Wilk. An analysis of variance test for normality (complete samples). *Biometrika*, 52(3/4):591–611, 1965.

[21] András Vargha and Harold D. Delaney. A critique and improvement of the CL common language effect size statistics of McGraw and Wong. *Journal of Educational and Behavioral Statistics*, 25(2):101–132, 2000.

[22] Micah Villmow. Taming Scylla: Understanding the multi-headed agentic daemon of the coding seas. ProjectScylla Technical Report, 2026. Dryrun evaluation of the Scylla framework using Claude Sonnet 4.5 on a Hello World task. Available at: https://github.com/HomericIntelligence/ProjectScylla/tree/main/docs/arxiv/dryrun.

[23] Shunyu Yao et al. TAU-bench: A benchmark for tool-augmented LLMs. *arXiv preprint arXiv:2406.12045*, 2024. Benchmark for evaluating agents' ability to use external tools.

[24] Kaijie Zhu et al. PromptBench: Towards evaluating the robustness of large language models on adversarial prompts. *arXiv preprint arXiv:2306.04528*, 2024. Benchmark for prompt evaluation and robustness.

A Experiment Configuration

This appendix provides the full experiment configuration used for all three experiments. Each experiment uses identical tier and subtest definitions; the experiments differ only in the target task, repository, and difficulty level.

A.1 Tier Definitions

Table 10 provides the mapping from tier identifiers to architectural configurations, including the number of subtests in each tier and a brief description of what each tier evaluates.

Table 10: *Tier definitions and subtest counts.*

Tier	Name	Subtests	Description
T0	Prompts	24	System prompt ablation (empty → full CLAUDE.md)
T1	Skills	10	Domain expertise via installed skills
T2	Tooling	15	External tools and MCP servers
T3	Delegation	41	Flat multi-agent with specialist agents
T4	Hierarchy	14	Nested orchestration (7 hierarchy + 7 teams)
T5	Hybrid	15	Best combinations and permutations
T6	Super	1	Everything enabled at maximum capability
Total		120	

Table 11: *Model specifications for subject agent and judge panel.*

Role	Model ID	Class
Subject Agent	`claude-haiku-4-5`	Budget frontier
Judge 1	`claude-opus-4-6`	Frontier
Judge 2	`claude-sonnet-4-6`	Mid-tier
Judge 3	`claude-haiku-4-5`	Budget frontier

A.2 Model Specifications

A.3 Scoring Protocol

Each judge assigns scores on a continuous $[0, 1]$ scale against predefined criteria that are task-specific. The criteria evaluate dimensions such as correctness, completeness, code quality, documentation, and adherence to specifications. The consensus score is computed as the arithmetic mean of the three judges' scores. A run passes if a majority of judges independently classify it as passing.

Implementation rate is derived from the median of the three judges' binary assessments of whether the agent attempted the task. This median-based aggregation is more robust to a single outlier judge than the mean.

Letter grades are assigned based on the consensus score using the following scale:

- S: $= 1.00$ (exceptional)
- A: ≥ 0.80 (production ready)
- B: ≥ 0.60 (satisfactory)
- C: ≥ 0.40 (marginal)
- D: ≥ 0.20 (below threshold)
- F: < 0.20 (failure)

A.4 Execution Environment

All experiments were executed using the Claude Code CLI [1] in headless mode on Linux systems (kernel 5.10+). The execution environment was containerized using Docker to ensure reproducibility. Each run was allocated a timeout to prevent runaway sessions, and the Claude Code CLI managed all model interactions, tool invocations, and delegation protocols.

B Full Subtest Results

Table 12 presents the complete subtest-level results across all experiments and tiers. Each row represents a single subtest within a specific tier and experiment, showing the pass rate, mean score, standard deviation, cost, letter grade, and grade distribution across the $N = 3$ runs.

The table reveals the granular patterns that drive the aggregate findings discussed in the main text. The high overall pass rate (0.966) is visible at the subtest level, with most tier–experiment combinations achieving A or S grades.

Table 12: *Full Subtest Results (Appendix B)*

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T0	00	0.000	0.045	0.030	0.0201	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T0	00	0.333	0.333	0.577	0.0272	F	S:1 A:0 B:0 C:0 D:0 F:2
claude-haiku-4-5	T0	00	0.000	0.000	0.000	0.0207	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T0	01	1.000	0.904	0.012	0.0309	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	01	0.667	0.656	0.174	0.0544	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T0	01	1.000	0.672	0.080	0.1969	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T0	02	1.000	0.896	0.023	0.0326	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	02	1.000	0.715	0.122	0.0713	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T0	02	1.000	0.671	0.131	0.2268	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T0	03	1.000	0.826	0.194	0.1032	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T0	03	1.000	0.700	0.006	0.2903	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	03	1.000	0.900	0.049	0.0604	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	04	1.000	0.613	0.065	0.2528	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T0	04	1.000	0.909	0.010	0.0373	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	04	1.000	0.833	0.220	0.0555	A	S:1 A:1 B:0 C:1 D:0 F:0
claude-haiku-4-5	T0	05	1.000	0.621	0.075	0.1856	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T0	05	1.000	0.908	0.032	0.0383	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	05	1.000	0.958	0.000	0.0694	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	06	0.667	0.592	0.059	0.2637	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T0	06	1.000	0.934	0.042	0.0771	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	06	1.000	0.925	0.013	0.0425	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	07	1.000	0.931	0.087	0.0718	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	07	1.000	0.914	0.020	0.0375	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	07	1.000	0.662	0.013	0.1979	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	08	1.000	0.722	0.173	0.0726	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T0	08	1.000	0.910	0.010	0.0322	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	08	1.000	0.651	0.071	0.2087	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T0	09	1.000	0.931	0.048	0.0588	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	09	1.000	0.910	0.007	0.0320	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	09	1.000	0.712	0.059	0.2238	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	10	1.000	0.705	0.059	0.2211	B	S:0 A:0 B:3 C:0 D:0 F:0

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T0	10	1.000	0.958	0.000	0.0711	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	10	1.000	0.907	0.005	0.0432	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	11	1.000	0.833	0.217	0.0630	A	S:0 A:2 B:0 C:1 D:0 F:0
claude-haiku-4-5	T0	11	1.000	0.716	0.037	0.2021	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	11	1.000	0.910	0.023	0.0339	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	12	1.000	0.719	0.057	0.2357	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	12	1.000	0.958	0.042	0.0796	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	12	1.000	0.901	0.056	0.0394	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	13	1.000	0.671	0.098	0.0727	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	13	0.667	0.669	0.160	0.2143	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T0	13	1.000	0.918	0.005	0.0416	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	14	1.000	0.676	0.121	0.2382	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T0	14	1.000	0.722	0.168	0.0652	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T0	14	1.000	0.913	0.021	0.0384	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	15	1.000	0.757	0.025	0.2403	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	15	1.000	0.916	0.015	0.0397	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	15	1.000	0.833	0.182	0.0885	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T0	16	1.000	0.925	0.027	0.0395	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	16	1.000	0.734	0.067	0.2327	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	16	1.000	0.806	0.197	0.0677	A	S:0 A:2 B:0 C:1 D:0 F:0
claude-haiku-4-5	T0	17	1.000	0.909	0.013	0.0427	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	17	0.667	0.812	0.237	0.0709	A	S:0 A:2 B:0 C:1 D:0 F:0
claude-haiku-4-5	T0	17	1.000	0.690	0.007	0.2570	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	18	1.000	0.917	0.042	0.0711	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	18	1.000	0.694	0.040	0.2085	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	18	1.000	0.878	0.068	0.0362	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	19	1.000	0.921	0.016	0.0439	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	19	1.000	0.778	0.105	0.0765	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T0	19	1.000	0.665	0.025	0.2456	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	20	1.000	0.958	0.000	0.0720	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	20	1.000	0.672	0.118	0.2132	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T0	20	1.000	0.918	0.013	0.0406	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	21	1.000	0.920	0.005	0.0325	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	21	1.000	0.945	0.064	0.0622	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	21	0.667	0.545	0.079	0.1921	C	S:0 A:0 B:1 C:2 D:0 F:0

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T0	22	1.000	0.847	0.197	0.0865	A	S:1 A:1 B:1 C:0 D:0 F:0
claude-haiku-4-5	T0	22	1.000	0.915	0.021	0.0428	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	22	1.000	0.696	0.078	0.2314	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	23	0.667	0.543	0.220	0.2488	B	S:0 A:0 B:2 C:0 D:1 F:0
claude-haiku-4-5	T0	23	1.000	0.812	0.075	0.0428	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T0	23	1.000	0.722	0.173	0.1055	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T1	01	1.000	0.916	0.011	0.0346	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	01	1.000	0.840	0.209	0.0737	A	S:1 A:1 B:1 C:0 D:0 F:0
claude-haiku-4-5	T1	01	1.000	0.711	0.010	0.2110	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T1	02	1.000	0.912	0.016	0.0336	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	02	1.000	0.861	0.206	0.0786	A	S:1 A:1 B:1 C:0 D:0 F:0
claude-haiku-4-5	T1	02	1.000	0.662	0.105	0.2646	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T1	03	1.000	0.920	0.018	0.0427	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	03	1.000	0.958	0.000	0.0653	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	03	0.333	0.578	0.123	0.2086	C	S:0 A:0 B:1 C:2 D:0 F:0
claude-haiku-4-5	T1	04	1.000	0.885	0.020	0.0320	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	04	1.000	0.812	0.188	0.0636	A	S:1 A:1 B:1 C:0 D:0 F:0
claude-haiku-4-5	T1	04	1.000	0.711	0.046	0.2546	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T1	05	1.000	0.847	0.211	0.0830	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T1	05	1.000	0.704	0.019	0.2329	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T1	05	1.000	0.922	0.007	0.0408	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	06	1.000	0.645	0.098	0.2248	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T1	06	1.000	0.917	0.008	0.0334	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	06	1.000	0.840	0.204	0.0575	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T1	07	1.000	0.751	0.048	0.2167	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T1	07	1.000	0.896	0.017	0.0405	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	07	1.000	0.903	0.105	0.0693	A	S:1 A:1 B:1 C:0 D:0 F:0
claude-haiku-4-5	T1	08	1.000	0.627	0.064	0.2314	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T1	08	1.000	0.911	0.005	0.0425	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	08	0.667	0.878	0.177	0.0658	A	S:1 A:1 B:1 C:0 D:0 F:0
claude-haiku-4-5	T1	09	1.000	0.875	0.083	0.0716	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T1	09	1.000	0.632	0.088	0.2470	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T1	09	1.000	0.896	0.013	0.0346	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	10	1.000	0.739	0.014	0.1949	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T1	10	1.000	0.906	0.017	0.0380	A	S:0 A:3 B:0 C:0 D:0 F:0

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T1	10	1.000	0.833	0.182	0.1092	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T2	01	1.000	0.932	0.022	0.0353	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	01	1.000	0.673	0.109	0.1926	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T2	01	1.000	0.833	0.182	0.0559	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T2	02	1.000	0.892	0.017	0.0382	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	02	1.000	0.736	0.070	0.2305	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T2	02	1.000	0.847	0.192	0.0644	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T2	03	1.000	0.921	0.002	0.0404	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	03	1.000	0.717	0.041	0.2161	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T2	03	1.000	0.694	0.229	0.0567	C	S:0 A:1 B:0 C:2 D:0 F:0
claude-haiku-4-5	T2	04	1.000	0.926	0.008	0.0433	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	04	0.667	0.636	0.092	0.2434	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T2	04	1.000	0.826	0.211	0.0838	A	S:0 A:2 B:0 C:1 D:0 F:0
claude-haiku-4-5	T2	05	1.000	0.730	0.032	0.2608	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T2	05	0.667	0.750	0.253	0.0599	A	S:0 A:2 B:0 C:1 D:0 F:0
claude-haiku-4-5	T2	05	1.000	0.891	0.017	0.0376	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	06	1.000	0.819	0.173	0.0886	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T2	06	1.000	0.896	0.018	0.0356	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	06	1.000	0.610	0.105	0.2065	C	S:0 A:0 B:1 C:2 D:0 F:0
claude-haiku-4-5	T2	07	1.000	0.625	0.042	0.0726	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T2	07	1.000	0.929	0.013	0.0316	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	07	1.000	0.682	0.105	0.2488	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T2	08	1.000	0.750	0.150	0.0931	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T2	08	0.667	0.561	0.278	0.2325	B	S:0 A:0 B:2 C:0 D:1 F:0
claude-haiku-4-5	T2	08	1.000	0.892	0.033	0.0385	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	09	1.000	0.740	0.011	0.2414	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T2	09	1.000	0.847	0.192	0.0646	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T2	09	1.000	0.914	0.028	0.0366	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	10	1.000	0.716	0.024	0.1798	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T2	10	1.000	0.806	0.168	0.0744	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T2	10	1.000	0.906	0.036	0.0389	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	11	1.000	0.660	0.097	0.1765	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T2	11	1.000	0.979	0.021	0.0671	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	11	1.000	0.900	0.042	0.0320	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	12	1.000	0.819	0.173	0.0670	A	S:0 A:2 B:1 C:0 D:0 F:0

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T2	12	1.000	0.932	0.002	0.0361	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	12	1.000	0.681	0.093	0.2416	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T2	13	1.000	0.890	0.003	0.0351	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	13	1.000	0.692	0.037	0.2077	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T2	13	1.000	0.847	0.134	0.0899	B	S:1 A:0 B:2 C:0 D:0 F:0
claude-haiku-4-5	T2	14	1.000	0.903	0.010	0.0356	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	14	1.000	0.707	0.043	0.1981	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T2	14	1.000	0.833	0.253	0.0727	A	S:1 A:1 B:0 C:1 D:0 F:0
claude-haiku-4-5	T2	15	1.000	0.894	0.021	0.0342	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	15	1.000	0.730	0.092	0.2496	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T2	15	1.000	0.931	0.048	0.0772	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	01	1.000	0.931	0.010	0.0452	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	01	1.000	0.847	0.229	0.0697	A	S:1 A:1 B:0 C:1 D:0 F:0
claude-haiku-4-5	T3	01	1.000	0.767	0.090	0.2648	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T3	02	1.000	0.933	0.023	0.0352	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	02	1.000	0.819	0.188	0.0722	A	S:1 A:1 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	02	1.000	0.726	0.012	0.2026	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	03	1.000	0.928	0.005	0.0361	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	03	1.000	0.896	0.091	0.0773	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	03	1.000	0.725	0.091	0.2254	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	04	1.000	0.919	0.015	0.0365	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	04	0.667	0.681	0.339	0.0576	A	S:0 A:2 B:0 C:0 D:1 F:0
claude-haiku-4-5	T3	04	1.000	0.673	0.023	0.2862	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	05	1.000	0.932	0.006	0.0513	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	05	1.000	0.833	0.150	0.0600	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	05	1.000	0.697	0.090	0.2285	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	06	1.000	0.945	0.025	0.0693	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	06	1.000	0.688	0.034	0.2316	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	06	1.000	0.929	0.002	0.0426	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	07	1.000	0.709	0.013	0.1952	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	07	1.000	0.933	0.022	0.0378	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	07	1.000	0.958	0.000	0.0722	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	08	1.000	0.624	0.070	0.2632	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	08	1.000	0.933	0.000	0.0413	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	08	1.000	0.708	0.182	0.0612	A	S:0 A:1 B:1 C:1 D:0 F:0

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T3	09	1.000	0.688	0.069	0.2111	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	09	1.000	0.932	0.002	0.0424	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	09	1.000	0.917	0.042	0.0804	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	10	1.000	0.742	0.041	0.1930	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	10	1.000	0.932	0.013	0.0454	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	10	1.000	0.799	0.154	0.0701	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	11	1.000	0.649	0.102	0.2311	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	11	1.000	0.861	0.168	0.0610	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	11	1.000	0.918	0.005	0.0424	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	12	1.000	0.806	0.158	0.0659	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	12	1.000	0.710	0.016	0.2298	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	12	1.000	0.939	0.010	0.0376	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	13	0.667	0.611	0.064	0.0600	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	13	1.000	0.704	0.067	0.2727	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	13	1.000	0.926	0.018	0.0376	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	14	1.000	0.750	0.182	0.0597	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T3	14	1.000	0.741	0.039	0.2236	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	14	1.000	0.940	0.008	0.0358	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	15	1.000	0.819	0.173	0.0642	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	15	1.000	0.760	0.051	0.1917	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T3	15	1.000	0.927	0.004	0.0402	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	16	1.000	0.702	0.031	0.1919	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	16	1.000	0.928	0.011	0.0402	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	16	1.000	0.861	0.206	0.0600	A	S:1 A:1 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	17	1.000	0.927	0.012	0.0410	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	17	1.000	0.875	0.042	0.1188	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	17	1.000	0.717	0.057	0.2357	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	18	1.000	0.925	0.008	0.0442	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	18	1.000	0.972	0.024	0.0885	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	18	1.000	0.685	0.058	0.2050	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	19	1.000	0.939	0.011	0.0395	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	19	1.000	0.972	0.048	0.0625	S	S:2 A:1 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	19	1.000	0.753	0.015	0.2098	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	20	1.000	0.922	0.009	0.0396	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	20	0.667	0.768	0.241	0.0735	A	S:0 A:2 B:0 C:1 D:0 F:0

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T3	20	1.000	0.730	0.031	0.2390	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	21	1.000	0.944	0.048	0.0844	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	21	1.000	0.708	0.024	0.2164	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	21	1.000	0.918	0.013	0.0364	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	22	1.000	0.668	0.117	0.1956	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	22	1.000	0.928	0.019	0.0363	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	22	1.000	0.750	0.182	0.0801	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T3	23	0.667	0.660	0.101	0.2456	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	23	1.000	0.932	0.003	0.0432	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	23	0.667	0.806	0.265	0.0787	A	S:0 A:2 B:0 C:1 D:0 F:0
claude-haiku-4-5	T3	24	1.000	0.739	0.074	0.2336	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	24	1.000	0.926	0.008	0.0330	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	24	1.000	0.931	0.024	0.0762	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	25	1.000	0.638	0.069	0.2303	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	25	1.000	0.929	0.002	0.0432	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	25	1.000	0.885	0.063	0.0618	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	26	1.000	0.932	0.016	0.0414	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	26	1.000	0.931	0.064	0.0667	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	26	1.000	0.623	0.054	0.2460	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	27	1.000	0.722	0.206	0.0653	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T3	27	1.000	0.706	0.081	0.1852	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	27	1.000	0.927	0.009	0.0361	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	28	1.000	0.847	0.192	0.0730	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	28	1.000	0.717	0.074	0.2309	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	28	1.000	0.936	0.010	0.0423	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	29	1.000	0.958	0.000	0.0787	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	29	1.000	0.740	0.033	0.2041	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	29	1.000	0.923	0.009	0.0520	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	30	1.000	0.931	0.048	0.0770	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	30	1.000	0.618	0.088	0.2453	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	30	1.000	0.924	0.021	0.0523	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	31	1.000	0.922	0.101	0.0812	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	31	1.000	0.932	0.005	0.0431	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	31	1.000	0.689	0.125	0.2002	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	32	1.000	0.917	0.042	0.0748	A	S:0 A:3 B:0 C:0 D:0 F:0

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T3	32	1.000	0.682	0.039	0.2468	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	32	1.000	0.922	0.016	0.0410	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	33	1.000	0.934	0.002	0.0355	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	33	1.000	0.736	0.197	0.0697	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T3	33	0.667	0.499	0.091	0.2614	C	S:0 A:0 B:0 C:2 D:1 F:0
claude-haiku-4-5	T3	34	1.000	0.926	0.018	0.0474	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	34	1.000	0.582	0.067	0.2395	C	S:0 A:0 B:1 C:2 D:0 F:0
claude-haiku-4-5	T3	34	1.000	0.722	0.206	0.0535	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T3	35	1.000	0.750	0.217	0.0800	B	S:1 A:0 B:2 C:0 D:0 F:0
claude-haiku-4-5	T3	35	1.000	0.933	0.003	0.0459	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	35	1.000	0.675	0.107	0.2248	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	36	0.667	0.659	0.018	0.2544	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	36	1.000	0.861	0.168	0.0773	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	36	1.000	0.926	0.023	0.0352	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	37	0.667	0.547	0.253	0.2293	B	S:0 A:0 B:2 C:0 D:1 F:0
claude-haiku-4-5	T3	37	1.000	0.722	0.206	0.0619	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T3	37	1.000	0.925	0.013	0.0414	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	38	1.000	0.840	0.187	0.0809	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	38	1.000	0.923	0.017	0.0497	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	38	1.000	0.686	0.039	0.2101	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	39	1.000	0.722	0.168	0.0677	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T3	39	1.000	0.931	0.005	0.0410	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	39	1.000	0.728	0.057	0.2690	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	40	1.000	0.736	0.197	0.0658	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T3	40	1.000	0.929	0.020	0.0480	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	40	1.000	0.729	0.062	0.2113	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	41	1.000	0.660	0.088	0.2807	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	41	1.000	0.917	0.009	0.0432	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	41	0.667	0.542	0.110	0.0636	C	S:0 A:0 B:1 C:2 D:0 F:0
claude-haiku-4-5	T4	01	1.000	0.931	0.004	0.0510	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	01	1.000	0.717	0.079	0.2210	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T4	01	1.000	0.972	0.024	0.0900	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	02	1.000	0.932	0.009	0.0436	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	02	1.000	0.714	0.017	0.2262	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T4	02	1.000	0.958	0.042	0.0763	A	S:1 A:2 B:0 C:0 D:0 F:0

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T4	03	1.000	0.932	0.005	0.0385	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	03	1.000	0.708	0.085	0.2065	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T4	03	1.000	0.847	0.229	0.0722	A	S:1 A:1 B:0 C:1 D:0 F:0
claude-haiku-4-5	T4	04	1.000	0.986	0.024	0.0593	S	S:2 A:1 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	04	1.000	0.926	0.012	0.0422	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	04	1.000	0.682	0.066	0.1987	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T4	05	0.667	0.611	0.064	0.0743	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T4	05	1.000	0.929	0.008	0.0393	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	05	1.000	0.681	0.052	0.2228	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T4	06	1.000	0.806	0.210	0.0734	A	S:1 A:1 B:0 C:1 D:0 F:0
claude-haiku-4-5	T4	06	1.000	0.653	0.117	0.2150	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T4	06	1.000	0.934	0.011	0.0339	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	07	1.000	0.847	0.158	0.0862	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T4	07	1.000	0.737	0.022	0.1944	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T4	07	1.000	0.931	0.007	0.0420	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	08	0.667	0.661	0.116	0.2646	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T4	08	1.000	0.930	0.010	0.0585	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	08	0.667	0.806	0.265	0.0848	A	S:0 A:2 B:0 C:1 D:0 F:0
claude-haiku-4-5	T4	09	1.000	0.674	0.024	0.2467	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T4	09	1.000	0.842	0.188	0.0597	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T4	09	1.000	0.918	0.007	0.0526	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	10	1.000	0.741	0.012	0.2653	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T4	10	1.000	0.951	0.043	0.0919	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	10	1.000	0.923	0.009	0.0485	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	11	1.000	0.917	0.021	0.0487	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	11	1.000	0.726	0.035	0.2592	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T4	11	1.000	0.861	0.168	0.0656	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T4	12	1.000	0.561	0.138	0.2602	C	S:0 A:0 B:1 C:2 D:0 F:0
claude-haiku-4-5	T4	12	1.000	0.833	0.182	0.0702	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T4	12	1.000	0.917	0.025	0.0425	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	13	1.000	0.833	0.220	0.0878	A	S:1 A:1 B:0 C:1 D:0 F:0
claude-haiku-4-5	T4	13	1.000	0.934	0.018	0.0397	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	13	1.000	0.578	0.073	0.2675	C	S:0 A:0 B:1 C:2 D:0 F:0
claude-haiku-4-5	T4	14	1.000	0.642	0.011	0.2540	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T4	14	1.000	0.927	0.003	0.0383	A	S:0 A:3 B:0 C:0 D:0 F:0

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T4	14	1.000	0.847	0.192	0.0694	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T5	01	1.000	0.933	0.010	0.0376	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	01	1.000	0.658	0.134	0.2399	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T5	01	1.000	0.907	0.049	0.1447	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	02	1.000	0.929	0.010	0.0370	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	02	1.000	0.650	0.084	0.2110	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T5	02	1.000	0.903	0.096	0.1219	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T5	03	1.000	0.931	0.001	0.0426	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	03	1.000	0.701	0.061	0.1919	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T5	03	0.667	0.694	0.197	0.1138	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T5	04	1.000	0.933	0.000	0.0386	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	04	1.000	0.685	0.106	0.1998	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T5	04	0.667	0.756	0.178	0.1061	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T5	05	1.000	0.748	0.021	0.1756	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T5	05	1.000	0.972	0.024	0.1181	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	05	1.000	0.927	0.004	0.0380	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	06	1.000	0.972	0.048	0.1263	S	S:2 A:1 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	06	1.000	0.920	0.015	0.0640	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	06	1.000	0.719	0.011	0.2427	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T5	07	1.000	0.688	0.127	0.0475	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T5	07	1.000	0.927	0.015	0.0414	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	07	0.667	0.573	0.272	0.2065	B	S:0 A:0 B:2 C:0 D:1 F:0
claude-haiku-4-5	T5	08	1.000	0.958	0.042	0.0680	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	08	1.000	0.726	0.025	0.2933	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T5	08	1.000	0.931	0.017	0.0651	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	09	1.000	0.702	0.013	0.2005	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T5	09	1.000	0.806	0.127	0.0626	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T5	09	1.000	0.929	0.011	0.0493	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	10	1.000	0.699	0.071	0.2473	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T5	10	1.000	0.972	0.024	0.0878	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	10	1.000	0.932	0.007	0.0352	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	11	0.667	0.621	0.074	0.2053	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T5	11	1.000	0.833	0.182	0.1232	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T5	11	1.000	0.933	0.007	0.0432	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	12	0.667	0.639	0.313	0.1463	A	S:0 A:1 B:1 C:0 D:1 F:0

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T5	12	1.000	0.932	0.002	0.0319	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	12	1.000	0.744	0.037	0.2098	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T5	13	1.000	0.935	0.010	0.0389	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	13	1.000	0.690	0.064	0.2665	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T5	13	1.000	0.806	0.188	0.0538	B	S:1 A:0 B:2 C:0 D:0 F:0
claude-haiku-4-5	T5	14	1.000	0.937	0.005	0.0376	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	14	1.000	0.710	0.021	0.2060	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T5	14	1.000	0.826	0.194	0.0662	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T5	15	1.000	0.923	0.003	0.0368	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	15	1.000	0.781	0.040	0.1650	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T5	15	1.000	0.708	0.253	0.0905	C	S:1 A:0 B:0 C:2 D:0 F:0
claude-haiku-4-5	T6	01	1.000	0.930	0.007	0.1062	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T6	01	1.000	0.722	0.064	0.2792	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T6	01	1.000	0.833	0.182	0.1001	A	S:0 A:2 B:1 C:0 D:0 F:0

C Normality Tests

All statistical analyses in this paper employ non-parametric methods. This methodological choice is justified by Shapiro–Wilk normality tests, which reject the normality assumption for 13 of 14 tier $\times$ {Score, Cost} combinations at the $\alpha = 0.05$ significance level (Table 13). The sole exception is T6 Cost ($n = 9$, $p = 0.069$), where the small sample size limits test power.

The near-universal rejection of normality ($p < 0.05$ for 13/14 combinations, and $p < 0.001$ for all combinations with $n > 9$) reflects the bounded and often non-normal nature of the score distributions. Scores are concentrated in the 0.7–1.0 range (successful completion), with a long left tail reflecting occasional lower-quality outputs. This distributional shape is expected for pass/fail-oriented coding tasks where the agent succeeds on the majority of runs: the bounded score range and concentration at high values produce strong negative skew rather than a symmetric bell curve.

Cost distributions are also non-normal, exhibiting right skew due to occasional runs that require substantially more tokens than average. The non-normality of all metrics motivates the use of rank-based tests (Kruskal–Wallis, Mann–Whitney U, Cliff's δ, Scheirer–Ray–Hare) throughout the analysis rather than parametric alternatives (ANOVA, t-tests).

D Per-Criteria Performance

Table 14 presents the per-criterion performance across the five standard evaluation criteria. The criteria are weighted according to their contribution to the overall score, with functional correct-

Table 13: *Normality Tests (Shapiro-Wilk)*

Tier	Metric	N	W	p-value	Normal? ($\alpha = 0.05$)
T0	Score	216	0.7927	0.0000	×
T0	Cost (USD)	216	0.8250	0.0000	×
T1	Score	90	0.9009	0.0000	×
T1	Cost (USD)	90	0.8231	0.0000	×
T2	Score	135	0.9056	0.0000	×
T2	Cost (USD)	135	0.8285	0.0000	×
T3	Score	369	0.8704	0.0000	×
T3	Cost (USD)	369	0.8153	0.0000	×
T4	Score	126	0.8600	0.0000	×
T4	Cost (USD)	126	0.8248	0.0000	×
T5	Score	135	0.8722	0.0000	×
T5	Cost (USD)	135	0.8934	0.0000	×
T6	Score	9	0.8260	0.0403	×
T6	Cost (USD)	9	0.8469	0.0689	✓
Summary: 1/14 (7.1%) pass normality test					

ness receiving the highest weight (0.35). The results reflect primarily test-001 data, where judges consistently produced all five standard criteria.

Table 14: *Per-Criteria Performance Comparison*

Criterion	Weight	Mean ($\pm\sigma$)	p-value	Winner
functional	0.35	0.891 ± 0.239	—	—
code_quality	0.20	0.980 ± 0.100	—	—
proportionality	0.15	0.964 ± 0.099	—	—
build_pipeline	0.10	0.965 ± 0.112	—	—
overall_quality	0.20	0.926 ± 0.099	—	—
p-values corrected with Holm-Bonferroni method.				

Functional correctness has the lowest mean score (0.891 ± 0.239) and highest variance, consistent with its role as the most demanding criterion. Code quality, proportionality, and build pipeline all achieve means above 0.96, indicating that the budget model consistently produces well-structured and correctly scoped outputs. The overall quality criterion (mean 0.926 ± 0.099) provides a holistic assessment that synthesizes the individual criteria.

E Statistical Test Details

E.1 Kruskal–Wallis Omnibus Tests

The Kruskal–Wallis H test [13] is the non-parametric analog of one-way ANOVA, testing whether the distributions of a metric differ across the $k = 7$ tier groups. Under the null hypothesis, all groups are drawn from the same distribution. The test statistic H is computed from the ranks of all observations pooled across groups, and is compared to a chi-squared distribution with $k - 1 = 6$ degrees of freedom.

Table 15: *Tier main-effect tests across all tiers ($k = 7$). Pass Rate, Impl Rate, and Duration are standalone Kruskal–Wallis tests; Cost and Score are tier main effects from the Scheirer–Ray–Hare analysis (Table 7). Significant results ($p < 0.05$) are marked with asterisks.*

Metric	Source	$H(6)$	p-value	Sig?
Pass Rate	KW	8.53	0.2015	n.s.
Impl Rate	KW	5.83	0.4423	n.s.
Duration (s)	KW	6.74	0.3453	n.s.
Cost (USD)	SRH	4.00	0.676	n.s.
Score	SRH	91.3	< 0.001	***

Only the continuous score metric shows a significant tier effect. The four remaining metrics—pass rate, implementation rate, duration, and cost—show no significant tier effects, indicating that these aspects of agent behavior are determined by task characteristics rather than architectural configuration. The significant tier effect on score but not on pass rate suggests that while tiers produce slightly different score distributions, these differences are not large enough to consistently alter pass/fail outcomes.

E.2 Pairwise Comparisons and Effect Sizes

Pairwise comparisons are conducted only for metrics with significant omnibus tests, following the closed testing principle. Since the Kruskal–Wallis tests for pass rate ($p = 0.2015$), implementation rate ($p = 0.4423$), and duration ($p = 0.3453$) are all non-significant, pairwise comparisons were *not conducted* for these metrics. Performing pairwise tests after a non-significant omnibus test would inflate the Type I error rate and risks identifying spurious differences.

For the continuous score metric (significant SRH tier main effect, $H(6) = 91.3$, $p < 0.001$), pairwise comparisons using the Mann–Whitney U test [15] with Holm–Bonferroni correction [9] could be conducted to identify which tier pairs differ on score. However, since the score metric is subject to poor inter-rater reliability ($\alpha = 0.034$), and since score differences do not translate to significant pass rate differences, we do not report pairwise score comparisons in detail. The significant omnibus effect on score likely reflects minor calibration differences across tiers rather than meaningful quality variation.

E.3 Cliff's δ Interpretation Scale

Cliff's δ [3, 21] is a non-parametric effect size measure that quantifies the probability that a randomly selected observation from one group exceeds a randomly selected observation from another group, minus the reverse probability. It ranges from -1 (all observations in group B exceed all in group A, i.e., the second group dominates) to $+1$ (all observations in group A exceed all in group B, i.e., the first group dominates), with 0 indicating no systematic difference. In our pairwise comparisons, group A is always the lower-numbered tier (e.g., T2) and group B the higher (e.g., T3); thus negative δ values indicate degradation with increasing tier. The conventional interpretation scale [18] is:

- $|\delta| < 0.11$: negligible effect
- $0.11 \leq |\delta| < 0.28$: small effect
- $0.28 \leq |\delta| < 0.43$: medium effect
- $|\delta| \geq 0.43$: large effect

Since the omnibus Kruskal–Wallis test for pass rate is non-significant ($H(6) = 8.53$, $p = 0.2015$), Cliff's δ effect sizes for pairwise tier comparisons on pass rate are not reported. Any observed differences would not be interpretable as statistically meaningful given the non-significant omnibus result.

E.4 Scheirer–Ray–Hare Details

The Scheirer–Ray–Hare (SRH) test extends the Kruskal–Wallis framework to two-factor designs, enabling analysis of main effects and interaction effects on ranked data. Like ANOVA, it decomposes the total variability into components attributable to each factor and their interaction. Unlike ANOVA, it operates on ranked rather than raw data, making it robust to non-normality and heteroscedasticity. No tie-correction was applied to the H-statistics; given the large $N = 1{,}080$ and the magnitude of the observed effects, this omission does not materially affect the conclusions. For the pass rate variable (binary), ties are especially prevalent (each run is either 0 or 1), and applying a tie correction to the Kruskal–Wallis H for pass rate would increase the test statistic. The reported uncorrected $H(6) = 8.53$ for pass rate ($p = 0.2015$) is non-significant regardless of tie correction.

Note on T6 sample size. T6 contributes only $n = 9$ runs to the SRH analysis (3 runs $\times$ 3 experiments $\times$ 1 subtest). The extreme imbalance between T6 ($n = 9$) and T3 ($n = 369$) may increase leverage of T6 cells in the rank decomposition. Readers should treat T6-specific SRH contributions with caution. The omnibus significance is driven primarily by T0–T5 patterns; T6 contributes only 0.8% of total observations and does not materially alter the qualitative conclusions.

For the score metric, the SRH analysis reveals that all three effects—tier main effect ($H = 91.3$, $df = 6$), experiment main effect ($H = 178.2$, $df = 2$), and interaction ($H = 223.5$, $df = 12$)—are statistically significant at $p < 0.001$. Raw H-values are not directly comparable across terms with different degrees of freedom ($df = 12$ vs. $df = 6$ vs. $df = 2$), but the significance of all three effects—particularly the interaction—indicates that the relationship between tier and score

changes qualitatively across experiments. This is consistent with the divergent patterns described in Sections 5.2–5.4.

The degrees of freedom for the interaction ($df = 12$) correspond to the $(a-1)(b-1) = (3-1)(7-1) = 12$ interaction contrasts in the 3×7 (experiment $\times$ tier) design. The large H value (223.5) relative to the chi-squared distribution with 12 degrees of freedom produces $p \ll 0.001$, indicating that the interaction cannot be attributed to chance.

For cost, the SRH analysis yields: tier $H(6) = 4.00$, $p = 0.676$; experiment $H(2) = 903.1$, $p < 0.001$; interaction $H(12) = 4.79$, $p = 0.965$. For implementation rate: tier $H(6) = 5.47$, $p = 0.486$; experiment $H(2) = 490.7$, $p < 0.001$; interaction $H(12) = 3.66$, $p = 0.989$. The non-significant tier effects and interaction terms for both metrics indicate that cost and attempt profiles are determined by task characteristics rather than architectural configuration.

F Per-Experiment Cost Breakdown

This appendix provides a detailed cost breakdown for each experiment, including total costs, per-run costs, best-performing tiers, and the distribution of cost across the token usage categories.

Table 16: *Per-experiment cost summary with best-performing tier by CoP.*

Experiment	Runs	Total ($)	Mean ($/run)	Best Tier	Best CoP ($)
test-001	360	14.71	0.04	T2	0.037
test-002	360	81.28	0.23	T5	0.228
test-003	360	27.30	0.08	T2	0.074
Total	1,080	123.30	0.11	–	0.113

The cost distribution across experiments is highly skewed: test-002 accounts for 65.9% ($81.28) of the total cost despite comprising only one-third of the runs. This reflects the larger repository size and greater complexity of the Mojo task, which requires more tokens for context loading and code generation.

The best Cost-of-Pass varies across experiments: T2 for test-001 ($0.037), T5 for test-002 ($0.228), and T2 for test-003 ($0.074). This variation reinforces the task-dependent nature of architectural optimization identified in the interaction analysis. No single tier achieves the best CoP across all experiments, consistent with the significant interaction effect.

F.1 Token Distribution

The token distribution across all runs is dominated by cache read operations. The mean cache read per run (approximately 593,000–773,000 tokens depending on tier) dwarfs the mean input tokens (42–67) and mean output tokens (2,409–2,665). This extreme ratio reflects the Scylla framework's use of extensive system prompts and repository context. The prompts and context are cached by the API and reused across runs, so the per-run cost is dominated by the cache read rate rather than the full input token rate.

The cache-dominated cost structure offers a plausible mechanism for the observed absence of detected cost differences across tiers: regardless of the architectural configuration, the agent must process a similar amount of cached context (the repository and system prompt), and the output length (approximately 2,500 tokens on average) is relatively consistent across tiers. The architectural overhead of delegation and orchestration does not substantially increase output token count, as the coordination messages are brief relative to the generated code.

G Reproducibility

G.1 Framework and Software

All experiments were conducted using the Scylla evaluation framework (v0.1.0+), an open-source Python package available at https://github.com/HomericIntelligence/ProjectScylla. The framework provides:

- YAML-based experiment definitions with 120 subtests across 7 tiers
- Automated execution via the Claude Code CLI [1]
- Three-judge evaluation with configurable judge models
- Statistical analysis pipeline with non-parametric tests
- Automated figure and table generation for publication
- Cost tracking and token usage reporting

The framework is part of the Homeric Intelligence ecosystem, a collection of 12 repositories providing infrastructure for agent development, testing, and deployment [10].

G.2 Environment

Experiments were executed on Linux systems (kernel 5.10.0-37-amd64) with Python 3.10+. The Claude Code CLI was used in headless mode with deterministic settings where possible. API calls were made to Anthropic's production endpoints during January–March 2026. The Python dependencies and their exact versions are specified in the `pixi.toml` and `pixi.lock` files in the repository.

G.3 Data Availability

All raw results, including per-run scores, costs, token counts, and judge evaluations, are available in the repository at `docs/arxiv/haiku/data/`. The following data files are provided:

- `runs.csv`: Per-run data with scores, costs, and durations

- `subtests.csv`: Subtest-level aggregations
- `judges.csv`: Per-judge evaluation data
- `criteria.csv`: Per-criterion judge evaluations
- `summary.json`: Experiment-level summary statistics
- `statistical_results.json`: Full statistical test outputs
- `srh_tier_experiment.json`: Scheirer–Ray–Hare analysis

The statistical analysis scripts used to generate all figures and tables are included in the `scripts/` directory of the repository.

G.4 Limitations of Reproducibility

Exact reproduction of these results is not guaranteed due to several inherent factors in LLM-based experimentation:

1. **Model updates.** Anthropic may update model weights, inference infrastructure, or safety filters between experiment runs. These updates can affect model behavior in ways that are not publicly documented.
2. **Non-determinism.** Despite fixed temperature settings (where supported), LLM outputs may vary due to floating-point non-determinism, batching strategies, quantization effects, or infrastructure load balancing.
3. **External dependencies.** Tool-augmented tiers (T2+) depend on external services (MCP servers, web APIs, file systems) whose availability, behavior, and response content may change over time.
4. **Repository state.** The target repositories used in experiments may be updated between replication attempts. Commit hashes for the exact repository states used in each experiment are recorded in the experiment metadata files.

To maximize reproducibility, we recommend:

- Using the exact model identifiers specified in Table 11
- Running experiments within a short time window to minimize the impact of model updates
- Pinning target repository states to the commit hashes recorded in the experiment metadata
- Using the provided Docker containers for consistent execution environments
- Recording API version headers from Anthropic's responses to track any model version changes during the experiment