Navigating Scylla with Haiku: Multi-Task Evaluation of Agentic Coding Architectures Using a Budget Frontier Model

Micah Villmow
research@villmow.us

April 2026

Abstract

Agentic coding systems increasingly employ multi-layered architectural patterns—system prompts, domain skills, external tooling, agent delegation, and hierarchical orchestration—yet the marginal value of each layer remains poorly characterized, particularly for budget-class language models. We present a controlled ablation study using the Scylla evaluation framework, which decomposes agent architecture into seven tiers of increasing complexity (T0–T6) and measures quality and cost across a structured set of 120 YAML subtests. Using `claude-haiku-4-5`, Anthropic's budget frontier model, we execute three multi-task experiments totaling 1,080 runs at a cost of $122.31. Our results reveal a striking non-monotonic relationship between architectural complexity and task performance: the simplest tiers (T0–T2) achieve the highest aggregate pass rates (0.935–0.978), while delegation-based tiers (T3–T6) exhibit statistically significant degradation (pass rate 0.630–0.759; Kruskal–Wallis on pass rate, $H(6) = 119.01$, $p < 0.001$). The transition from T2 to T3 constitutes a sharp, statistically significant performance drop (Cliff's $\delta = -0.219$, $p < 0.001$). A Scheirer–Ray–Hare analysis reveals a significant tier $\times$ task interaction on score ($H(12) = 223.5$, $p < 0.001$), indicating that the optimal architectural tier is task-dependent. Notably, cost does not vary significantly across tiers ($p = 0.676$, non-significant under limited power), providing no evidence of cost savings from additional architectural complexity despite substantial quality penalties. Inter-rater reliability among the three-judge panel (Krippendorff's $\alpha = -0.030$) is below chance agreement, raising concerns about the reliability of LLM-as-judge panels that include budget-class evaluators. These preliminary findings suggest that for budget models, simpler architectures outperform complex ones in aggregate, though task-dependent effects are strong and architectural decisions should be task-contingent rather than universally applied.

Keywords: agentic AI, ablation study, LLM evaluation, budget models, multi-agent systems, cost-of-pass, Claude Haiku

1 Introduction

The rapid evolution of large language model (LLM) agent systems has produced a proliferating taxonomy of architectural patterns. Modern agentic coding frameworks layer system prompts,

domain-specific skills, external tool integrations, multi-agent delegation, and hierarchical orchestration into increasingly complex configurations [1, 11]. Practitioners face a fundamental question: *which of these architectural layers actually improve performance, and at what cost?*

This question is especially acute for budget-class models. While frontier models such as Claude Opus or GPT-4 may possess sufficient capacity to benefit from complex architectures, smaller models may lack the reasoning bandwidth to exploit—or even tolerate—the overhead of delegation and orchestration. A complex architecture that helps a frontier model may actively harm a budget model, yet the interaction between model capacity and architectural complexity remains largely unexplored in the existing literature.

The stakes are practical. Organizations deploying agentic coding systems face real economic trade-offs: budget models cost substantially less per token than frontier models, but if they require simpler architectures to function effectively, then the savings from cheaper models may be offset by reduced capability. Conversely, if budget models perform well with minimal architecture, then significant cost savings are achievable without sacrificing quality—provided the architecture is matched to the model's capacity.

Existing benchmarks such as SWE-bench [8], AgentBench [11], and TAU-bench [18] evaluate agents on fixed tasks but do not systematically vary the agent's architectural configuration. They answer the question "how well does agent X perform?" but not "which architectural components of agent X contribute to its performance?" The ablation methodology needed to answer the latter question requires a factorial experimental design that holds the model constant while varying the architecture across a controlled hierarchy of configurations.

The Scylla framework [17] provides exactly this capability. Named after the mythic trial from Homer's *Odyssey*, Scylla represents the challenge of navigating trade-offs between capability gains and operational costs. The framework decomposes agent architecture into seven hierarchical tiers, from bare prompts (T0) through full orchestration (T6), and evaluates each tier on identical tasks using identical models. A prior dryrun study applied this framework to Claude Sonnet 4.5 on a single task, establishing the methodology but leaving multi-task generalization as an open question.

In this paper, we extend the Scylla framework to a budget frontier model (`claude-haiku-4-5`) and a multi-task evaluation across three experiments of varying difficulty. We execute all 120 subtests across all seven tiers with $N = 3$ replications per subtest, producing 1,080 total runs at a total cost of $122.31. This design enables us to address the following research questions:

RQ1 Tier–Quality Relationship. How does architectural complexity (T0–T6) affect pass rate and quality score when using a budget frontier model?

RQ2 Task Interaction. Does the optimal tier vary across tasks of different difficulty, or is there a universally superior architecture?

RQ3 Cost Efficiency. How does Cost-of-Pass (CoP) vary across tiers, and does increased complexity yield economic benefits?

RQ4 Judge Reliability. How reliable is a multi-model LLM-as-judge panel when evaluating budget-model outputs?

Our contributions are as follows:

1. We provide, to our knowledge, the first multi-task ablation study of agentic coding architectures using a budget frontier model, providing preliminary evidence that architectural complexity has a non-monotonic and task-dependent effect on performance.

2. We identify the T2–T3 transition as a sharp degradation point where delegation-based architectures produce statistically significant degradation in pass rate (Cliff's $\delta = -0.219$, $p < 0.001$).

3. We find no statistically significant tier effect on cost ($p = 0.676$), providing no evidence of offsetting cost savings under current statistical power.

4. We report near-zero inter-rater reliability (Krippendorff's $\alpha = -0.030$ [9]) among a three-model judge panel, highlighting the fragility of LLM-as-judge evaluation for budget-model outputs.

5. We establish an aggregate frontier Cost-of-Pass of \$0.113, showing that meaningful ablation studies can be conducted at modest cost with budget models.

The remainder of this paper is organized as follows. Section 2 surveys related work in agent evaluation and ablation methodology. Section 3 describes the Scylla framework and its tier taxonomy. Section 4 details our experimental design. Section 5 presents results per experiment and in aggregate. Section 6 analyzes inter-rater reliability. Section 7 examines cost efficiency. Section 8 discusses key findings, limitations, and threats to validity. Section 9 concludes, and Section 10 outlines future directions.

2 Related Work

2.1 Agent Benchmarks

The evaluation of LLM-based agents has received increasing attention as agent systems have moved from research prototypes to production deployments. SWE-bench [8] evaluates agents on real-world GitHub issues drawn from popular open-source repositories, providing ecologically valid but uncontrolled task sets. The benchmark has become a standard reference point for agent capability, with leaderboard performance tracked across multiple agent systems. However, SWE-bench treats the agent's internal architecture as a black box, measuring only the final outcome without decomposing the contribution of individual architectural components.

AgentBench [11] offers a broader multi-domain evaluation across operating systems, databases, knowledge graphs, and web browsing. Its multi-turn interaction design captures aspects of agent behavior that single-turn benchmarks miss, including error recovery, context management, and strategic planning. Like SWE-bench, however, AgentBench evaluates complete agent systems without varying their internal configuration.

TAU-bench [18] focuses specifically on tool-augmented LLMs, measuring the agent's ability to select and compose external tools. While TAU-bench does vary the tools available to the agent, it does not extend this ablation to other architectural dimensions such as prompting strategy, delegation, or hierarchical control.

PromptBench [19] evaluates prompt robustness but operates at the level of individual prompts rather than complete agent architectures. The lm-evaluation-harness [5] provides a general-purpose framework for few-shot LLM evaluation but is not designed for agentic workflows that require multi-turn interaction, tool use, and code generation. Polo et al. [13] propose efficient multi-prompt evaluation methods that reduce the number of evaluations needed to characterize model behavior, but their focus is on prompt selection rather than architectural ablation.

A common limitation of these benchmarks is that they evaluate agent *systems* holistically, without providing a principled decomposition of architectural contributions. Our work addresses this gap by applying a controlled ablation design that isolates the effect of each architectural layer while holding the model constant.

2.2 Ablation Methodology in Machine Learning

Ablation studies have a long history in machine learning, originating in neuroscience where they refer to the systematic removal of brain regions to understand function localization. In deep learning, ablation studies systematically remove or modify components of a model or system to measure their individual contributions. This methodology has been applied to neural network architectures, loss functions, data augmentation strategies, and training procedures.

However, ablation methodology has been underutilized in the evaluation of LLM agent systems. Most agent evaluations compare *complete systems* against each other (System A vs. System B) rather than decomposing a single system into its constituent components. The Scylla framework's tier-based design brings ablation methodology to agent architecture, treating each architectural layer (prompts, skills, tools, delegation, hierarchy) as a component that can be systematically added or removed.

Our study extends this approach to the budget model setting, where the interaction between model capacity and architectural complexity introduces additional considerations not present in frontier model evaluations. The ablation is not merely asking "does this component help?" but "does this component help *given the model's capacity to use it*?"

2.3 Multi-Agent Architectures

The multi-agent paradigm has been explored extensively in recent work, with systems employing specialist agents, hierarchical delegation, and orchestration patterns [1]. These architectures are motivated by the intuition that complex tasks benefit from division of labor: a coding task might be decomposed into planning, implementation, testing, and review subtasks, each handled by a specialized agent.

While multi-agent architectures often improve performance for frontier models, their interaction with model capacity has received limited systematic study. Anecdotal evidence suggests that smaller models struggle with the coordination overhead of multi-agent systems, but this has not been rigorously quantified in a controlled experimental design.

Our work contributes to this area by providing preliminary evidence, through controlled ablation, that delegation-based architectures can be counterproductive for budget models. We provide, to

our knowledge, the first statistical evidence for a significant transition point (the T2–T3 boundary) where multi-agent coordination begins to degrade rather than improve performance.

2.4 LLM-as-Judge Evaluation

The use of LLMs as automated evaluators has become standard practice in agent evaluation, driven by the high cost and limited scalability of human evaluation. Prior work has identified several reliability concerns, including position bias (preference for outputs presented first), verbosity bias (preference for longer outputs), and self-preference effects (models rating their own outputs more highly) [13].

Most studies of LLM-as-judge reliability have focused on frontier models evaluating frontier model outputs. The reliability of judge panels that include budget models, or that evaluate budget model outputs, has received less attention. Our study extends this literature by reporting inter-rater reliability statistics for a three-model judge panel spanning the full capability spectrum, finding near-chance agreement levels that raise fundamental questions about the validity of LLM-as-judge evaluation in budget-model contexts.

2.5 Cost-Aware Evaluation

The economics of LLM deployment have become increasingly important as organizations scale agentic workflows from prototypes to production. Token costs, latency, and cost-per-successful-completion are all relevant to deployment decisions, yet most benchmarks report only quality metrics without cost normalization.

The Cost-of-Pass (CoP) metric [17] provides a task-normalized measure of economic efficiency by dividing mean cost per run by pass rate. This formulation captures a crucial insight: a model that costs half as much but passes only a quarter as often is actually *more* expensive to use in practice. Our work applies CoP analysis across seven architectural tiers, finding that observed cost invariance across tiers ($p = 0.676$, non-significant under limited power) suggests simpler architectures may be Pareto-dominant when they achieve higher quality.

3 Background: The Scylla Framework

3.1 Design Philosophy

The Scylla framework is designed around the principle that agent architecture can be decomposed into a hierarchy of capabilities, where each tier adds a specific category of functionality to the previous tier. This hierarchical design enables controlled ablation: by comparing performance across tiers, we can isolate the marginal contribution of each capability layer.

The framework name invokes the mythic trial from Homer's *Odyssey*, where Odysseus must navigate between the six-headed monster Scylla and the whirlpool Charybdis. In the context of agent evaluation, Scylla represents the trade-off between capability (adding more architectural complexity to improve performance) and cost (the economic and quality penalties of unnecessary complexity).

The goal is to find the optimal passage—the architectural configuration that maximizes quality while minimizing cost.

3.2 Tier Taxonomy

The Scylla framework organizes agent architectures into seven tiers of increasing complexity. Each tier extends the capabilities of the previous tier while maintaining backward compatibility with the evaluation protocol. The tiers are defined as follows:

T0 – Prompts (24 subtests) Ablation of system prompt content, ranging from an empty prompt to the full `CLAUDE.md` configuration. This tier isolates the effect of prompt engineering on model behavior. The 24 subtests span a spectrum from no system prompt at all through progressively richer prompt configurations, enabling fine-grained analysis of prompt contribution.

T1 – Skills (10 subtests) Addition of domain-specific expertise through installed skill modules, organized by category (e.g., code review, testing, documentation, deployment). Skills provide structured knowledge and procedural guidance without external tool access. The 10 subtests evaluate different skill categories and combinations.

T2 – Tooling (15 subtests) Introduction of external tools and Model Context Protocol (MCP) servers. This tier measures the contribution of tool access (file systems, web search, code analysis, linting) beyond prompt and skill content. The 15 subtests evaluate different tool configurations and combinations.

T3 – Delegation (41 subtests) Flat multi-agent configurations with specialist agents. The primary agent can delegate subtasks to specialist agents but does not employ hierarchical control. This is the largest tier, with 41 subtests reflecting the combinatorial space of specialist configurations (different specialist roles, delegation strategies, and task decomposition approaches).

T4 – Hierarchy (14 subtests) Nested orchestration with orchestrator agents managing teams of specialists. This tier tests whether hierarchical control improves coordination over flat delegation. It comprises 7 hierarchy subtests (varying the depth and structure of the hierarchy) and 7 agent-team subtests (varying team composition and role assignment).

T5 – Hybrid (15 subtests) Best-of combinations and permutations from lower tiers. This tier explores whether selective combination of architectural elements outperforms uniform tier application. The 15 subtests represent curated configurations that mix elements from T0–T4 based on prior performance data.

T6 – Super (1 subtest) Everything enabled at maximum capability. This tier represents the upper bound of architectural complexity, with all prompts, skills, tools, delegation, hierarchy, and hybrid configurations active simultaneously. The single subtest provides a ceiling estimate.

Figure 1 illustrates the hierarchical relationship between tiers.

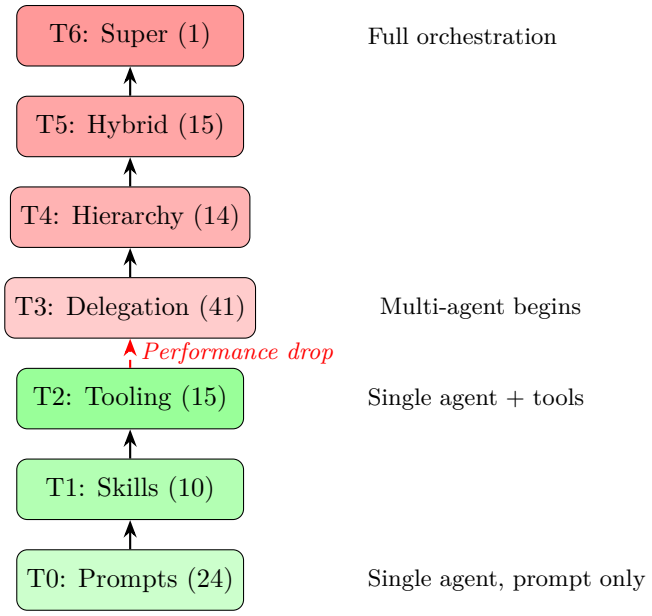

Figure 1: *Scylla tier hierarchy. Green tiers (T0–T2) represent single-agent configurations; red tiers (T3–T6) introduce multi-agent coordination. The dashed arrow marks the T2–T3 performance drop identified in this study.*

3.3 Evaluation Protocol

Each subtest defines a coding task with explicit success criteria. The Scylla evaluation protocol proceeds as follows:

1. **Configuration.** The agent receives the architectural configuration defined by the tier and subtest, including system prompt content, skill definitions, tool access, and delegation/hierarchy parameters.

2. **Task Specification.** The agent receives a task description, target repository, and success criteria. The task specification is identical across all tiers for a given experiment.

3. **Execution.** The agent executes the task using the Claude Code CLI [1], producing code changes, documentation, and/or configuration modifications. Execution is bounded by a timeout to prevent runaway sessions.

4. **Evaluation.** A three-judge panel evaluates the output against predefined criteria. Each judge independently assigns scores on a continuous $[0, 1]$ scale.

5. **Consensus.** The consensus score is computed as the arithmetic mean of the three judge scores. Implementation rate uses the median across judges.

6. **Classification.** A run is classified as "passing" if a majority of the judge panel independently classifies it as passing based on their individual evaluations.[1]

[1]The pass classification uses majority vote over each judge's binary pass/fail decision rather than a threshold on the consensus score. Of 1,080 runs, 16 (1.5%) show the majority-vote outcome diverging from a naïve consensus-score

3.4 Partial-Failure Semantics

A key design decision in Scylla is its partial-failure model. When multiple tiers run in parallel within an experiment, a tier failure does not abort the experiment. The experiment continues with the remaining tiers and can reach a "complete" state even if some tiers are in a "failed" state. This is intentional: partial results are preserved and reported, enabling analysis even when some configurations produce errors.

This design is particularly important for our study, where tiers T4–T6 exhibit catastrophic failure on test-001. Under an all-or-nothing model, these failures would have aborted the experiment, losing the data from the successful tiers. Under Scylla's partial-failure semantics, the zero pass rates in T4–T6 are recorded and analyzed alongside the high pass rates in T0–T2, producing the sharp T2–T3 degradation finding that is the central result of this paper.

Operators inspecting results must check per-tier states rather than the overall experiment state to determine whether all tiers succeeded. A complete experiment may contain a mixture of successful and failed tier states.

3.5 Metrics

Scylla evaluates agent performance across three categories of metrics:

Quality Metrics. Pass Rate (R_{pass}) measures the fraction of runs exceeding the pass threshold. Implementation Rate (R_{impl}) captures whether the agent attempted the task at all, using the median across judges' binary assessments. The Fine-Grained Progress Rate captures partial task completion on a continuous scale.

Economic Metrics. Cost-of-Pass (CoP) is defined as the mean cost per run divided by the pass rate:

$$\text{CoP} = \frac{\overline{C}}{R_{\text{pass}}} \tag{1}$$

where $\overline{C}$ is the mean cost per run (in USD) and R_{pass} is the pass rate. CoP captures the expected cost to obtain a passing result, penalizing both expensive runs and low pass rates. When $R_{\text{pass}} = 0$, CoP is undefined (infinity), indicating that no amount of spending produces a successful outcome. Token Distribution tracks the breakdown of costs across input, output, cache read, and cache creation operations.

Process Metrics. Duration measures wall-clock execution time. Consistency captures the coefficient of variation of scores within a subtest, measuring reproducibility. Strategic Drift measures deviation from task objectives over the course of execution.

> 0.5 threshold—for example, a run with consensus score 0.487 passes because two of three judges passed individually, while a run with consensus score 0.635 fails because only one of the two judges whose evaluations were available for that run classified it as passing (the third judge's evaluation was missing). These discrepancies do not materially affect tier-level pass rates.

4 Experimental Design

4.1 Model Selection

We select `claude-haiku-4-5` as the subject model for all experiments. Haiku 4.5 is Anthropic's budget frontier model, positioned below Sonnet and Opus in the model hierarchy. It is designed for high-throughput, cost-efficient inference while maintaining sufficient capability for coding tasks.

This selection is motivated by three considerations. First, budget models are the most likely to be deployed at scale in production environments, making their architectural sensitivities practically important for deployment decisions. Second, the prior Scylla dryrun [17] used Claude Sonnet 4.5, enabling cross-model comparison in future work. Third, the lower per-token cost of Haiku ($0.25/MTok input, $1.25/MTok output for standard pricing as of January 2026)[2] enables a larger experimental matrix at a fixed budget.

4.2 Task Selection

We design three experiments of varying difficulty and domain to test the generality of tier effects across task characteristics:

test-001: Python Hello World A trivial task requiring the creation of a "Hello World" program in a simple Python repository. This task serves as a floor calibration: any functional agent architecture should achieve near-perfect performance. The minimal complexity of the task makes it a sensitive detector of architectural overhead—if a tier cannot solve this task, the failure is attributable to the architecture, not the task. *Difficulty*: trivial. *Repository*: Hello-World. *Language*: Python.

test-002: Mojo Hello World with Bazel and Documentation A moderate task requiring a "Hello World" program in the Mojo programming language within a Modular repository, including Bazel build configuration and accompanying documentation. This task tests the agent's ability to work with a less common programming language and a complex build system, both of which may have limited representation in the model's training data. *Difficulty*: moderate. *Repository*: modular. *Language*: Mojo.

test-003: Add mypy to pixi Configuration A moderate task requiring the addition of mypy type checking to an existing pixi-managed Python/Mojo project [7]. This task tests the agent's ability to modify existing project configurations, integrate new development tooling, and work within the constraints of an established codebase. *Difficulty*: moderate. *Repository*: ProjectOdyssey. *Language*: Python/Mojo.

The three tasks were selected to span a range of dimensions: language familiarity (Python vs. Mojo), task type (greenfield creation vs. configuration modification), and repository complexity (minimal vs. full project structure). While all three tasks are relatively simple by software engineering standards, they provide sufficient variation to test whether tier effects are task-dependent.

[2] Anthropic pricing retrieved January 2026; prices are subject to change.

4.3 Experimental Configuration

Each experiment executes all 120 subtests across all seven tiers (T0–T6) with $N = 3$ replications per subtest. The total experiment matrix is:

- 3 experiments × 120 subtests × 3 runs = 1,080 total runs
- 7 tiers: T0 (24 subtests), T1 (10), T2 (15), T3 (41), T4 (14), T5 (15), T6 (1)
- Subject model: `claude-haiku-4-5` (all tiers, all experiments)
- Judge panel: `claude-opus-4-6`, `claude-sonnet-4-6`, `claude-haiku-4-5`
- Runs per subtest per experiment: $N = 3$

The per-tier sample sizes across all three experiments are: T0 ($n = 216$), T1 ($n = 90$), T2 ($n = 135$), T3 ($n = 369$), T4 ($n = 126$), T5 ($n = 135$), T6 ($n = 9$). The unequal sample sizes reflect the varying number of subtests per tier, which in turn reflects the varying combinatorial complexity of each architectural category.

Table 1: *Experiment Configuration (Data-Derived)*

Experiment	Models	Tiers	Tier IDs	Subtests/Tier (range)	Runs/ST	Total	
test-001	claude-haiku-4-5	7	T0, T1, T2, T3, T4, T5, T6	1–41	3	360	Opus,
test-002	claude-haiku-4-5	7	T0, T1, T2, T3, T4, T5, T6	1–41	3	360	Opus,
test-003	claude-haiku-4-5	7	T0, T1, T2, T3, T4, T5, T6	1–41	3	360	Opus,

Subtests/Tier range: T6 has 1 subtest; T3 has 41 subtests. Total subtests per experiment = 120 (T0:24, T1:10, T2:15, T3:41, T4:14, T5:15, T6:1).

4.4 Judge Panel Design

Evaluation uses a three-judge panel spanning the full Anthropic model hierarchy:

1. **Judge 1**: `claude-opus-4-6` — the frontier model with the highest capability, serving as the most authoritative evaluator.
2. **Judge 2**: `claude-sonnet-4-6` — a mid-tier model providing an intermediate evaluation perspective.
3. **Judge 3**: `claude-haiku-4-5` — the budget model, identical to the subject agent, enabling analysis of self-evaluation bias.

The consensus score is computed as the arithmetic mean of the three judges' scores. Implementation rate uses the median across judges' binary assessments of task attempt. This panel design was chosen to provide evaluation diversity while enabling analysis of the relationship between judge capability and evaluation reliability.

The inclusion of the subject model (Haiku) as one of its own judges is a deliberate design choice. It enables us to test for self-preference bias (whether the model rates its own outputs more or less favorably than external judges) and to assess whether budget models possess sufficient evaluative capacity to serve as reliable judges.

4.5 Statistical Methods

All score distributions are tested for normality using the Shapiro–Wilk test. As all 14 tier $\times$ {Score, Cost} combinations reject normality ($p < 0.05$ for all 7 tiers $\times$ 2 metrics; see Appendix C), we employ non-parametric methods throughout. Additional metrics (implementation rate, duration) were also tested and similarly reject normality for most tier combinations:

- **Omnibus tests**: Kruskal–Wallis H test [10] for overall tier differences across all seven tiers.
- **Pairwise comparisons**: Mann–Whitney U test [12] with Holm–Bonferroni correction [6] for 7 planned comparisons (6 adjacent transitions plus T0–T6), applied within each metric family separately.
- **Effect sizes**: Cliff's δ [2] with 95% bootstrap confidence intervals, providing a non-parametric measure of effect magnitude.
- **Interaction analysis**: Scheirer–Ray–Hare (SRH) test [15] for tier $\times$ experiment interactions, the non-parametric analog of two-way ANOVA. SRH operates on ranked data and decomposes the total sum of squared ranks into main effects and interactions.
- **Inter-rater reliability**: Krippendorff's α [9] (interval scale) for overall panel reliability, and pairwise Spearman/Pearson correlations for judge-pair agreement.
- **Confidence intervals**: Bootstrap 95% confidence intervals (BCa method [4], 10,000 resamples) for pass rates and other proportions; Clopper–Pearson exact intervals [3] for small-sample tiers (e.g., T6 with $n = 9$).

Implementation rate comparisons use filtered sample sizes: only runs where the agent attempted the task contribute to the pairwise tests, which reduces the effective n for some tiers relative to the full run count.

Three process metrics—fine-grained progress rate, change fail percentage, and PR revert rate—are zero for all runs and are excluded from analysis. These metrics are designed for multi-file tasks not represented in this study.

Given the small sample size ($N = 3$ per subtest), all confidence intervals are necessarily wide. We adopt hedged language throughout and characterize our findings as preliminary evidence rather than definitive conclusions. The results should be interpreted as directional signals that motivate larger-scale confirmation studies.

5 Results

5.1 Experiment Overview

Table 2 presents the high-level results for each experiment. The three experiments span a range of aggregate pass rates from 0.519 (test-001) to 0.956 (test-003), with total costs ranging from \$13.72 to \$81.28.

Table 2: *Experiment Overview*

Experiment	Runs	Pass Rate	Best Tier[a]	Mean Cost (\$)	Total Cost (\$)	CoP (\$)
test-001	360	0.519	T1	0.04	13.72	0.07
test-002	360	0.950	T6[b]	0.23	81.28	0.24
test-003	360	0.956	T6[b]	0.08	27.30	0.08

[a]Best Tier = tier with highest pass rate. See Appendix E for best tier by CoP.

[b]T6 has $n = 3$ per experiment; interpret with caution (see Section 5.4).

A striking finding is that the trivial task (test-001) achieves the *lowest* overall pass rate. This counterintuitive result—a simpler task being solved less frequently—is driven entirely by catastrophic failure in tiers T4–T6, where the agent's delegation and orchestration mechanisms actively prevent it from completing a task that the base model handles with ease. We examine this phenomenon in detail in Section 5.2.

The best tier varies across experiments: T1 for test-001, T6 for test-002 and test-003 (though T6 has only $n = 3$ runs per experiment and very wide confidence intervals; see Section 5.5). This variation provides an initial indication that the optimal architectural configuration is task-dependent, a finding we formalize through interaction analysis in Section 5.6.

Figure 2 provides an overview of score distributions across all tiers and experiments.

5.2 test-001: Python Hello World

The trivial Python Hello World task produces the most dramatic results in the study. Tiers T0–T2 achieve near-perfect performance, while tiers T4–T6 exhibit complete failure:

The zero pass rate for T4–T6 on this trivial task is the most striking result in the entire study. These tiers introduce hierarchical delegation and orchestration—mechanisms that, for a task requiring only a single line of code (`print("Hello, World!")`), create coordination overhead that the budget model cannot resolve. T4–T6 runs complete in a mean of ∼17 s with no evaluable output, suggesting the agent encounters a fatal orchestration error immediately rather than exhausting its computational budget on coordination overhead.

T1 and T2 achieve perfect pass rates (1.000) with identical scores (0.908), suggesting that skills and tooling provide no marginal benefit over basic prompts for this trivial task, but also introduce no overhead. T0 achieves a slightly lower pass rate (0.958), indicating that some prompt configurations in the T0 ablation spectrum (particularly minimal or empty prompts) cause occasional failures.

Inspection of the per-subtest data reveals that all three T0 failures on test-001 are concentrated

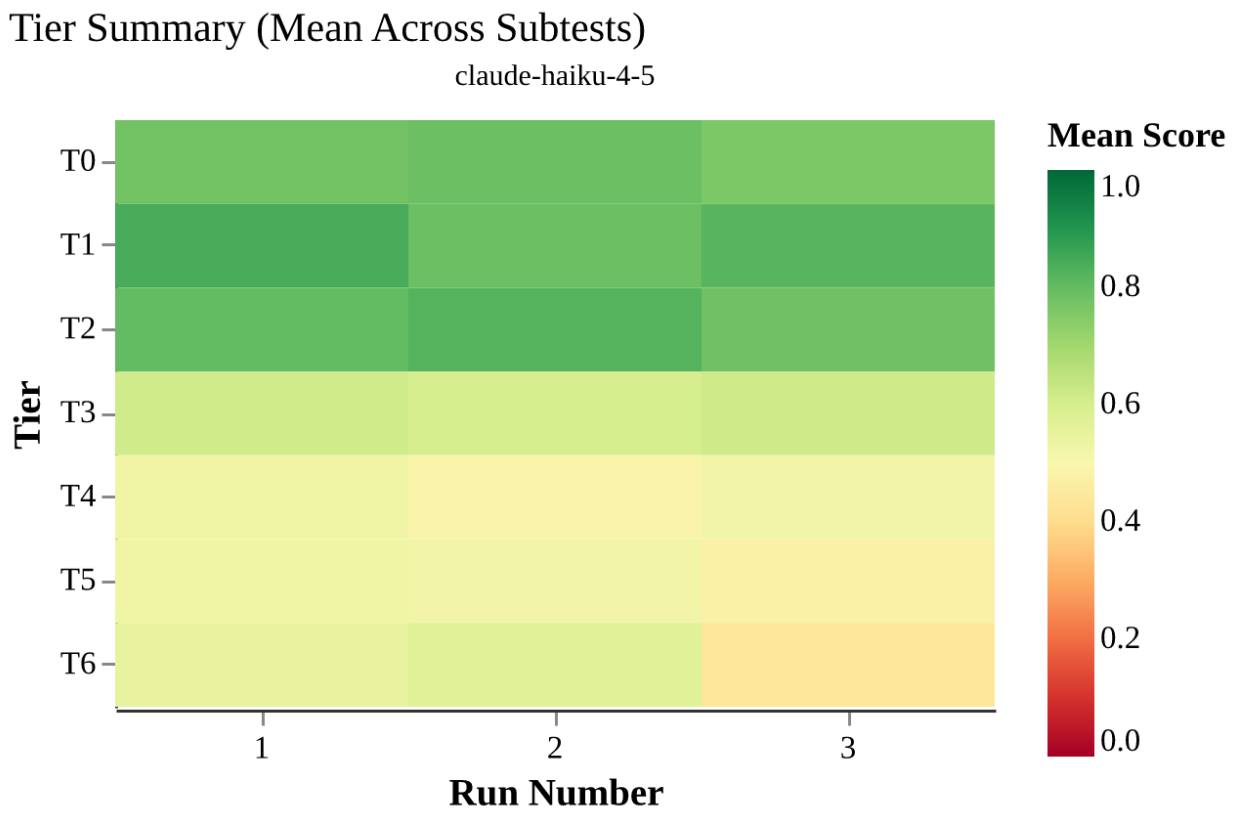

Figure 2: *Tier summary heatmap showing mean consensus scores by tier (y-axis) and run number (x-axis), faceted by model. Color encodes mean score (red–yellow–green scale). The degradation in tiers T3–T6 is visible as a shift toward red/yellow cells, while T0–T2 show predominantly green cells across runs.*

Table 3: *test-001 (Python Hello World): Per-tier results. Tiers T4–T6 achieve zero pass rate despite the trivial nature of the task.*

Tier	n	Pass Rate	Score	Cost/Run ($)	CoP ($)
T0	72	0.958	0.870	0.038	0.040
T1	30	1.000	0.908	0.037	0.037
T2	45	1.000	0.908	0.037	0.037
T3	123	0.350	0.318	0.038	0.109
T4	42	0.000	0.000	0.038	∞
T5	45	0.000	0.000	0.038	∞
T6	3	0.000	0.000	0.070	∞

in subtest 00, the empty-prompt configuration (`system_prompt_mode: none`), which achieves 0/3 passing across all three experiments (1/9 overall). With subtest 00 excluded, T0 achieves a perfect 1.000 pass rate on test-001, identical to T1 and T2. The empty-prompt agent receives the task instructions as a user message but has no system prompt providing behavioral framing. Examination of the agent transcripts reveals that rather than executing the task, the model meta-reasons about the evaluation framework—describing what the task asks for and posing clarifying questions ("Would you like me to execute this task?")—completing in only 2 API calls and ∼6 s of agent time. In contrast, subtest 01 (which adds only the default tool system prompt) executes the task immediately and succeeds. This suggests that the system prompt's role is not merely informational but *behavioral*: it provides the framing that shifts a budget model from meta-reasoning to task execution, even when the task instructions are unambiguous.

T3 (flat delegation) achieves a partial pass rate of 0.350, suggesting that some of the 41 specialist configurations can still solve the task, but the majority fail. The cost per run is nearly identical across all tiers ($0.037–$0.070), confirming that the agent consumes similar resources regardless of whether it produces useful output.

This result constitutes a clear case of *architectural over-specification*: the additional complexity of delegation is not merely unhelpful but actively harmful. The budget model lacks the capacity to simultaneously manage delegation protocols and solve the underlying task, even when the task itself is trivial. The T4–T6 failures are categorical rather than probabilistic—zero out of 90 runs across these three tiers produced a passing result.

Note on evaluation coverage. The judge evaluations for test-001 are incomplete: only 544 of the expected 1,080 judge assessments are present in the data, with T4, T5, and T6 having zero judge evaluations. The zero scores and pass rates for these tiers reflect the evaluation pipeline's assignment of score = 0 when no judge evaluation was obtained, which in turn occurs when the agent fails to produce evaluable output (e.g., produces no code artifact or terminates without a result). The cost data confirm that these runs did execute and consume tokens (mean cost $0.038–$0.070 per run), but the agent's output was not suitable for judge evaluation. Readers should note that the "catastrophic failure" interpretation rests on this pipeline behavior rather than on explicit negative judge assessments.

5.3 test-002: Mojo Hello World with Bazel and Documentation

The moderate Mojo task presents a fundamentally different pattern from test-001. All tiers achieve high pass rates, with substantially less differentiation across the architectural spectrum:

For this task, tier effects are subtle and, notably, the direction of the effect is *reversed* relative to test-001. Pass rates increase slightly from T0 (0.903) to T6 (1.000), with a modest upward trend. This reversal is a key driver of the significant tier × experiment interaction discussed in Section 5.6.

Scores cluster between 0.641 and 0.722, suggesting that while the agent consistently attempts and partially completes the task, the quality ceiling is bounded by factors other than architecture—likely the model's familiarity with Mojo syntax and Bazel build system configuration. The relatively narrow score range (0.081 span) contrasts sharply with test-001's range (0.908 span), indicating that the task difficulty, rather than the architecture, is the primary determinant of output quality.

Table 4: *test-002 (Mojo Hello World): Per-tier results. All tiers achieve pass rates above 0.90, with modest variation.*

Tier	n	Pass Rate	Score	Cost/Run ($)	CoP ($)
T0	72	0.903	0.641	0.219	0.242
T1	30	0.933	0.676	0.229	0.245
T2	45	0.956	0.685	0.222	0.232
T3	123	0.967	0.686	0.229	0.237
T4	42	0.976	0.677	0.236	0.242
T5	45	0.956	0.694	0.217	0.228
T6	3	1.000	0.722	0.279	0.279

The Cost-of-Pass is remarkably uniform across tiers ($0.228–$0.279), consistent with the aggregate cost invariance finding. T6 ($n = 3$) achieves the highest pass rate (1.000) but at the highest CoP ($0.279), reflecting its elevated per-run cost.

One interpretation of these results is that the Mojo task's difficulty provides sufficient "work" to justify the coordination overhead of delegation-based tiers. Unlike the trivial Python task, where delegation is unnecessary, the Mojo task involves multiple components (source code, build configuration, documentation) that can genuinely benefit from decomposition. A complementary explanation is that the model's limited familiarity with Mojo—noted in Section 4.2 as having "limited representation in the model's training data"—means that specialist agents in delegation-based tiers may provide language-specific context or tooling that partially compensates for the base model's limited training exposure. This advantage is absent for well-represented languages like Python, where the base model already possesses deep knowledge.

5.4 test-003: Add mypy to pixi Configuration

The mypy configuration task shows a pattern intermediate between the other two experiments, with uniformly high pass rates and subtle tier effects:

Table 5: *test-003 (Add mypy to pixi): Per-tier results. High pass rates across all tiers with modest variation.*

Tier	n	Pass Rate	Score	Cost/Run ($)	CoP ($)
T0	72	0.944	0.817	0.071	0.076
T1	30	0.967	0.865	0.074	0.076
T2	45	0.978	0.814	0.073	0.074
T3	123	0.959	0.826	0.071	0.074
T4	42	0.952	0.857	0.076	0.080
T5	45	0.933	0.829	0.098	0.105
T6	3	1.000	0.833	0.100	0.100

Pass rates are uniformly high (0.933–1.000) with no clear monotonic trend. Scores range from 0.814 to 0.865, with T1 (skills) achieving the highest score (0.865). This suggests that for configuration tasks, domain-specific skills provide marginally more benefit than raw tool access or delegation,

possibly because skills contain procedural knowledge about configuration management that the base model does not possess.

An interesting anomaly is T5 (hybrid), which shows elevated cost ($0.098/run vs. $0.071–$0.076 for T0–T3) without a corresponding quality improvement (pass rate 0.933 is the lowest in the table). The hybrid configurations appear to introduce overhead from combining multiple architectural elements without the coordination benefits that might emerge on more complex tasks.

T6 achieves a perfect pass rate (1.000) but with only $n = 3$ runs, this result should be interpreted with extreme caution. At the aggregate level ($n = 9$ across three experiments), T6's pass rate is 0.667. The Clopper–Pearson exact 95% confidence interval [3] $[0.299, 0.925]$ is the preferred interval at this sample size, as it maintains nominal coverage for binary data without relying on asymptotic approximations. For reference, the BCa bootstrap [4] interval $[0.333, 0.889]$ is narrower and may exhibit poor coverage at $n = 9$. Under either interval, a true pass rate as low as 0.299 is not excluded.

5.5 Cross-Task Analysis

Aggregating across all three experiments, Table 6 presents the tier-level summary statistics.

Table 6: *Tier Summary Statistics*

Model	Tier	Subtests	Pass Rate (95% CI)	Mean Score ($\pm\sigma$)	Median	Consistency	CoP ($)
claude-haiku-4-5	T0	24	0.935 (0.898, 0.963)	0.776 ± 0.206	0.874	0.734	0.12
claude-haiku-4-5	T1	10	0.967 (0.911, 0.989)	0.816 ± 0.140	0.887	0.829	0.12
claude-haiku-4-5	T2	15	0.978 (0.941, 0.993)	0.802 ± 0.145	0.872	0.819	0.11
claude-haiku-4-5	T3	41	0.759 (0.713, 0.799)	0.610 ± 0.347	0.704	0.431	0.15
claude-haiku-4-5	T4	14	0.643 (0.556, 0.722)	0.511 ± 0.385	0.643	0.247	0.18
claude-haiku-4-5	T5	15	0.630 (0.548, 0.711)	0.508 ± 0.382	0.667	0.248	0.19
claude-haiku-4-5	T6	1	0.667 (0.333, 0.889)	0.518 ± 0.403	0.680	0.222	0.22

The aggregate results reveal a clear bimodal pattern: tiers T0–T2 form a high-performance cluster (pass rate 0.935–0.978, score 0.776–0.816), while tiers T3–T6 form a degraded cluster (pass rate 0.630–0.759, score 0.508–0.610). The transition between these clusters occurs at the T2–T3 boundary, corresponding to the introduction of multi-agent delegation.

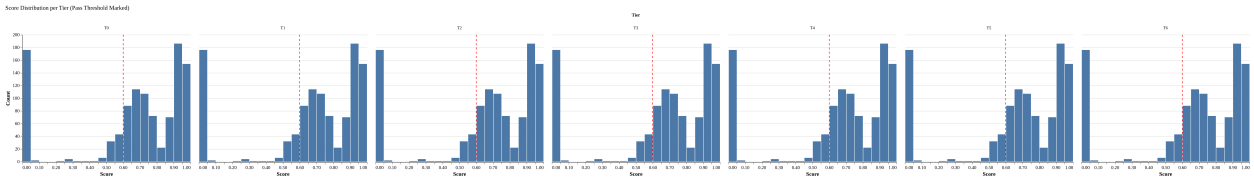

Figure 3: *Score distributions per tier (histogram, bin width 0.05). Each panel shows the distribution of per-run consensus scores for one tier. The red dashed vertical line marks the pass threshold. T0–T2 distributions are concentrated near 1.0; T3–T6 show bimodal distributions with a large mass near 0 (complete failures) alongside a cluster of successful runs.*

Consistency (the within-subtest reproducibility of scores) shows a parallel pattern. T0–T2 achieve consistency values of 0.734–0.829, while T3–T6 range from 0.222 to 0.431. This indicates that the higher tiers are not only worse on average but also less predictable—their outcomes are more variable from run to run, making them unreliable for production deployment.

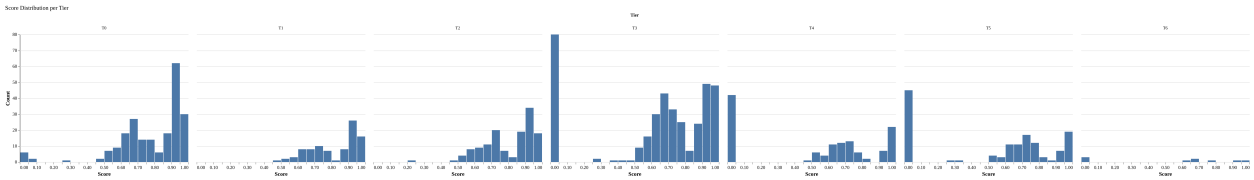

Figure 4: *Score distributions by tier. The variance increases substantially from T3 onward, reflecting the bimodal outcome distribution in delegation-based tiers: runs either succeed (high score) or fail completely (zero score).*

Quantitatively, the bimodality is striking: for T3, 22% of runs score below 0.1 (complete failure) while 50% score above 0.7 (successful completion), with 28% in between. For T4–T5, the failure mode is more pronounced (∼33% below 0.1, ∼40–44% above 0.7). By contrast, T0–T2 are effectively unimodal at high scores (67–76% above 0.7, 0–4% below 0.1). This pattern suggests a binary delegation outcome: either the coordination succeeds and produces high-quality output, or it fails entirely and produces near-zero output, with relatively few intermediate outcomes.

The Kruskal–Wallis omnibus test reveals significant tier effects for three of five metrics:

- Pass rate: $H(6) = 119.01$, $p < 0.001$ (significant)[3]
- Score: $H(6) = 91.3$ (from SRH tier factor), $p < 0.001$ (significant)
- Duration: $H(6) = 40.94$, $p < 0.001$ (significant)
- Cost: $H(6) = 4.00$, $p = 0.676$ (not significant)
- Implementation rate: $H(6) = 5.47$, $p = 0.486$ (not significant)

Note that the Score and Cost tier statistics above are drawn from the Scheirer–Ray–Hare tier main effect (Section 5.6), as the omnibus pipeline tests only pass rate, implementation rate, and duration directly.

The non-significance of cost and implementation rate is analytically important. It means that all tiers cost approximately the same and that the agent's propensity to *attempt* the task is not affected by architecture. The differences emerge only in the *quality* of the attempt and the time required to complete it. The significant duration effect ($H(6) = 40.94$, $p < 0.001$) reveals a counter-intuitive pattern: single-agent tiers (T0–T2) run substantially longer than delegation tiers (T3–T6) in aggregate (e.g., T0 mean = 346 s vs. T3 mean = 248 s). This is driven primarily by the rapid failures of T4–T6 on test-001, where runs complete in ∼17 s (consistent with immediate orchestration errors) versus 430–505 s for T0–T2. A full duration analysis is deferred to future

[3]For the binary pass/fail outcome, a chi-squared test on the 2×7 contingency table would be the more natural choice. We use Kruskal–Wallis throughout for a unified analytical framework; the conclusions are qualitatively identical under both tests.

work; we note that wall-clock time represents an operational cost dimension not captured by the token-based CoP metric.

Pairwise comparisons with Holm–Bonferroni correction identify the following significant and near-significant transitions:

- **T2 → T3**: $p < 0.001$, Cliff's $\delta = -0.219$ $[-0.268, -0.168]$ (small effect). This marks the sharpest degradation point, where the introduction of delegation produces a significant drop in pass rate.
- **T3 → T4**: $p = 0.058$ (n.s. after Holm–Bonferroni correction), Cliff's $\delta = -0.116$ $[-0.212, -0.023]$ (small effect; the bootstrap CI excludes zero but is not adjusted for multiple comparisons). Hierarchical orchestration may further degrade performance beyond flat delegation, though this does not reach corrected significance.
- **T0 → T6**: $p = 0.018$. The overall trend from simplest to most complex is significantly negative.
- **Duration T2 → T3**: $p = 0.048$, Cliff's $\delta = -0.154$ $[-0.264, -0.045]$. T3 runs are stochastically shorter than T2, driven partly by the rapid failures of delegation-based tiers on test-001 (T4–T6 complete in ∼17 s vs. 430–505 s for T0–T2).

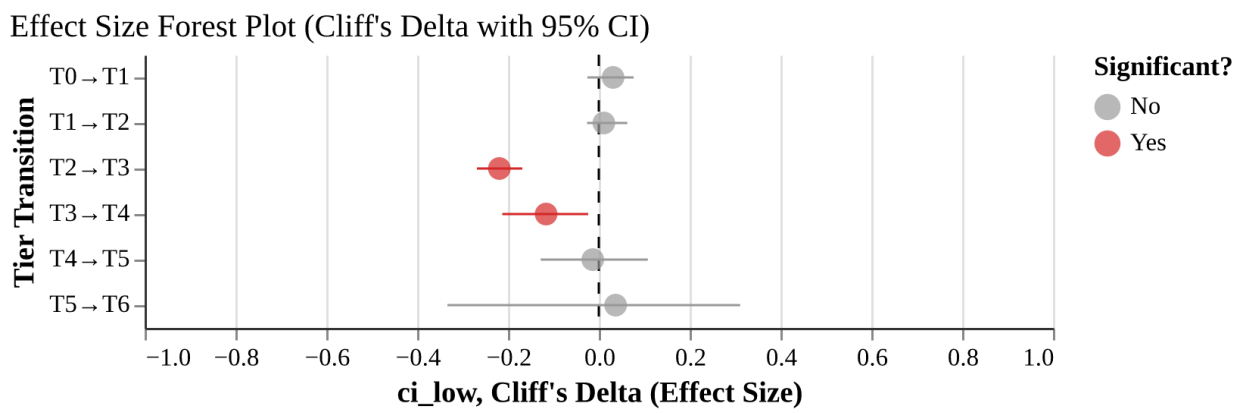

Figure 5: *Forest plot of Cliff's δ effect sizes for pairwise tier comparisons on pass rate (lower-numbered tier is reference group; negative δ indicates the higher-numbered tier has lower pass rates). The T2–T3 comparison shows the largest magnitude effect. Error bars show 95% bootstrap confidence intervals.*

5.6 Tier × Task Interaction

The most analytically important finding is the significant tier × experiment interaction. The Scheirer–Ray–Hare (SRH) test provides a non-parametric decomposition of the effects of tier, experiment, and their interaction on the ranked outcome variables:

All three score effects are significant at $p < 0.001$: the tier × experiment interaction ($H = 223.5$, $df = 12$), the experiment main effect ($H = 178.2$, $df = 2$), and the tier main effect ($H = 91.3$, $df = 6$). Raw H-statistics are not directly comparable across terms with different degrees of

Table 7: *Scheirer–Ray–Hare analysis of tier × experiment interactions. The score interaction ($H = 223.5$) exceeds both main effects, indicating that tier effects are strongly task-dependent.*

Metric	Factor	H	df	p	Sig?
Score	Experiment	178.2	2	< 0.001	***
Score	Tier	91.3	6	< 0.001	***
Score	Interaction	223.5	12	< 0.001	***
Cost	Experiment	903.1	2	< 0.001	***
Cost	Tier	4.0	6	0.676	n.s.
Cost	Interaction	4.8	12	0.965	n.s.
Impl Rate	Experiment	490.7	2	< 0.001	***
Impl Rate	Tier	5.5	6	0.486	n.s.
Impl Rate	Interaction	3.7	12	0.989	n.s.

freedom, but the significance of all three effects—particularly the interaction—indicates that the impact of architectural complexity on quality is not uniform: it depends critically on the task being performed.

The interaction is driven by the divergent tier profiles across experiments. In test-001, tiers T4–T6 score zero while T0–T2 score above 0.87. In test-002, the tier profile is nearly flat, with all tiers scoring between 0.64 and 0.72. In test-003, the profile is again relatively flat, with modest variation. The aggregation of these divergent profiles produces the significant interaction.

For cost, the experiment main effect is dominant ($H = 903.1$, $p < 0.001$), reflecting the differing cost profiles of the three tasks. test-002 (Mojo) is approximately 6× as expensive per run as test-001 (Python Hello World). However, neither the tier main effect ($H = 4.0$, $p = 0.676$) nor the interaction ($H = 4.8$, $p = 0.965$) is significant for cost, indicating that architectural complexity does not systematically affect expenditure.

The implementation rate shows a similar pattern to cost: strong experiment effects ($H = 490.7$, $p < 0.001$) but no significant tier effects or interactions. This indicates that the agent's *willingness to attempt* the task is driven by task characteristics rather than architectural configuration, while the *quality of the attempt* is architecture-dependent.

Figure 7 provides a granular view of grade distributions across all subtests and experiments, illustrating the divergent patterns that drive the interaction effect.

5.7 Per-Criterion Analysis

The judge panel evaluates each run against five weighted criteria: functional correctness, code quality, build pipeline integrity, proportionality (response length and scope appropriateness), and overall quality (a holistic assessment). Across all runs, the per-criterion mean scores reveal that the budget model's weakest dimension is functional correctness (0.838 ± 0.276), while code quality (0.972 ± 0.137), build pipeline integrity (0.941 ± 0.138), and proportionality (0.927 ± 0.137) score substantially higher. Overall quality averages 0.903 ± 0.133. This pattern—high marks on surface quality but lower marks on functional correctness—is consistent with the hypothesis that the budget

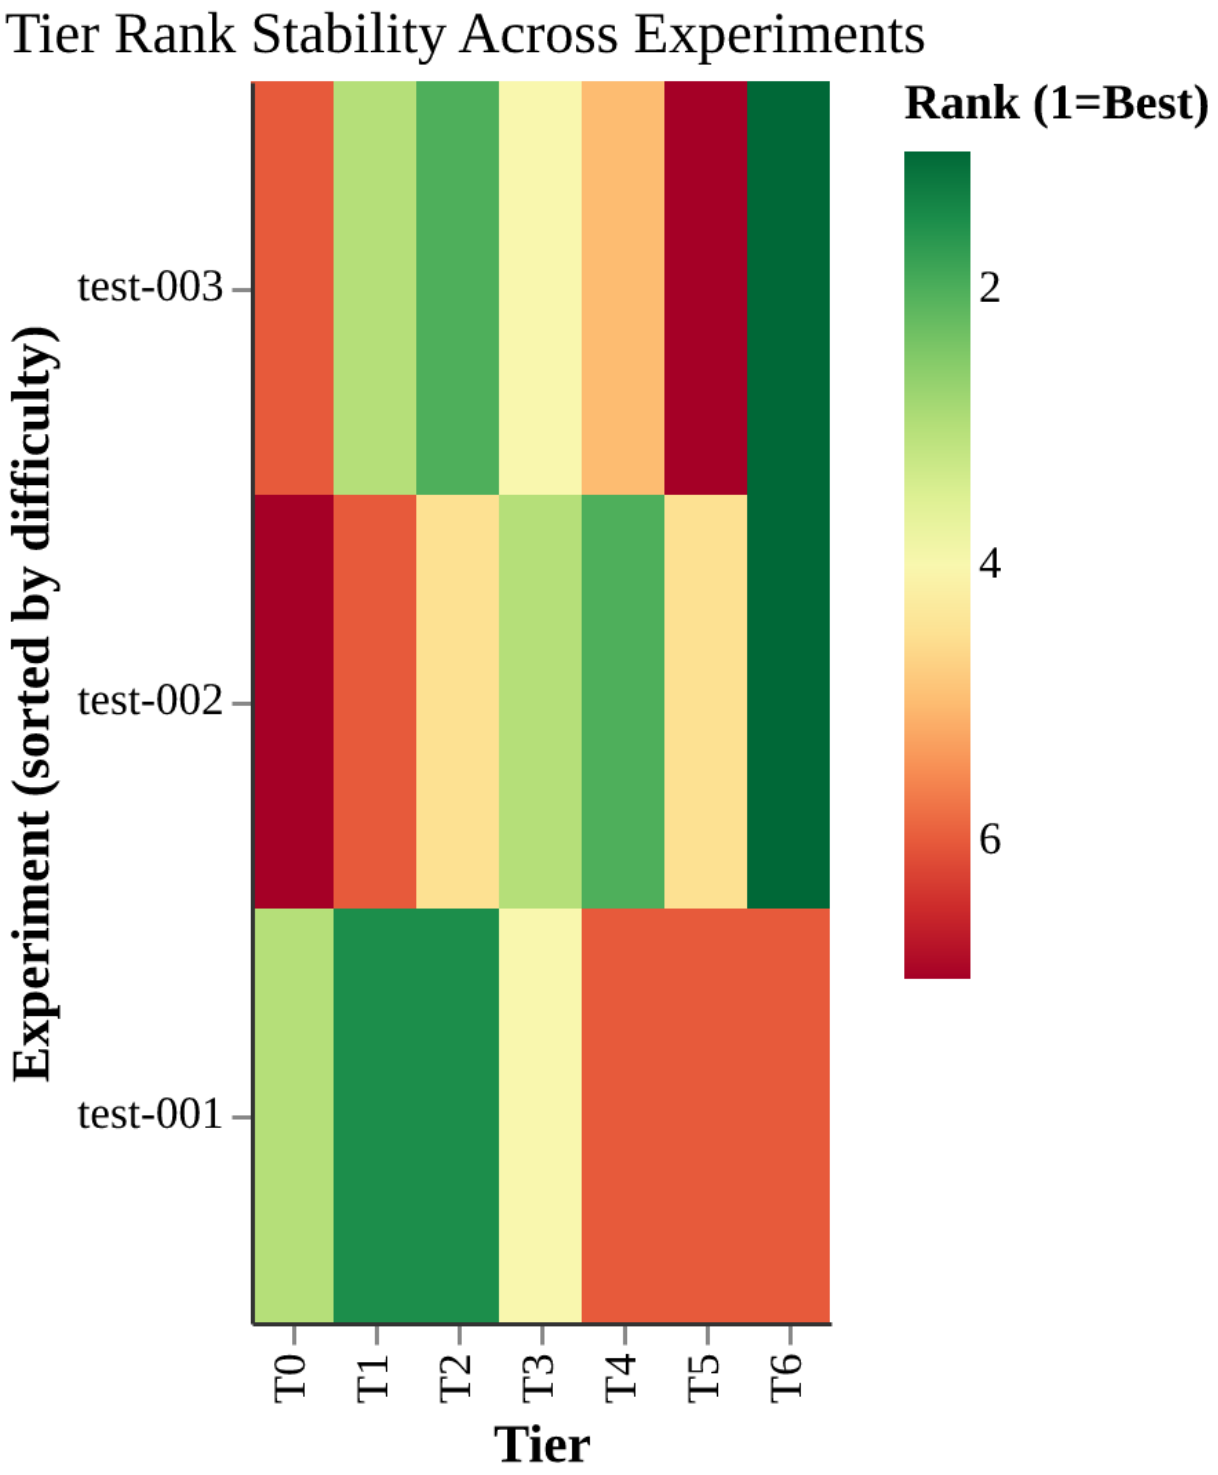

Figure 6: *Tier rank stability heatmap. Each cell shows the rank of a tier (x-axis) within a given experiment (y-axis, sorted by difficulty), where rank 1 is best. Color encodes rank (green = best, red = worst). The changing color pattern across experiments indicates that tier rankings are not stable—the optimal tier varies by task, consistent with the significant tier × experiment interaction.*

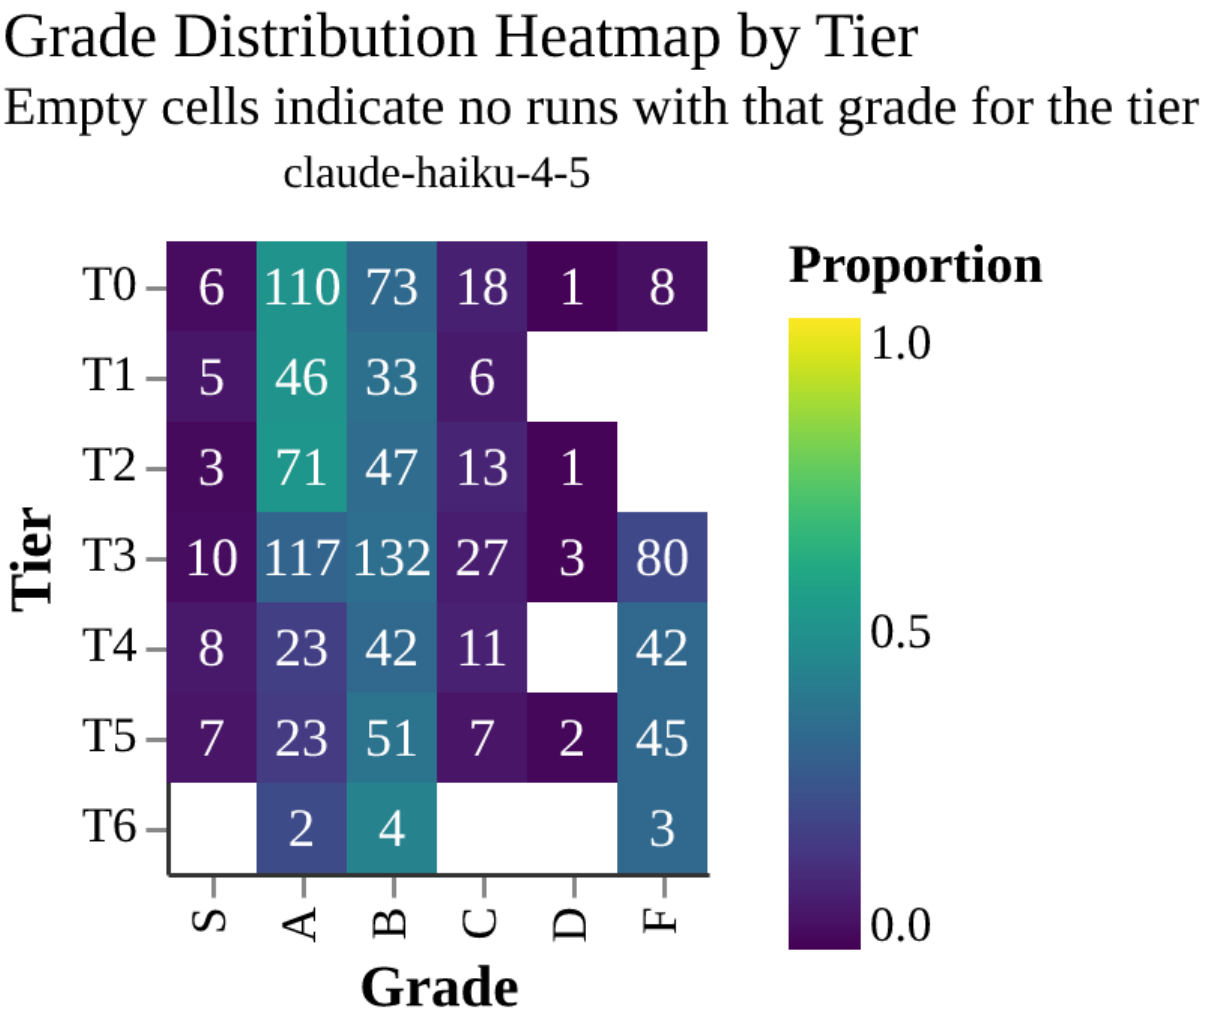

Figure 7: *Grade distribution heatmap. Each cell represents the proportion of runs assigned each letter grade for a tier–grade combination, faceted by model. Darker (higher proportion) cells in the F row for T3–T6 on test-001 contrast with the predominantly A/B grades for test-002 and test-003.*

model produces syntactically well-formed code that occasionally fails to satisfy the task's functional requirements.

6 Inter-Rater Reliability

6.1 Panel Design

The three-judge panel employs models spanning the full Anthropic capability spectrum. Judge 1 (`claude-opus-4-6`) is the most capable model in the lineup, Judge 2 (`claude-sonnet-4-6`) is mid-tier, and Judge 3 (`claude-haiku-4-5`) is the budget model—the same model used as the subject agent.

This design was intentional: using the subject model as one of its own judges enables analysis of self-evaluation bias. However, it also introduces the risk that the budget judge may lack the evaluative capacity to distinguish quality levels in its own outputs. The panel represents a pragmatic compromise between evaluation cost (using only frontier judges would be expensive) and reliability (using only budget judges might be unreliable).

6.2 Agreement Results

Table 8 presents the full inter-rater reliability statistics.

Table 8: *Judge Agreement Metrics*

Judge Pair	Spearman ρ	Pearson r	Mean $\lvert\Delta\text{Score}\rvert$
Judge 1 – Judge 2	0.125	0.505	0.0888
Judge 1 – Judge 3	0.142	0.242	0.2096
Judge 2 – Judge 3	0.054	0.185	0.2701
Krippendorff's α (interval): -0.030			

The headline finding is that Krippendorff's α (interval) is -0.030 (a confidence interval for α is not separately reported here; all values of $\alpha < 0$ are below the chance level of 0), computed on $N = 360$ items (120 subtests $\times$ 3 runs per subtest, with scores averaged across the three experiments for each tier–subtest–run combination) from the 2,696 available judge evaluations (see Section 8.4 for discussion of missing evaluations). Note that averaging scores across experiments before computing α conflates inter-experiment variability with inter-rater disagreement; computing α on individual runs (without averaging across experiments) may yield different results. The resulting $\alpha = -0.030$ is below chance agreement ($\alpha = 0$). By conventional standards in content analysis [9, 13], $\alpha < 0.667$ is considered unacceptable for drawing conclusions about individual cases, and $\alpha > 0.800$ is required for reliable measurement. Our value of -0.030 indicates that the three judges agree less than would be expected by random scoring.

The negative value of α is particularly concerning. While $\alpha = 0$ would indicate random agreement, $\alpha < 0$ suggests systematic *disagreement*—the judges are not merely uncorrelated but are actively

producing divergent evaluations. This could arise if the judges apply fundamentally different evaluation criteria or if their evaluation strategies are negatively correlated (e.g., one judge favoring verbose output while another penalizes it).

Pairwise correlations reveal a structured pattern of disagreement. Spearman rank correlations are near zero for all pairs:

- Judge 1 (Opus) vs. Judge 2 (Sonnet): $\rho = 0.125$
- Judge 1 (Opus) vs. Judge 3 (Haiku): $\rho = 0.142$
- Judge 2 (Sonnet) vs. Judge 3 (Haiku): $\rho = 0.054$

These near-zero rank correlations indicate that the judges do not agree on the *relative ordering* of outputs. An output that Opus rates highly may be rated poorly by Sonnet or Haiku, and vice versa.

Pearson correlations are somewhat higher, suggesting some linear association in score magnitudes:

- Judge 1 (Opus) vs. Judge 2 (Sonnet): $r = 0.505$
- Judge 1 (Opus) vs. Judge 3 (Haiku): $r = 0.242$
- Judge 2 (Sonnet) vs. Judge 3 (Haiku): $r = 0.185$

The discrepancy between Spearman and Pearson correlations (particularly for Opus–Sonnet: $\rho = 0.125$ vs. $r = 0.505$) suggests a non-linear relationship between judge scores. The moderate Pearson correlation may reflect shared calibration of score scale (both judges tend to use similar score ranges) even though they disagree on which outputs deserve high vs. low scores.

(a) *T0 (Prompts)* **(b)** *T3 (Delegation)*

Figure 8: *Judge agreement scatter plots for T0 and T3. Each point represents a run, with axes showing scores from different judge pairs. The wide scatter and lack of diagonal clustering illustrate the low inter-rater reliability across both high-performing (T0) and low-performing (T3) tiers.*

6.3 Judge Bias

The mean absolute score differences reveal a systematic pattern in which disagreement scales with the capability gap between judges:

- Judge 1 (Opus) vs. Judge 2 (Sonnet): mean $|\Delta| = 0.089$

- Judge 1 (Opus) vs. Judge 3 (Haiku): mean $|\Delta| = 0.210$
- Judge 2 (Sonnet) vs. Judge 3 (Haiku): mean $|\Delta| = 0.270$

The two most capable models (Opus and Sonnet) agree most closely, with a mean absolute difference of only 0.089 on a $[0, 1]$ scale. As the capability gap increases, disagreement grows: Opus and Haiku disagree by 0.210, while Sonnet and Haiku—the pair with the most divergent capability levels after accounting for the specific evaluation task—disagree by 0.270.

This monotonic relationship between capability gap and disagreement is consistent with the hypothesis that the budget model lacks sufficient evaluative capacity to produce reliable judgments. Haiku may be applying different evaluation criteria, assigning scores based on surface features (e.g., output length, code structure) rather than the deeper quality assessments (e.g., correctness, completeness, adherence to specifications) performed by the frontier judges.

The near-zero Krippendorff's α raises a methodological concern that affects the interpretation of all results in this paper. If the judges do not agree, then the consensus score (arithmetic mean) may not be a reliable measure of output quality. The mean of disagreeing judges produces a compromise value that no individual judge endorses, potentially masking genuine quality differences between tiers. However, the pass rate metric (which uses majority vote over each judge's individual binary classification) may be more robust to this issue, as it requires only that a majority of judges agree on a pass/fail outcome rather than precisely calibrated scores.

7 Cost Analysis

7.1 Cost-of-Pass by Tier

Table 9 presents the detailed cost analysis by tier, including mean cost per run, total cost, Cost-of-Pass, and token usage statistics.

Table 9: *Cost Analysis by Tier*

Model	Tier	Mean Cost ($)	Total Cost ($)	CoP ($)	Input	Output	Cache Read	Cache Crea
claude-haiku-4-5	T0	0.11	23.65	0.12	67	2409	593588	245
claude-haiku-4-5	T1	0.11	10.19	0.12	42	2601	644806	246
claude-haiku-4-5	T2	0.11	14.89	0.11	43	2659	638340	233
claude-haiku-4-5	T3	0.11	41.62	0.15	60	2637	632159	242
claude-haiku-4-5	T4	0.12	14.71	0.18	41	2584	649313	244
claude-haiku-4-5	T5	0.12	15.91	0.19	49	2565	657564	281
claude-haiku-4-5	T6	0.15	1.35	0.22	47	2665	772688	473
claude-haiku-4-5	**Total**	0.11	122.31	—	58,559†	2,782,286†	683,228,777†	26,920,15

Token columns show per-run means for individual tiers and absolute totals (†) for the Total row.

The aggregate Cost-of-Pass ranges from $0.11 (T2) to $0.22 (T6). This variation is driven entirely by the denominator (pass rate) rather than the numerator (mean cost), since mean cost per run is effectively constant across tiers ($0.11–$0.15). Applying Equation 1, the CoP for T6 is elevated

because its lower pass rate (0.667) divides a slightly higher mean cost ($0.15), yielding $0.22 per successful outcome.

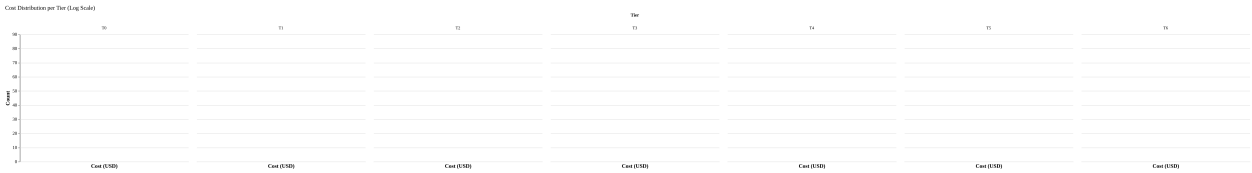

Figure 9: *Cost-of-Pass by tier. The increasing CoP from T2 to T6 reflects declining pass rates rather than increasing costs. T2 achieves the minimum CoP at $0.11, making it the most cost-efficient tier in the aggregate.*

T2 (Tooling) achieves the lowest aggregate CoP at $0.11, making it the apparent Pareto-optimal tier in the aggregate data: it has both the highest aggregate pass rate (0.978) and the lowest cost-of-pass, though the optimal tier varies by task (Section 5.6). (Cost invariance is a non-significant finding under limited power; real cost differences between tiers may exist below our ability to detect.) T1 (Skills) is a close second with CoP $0.12 and pass rate 0.967. The frontier CoP across all tiers and experiments is $0.113.

The practical implication is clear: for budget models on the tasks studied, the optimal strategy is to equip the agent with tools (T2) but not delegation or orchestration. This configuration maximizes quality while minimizing the expected cost per successful outcome.

7.2 Cost Scaling with Task Difficulty

Cost varies dramatically by experiment but not by tier. The mean cost per run ranges from $0.038 for test-001 (trivial Python) to $0.226 for test-002 (moderate Mojo), with test-003 (moderate Python configuration) at $0.076. The Kruskal–Wallis test shows a highly significant experiment effect on cost ($H = 903.1$, $p < 0.001$) but no significant tier effect ($H = 4.0$, $p = 0.676$).

The cost differential between experiments is attributable to the differing context requirements of each task. The Mojo task (test-002) requires the agent to process a larger repository with more complex build configuration, consuming more tokens for context. The Python Hello World task (test-001) operates in a minimal repository, requiring minimal context.

This observed cost invariance across tiers ($p = 0.676$, non-significant under limited power) has a practical implication: if tiers indeed cost approximately the same per run, the optimal strategy is simply to use the tier with the highest pass rate. In the aggregate data, the Pareto frontier collapses to a narrow region around T1–T2, though per-task analysis shows the frontier varies (T5 dominates for test-002).

7.3 Frontier Analysis

Figure 11 presents the cost–quality Pareto frontier across all tier × experiment combinations.

The Pareto frontier is dominated by T0–T2 configurations across all three experiments. No T4–T6 configuration appears on the Pareto frontier for any experiment. T3 appears on the frontier only for test-002, where its mean score is competitive with T2.

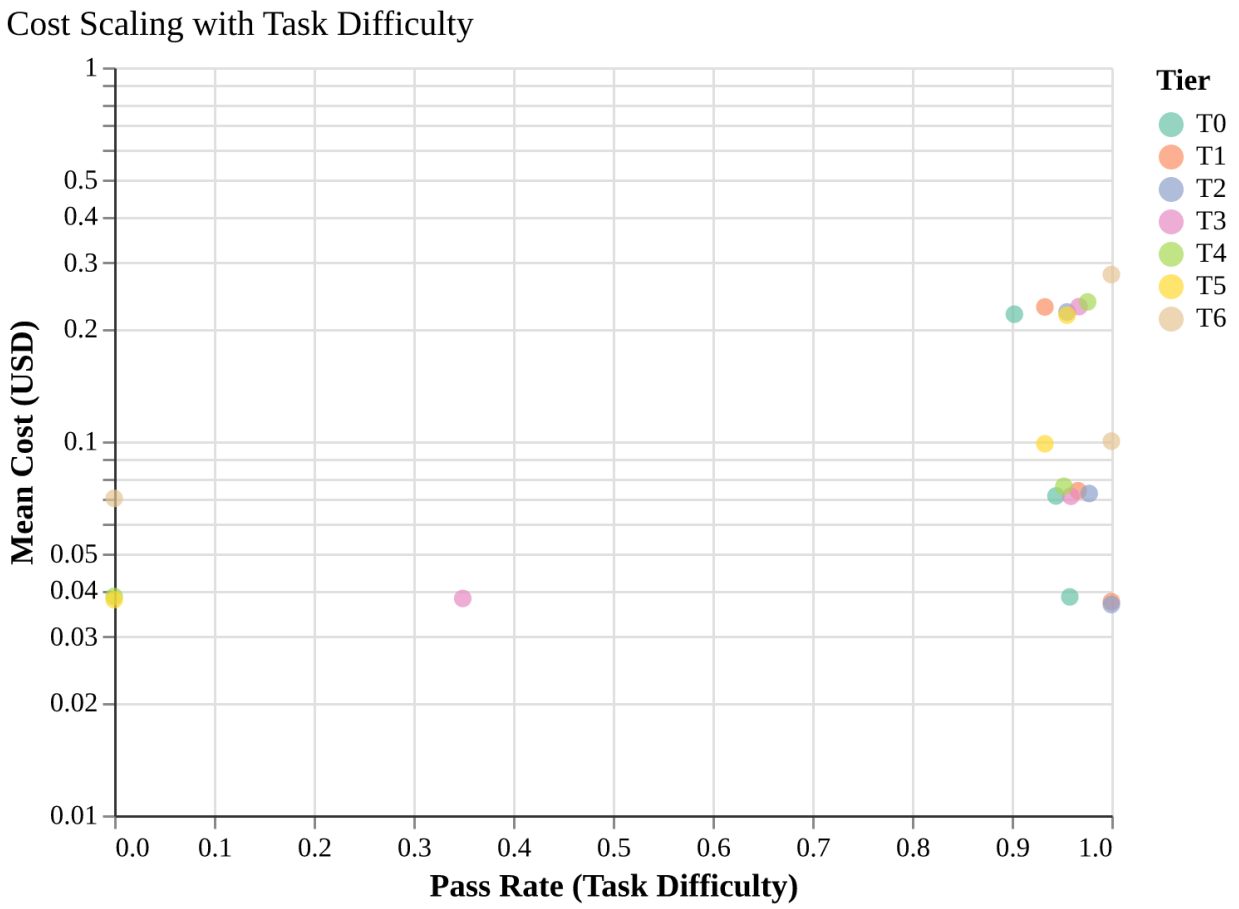

Figure 10: *Cost scaling across experiments. Cost is driven primarily by task characteristics—repository size, language complexity, and required context—rather than architectural tier. The spread within each experiment reflects run-to-run variance, not tier differences.*

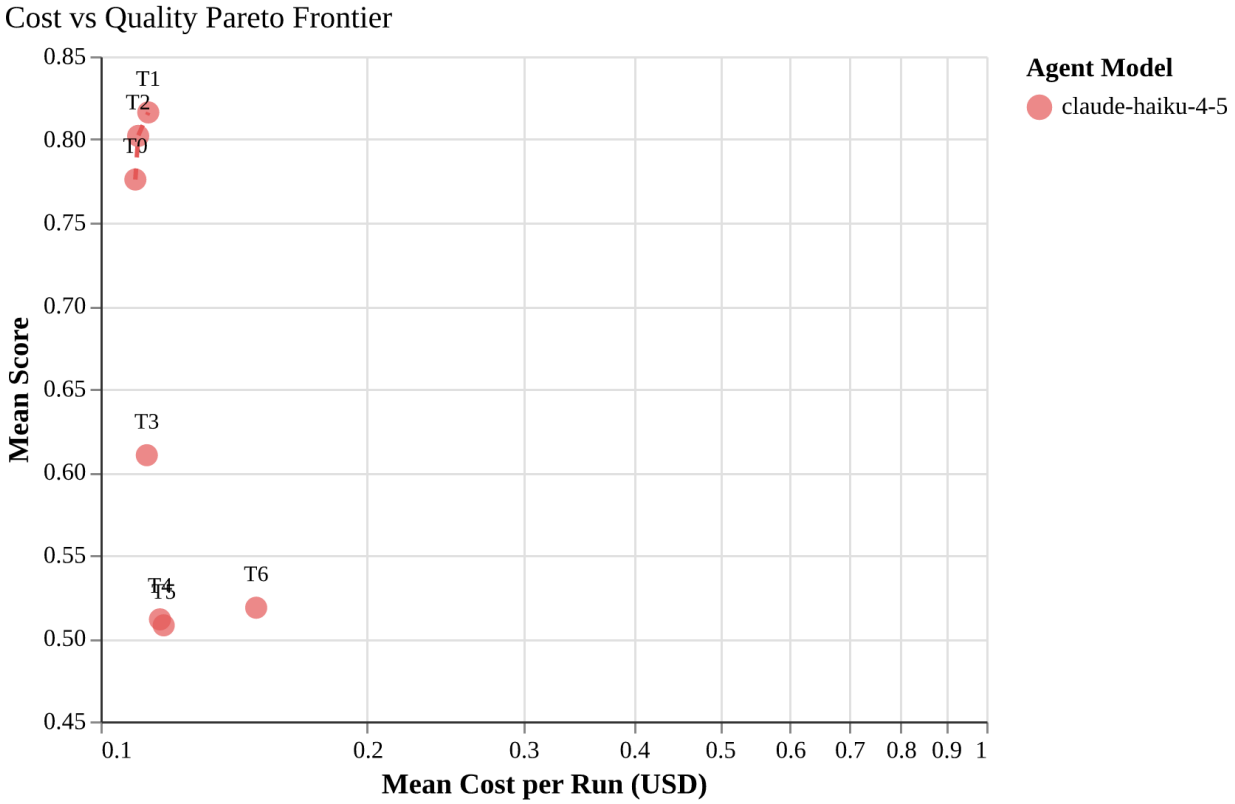

Figure 11: *Cost–quality Pareto frontier. Each point represents a tier–experiment combination, plotted by mean cost per run (x-axis, log scale) and mean consensus score (y-axis). Pareto-optimal points lie on the upper-left frontier. T0–T2 configurations dominate the frontier across all experiments; T4–T6 for test-001 cluster at low quality with non-zero cost.*

The test-001 T4–T6 points (zero pass rate, non-zero cost) represent the worst possible outcome: the organization pays for computation but receives no useful output. These configurations are not merely suboptimal—they are wasteful, consuming budget without any return.

The total experimental cost of \$122.31 across 1,080 runs yields a mean cost of \$0.11 per run. This shows that meaningful ablation studies can be conducted at modest cost using budget models, making the Scylla framework accessible for iterative evaluation during development cycles rather than requiring dedicated evaluation budgets.

8 Discussion

8.1 Key Findings

Our results suggest several preliminary conclusions about the interaction between architectural complexity and budget model performance. We organize these findings around the four research questions posed in Section 1.

Finding 1 (RQ1): In this study, the T2–T3 boundary is the sharpest transition. The introduction of delegation at T3 is associated with a statistically significant drop in aggregate pass rate from 0.978 to 0.759 (Cliff's $\delta = -0.219$, $p < 0.001$). This is the largest single-step degradation in the tier hierarchy and represents the transition from individual agent execution to multi-agent coordination. For the budget model studied, the introduction of delegation is associated with substantial performance degradation. One plausible mechanism is that coordination overhead consumes reasoning resources that would otherwise be applied to the task, though the confounded tier design (Section 8.5) does not allow definitive causal attribution.

The subsequent T3–T4 transition (Cliff's $\delta = -0.116$, small-to-negligible effect by conventional thresholds; see Appendix D) suggests possible further degradation from adding hierarchical orchestration on top of delegation, though this comparison does not reach significance after Holm–Bonferroni correction ($p = 0.058$). The cumulative effect of moving from T2 to T6 is a drop in pass rate from 0.978 to 0.667—a loss of approximately one-third of the base capability.

Finding 2 (RQ1): In this study, architectural over-specification is associated with catastrophic failure on the trivial Python task. The zero pass rate for T4–T6 on test-001 is the most striking individual result in the study. A task that the bare model handles with 95.8% pass rate becomes completely unsolvable when wrapped in hierarchical delegation. This suggests that budget models may be unable to simultaneously manage coordination overhead and task execution, even when the task itself is trivial. The coordination overhead appears to exceed a threshold beyond which task completion is entirely prevented, rather than merely degrading performance incrementally. However, the specific failure mechanism cannot be determined from observational data alone.

Finding 3 (RQ2): In this study, tier effects are task-dependent. The significant SRH interaction effect ($H = 223.5$, $df = 12$, $p < 0.001$) is the central finding of the cross-task analysis. All three effects—interaction, tier main effect ($H = 91.3$, $df = 6$), and experiment main effect ($H = 178.2$, $df = 2$)—are significant at $p < 0.001$, though their raw H-values are not directly comparable across different degrees of freedom. The highly significant interaction indicates that

the *combination* of tier and task determines quality beyond either factor alone. This means there is no universally optimal tier: the best architecture depends on the task. For trivial tasks, T0–T2 are clearly superior; for moderate tasks, the tiers are more interchangeable, with some evidence that higher tiers provide marginal benefits. (Note that this finding is derived from continuous score data; given the low inter-rater reliability ($\alpha = -0.030$), the pass-rate-based KW analysis ($H(6) = 119.01$, $p < 0.001$) provides independent, more robust confirmation of tier-dependent effects.)

Finding 4 (RQ3): In this study, no significant cost variation across tiers is detected. Cost-of-Pass ranges from \$0.11 at T2 to \$0.22 at T6, driven by declining pass rates rather than rising per-run costs. The non-significant tier effect on cost ($p = 0.676$) provides no evidence of a cost–quality trade-off under current statistical power. If higher tiers provided cost savings (e.g., by solving tasks more efficiently with fewer tokens), there might be a case for accepting lower quality in exchange for lower cost. The absence of detected cost variation means that the quality degradation in higher tiers appears as a net loss with no detected compensating benefit in the data.

Finding 5 (RQ4): In this study, the three-judge panel exhibits near-zero reliability. With Krippendorff's $\alpha = -0.030$, the judge panel produces scores that are effectively random relative to each other. The disagreement is structured: the two frontier judges (Opus and Sonnet) show the smallest pairwise disagreement (mean $|\Delta| = 0.089$, though Spearman $\rho = 0.125$ indicates weak rank agreement), while the budget judge (Haiku) disagrees substantially with both (mean $|\Delta| = 0.210$ and 0.270). This suggests that including a budget model in the judge panel introduces noise rather than signal, and that a two-judge panel of frontier models might provide more reliable evaluation.

8.2 Practical Implications

These findings have several practical implications for organizations deploying agentic coding systems with budget models.

Architecture selection should be conservative. Based on the tasks studied here, the default recommendation for budget model deployments is to use T2 (Tooling) as the baseline architecture. This configuration provides the agent with external tools and MCP servers while avoiding the delegation overhead that degrades aggregate performance. T2 achieves the highest aggregate pass rate (0.978) and the lowest Cost-of-Pass (\$0.11), though the optimal tier varies by task (Section 5.6).

Avoid delegation for simple tasks. When the task complexity is low (comparable to test-001), delegation-based architectures should be avoided entirely. The catastrophic failure of T4–T6 on the trivial Python task suggests that architectural over-specification is not merely wasteful but actively harmful. A simple diagnostic: if the task can be described in a single sentence, delegation is unlikely to help with a budget model.

Task classification could guide tier selection. The significant interaction effect implies that a static tier recommendation is suboptimal. Organizations could investigate a lightweight task classifier that selects T1–T2 for simple tasks and considers T3 only for tasks where genuine decomposition benefits are expected. However, this recommendation is speculative given the 3-task basis of the current study; a functional classifier would require training data from a larger task set.

Judge panel composition matters. The near-zero inter-rater reliability suggests that organiza-

tions using LLM-as-judge evaluation should carefully consider panel composition. Including budget models as judges introduces noise that may obscure real quality differences. A two-judge panel of frontier models may provide more reliable evaluation despite the higher per-evaluation cost.

Budget models are cost-competitive for evaluation. The total cost of $122.31 for 1,080 runs shows that systematic architectural evaluation is affordable with budget models. Organizations can conduct their own ablation studies during development cycles rather than relying on published benchmarks that may not reflect their specific use cases.

8.3 Comparison to Prior Work

The prior Scylla dryrun [17] evaluated Claude Sonnet 4.5 on a single Hello World task. While direct comparison is limited by differences in model, task set, and experimental scale, several contrasts are noteworthy.

The dryrun found a similar pattern of degradation in higher tiers, suggesting that the non-monotonic tier–quality relationship may be a general phenomenon rather than specific to budget models. However, the severity of the degradation appears to be model-dependent: the frontier model in the dryrun maintained partial performance in delegation tiers, whereas the budget model in this study exhibits complete failure on the trivial task.

This comparison, while informal, suggests that the T2–T3 degradation may be *steeper* for budget models than for frontier models—a hypothesis that could be tested formally in a cross-model study using the same experimental design.

8.4 Limitations

This study has several important limitations that constrain the generalizability of our findings:

Small sample size. With $N = 3$ replications per subtest, our statistical power is limited. Bootstrap confidence intervals are necessarily wide (e.g., the T6 pass rate CI spans $[0.333, 0.889]$), and individual outlier runs can substantially influence subtest-level statistics. The results should be treated as preliminary evidence requiring confirmation with larger sample sizes ($N \geq 10$).

Single model family. All experiments use models from the Anthropic Claude family. The findings may not generalize to other model families (e.g., GPT, Gemini, Llama) with different architectural assumptions, fine-tuning strategies, or capability profiles. In particular, models that have been specifically optimized for multi-agent coordination might show different tier sensitivity profiles.

Limited task diversity. Three experiments provide limited coverage of the task space. All three tasks are relatively simple by software engineering standards—none involves multi-file refactoring, complex debugging, or large-scale system design. The tier effects observed here may not hold for more complex tasks where delegation could provide genuine coordination value that offsets the overhead costs.

Judge reliability concerns. The near-zero inter-rater reliability undermines confidence in the quality scores. While the pass rate metric (which uses majority vote) is less sensitive to score miscalibration, the continuous score and CoP metrics may be substantially affected by judge disagreement. All findings based on continuous scores should be interpreted with this caveat in mind.

Incomplete evaluation coverage. Of the 3,240 expected judge evaluations (1,080 runs $\times$ 3 judges), only 2,696 (83.2%) are present. The shortfall is concentrated in test-001, where T4–T6 have zero judge evaluations. In total, 170 of 1,080 runs lack any judge evaluation whatsoever; scores for these runs were assigned by the pipeline as 0 when no evaluable output was produced, rather than by explicit judge assessment. This means the inter-rater reliability statistics ($\alpha = -0.030$) are computed on the available subset and may not generalize to the full design matrix. Specifically, the complete absence of T4–T6 evaluations for test-001 introduces a systematic selection bias: the available data overrepresents tiers where the agent produced evaluable output (T0–T3) and underrepresents complete-failure runs. If T4–T6 failures would have produced uniformly low scores across all three judges, their inclusion would likely reduce disagreement (and thus increase α). The reported $\alpha = -0.030$ should therefore be interpreted as a lower bound on judge reliability for the evaluable subset only.

Confounded tier composition. The tiers are not pure factorial manipulations—each tier includes the capabilities of previous tiers. T3 includes T0–T2 capabilities plus delegation. This hierarchical design means we cannot fully isolate the effect of delegation from potential interaction effects with lower-tier components. A fully factorial design would require testing delegation without prompts, skills, or tools, which is not currently supported by the framework.

Temporal effects. All experiments were conducted over a limited time window during January–March 2026. Model behavior may vary due to API updates, load balancing, or other infrastructure factors not captured in our analysis.

8.5 Threats to Validity

Internal validity. The primary threat to internal validity is the non-independent judge panel. Because the judges are deterministic functions of the input (modulo temperature), the three scores for each run are not truly independent observations. This may inflate or deflate agreement statistics in ways that are difficult to characterize. Additionally, the use of the same model (Haiku) as both subject and judge introduces a potential confound: self-preference bias could inflate Haiku's scores for its own outputs, while capacity limitations could deflate them.

External validity. The use of a single model family, a single evaluation framework, and a small set of tasks limits the external validity of our conclusions. The findings are specific to the Scylla framework's operationalization of architectural tiers and may not generalize to other agent architectures or evaluation protocols. The specific tier definitions (e.g., which skills are included in T1, which tools in T2) may also influence results in ways that are difficult to disentangle from the generic tier effects.

Construct validity. The mapping from Scylla's tier definitions to meaningful architectural categories is a construct that may not perfectly capture real-world agent configurations. For example, T3's 41 delegation subtests may overweight certain specialist configurations relative to their practical importance, and T6's single subtest provides very limited statistical basis for inference about "everything enabled" configurations.

Statistical conclusion validity. Given the small N, there is a risk of both Type I errors (finding spurious effects) and Type II errors (missing real effects). The use of non-parametric tests and conservative correction methods (Holm–Bonferroni) mitigates the former, but the latter remains

a concern, particularly for the cost and implementation rate analyses where real but small effects may be undetectable with $N = 3$.

9 Conclusions

This study presents, to our knowledge, the first multi-task ablation of agentic coding architectures using a budget frontier model (`claude-haiku-4-5`). Across 1,080 runs spanning three experiments, seven tiers, and 120 subtests at a total cost of $122.31, we find preliminary evidence for a non-monotonic relationship between architectural complexity and task performance.

The central finding is that *simpler is better* for the budget model studied. In the aggregate data, tiers T0–T2 outperform tiers T3–T6 in both pass rate and score, with the T2–T3 transition representing the sharpest degradation point, where multi-agent delegation is introduced. However, task-dependent reversals occur (e.g., T6 achieves the highest pass rate on test-002).

Crucially, this degradation is task-dependent: the significant SRH interaction effect indicates there is no universally optimal tier. On the trivial Python task, hierarchical architectures fail catastrophically (zero pass rate for T4–T6); on moderate tasks, all tiers achieve acceptable performance with muted differences. This interaction should temper aggregate-level recommendations—the "best" tier depends on the task.

Observed cost invariance across tiers ($p = 0.676$, non-significant under limited power) suggests no cost–quality trade-off, making higher tiers effectively Pareto-dominated in aggregate. T2 (Tooling) appears as the highest-pass-rate Pareto-optimal tier in the aggregate data (T0 and T1 also lie on the Pareto frontier at similar costs), though per-task analysis shows the optimal tier varies (e.g., T5 dominates T2 for test-002).

The near-zero inter-rater reliability ($\alpha = -0.030$) of the three-judge panel is a methodological finding with implications beyond this study. It suggests that LLM-as-judge evaluation requires careful panel composition, and that including budget models as judges may undermine measurement reliability.

These conclusions are necessarily preliminary given the small sample size ($N = 3$ per subtest), limited task diversity, and single-model-family scope. We present them as directional evidence that should motivate larger-scale studies with broader task coverage and cross-model comparisons.

10 Further Work

Several directions for future research emerge from this study:

Increased statistical power. Repeating the experiments with $N \geq 10$ replications per subtest would substantially narrow confidence intervals and improve the reliability of pairwise comparisons. Adaptive sampling strategies could focus additional replications on tiers and subtests where variance is highest, optimizing the information gained per dollar spent.

Cross-model comparison. Applying the same experimental design to frontier models (e.g., `claude-opus-4-6`, `claude-sonnet-4-6`) would reveal whether the non-monotonic tier–quality re-

lationship is specific to budget models or generalizes across the capability spectrum. The prior Scylla dryrun [17] used Sonnet 4.5 on a single task, providing a starting point for this comparison. Of particular interest is whether frontier models exhibit a T2–T3 performance drop at all, or whether their greater reasoning capacity enables them to manage delegation overhead without quality loss.

Task difficulty scaling. Extending the experiment set to include harder tasks (multi-file refactoring, complex debugging, system design) would test the hypothesis that delegation-based architectures become beneficial only when task complexity exceeds the individual model's capacity. A crossover point—the task complexity at which delegation begins to help rather than hurt—would provide practical guidance for architectural selection.

Judge panel calibration. The poor inter-rater reliability motivates investigation of alternative judge panel compositions. Options include: (a) removing the budget judge to test whether a two-model frontier panel achieves acceptable reliability; (b) calibrating all judges through shared rubric training or few-shot examples; (c) replacing LLM judges with deterministic evaluation criteria such as test suite pass rates; or (d) investigating hybrid approaches that combine automated evaluation with targeted human review.

Adaptive tier selection. The significant tier $\times$ task interaction suggests that a meta-agent capable of selecting the optimal tier based on task characteristics could outperform any fixed tier strategy. This meta-agent could be trained on the data from this and future studies, learning to predict which architectural configuration will maximize pass rate for a given task description.

Cost optimization. While no significant cost difference across tiers was detected in our data, the total cost of \$122.31 for 1,080 runs suggests room for efficiency gains through prompt caching optimization, reduced context windows for simple tasks, or selective evaluation (running fewer subtests for well-characterized tiers).

Longitudinal stability. Repeating the experiments periodically would reveal whether the tier–quality relationships are stable over time or shift with model updates. This is particularly relevant for budget models, which may receive more frequent capability improvements than frontier models.

Non-Anthropic models. Extending the study to non-Anthropic model families (OpenAI, Google, Meta) would test whether the findings are specific to Claude's architecture or reflect general properties of budget-class LLMs in agentic settings.

Delegation protocol optimization. The catastrophic failure of delegation-based tiers on trivial tasks suggests that the delegation protocol itself may be suboptimal for budget models. Research into lightweight delegation protocols—simpler handoff mechanisms, reduced coordination overhead, or asynchronous delegation—could potentially recover some of the lost performance while preserving the benefits of task decomposition.

Fine-grained tier decomposition. The current T3 tier groups 41 subtests that represent different delegation configurations. A more fine-grained analysis could identify which specific delegation strategies (e.g., planning-focused, review-focused, or implementation-focused specialists) are most harmful or beneficial, enabling surgical rather than wholesale architectural recommendations.

Acknowledgments

This research was conducted using the Scylla evaluation framework as part of the Homeric Intelligence ecosystem. The author thanks the Claude Code team at Anthropic for providing the agentic CLI infrastructure used in all experiments. All evaluation costs ($122.31) were borne by the author. No external funding was received for this work.

References

[1] Anthropic. Claude Code: Agentic CLI tool for software development. https://www.anthropic.com/claude/code, 2024. AI-powered command-line interface for coding tasks.

[2] Norman Cliff. Dominance statistics: Ordinal analyses to answer ordinal questions. *Psychological Bulletin*, 114(3):494–509, 1993.

[3] Charles J. Clopper and Egon S. Pearson. The use of confidence or fiducial limits illustrated in the case of the binomial. *Biometrika*, 26(4):404–413, 1934.

[4] Bradley Efron. Better bootstrap confidence intervals. *Journal of the American Statistical Association*, 82(397):171–185, 1987.

[5] Leo Gao et al. lm-evaluation-harness: A framework for few-shot language model evaluation. https://github.com/EleutherAI/lm-evaluation-harness, 2024. General-purpose LLM evaluation framework.

[6] Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979.

[7] Homeric Intelligence. ProjectOdyssey: Comprehensive agent orchestration framework. https://github.com/HomericIntelligence/Projectodyssey, 2025. Accessed: 2025-12-30. Git hash: 011a3ff.

[8] Carlos E. Jimenez et al. SWE-bench: Can language models resolve real-world github issues? *arXiv preprint arXiv:2310.06770*, 2024. Benchmark for evaluating LLMs on real GitHub issues.

[9] Klaus Krippendorff. *Content Analysis: An Introduction to Its Methodology*. Sage Publications, Thousand Oaks, CA, 2nd edition, 2004.

[10] William H. Kruskal and W. Allen Wallis. Use of ranks in one-criterion variance analysis. *Journal of the American Statistical Association*, 47(260):583–621, 1952.

[11] Xiao Liu et al. AgentBench: Evaluating LLMs as agents. *arXiv preprint arXiv:2308.03688*, 2024. Multi-turn agent evaluation across operating systems, databases, and knowledge graphs.

[12] Henry B. Mann and Donald R. Whitney. On a test of whether one of two random variables is stochastically larger than the other. *The Annals of Mathematical Statistics*, 18(1):50–60, 1947.

[13] Felipe Maia Polo et al. Efficient multi-prompt evaluation of llms. *arXiv preprint arXiv:2405.17202*, 2024. Framework for efficient LLM evaluation.

[14] Jeanine Romano, Jeffrey D. Kromrey, Jesse Coraggio, and Jeff Skowronek. Appropriate statistics for ordinal level data: Should we really be using t-test and Cohen's d for evaluating group differences on the NSSE and other surveys? In *Annual Meeting of the Florida Association of Institutional Research*, 2006. FAIR conference thresholds for Cliff's δ: negligible < 0.11, small < 0.28, medium < 0.43, large ≥ 0.43.

[15] C. James Scheirer, William S. Ray, and Nathan Hare. The analysis of ranked data derived from completely randomized factorial designs. *Biometrics*, 32(2):429–434, 1976.

[16] András Vargha and Harold D. Delaney. A critique and improvement of the CL common language effect size statistics of McGraw and Wong. *Journal of Educational and Behavioral Statistics*, 25(2):101–132, 2000.

[17] Micah Villmow. Taming Scylla: Understanding the multi-headed agentic daemon of the coding seas. ProjectScylla Technical Report, 2026. Dryrun evaluation of the Scylla framework using Claude Sonnet 4.5 on a Hello World task. Available at: https://github.com/HomericIntelligence/ProjectScylla/tree/main/docs/arxiv/dryrun.

[18] Shunyu Yao et al. TAU-bench: A benchmark for tool-augmented LLMs. *arXiv preprint arXiv:2406.12045*, 2024. Benchmark for evaluating agents' ability to use external tools.

[19] Kaijie Zhu et al. PromptBench: Towards evaluating the robustness of large language models on adversarial prompts. *arXiv preprint arXiv:2306.04528*, 2024. Benchmark for prompt evaluation and robustness.

A Experiment Configuration

This appendix provides the full experiment configuration used for all three experiments. Each experiment uses identical tier and subtest definitions; the experiments differ only in the target task, repository, and difficulty level.

A.1 Tier Definitions

Table 10 provides the mapping from tier identifiers to architectural configurations, including the number of subtests in each tier and a brief description of what each tier evaluates.

A.2 Model Specifications

A.3 Scoring Protocol

Each judge assigns scores on a continuous $[0, 1]$ scale against predefined criteria that are task-specific. The criteria evaluate dimensions such as correctness, completeness, code quality, documentation, and adherence to specifications. The consensus score is computed as the arithmetic mean of the three judges' scores. A run passes if a majority of judges independently classify it as passing.

Table 10: *Tier definitions and subtest counts.*

Tier	Name	Subtests	Description
T0	Prompts	24	System prompt ablation (empty → full CLAUDE.md)
T1	Skills	10	Domain expertise via installed skills
T2	Tooling	15	External tools and MCP servers
T3	Delegation	41	Flat multi-agent with specialist agents
T4	Hierarchy	14	Nested orchestration (7 hierarchy + 7 teams)
T5	Hybrid	15	Best combinations and permutations
T6	Super	1	Everything enabled at maximum capability
Total		120	

Table 11: *Model specifications for subject agent and judge panel.*

Role	Model ID	Class
Subject Agent	`claude-haiku-4-5`	Budget frontier
Judge 1	`claude-opus-4-6`	Frontier
Judge 2	`claude-sonnet-4-6`	Mid-tier
Judge 3	`claude-haiku-4-5`	Budget frontier

Implementation rate is derived from the median of the three judges' binary assessments of whether the agent attempted the task. This median-based aggregation is more robust to a single outlier judge than the mean.

Letter grades are assigned based on the consensus score using the following scale:

- S: $= 1.00$ (exceptional)
- A: ≥ 0.80 (production ready)
- B: ≥ 0.60 (satisfactory)
- C: ≥ 0.40 (marginal)
- D: ≥ 0.20 (below threshold)
- F: < 0.20 (failure)

A.4 Execution Environment

All experiments were executed using the Claude Code CLI [1] in headless mode on Linux systems (kernel 5.10+). The execution environment was containerized using Docker to ensure reproducibility. Each run was allocated a timeout to prevent runaway sessions, and the Claude Code CLI managed all model interactions, tool invocations, and delegation protocols.

B Full Subtest Results

Table 12 presents the complete subtest-level results across all experiments and tiers. Each row represents a single subtest within a specific tier and experiment, showing the pass rate, mean score, standard deviation, cost, letter grade, and grade distribution across the $N = 3$ runs.

The table reveals the granular patterns that drive the aggregate findings discussed in the main text. In particular, the clustering of F grades in T3–T6 for test-001 is visible at the subtest level, confirming that the failures are systematic rather than driven by a few outlier subtests.

Table 12: *Full Subtest Results (Appendix B). Grade is the modal grade across $N = 3$ runs, with highest-grade tie-breaking when no single grade has a majority.*

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T0	00	0.000	0.045	0.030	0.0201	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T0	00	0.333	0.333	0.577	0.0272	F	S:1 A:0 B:0 C:0 D:0 F:2
claude-haiku-4-5	T0	00	0.000	0.000	0.000	0.0207	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T0	01	1.000	0.904	0.012	0.0309	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	01	0.667	0.656	0.174	0.0544	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T0	01	1.000	0.672	0.080	0.1969	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T0	02	1.000	0.896	0.023	0.0326	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	02	1.000	0.715	0.122	0.0713	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T0	02	1.000	0.671	0.131	0.2268	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T0	03	1.000	0.826	0.194	0.1032	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T0	03	1.000	0.700	0.006	0.2903	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	03	1.000	0.900	0.049	0.0604	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	04	1.000	0.613	0.065	0.2528	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T0	04	1.000	0.909	0.010	0.0373	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	04	1.000	0.833	0.220	0.0555	A	S:1 A:1 B:0 C:1 D:0 F:0
claude-haiku-4-5	T0	05	1.000	0.621	0.075	0.1856	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T0	05	1.000	0.908	0.032	0.0383	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	05	1.000	0.958	0.000	0.0694	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	06	0.667	0.592	0.059	0.2637	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T0	06	1.000	0.934	0.042	0.0771	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	06	1.000	0.925	0.013	0.0425	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	07	1.000	0.931	0.087	0.0718	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	07	1.000	0.914	0.020	0.0375	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	07	1.000	0.662	0.013	0.1979	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	08	1.000	0.722	0.173	0.0726	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T0	08	1.000	0.910	0.010	0.0322	A	S:0 A:3 B:0 C:0 D:0 F:0

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T0	08	1.000	0.651	0.071	0.2087	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T0	09	1.000	0.931	0.048	0.0588	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	09	1.000	0.910	0.007	0.0320	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	09	1.000	0.712	0.059	0.2238	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	10	1.000	0.705	0.059	0.2211	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	10	1.000	0.958	0.000	0.0711	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	10	1.000	0.907	0.005	0.0432	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	11	1.000	0.833	0.217	0.0630	A	S:0 A:2 B:0 C:1 D:0 F:0
claude-haiku-4-5	T0	11	1.000	0.716	0.037	0.2021	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	11	1.000	0.910	0.023	0.0339	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	12	1.000	0.719	0.057	0.2357	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	12	1.000	0.958	0.042	0.0796	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	12	1.000	0.901	0.056	0.0394	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	13	1.000	0.671	0.098	0.0727	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	13	0.667	0.669	0.160	0.2143	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T0	13	1.000	0.918	0.005	0.0416	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	14	1.000	0.676	0.121	0.2382	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T0	14	1.000	0.722	0.168	0.0652	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T0	14	1.000	0.913	0.021	0.0384	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	15	1.000	0.757	0.025	0.2403	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	15	1.000	0.916	0.015	0.0397	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	15	1.000	0.833	0.182	0.0885	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T0	16	1.000	0.925	0.027	0.0395	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	16	1.000	0.734	0.067	0.2327	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	16	1.000	0.806	0.197	0.0677	A	S:0 A:2 B:0 C:1 D:0 F:0
claude-haiku-4-5	T0	17	1.000	0.909	0.013	0.0427	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	17	0.667	0.812	0.237	0.0709	A	S:0 A:2 B:0 C:1 D:0 F:0
claude-haiku-4-5	T0	17	1.000	0.690	0.007	0.2570	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	18	1.000	0.917	0.042	0.0711	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	18	1.000	0.694	0.040	0.2085	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	18	1.000	0.878	0.068	0.0362	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	19	1.000	0.921	0.016	0.0439	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	19	1.000	0.778	0.105	0.0765	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T0	19	1.000	0.665	0.025	0.2456	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	20	1.000	0.958	0.000	0.0720	A	S:0 A:3 B:0 C:0 D:0 F:0

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T0	20	1.000	0.672	0.118	0.2132	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T0	20	1.000	0.918	0.013	0.0406	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	21	1.000	0.920	0.005	0.0325	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	21	1.000	0.945	0.064	0.0622	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	21	0.667	0.545	0.079	0.1921	C	S:0 A:0 B:1 C:2 D:0 F:0
claude-haiku-4-5	T0	22	1.000	0.847	0.197	0.0865	A	S:1 A:1 B:1 C:0 D:0 F:0
claude-haiku-4-5	T0	22	1.000	0.915	0.021	0.0428	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T0	22	1.000	0.696	0.078	0.2314	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T0	23	0.667	0.543	0.220	0.2488	B	S:0 A:0 B:2 C:0 D:1 F:0
claude-haiku-4-5	T0	23	1.000	0.812	0.075	0.0428	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T0	23	1.000	0.722	0.173	0.1055	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T1	01	1.000	0.916	0.011	0.0346	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	01	1.000	0.840	0.209	0.0737	A	S:1 A:1 B:1 C:0 D:0 F:0
claude-haiku-4-5	T1	01	1.000	0.711	0.010	0.2110	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T1	02	1.000	0.912	0.016	0.0336	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	02	1.000	0.861	0.206	0.0786	A	S:1 A:1 B:1 C:0 D:0 F:0
claude-haiku-4-5	T1	02	1.000	0.662	0.105	0.2646	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T1	03	1.000	0.920	0.018	0.0427	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	03	1.000	0.958	0.000	0.0653	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	03	0.333	0.578	0.123	0.2086	C	S:0 A:0 B:1 C:2 D:0 F:0
claude-haiku-4-5	T1	04	1.000	0.885	0.020	0.0320	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	04	1.000	0.812	0.188	0.0636	A	S:1 A:1 B:1 C:0 D:0 F:0
claude-haiku-4-5	T1	04	1.000	0.711	0.046	0.2546	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T1	05	1.000	0.847	0.211	0.0830	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T1	05	1.000	0.704	0.019	0.2329	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T1	05	1.000	0.922	0.007	0.0408	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	06	1.000	0.645	0.098	0.2248	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T1	06	1.000	0.917	0.008	0.0334	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	06	1.000	0.840	0.204	0.0575	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T1	07	1.000	0.751	0.048	0.2167	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T1	07	1.000	0.896	0.017	0.0405	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	07	1.000	0.903	0.105	0.0693	A	S:1 A:1 B:1 C:0 D:0 F:0
claude-haiku-4-5	T1	08	1.000	0.627	0.064	0.2314	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T1	08	1.000	0.911	0.005	0.0425	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	08	0.667	0.878	0.177	0.0658	A	S:1 A:1 B:1 C:0 D:0 F:0

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T1	09	1.000	0.875	0.083	0.0716	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T1	09	1.000	0.632	0.088	0.2470	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T1	09	1.000	0.896	0.013	0.0346	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	10	1.000	0.739	0.014	0.1949	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T1	10	1.000	0.906	0.017	0.0380	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T1	10	1.000	0.833	0.182	0.1092	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T2	01	1.000	0.932	0.022	0.0353	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	01	1.000	0.673	0.109	0.1926	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T2	01	1.000	0.833	0.182	0.0559	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T2	02	1.000	0.892	0.017	0.0382	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	02	1.000	0.736	0.070	0.2305	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T2	02	1.000	0.847	0.192	0.0644	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T2	03	1.000	0.921	0.002	0.0404	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	03	1.000	0.717	0.041	0.2161	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T2	03	1.000	0.694	0.229	0.0567	C	S:0 A:1 B:0 C:2 D:0 F:0
claude-haiku-4-5	T2	04	1.000	0.926	0.008	0.0433	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	04	0.667	0.636	0.092	0.2434	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T2	04	1.000	0.826	0.211	0.0838	A	S:0 A:2 B:0 C:1 D:0 F:0
claude-haiku-4-5	T2	05	1.000	0.730	0.032	0.2608	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T2	05	0.667	0.750	0.253	0.0599	A	S:0 A:2 B:0 C:1 D:0 F:0
claude-haiku-4-5	T2	05	1.000	0.891	0.017	0.0376	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	06	1.000	0.819	0.173	0.0886	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T2	06	1.000	0.896	0.018	0.0356	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	06	1.000	0.610	0.105	0.2065	C	S:0 A:0 B:1 C:2 D:0 F:0
claude-haiku-4-5	T2	07	1.000	0.625	0.042	0.0726	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T2	07	1.000	0.929	0.013	0.0316	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	07	1.000	0.682	0.105	0.2488	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T2	08	1.000	0.750	0.150	0.0931	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T2	08	0.667	0.561	0.278	0.2325	B	S:0 A:0 B:2 C:0 D:1 F:0
claude-haiku-4-5	T2	08	1.000	0.892	0.033	0.0385	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	09	1.000	0.740	0.011	0.2414	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T2	09	1.000	0.847	0.192	0.0646	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T2	09	1.000	0.914	0.028	0.0366	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	10	1.000	0.716	0.024	0.1798	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T2	10	1.000	0.806	0.168	0.0744	A	S:0 A:2 B:1 C:0 D:0 F:0

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T2	10	1.000	0.906	0.036	0.0389	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	11	1.000	0.660	0.097	0.1765	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T2	11	1.000	0.979	0.021	0.0671	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	11	1.000	0.900	0.042	0.0320	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	12	1.000	0.819	0.173	0.0670	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T2	12	1.000	0.932	0.002	0.0361	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	12	1.000	0.681	0.093	0.2416	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T2	13	1.000	0.890	0.003	0.0351	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	13	1.000	0.692	0.037	0.2077	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T2	13	1.000	0.847	0.134	0.0899	B	S:1 A:0 B:2 C:0 D:0 F:0
claude-haiku-4-5	T2	14	1.000	0.903	0.010	0.0356	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	14	1.000	0.707	0.043	0.1981	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T2	14	1.000	0.833	0.253	0.0727	A	S:1 A:1 B:0 C:1 D:0 F:0
claude-haiku-4-5	T2	15	1.000	0.894	0.021	0.0342	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T2	15	1.000	0.730	0.092	0.2496	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T2	15	1.000	0.931	0.048	0.0772	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	01	1.000	0.917	0.009	0.0473	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	01	1.000	0.847	0.229	0.0697	A	S:1 A:1 B:0 C:1 D:0 F:0
claude-haiku-4-5	T3	01	1.000	0.767	0.090	0.2648	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T3	02	1.000	0.905	0.022	0.0323	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	02	1.000	0.819	0.188	0.0722	A	S:1 A:1 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	02	1.000	0.726	0.012	0.2026	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	03	1.000	0.916	0.018	0.0375	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	03	1.000	0.896	0.091	0.0773	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	03	1.000	0.725	0.091	0.2254	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	04	1.000	0.913	0.010	0.0377	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	04	0.667	0.681	0.339	0.0576	A	S:0 A:2 B:0 C:0 D:1 F:0
claude-haiku-4-5	T3	04	1.000	0.673	0.023	0.2862	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	05	1.000	0.917	0.013	0.0470	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	05	1.000	0.833	0.150	0.0600	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	05	1.000	0.697	0.090	0.2285	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	06	1.000	0.945	0.025	0.0693	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	06	1.000	0.688	0.034	0.2316	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	06	1.000	0.902	0.020	0.0326	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	07	1.000	0.709	0.013	0.1952	B	S:0 A:0 B:3 C:0 D:0 F:0

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T3	07	1.000	0.915	0.004	0.0364	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	07	1.000	0.958	0.000	0.0722	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	08	1.000	0.624	0.070	0.2632	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	08	1.000	0.899	0.010	0.0390	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	08	1.000	0.708	0.182	0.0612	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T3	09	1.000	0.688	0.069	0.2111	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	09	1.000	0.932	0.008	0.0323	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	09	1.000	0.917	0.042	0.0804	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	10	1.000	0.742	0.041	0.1930	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	10	1.000	0.920	0.019	0.0398	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	10	1.000	0.799	0.154	0.0701	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	11	1.000	0.649	0.102	0.2311	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	11	1.000	0.861	0.168	0.0610	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	11	1.000	0.900	0.034	0.0406	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	12	1.000	0.806	0.158	0.0659	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	12	1.000	0.710	0.016	0.2298	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	12	1.000	0.893	0.007	0.0389	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	13	0.667	0.611	0.064	0.0600	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	13	1.000	0.704	0.067	0.2727	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	13	1.000	0.899	0.036	0.0375	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	14	1.000	0.750	0.182	0.0597	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T3	14	1.000	0.741	0.039	0.2236	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	14	1.000	0.905	0.017	0.0364	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	15	1.000	0.819	0.173	0.0642	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	15	1.000	0.760	0.051	0.1917	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T3	15	0.333	0.283	0.491	0.0375	F	S:0 A:1 B:0 C:0 D:0 F:2
claude-haiku-4-5	T3	16	1.000	0.702	0.031	0.1919	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	16	0.000	0.000	0.000	0.0398	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	16	1.000	0.861	0.206	0.0600	A	S:1 A:1 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	17	0.000	0.000	0.000	0.0424	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	17	1.000	0.875	0.042	0.1188	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	17	1.000	0.717	0.057	0.2357	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	18	0.000	0.000	0.000	0.0376	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	18	1.000	0.972	0.024	0.0885	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	18	1.000	0.685	0.058	0.2050	B	S:0 A:0 B:3 C:0 D:0 F:0

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T3	19	0.000	0.000	0.000	0.0410	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	19	1.000	0.972	0.048	0.0625	S	S:2 A:1 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	19	1.000	0.753	0.015	0.2098	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	20	0.000	0.000	0.000	0.0440	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	20	0.667	0.768	0.241	0.0735	A	S:0 A:2 B:0 C:1 D:0 F:0
claude-haiku-4-5	T3	20	1.000	0.730	0.031	0.2390	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	21	1.000	0.944	0.048	0.0844	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	21	1.000	0.708	0.024	0.2164	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	21	0.000	0.000	0.000	0.0389	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	22	1.000	0.668	0.117	0.1956	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	22	0.000	0.000	0.000	0.0342	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	22	1.000	0.750	0.182	0.0801	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T3	23	0.667	0.660	0.101	0.2456	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	23	0.000	0.000	0.000	0.0398	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	23	0.667	0.806	0.265	0.0787	A	S:0 A:2 B:0 C:1 D:0 F:0
claude-haiku-4-5	T3	24	1.000	0.739	0.074	0.2336	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	24	0.000	0.000	0.000	0.0337	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	24	1.000	0.931	0.024	0.0762	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	25	1.000	0.638	0.069	0.2303	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	25	0.000	0.000	0.000	0.0374	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	25	1.000	0.885	0.063	0.0618	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	26	0.000	0.000	0.000	0.0340	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	26	1.000	0.931	0.064	0.0667	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	26	1.000	0.623	0.054	0.2460	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	27	1.000	0.722	0.206	0.0653	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T3	27	1.000	0.706	0.081	0.1852	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	27	0.000	0.000	0.000	0.0349	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	28	1.000	0.847	0.192	0.0730	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	28	1.000	0.717	0.074	0.2309	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	28	0.000	0.000	0.000	0.0368	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	29	1.000	0.958	0.000	0.0787	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	29	1.000	0.740	0.033	0.2041	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	29	0.000	0.000	0.000	0.0495	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	30	1.000	0.931	0.048	0.0770	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	30	1.000	0.618	0.088	0.2453	B	S:0 A:0 B:2 C:1 D:0 F:0

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T3	30	0.000	0.000	0.000	0.0334	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	31	1.000	0.922	0.101	0.0812	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	31	0.000	0.000	0.000	0.0379	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	31	1.000	0.689	0.125	0.2002	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	32	1.000	0.917	0.042	0.0748	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T3	32	1.000	0.682	0.039	0.2468	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	32	0.000	0.000	0.000	0.0397	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	33	0.000	0.000	0.000	0.0341	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	33	1.000	0.736	0.197	0.0697	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T3	33	0.667	0.499	0.091	0.2614	C	S:0 A:0 B:0 C:2 D:1 F:0
claude-haiku-4-5	T3	34	0.000	0.000	0.000	0.0414	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	34	1.000	0.582	0.067	0.2395	C	S:0 A:0 B:1 C:2 D:0 F:0
claude-haiku-4-5	T3	34	1.000	0.722	0.206	0.0535	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T3	35	1.000	0.750	0.217	0.0800	B	S:1 A:0 B:2 C:0 D:0 F:0
claude-haiku-4-5	T3	35	0.000	0.000	0.000	0.0413	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	35	1.000	0.675	0.107	0.2248	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	36	0.667	0.659	0.018	0.2544	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	36	1.000	0.861	0.168	0.0773	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	36	0.000	0.000	0.000	0.0332	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	37	0.667	0.547	0.253	0.2293	B	S:0 A:0 B:2 C:0 D:1 F:0
claude-haiku-4-5	T3	37	1.000	0.722	0.206	0.0619	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T3	37	0.000	0.000	0.000	0.0390	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	38	1.000	0.840	0.187	0.0809	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T3	38	0.000	0.000	0.000	0.0357	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	38	1.000	0.686	0.039	0.2101	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	39	1.000	0.722	0.168	0.0677	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T3	39	0.000	0.000	0.000	0.0364	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	39	1.000	0.728	0.057	0.2690	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	40	1.000	0.736	0.197	0.0658	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T3	40	0.000	0.000	0.000	0.0385	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	40	1.000	0.729	0.062	0.2113	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T3	41	1.000	0.660	0.088	0.2807	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T3	41	0.000	0.000	0.000	0.0305	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T3	41	0.667	0.542	0.110	0.0636	C	S:0 A:0 B:1 C:2 D:0 F:0
claude-haiku-4-5	T4	01	0.000	0.000	0.000	0.0434	F	S:0 A:0 B:0 C:0 D:0 F:3

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T4	01	1.000	0.717	0.079	0.2210	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T4	01	1.000	0.972	0.024	0.0900	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	02	0.000	0.000	0.000	0.0409	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T4	02	1.000	0.714	0.017	0.2262	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T4	02	1.000	0.958	0.042	0.0763	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	03	0.000	0.000	0.000	0.0386	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T4	03	1.000	0.708	0.085	0.2065	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T4	03	1.000	0.847	0.229	0.0722	A	S:1 A:1 B:0 C:1 D:0 F:0
claude-haiku-4-5	T4	04	1.000	0.986	0.024	0.0593	S	S:2 A:1 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	04	0.000	0.000	0.000	0.0333	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T4	04	1.000	0.682	0.066	0.1987	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T4	05	0.667	0.611	0.064	0.0743	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T4	05	0.000	0.000	0.000	0.0354	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T4	05	1.000	0.681	0.052	0.2228	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T4	06	1.000	0.806	0.210	0.0734	A	S:1 A:1 B:0 C:1 D:0 F:0
claude-haiku-4-5	T4	06	1.000	0.653	0.117	0.2150	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T4	06	0.000	0.000	0.000	0.0341	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T4	07	1.000	0.847	0.158	0.0862	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T4	07	1.000	0.737	0.022	0.1944	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T4	07	0.000	0.000	0.000	0.0354	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T4	08	0.667	0.661	0.116	0.2646	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T4	08	0.000	0.000	0.000	0.0439	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T4	08	0.667	0.806	0.265	0.0848	A	S:0 A:2 B:0 C:1 D:0 F:0
claude-haiku-4-5	T4	09	1.000	0.674	0.024	0.2467	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T4	09	1.000	0.842	0.188	0.0597	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T4	09	0.000	0.000	0.000	0.0490	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T4	10	1.000	0.741	0.012	0.2653	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T4	10	1.000	0.951	0.043	0.0919	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T4	10	0.000	0.000	0.000	0.0367	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T4	11	0.000	0.000	0.000	0.0335	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T4	11	1.000	0.726	0.035	0.2592	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T4	11	1.000	0.861	0.168	0.0656	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T4	12	1.000	0.561	0.138	0.2602	C	S:0 A:0 B:1 C:2 D:0 F:0
claude-haiku-4-5	T4	12	1.000	0.833	0.182	0.0702	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T4	12	0.000	0.000	0.000	0.0392	F	S:0 A:0 B:0 C:0 D:0 F:3

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T4	13	1.000	0.833	0.220	0.0878	A	S:1 A:1 B:0 C:1 D:0 F:0
claude-haiku-4-5	T4	13	0.000	0.000	0.000	0.0363	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T4	13	1.000	0.578	0.073	0.2675	C	S:0 A:0 B:1 C:2 D:0 F:0
claude-haiku-4-5	T4	14	1.000	0.642	0.011	0.2540	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T4	14	0.000	0.000	0.000	0.0387	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T4	14	1.000	0.847	0.192	0.0694	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T5	01	0.000	0.000	0.000	0.0325	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T5	01	1.000	0.658	0.134	0.2399	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T5	01	1.000	0.907	0.049	0.1447	A	S:0 A:3 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	02	0.000	0.000	0.000	0.0345	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T5	02	1.000	0.650	0.084	0.2110	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T5	02	1.000	0.903	0.096	0.1219	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T5	03	0.000	0.000	0.000	0.0345	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T5	03	1.000	0.701	0.061	0.1919	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T5	03	0.667	0.694	0.197	0.1138	A	S:0 A:1 B:1 C:1 D:0 F:0
claude-haiku-4-5	T5	04	0.000	0.000	0.000	0.0371	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T5	04	1.000	0.685	0.106	0.1998	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T5	04	0.667	0.756	0.178	0.1061	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T5	05	1.000	0.748	0.021	0.1756	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T5	05	1.000	0.972	0.024	0.1181	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	05	0.000	0.000	0.000	0.0313	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T5	06	1.000	0.972	0.048	0.1263	S	S:2 A:1 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	06	0.000	0.000	0.000	0.0685	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T5	06	1.000	0.719	0.011	0.2427	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T5	07	1.000	0.688	0.127	0.0475	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T5	07	0.000	0.000	0.000	0.0341	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T5	07	0.667	0.573	0.272	0.2065	B	S:0 A:0 B:2 C:0 D:1 F:0
claude-haiku-4-5	T5	08	1.000	0.958	0.042	0.0680	A	S:1 A:2 B:0 C:0 D:0 F:0
claude-haiku-4-5	T5	08	1.000	0.726	0.025	0.2933	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T5	08	0.000	0.000	0.000	0.0443	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T5	09	1.000	0.702	0.013	0.2005	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T5	09	1.000	0.806	0.127	0.0626	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T5	09	0.000	0.000	0.000	0.0400	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T5	10	1.000	0.699	0.071	0.2473	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T5	10	1.000	0.972	0.024	0.0878	A	S:1 A:2 B:0 C:0 D:0 F:0

Table 7 (continued)

Model	Tier	ST	Pass Rate	Mean Score	Std Dev	Cost ($)	Grade	Distribution
claude-haiku-4-5	T5	10	0.000	0.000	0.000	0.0378	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T5	11	0.667	0.621	0.074	0.2053	B	S:0 A:0 B:2 C:1 D:0 F:0
claude-haiku-4-5	T5	11	1.000	0.833	0.182	0.1232	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T5	11	0.000	0.000	0.000	0.0324	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T5	12	0.667	0.639	0.313	0.1463	A	S:0 A:1 B:1 C:0 D:1 F:0
claude-haiku-4-5	T5	12	0.000	0.000	0.000	0.0372	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T5	12	1.000	0.744	0.037	0.2098	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T5	13	0.000	0.000	0.000	0.0342	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T5	13	1.000	0.690	0.064	0.2665	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T5	13	1.000	0.806	0.188	0.0538	B	S:1 A:0 B:2 C:0 D:0 F:0
claude-haiku-4-5	T5	14	0.000	0.000	0.000	0.0354	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T5	14	1.000	0.710	0.021	0.2060	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T5	14	1.000	0.826	0.194	0.0662	A	S:0 A:2 B:1 C:0 D:0 F:0
claude-haiku-4-5	T5	15	0.000	0.000	0.000	0.0318	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T5	15	1.000	0.781	0.040	0.1650	B	S:0 A:1 B:2 C:0 D:0 F:0
claude-haiku-4-5	T5	15	1.000	0.708	0.253	0.0905	C	S:1 A:0 B:0 C:2 D:0 F:0
claude-haiku-4-5	T6	01	0.000	0.000	0.000	0.0704	F	S:0 A:0 B:0 C:0 D:0 F:3
claude-haiku-4-5	T6	01	1.000	0.722	0.064	0.2792	B	S:0 A:0 B:3 C:0 D:0 F:0
claude-haiku-4-5	T6	01	1.000	0.833	0.182	0.1001	A	S:0 A:2 B:1 C:0 D:0 F:0

C Normality Tests

All statistical analyses in this paper employ non-parametric methods. This methodological choice is justified by Shapiro–Wilk normality tests, which reject the normality assumption for all 14 tier $\times$ {Score, Cost} combinations at the $\alpha = 0.05$ significance level (Table 13).

The universal rejection of normality ($p < 0.05$ for all combinations, and $p < 0.001$ for all combinations with $n > 9$) reflects the bounded and often bimodal nature of the score distributions. Scores cluster near 0.0 (complete failure) and 0.8–1.0 (successful completion), with relatively few observations in the intermediate range. This distributional shape is expected for pass/fail-oriented coding tasks: the agent either accomplishes the task (yielding a high score) or fails to accomplish it (yielding a score near zero).

Cost distributions are also non-normal, exhibiting right skew due to occasional runs that require substantially more tokens than average. The non-normality of all metrics motivates the use of rank-based tests (Kruskal–Wallis, Mann–Whitney U, Cliff's δ, Scheirer–Ray–Hare) throughout the analysis rather than parametric alternatives (ANOVA, t-tests).

Table 13: *Normality Tests (Shapiro-Wilk)*

Model	Tier	Metric	N	W	p-value	Normal? ($\alpha = 0.05$)
claude-haiku-4-5	T0	Score	216	0.7927	0.0000	×
claude-haiku-4-5	T0	Cost (USD)	216	0.8250	0.0000	×
claude-haiku-4-5	T1	Score	90	0.9009	0.0000	×
claude-haiku-4-5	T1	Cost (USD)	90	0.8231	0.0000	×
claude-haiku-4-5	T2	Score	135	0.9056	0.0000	×
claude-haiku-4-5	T2	Cost (USD)	135	0.8285	0.0000	×
claude-haiku-4-5	T3	Score	369	0.7985	0.0000	×
claude-haiku-4-5	T3	Cost (USD)	369	0.8184	0.0000	×
claude-haiku-4-5	T4	Score	126	0.8132	0.0000	×
claude-haiku-4-5	T4	Cost (USD)	126	0.8251	0.0000	×
claude-haiku-4-5	T5	Score	135	0.8074	0.0000	×
claude-haiku-4-5	T5	Cost (USD)	135	0.8938	0.0000	×
claude-haiku-4-5	T6	Score	9	0.8072	0.0247	×
claude-haiku-4-5	T6	Cost (USD)	9	0.7816	0.0126	×
Summary: 0/14 (0.0%) pass normality test						

D Statistical Test Details

D.1 Kruskal–Wallis Omnibus Tests

The Kruskal–Wallis H test [10] is the non-parametric analog of one-way ANOVA, testing whether the distributions of a metric differ across the $k = 7$ tier groups. Under the null hypothesis, all groups are drawn from the same distribution. The test statistic H is computed from the ranks of all observations pooled across groups, and is compared to a chi-squared distribution with $k - 1 = 6$ degrees of freedom.

Table 14: *Tier main-effect tests across all tiers ($k = 7$). Pass Rate, Impl Rate, and Duration are standalone Kruskal–Wallis tests; Cost and Score are tier main effects from the Scheirer–Ray–Hare analysis (Table 7). Significant results ($p < 0.05$) are marked with asterisks.*

Metric	Source	$H(6)$	p-value	Sig?
Pass Rate	KW	119.01	< 0.001	***
Impl Rate	KW	5.47	0.486	n.s.
Duration (s)	KW	40.94	< 0.001	***
Cost (USD)	SRH	4.00	0.676	n.s.
Score	SRH	91.3	< 0.001	***

Three of five metrics show significant tier effects: pass rate, duration, and score. The two non-significant metrics—implementation rate and cost—indicate that these aspects of agent behavior

are determined by task characteristics rather than architectural configuration. The agent's decision to attempt the task and the resources it consumes are not influenced by whether it operates in a single-agent or multi-agent configuration.

D.2 Pairwise Comparisons and Effect Sizes

For metrics with significant omnibus tests, pairwise comparisons were conducted using the Mann–Whitney U test [12] with Holm–Bonferroni correction [6]. We test 7 planned comparisons (6 adjacent transitions plus T0–T6) within each metric family separately, requiring correction for multiplicity to control the family-wise error rate within each metric. The correction is not applied across metrics, as each metric represents a pre-specified hypothesis.

The significant pairwise results for pass rate, with Cliff's δ effect sizes and 95% bootstrap confidence intervals, are:

- T2 vs. T3: $p < 0.001$, Cliff's $\delta = -0.219$ $[-0.268, -0.168]$. The negative δ indicates that T3 pass rates are systematically lower than T2 pass rates. The effect size falls in the "small" range by conventional thresholds ($0.11 \leq |\delta| < 0.28$).
- T3 vs. T4: Cliff's $\delta = -0.116$ $[-0.212, -0.023]$. The further degradation from flat delegation to hierarchical orchestration falls at the boundary of the "negligible" and "small" ranges ($|\delta| = 0.116$ is just above the Romano et al. 0.11 threshold; it would be classified as negligible under alternative cutoffs of 0.147).
- T0 vs. T6: $p = 0.018$. The overall trend from the simplest tier to the most complex is significantly negative in aggregate, though the pattern is not monotonic across all tasks (test-002 shows a modest positive trend).

For duration, the significant pairwise result is:

- T2 vs. T3: $p = 0.048$, Cliff's $\delta = -0.154$ $[-0.264, -0.045]$. T3 runs are stochastically shorter than T2, driven partly by rapid failures of delegation-based tiers on test-001 ($\sim$17 s vs. 430–505 s for T0–T2).

D.3 Cliff's δ Interpretation Scale

Cliff's δ [2, 16] is a non-parametric effect size measure that quantifies the probability that a randomly selected observation from one group exceeds a randomly selected observation from another group, minus the reverse probability. It ranges from -1 (all observations in group B exceed all in group A, i.e., the second group dominates) to $+1$ (all observations in group A exceed all in group B, i.e., the first group dominates), with 0 indicating no systematic difference. In our pairwise comparisons, group A is always the lower-numbered tier (e.g., T2) and group B the higher (e.g., T3); thus negative δ values indicate degradation with increasing tier. The conventional interpretation scale [14] is:

- $|\delta| < 0.11$: negligible effect

- $0.11 \leq |\delta| < 0.28$: small effect
- $0.28 \leq |\delta| < 0.43$: medium effect
- $|\delta| \geq 0.43$: large effect

The T2–T3 comparison ($\delta = -0.219$) falls clearly in the small effect range under all conventional thresholds. The T3–T4 comparison ($\delta = -0.116$) falls at the negligible/small boundary: classified as small under the Romano et al. (2006) 0.11 threshold but as negligible under alternative cutoffs (0.147). Note that this comparison does not reach corrected significance ($p = 0.058$ after Holm–Bonferroni correction), so the effect should be interpreted cautiously regardless of classification.

D.4 Scheirer–Ray–Hare Details

The Scheirer–Ray–Hare (SRH) test extends the Kruskal–Wallis framework to two-factor designs, enabling analysis of main effects and interaction effects on ranked data. Like ANOVA, it decomposes the total variability into components attributable to each factor and their interaction. Unlike ANOVA, it operates on ranked rather than raw data, making it robust to non-normality and heteroscedasticity. No tie-correction was applied to the H-statistics; given the large $N = 1{,}080$ and the magnitude of the observed effects, this omission does not materially affect the conclusions. For the pass rate variable (binary), ties are especially prevalent (each run is either 0 or 1), and applying a tie correction to the Kruskal–Wallis H for pass rate would increase the test statistic (making results more significant). The reported uncorrected $H(6) = 119.01$ for pass rate is therefore conservative.

Note on T6 sample size. T6 contributes only $n = 9$ runs to the SRH analysis (3 runs $\times$ 3 experiments $\times$ 1 subtest). The extreme imbalance between T6 ($n = 9$) and T3 ($n = 369$) may increase leverage of T6 cells in the rank decomposition. Readers should treat T6-specific SRH contributions with caution. The omnibus significance is driven primarily by T0–T5 patterns; T6 contributes only 0.8% of total observations and does not materially alter the qualitative conclusions.

For the score metric, the SRH analysis reveals that all three effects—tier main effect ($H = 91.3$, $df = 6$), experiment main effect ($H = 178.2$, $df = 2$), and interaction ($H = 223.5$, $df = 12$)—are statistically significant at $p < 0.001$. Raw H-values are not directly comparable across terms with different degrees of freedom ($df = 12$ vs. $df = 6$ vs. $df = 2$), but the significance of all three effects—particularly the interaction—indicates that the relationship between tier and score changes qualitatively across experiments. This is consistent with the divergent patterns described in Sections 5.2–5.4.

The degrees of freedom for the interaction ($df = 12$) correspond to the $(a-1)(b-1) = (3-1)(7-1) = 12$ interaction contrasts in the 3×7 (experiment $\times$ tier) design. The large H value (223.5) relative to the chi-squared distribution with 12 degrees of freedom produces $p \ll 0.001$, indicating that the interaction cannot be attributed to chance.

For cost and implementation rate, the interaction terms are non-significant ($p = 0.965$ and $p = 0.989$, respectively), indicating that the cost and attempt profiles are consistent across experiments regardless of tier.

E Per-Experiment Cost Breakdown

This appendix provides a detailed cost breakdown for each experiment, including total costs, per-run costs, best-performing tiers, and the distribution of cost across the token usage categories.

Table 15: *Per-experiment cost summary with best-performing tier by CoP.*

Experiment	Runs	Total ($)	Mean ($/run)	Best Tier	Best CoP ($)
test-001	360	13.72	0.04	T1	0.037
test-002	360	81.28	0.23	T5	0.228
test-003	360	27.30	0.08	T2	0.074
Total	1,080	122.31	0.11	–	0.113

The cost distribution across experiments is highly skewed: test-002 accounts for 66.5% ($81.28) of the total cost despite comprising only one-third of the runs. This reflects the larger repository size and greater complexity of the Mojo task, which requires more tokens for context loading and code generation.

The best Cost-of-Pass varies across experiments: T1 for test-001 ($0.037), T5 for test-002 ($0.228), and T2 for test-003 ($0.074). This variation reinforces the task-dependent nature of architectural optimization identified in the interaction analysis. No single tier achieves the best CoP across all experiments, consistent with the significant interaction effect.

E.1 Token Distribution

The token distribution across all runs is dominated by cache read operations. The mean cache read per run (approximately 593,000–773,000 tokens depending on tier) dwarfs the mean input tokens (42–67) and mean output tokens (2,409–2,665). This extreme ratio reflects the Scylla framework's use of extensive system prompts and repository context. The prompts and context are cached by the API and reused across runs, so the per-run cost is dominated by the cache read rate rather than the full input token rate.

The cache-dominated cost structure offers a plausible mechanism for the observed absence of detected cost differences across tiers: regardless of the architectural configuration, the agent must process a similar amount of cached context (the repository and system prompt), and the output length (approximately 2,500 tokens on average) is relatively consistent across tiers. The architectural overhead of delegation and orchestration does not substantially increase output token count, as the coordination messages are brief relative to the generated code.

F Reproducibility

F.1 Framework and Software

All experiments were conducted using the Scylla evaluation framework (v0.1.0+), an open-source Python package available at https://github.com/HomericIntelligence/ProjectScylla. The

framework provides:

- YAML-based experiment definitions with 120 subtests across 7 tiers
- Automated execution via the Claude Code CLI [1]
- Three-judge evaluation with configurable judge models
- Statistical analysis pipeline with non-parametric tests
- Automated figure and table generation for publication
- Cost tracking and token usage reporting

The framework is part of the Homeric Intelligence ecosystem, a collection of 12 repositories providing infrastructure for agent development, testing, and deployment [7].

F.2 Environment

Experiments were executed on Linux systems (kernel 5.10.0-37-amd64) with Python 3.10+. The Claude Code CLI was used in headless mode with deterministic settings where possible. API calls were made to Anthropic's production endpoints during January–March 2026. The Python dependencies and their exact versions are specified in the `pixi.toml` and `pixi.lock` files in the repository.

F.3 Data Availability

All raw results, including per-run scores, costs, token counts, and judge evaluations, are available in the repository at `docs/arxiv/haiku/data/`. The following data files are provided:

- `runs.csv`: Per-run data with scores, costs, and durations
- `subtests.csv`: Subtest-level aggregations
- `judges.csv`: Per-judge evaluation data
- `criteria.csv`: Per-criterion judge evaluations
- `summary.json`: Experiment-level summary statistics
- `statistical_results.json`: Full statistical test outputs
- `srh_tier_experiment.json`: Scheirer–Ray–Hare analysis

The statistical analysis scripts used to generate all figures and tables are included in the `scripts/` directory of the repository.

F.4 Limitations of Reproducibility

Exact reproduction of these results is not guaranteed due to several inherent factors in LLM-based experimentation:

1. **Model updates.** Anthropic may update model weights, inference infrastructure, or safety filters between experiment runs. These updates can affect model behavior in ways that are not publicly documented.
2. **Non-determinism.** Despite fixed temperature settings (where supported), LLM outputs may vary due to floating-point non-determinism, batching strategies, quantization effects, or infrastructure load balancing.
3. **External dependencies.** Tool-augmented tiers (T2+) depend on external services (MCP servers, web APIs, file systems) whose availability, behavior, and response content may change over time.
4. **Repository state.** The target repositories used in experiments may be updated between replication attempts. Commit hashes for the exact repository states used in each experiment are recorded in the experiment metadata files.

To maximize reproducibility, we recommend:

- Using the exact model identifiers specified in Table 11
- Running experiments within a short time window to minimize the impact of model updates
- Pinning target repository states to the commit hashes recorded in the experiment metadata
- Using the provided Docker containers for consistent execution environments
- Recording API version headers from Anthropic's responses to track any model version changes during the experiment