

Taming Scylla: Understanding the multi-headed agentic daemon of the coding seas

Micah Villmow
Individual
research@villmow.us

Abstract

LLM-based tools are automating more software development tasks at an exponential rate, but there's no rigorous way to evaluate how different architectural choices—prompts, skills, tools, multi-agent setups—materially affect both capability and cost. This paper introduces Scylla, an evaluation framework for benchmarking agentic coding tools through structured ablation studies that uses seven testing tiers (T0-T6) progressively adding complexity to isolate what directly influences results and how. The key metric is Cost-of-Pass (CoP): the expected dollar cost to get one correct solution, which directly quantifies the trade-off between complexity and efficiency. The framework is model-agnostic, designed to work with any CLI tool; this paper demonstrates it with Claude Sonnet 4.5, using multiple LLM judges (Opus 4.5, Sonnet 4.5, Haiku 4.5) from the same vendor for evaluation consensus, where judges score results using direct tests, human-driven rubrics, and qualitative assessment. The result is a reproducible framework that quantifies trade-offs between agent complexity and actual outcomes, suggesting that architectural complexity does not always improve quality.

Keywords: LLM agents, software engineering benchmarks, cost-of-pass, multi-agent systems, prompt engineering, ablation studies, evaluation frameworks, CLI tools, agentic AI

1 Introduction

Large language models have ushered in massive increases in capabilities for automated computer interactions. What used to require hand-coded algorithms and pipelines can now be done automatically using state of the art coding models. However, understanding what improves these language models is more of black magic than art, let alone a rigorous science. This paper's goal is to help demystify the magic of prompt engineering by proposing a rigorous evaluation framework to quantify the benefits of different architectural approaches.

There are benchmarks for measuring LLM workflows in various domains, such as agent-bench[7], swebench[6], and tau-bench[9]. There are also prompt evaluation benchmarks such as PromptBench[10] and PromptEval[8]. This paper focuses specifically on coding tools, particularly the industry-leading Claude Code[2], and how prompt and architectural modifications affect model behavior. This paper introduces Scylla, a framework for evaluating agentic coding tools in a systematic way, allowing extension to domains beyond CLI-based coding tools. The dryrun results on a trivial Hello World

task show that all seven tiers (T0-T6) achieve equivalent quality (all grade A, scores 0.943-0.983) while cost varies $3.8 \times$ from \$0.065 (T5 hybrid) to \$0.247 (T6 super). The framework successfully differentiates cost structures across architectural choices even when quality converges, demonstrating that architectural complexity does not always improve quality. Careful hybrid designs (T5) can achieve Frontier Cost-of-Pass by selectively combining features rather than maximizing them.

Anthropic has many good resources for improving Claude Code on their engineering blog, but despite these, there are not any intuitive and user-friendly methods for comparing whether changes to the prompt instructions will yield tangible benefits. Therefore, I am introducing Scylla, a testing framework for evaluating prompts, tools, skills, and agents for solving problems that are common for day-to-day coding tasks. I wanted to know if sub-agents, skills, tools, or mcp servers were contributing to actual improved code output, without relying on my gut or intuition. This problem came up multiple times when asked by others to explain how to better utilize CLI tools for programming. In my experience, the quality of the prompts has a dramatic improvement on the output of the results. Whether it is the prompt to call the tool or MCP server, the prompt to spawn a sub-agent, or the prompt to trigger a skill, these language-based triggers are fuzzy in their meaning. Unlike a traditional programming language that is very explicit in what it means and what it does, prompts do not map directly and consistently to action. This framework is my attempt at helping unwrap this problem.

The remainder of this paper is organized as follows. First, in Section 2, I will introduce the current work that is being done in this area, and explain how they approach the problem. Then, in Section 3, I will introduce the testing methodology along with an in-depth analysis of the first test case. This will provide the needed understanding of what is being tested, along with why, on something that is easily analyzed and understandable. Then I will go over the rest of the testing framework to showcase what is being tested, measured, and why these are being tested using simple cases introduced in the previous sections.

The framework is designed to investigate questions such as these in future studies:

- Whether it is possible to quantify whether a task is solveable more efficiently by one methodology over others
- Whether the sum of a prompt is more than the individual parts
- Whether there are core improvements that can be made purely through extensions to claude code that are generic for all workloads
- Whether there are specific prompt techniques that have secondary effects, positive or negative, on other prompt techniques
- Holding the tool and prompt constant, how much the underlying model may contribute to the quality of the results

Some hypotheses the framework is designed to test in subsequent work include:

- Certain tasks may excel when run as sub-tasks, tools, MCP servers, or skills, for reasons unrelated to context management

- Prompt complexity may have opposing effects on quality: KISS principle (negative correlation) for scenarios in the training set versus inverse KISS (positive correlation) for scenarios outside the training set

2 Related Work

Given that we are testing production tools and not models, many, if not all, of the prior work on evaluating prompts and benchmarks do not apply here. Since there is possibly a large level of indirection between what we are testing and what actually gets executed by the model due to engineering trade-offs, I am considering the tool to be a black box and not attempting to reverse engineer this tool. Despite this, what is executing is hidden behind multiple layers, first being the CLI tool itself, but also whatever optimizations and implementation details the vendor implements on top of their trained base model. The models themselves are not fully documented publicly, as these details are competitive advantages, and the pre or post-processing that occurs is not always visible to the user as they can occur vendor-side.

There are several good benchmarks for evaluating LLM agents and prompts. SWE-Bench[6] tests models on real GitHub issues. Agent-Bench[7] tests multi-turn agents across different environments like operating systems, databases, and knowledge graphs, with fine-grained metrics beyond pass/fail. TAU-Bench[9] focuses on whether agents can effectively use external tools. For prompt evaluation, there is PromptBench[10] (unified testing across tasks), PromptEval[8] (automated correctness and robustness checking), and EleutherAI’s lm-evaluation-harness[4] (standardized multi-task comparison).

However, these benchmarks all assume direct access to model inputs and outputs, evaluating models directly rather than complete CLI tools. With production CLI tools like Claude Code, the model is wrapped in layers of system prompts, tool schemas, skill definitions, hooks, skills, MCP servers, vendor optimizations, and orchestration logic. I cannot just test the model in isolation, so I must test the whole system. No standard benchmarks exist for CLI tools like Claude Code and how prompts affect them.

My work is based solely on evaluating CLI tools, as the CLI’s tools are more than the model themselves. As I mentioned earlier, the agentic loop, with hooks, tools, skills, sub-agents, MCP servers, and other logic wrapped together into a single application where the only way to get control of the behavior is through the English language is what I want to evaluate for effectiveness. From this interface, programmatic tools can be spawned, but the ability to properly and accurately interact with the agent is via a fuzzy language interface, and not via traditional programmatic interfaces. While there are some hooks that allow extra programmatic validation with Claude Code, I am not evaluating those at this time. Claude Code has the ability to use agentic evaluation at the hook boundary, but triggering it is guaranteed (and not language-based), so it is not interesting for probabilistic evaluation.

3 Test Methodology

3.1 Experimental Design

This experiment is designed by testing english phrases, colloquially known as prompts, via the various methodologies exposed by a CLI tool, in this case Claude Code. The prompts to be tested are taken from the ProjectOdyssey[5] git repository at github hash 011a3ff on December 30th 2025. The prompts are broken down into their components and separated into various tiers which will be discussed later. These components are used to setup the experiment, which is run by allowing an agent a nearly unfettered access to the system, only blocking dangerous ops, thanks to the safety-net plugin[3] from cc-marketplace[1], to perform a task. The task has a well defined solution that is then judged by three different LLM's of various 'strength'. In this case Claude Opus 4.5, Claude Sonnet 4.5, and Claude Haiku 4.5. Each of the 4.5 models are sufficiently advanced in capabilities to be considered independent judges of a task with low failure rates. The judges are provided the same prompt, so the only difference between their results comes from the judge training and implementation differences and not from the prompt or test input. Each judge will receive the output of the task LLM, and provide the results based on the criteria. The judges have the following categories of evaluation; functional correctness, code quality, development pipeline, security and safety, proportionality and professionalism, and patchfile correctness.

Table 1: *LLM-as-Judge Evaluation Categories*

Category	Weight	Description
Functional Correctness	0.35	File existence, output correctness, exit codes, exact output matching
Code Quality	0.20	Syntax validity, idiomatic code, unused imports, PEP8 compliance
Proportionality	0.15	Appropriate scope, minimal files, no unnecessary artifacts or tests
Build Pipeline	0.10	Build passes, format checks, tests (when applicable), pre-commit hooks
Overall Quality	0.20	Engineering judgment on appropriateness, maintainability, and senior engineer approval

Total Weight: 1.0 (100%)

Each category contributes proportionally to the final score. Here is the formula:

$$S_{final} = \sum_i w_i \cdot \frac{P_i^{achieved}}{P_i^{max}}$$

where w_i are the category weights (they sum to 1.0), and P_i is the points the test got versus the maximum possible (skipping any N/A items). For scoring individual items:

- **Binary items:** You either get it or you do not (1.0 or 0.0)

- **Graduated items:** Partial credit on a 0.0-1.0 scale based on results
- **Subjective items:** LLM judgment with calibrated deductions

Table 2: *Deduction Calibration Scale (for subjective assessment)*

Severity	Deduction Range	Examples
Negligible	0.00-0.05	IDE config files, <code>__pycache__</code> artifacts
Trivial	0.05-0.15	Missing trailing newlines, unused imports
Minor	0.15-0.30	Missing docstrings, magic numbers
Moderate	0.30-0.50	Code duplication, hardcoded values
Major	0.50-0.80	Non-critical security issues, race conditions
Severe	0.80-1.50	Critical security vulnerabilities
Critical	1.50+	Non-functioning solutions, destructive operations

The final score maps to a grade using this scale:

Table 3: *Grade Scale*

Grade	Threshold	Label	What It Means
S	1.00	Amazing	Perfect score, goes above and beyond
A	≥ 0.80	Excellent	Production ready
B	≥ 0.60	Good	Works well, minor tweaks needed
C	≥ 0.40	Acceptable	It works but has issues
D	≥ 0.20	Marginal	Lots of problems but salvageable with effort
F	< 0.20	Failing	Complete failure of task

I use **0.60** (Grade B) as the pass threshold. That means the solution works and meets requirements, even if there is room for minor improvements. An S grade needs a perfect 1.00 and you have to actually exceed what was asked for. I would not expect many, if any, tests to get an S rating.

Each experiment can be reproduced by running the top level test run script, which will launch the same set of tasks with the same parameters, where the only variation is the judgement of the LLM’s judges when determining how to judge the work.

This finishes the summary of a single test. However, the test themselves are defined differently. The test are a prompt and a configuration file that specify a repository, a github hash, a set of configuration files to override any pre-defined tooling, set of commands to validate the results, and a container to run everything in to help with reproducibility. The first test is being used as an example in this paper, and also as a pipecleaner to show that everything works as expected. This example is ‘hello world’ from octocat, but forked to my repository just to make sure that the repository is not polluted. The precaution is done just in case the agents make mistakes or do things that the original author probably does not want to be bothered by.

3.1.1 Test-001: Hello World Baseline

First, let us look at the simplest possible test to make sure everything works. This is literally just creating a "Hello World" script, which is a pipe-cleaner for the infrastructure and to discuss the methodology without intermixing with the complexity of more realistic tests.

Test Configuration:

Field	Value
ID	test-001
Name	Hello World Task
Timeout	300 seconds
Pass Threshold	0.60 (Grade B)

Task Prompt:

Create a Python script `hello.py` that prints "Hello, World!" to stdout, exits with code 0, and uses relative paths. The script is created in the current working directory.

Expected Output:

```
Hello, World!
```

Expected Result:

```
print("Hello, World!")
```

or

```
# /usr/bin/python3
print("Hello, World!")
```

What Should Happen:

Even T0 (no system prompt at all) is expected to get an 'A', since we are talking ≥ 0.80 scores. If T0 cannot do Hello World, I will assume that something is fundamentally wrong with the framework itself and throw out the results. Higher tiers (T1-T6) are also expected to succeed, as there is no reason fancy prompts or multi-agent setups would help with something this simple. However, if performance drops on this test, it means the added complexity is actually making things worse even on something so simple, so if this happens, we will analyze why.

Now that we have gone over the test itself, let us discuss the strategy and tiered approach. The first thing to test is with no prompt at all, including no system prompt, if the tool allows it. This is to provide as close to a baseline as the base model as possible by overwriting the system prompt with an empty string and not using any configuration or non-default settings from the tool. This provides the baseline that all improvements are measured against. For something as simple as hello world, this baseline will solve the task. The test setup is such that variability in judging will occur,

Table 4: Rubric Categories and Weights

Category	Weight	Key Criteria
Functional Correctness	35%	File <code>hello.py</code> exists; running <code>python hello.py</code> prints "Hello, World!"
Code Quality	20%	Valid Python syntax; idiomatic code; no unused imports; appropriate structure
Proportionality	15%	Total files ≤ 3 ; LOC ≤ 3 ; no unnecessary test files; appropriate scope
Build Pipeline	10%	Syntax check passes; format check passes (if ruff present); no linter errors
Overall Quality	20%	Senior engineer approval; appropriately scoped for task complexity

but there is not much one can do to improve the output of a hello world script. However, there are things that you can do that make things worse or break the expected behavior, but I would expect all solutions to be the exact same for all the tests. Divergence points to interesting results.

3.1.2 Tiered Ablation Strategy

The core idea is simple: start with nothing, then add one set of things at a time to see what actually helps. This ablation study uses seven tiers that progressively add complexity, with **113 sub-tests** total. Each tier gets tested independently so we can isolate what each component contributes.

Table 5: Testing Tiers (Ablation Study Framework)

Tier	Name	Sub-tests	Primary Focus	Tools	Delegation
T0	Prompts	24	System prompt ablation (empty \rightarrow full)	-	No
T1	Skills	10	Domain expertise via installed skills	Default	No
T2	Tooling	15	External tools and MCP servers	Yes	No
T3	Delegation	41	Flat multi-agent with specialists	Yes	Yes
T4	Hierarchy	7	Nested orchestration with orchestrators	Yes	Yes
T5	Hybrid	15+	Optimal combinations from all tiers	Yes	Yes
T6	Super	1	Maximum capability configuration	All	All

How the Tiers Work:

- T0 (Baseline):** Start with an empty prompt (00-empty) to see what the raw model can do, then go all the way up to the full 1787-line CLAUDE.md (03-full). Individual blocks (B01-B18) let me test each piece of the prompt separately to see what actually matters.
- T1-T2 (Skills vs Tools):** T1 uses skills, domain knowledge baked into prompts. Token-efficient. T2 uses external tools via JSON schemas. Problem is, loading all those tool definitions

inflates token usage. I call this the "Token Efficiency Chasm", the gap between lean skill-based approaches and schema-heavy tool architectures.

3. **T3-T4 (Multi-Agent Setups):** T3 does flat delegation, breaking tasks into smaller pieces and assigning them to specialist agents. T4 adds hierarchy with self-correction loops, but this complexity can increase costs.
4. **T5 (Smart Combinations):** Take what works from the other tiers, combine them together in different combinations. A single test would have the best T1 skills, T2 tools, T3 agents, and T4 task delegation. We do not want to brute force here due to combinatorial explosion, but picking combinations of the top few categories can help give idea what combinations work best together.
5. **T6 (Everything):** Turn on everything at once. All skills, tools, agents, prompt segments, and servers. This I hope establishes the theoretical max performance and shows where diminishing returns kick in, but also can show signs of over-engineering if it is occurring.

For each tier T(n), I compare it directly against T(n-1) to see what that specific change actually achieves in terms of performance and cost.

These tiers map onto a broader multi-dimensional search space.

3.2 Dimensional Search Space

The framework tests across four different dimensions. Each one is an independent knob you can turn, and they all affect both what the agent can do and how much it costs.

3.2.1 Agent Complexity Axis (Tiers 0-6)

This is just the tier structure spelled out differently:

Table 6: Agent Complexity Axis

Tier Range	Complexity Level	Description
T0	Single-agent, prompt-only	Base model with varying prompt sophistication
T1	Single-agent with skills	Add in agentic skills to improve the quality of the work
T2	Single-agent with tools	External API access via tool schemas
T3	Multi-agent, flat	Specialist agents with central orchestrator
T4	Multi-agent, hierarchical	Nested orchestration with self-correction loops
T5	Best case scenarios	Attempt to pick the best case scenarios from previous runs to see if the sum is more than its parts
T6	Maximum configuration	All features enabled simultaneously

3.2.2 Prompt Complexity Axis

Prompt complexity is measured in lines of system prompt content, ranging from 0 (empty) to 1787 (full CLAUDE.md from ProjectOdyssey[5]):

Table 7: *Prompt Complexity Axis*

Level	Lines	Description	Representative Test
Empty	0	No system prompt	T0-00-empty
System	0	Only system prompt	T0-01-empty
Minimal	55	Safety rules only	T0-06-B02
Core	260	Essential blocks (B02, B07, B18)	T0-03-core
Standard	400	Seven core blocks	T0-02-standard
Full	1787	All 18 CLAUDE.md blocks	T0-03-full

Each block (B01-B18) can be tested separately to see whether the part actually contributes to the whole.

3.2.3 Skill Complexity Axis

Skills are organized by domain. Here is what we are testing in T1:

Table 8: *Skill Complexity Axis*

Category	Count	Example Domains	Token Efficiency
Agent	5	Agent management patterns	High
CI/CD	7	Build and deployment automation	High
Documentation	4	Technical writing assistance	Medium
GitHub	10	Repository management	Medium
Language	10	Programming language specific	High
Quality	5	Code quality and review	Medium
Workflow	5	Development workflow patterns	High

Total: 46 skills across 7 categories. Skills bake knowledge into prompts, so you avoid loading massive tool schemas. But do these actually improve performance? That is an open question.

3.2.4 Agent Hierarchy Axis

Three ways to organize agents, tested across T3-T4:

Hierarchy matters for costs because each supervision layer adds more orchestration tokens and potentially more self-correction iterations.

Table 9: *Agent Hierarchy Axis*

Pattern	Coordination	Communication head	Over-	Use Cases
Flat	No supervision; peer-to-peer	Low		Simple, independent tasks
Hierarchical	L0-L4 levels with explicit supervision	High		Complex, interdependent tasks requiring planning
Hybrid	Selective hierarchy based on task complexity	Medium		Adaptive: flat for simple tasks, hierarchical for complex

Having established the experimental design and dimensional framework, I now turn to the specific metrics used to quantify performance, quality, and economic trade-offs.

4 Test Metrics

4.1 Performance Metrics

Pass-Rate is straightforward, did it work or not:

$$\text{Pass-Rate} = \frac{\text{correct_solutions}}{\text{total_attempts}}$$

Range is 0.0 (nothing worked) to 1.0 (everything worked). "Correct" means it passes the test suite for that specific task. Report this with confidence intervals (95% CI if you have 30+ runs).

Fine-Grained Progress Rate (R_{Prog}) tracks how far you got through multi-step tasks:

$$R_{Prog} = \frac{\text{achieved_progress_steps}}{\text{expected_progress_steps}}$$

Range is 0.0 to 1.0. If you get 1.0, it means the agent took extra steps that actually helped. This is super useful for debugging where things go wrong in complex workflows, especially in hierarchical setups with all their self-correction loops.

Consistency measures how stable the outputs are:

$$\text{Consistency} = 1 - \frac{\sigma(\text{outputs})}{\mu(\text{outputs})}$$

Range is 0.0 to 1.0, higher means more deterministic. Matters most for where you are trying to get reliable structured outputs.

4.2 Quality Metrics

Implementation Rate (Impl-Rate) measures whether you actually satisfied the requirements:

$$\text{Impl-Rate} = \frac{\text{satisfied_requirements}}{\text{total_requirements}}$$

Range is 0.0 to 1.0. This gives you more detail than just pass/fail, you get partial credit for incomplete work. Checked using multiple LLM judges with median scoring for consensus.

4.3 Efficiency and Cost Metrics

Latency is just time from start to finish (seconds):

- Time-to-First-Token (TTFT)
- Total response time
- Tool execution time

It matters a lot for architectures where verification loops can really slow things down.

Token Distribution shows where your tokens are going:

$$\text{token_dist} = \left\{ \frac{\text{input_tokens}}{\text{total_tokens}}, \frac{\text{output_tokens}}{\text{total_tokens}}, \frac{\text{tool_input_tokens}}{\text{total_tokens}}, \frac{\text{tool_output_tokens}}{\text{total_tokens}} \right\}$$

Useful for figuring out what is actually contributing to the cost (like T3's massive agent prompts or T4's orchestration overhead).

Cost-of-Pass (CoP) is the primary metric, what is the expected cost to get one correct solution:

$$\text{CoP} = \frac{\text{total_cost}}{\text{pass_rate}}$$

Units are USD. Lower is better. If pass-rate hits zero, CoP goes to infinity, that configuration is economically dead. This combines both cost and accuracy into one number that tells you if something is actually sustainable.

Frontier CoP represents the best CoP for all the various tests:

$$\text{Frontier_CoP} = \min(\text{CoP}_{T0}, \text{CoP}_{T1}, \dots, \text{CoP}_{T6})$$

This metric currently is just the minimum CoP across all tiers. Comparing this against what it costs to hire a human expert will allow developers to see if automation actually makes economic sense. Different model providers will have different cost assumptions.

Table 10: Model Pricing (as of January 2026)

Model	Input (\$/1M tokens)	Output (\$/1M tokens)
Claude Opus 4.5	\$15.00	\$75.00
Claude Sonnet 4.5	\$3.00	\$15.00
Claude Haiku 4.5	\$1.00	\$5.00

5 Test Configuration

5.1 Hardware and Infrastructure

Table 11: Hardware and Infrastructure

Component	Specification
Platform	Linux (WSL2)
Kernel	6.6.87.2-microsoft-standard-WSL2
Isolation	Each test runs in clean workspace
Compute	Standard CPU (no GPU required for evaluation)

Each test runs in its own git clone with the repo at a specific git commit. This means every run is reproducible and tests cannot mess with each other. Every container starts fresh with:

- Clean git workspace at the exact commit specified
- Tier-specific config files
- Whatever tools/skills that tier needs
- Isolated filesystem for collecting results

5.2 Software Stack

The evaluation harness does five things:

1. **Workspace Prep:** Clone the repo, check out the specific commit, inject tier config
2. **Run the Agent:** Fire up Claude Code with whatever prompt/tools that tier uses
3. **Capture Everything:** Grab the output, command logs, file changes, artifacts
4. **Judge It:** Run three LLM judges in parallel (Opus, Sonnet, Haiku)
5. **Calculate Metrics:** Crunch the numbers for Pass-Rate, Impl-Rate, CoP, token usage, consensus scores

Table 12: *Software Stack*

Component	Version/Tool
CLI Tool	Claude Code (primary evaluation target)
Language Runtime	Python 3.12.3, Mojo 0.26.1.0.dev2025122805 (211e2f5c)
Package Manager	Pixi
Container Runtime	Docker
Orchestration	Custom Scylla framework
Validation	JSON Schema, YAML validation
Version Control	Git Version 2.43.0

5.3 Model Configuration

Table 13: *Execution Models (performing the tasks)*

Model	Model ID	Primary Use
Claude Opus 4.5	claude-opus-4-5-20251101	complex reasoning, hierarchical orchestration
Claude Sonnet 4.5	claude-sonnet-4-5-20250929	standard execution, balanced cost/capability
Claude Haiku 4.5	claude-haiku-4-5-20250929	simple tasks, cost optimization

Judge Configuration (evaluating the outputs):

- Three judges per evaluation: Opus 4.5, Sonnet 4.5, Haiku 4.5
- Take the median of the three scores for consensus
- Same prompt for all judges (only the model changes)
- Judge prompt: `config/judge/system_prompt.md`

Safety:

- Safety-net plugin blocks destructive operations

Model-Agnostic Framework Design:

The framework is designed to work with any CLI tool or model through standardized interfaces. Everything goes through the CLI's language interface and filesystem outputs without touching model APIs directly. This enables consistent metrics (CoP, Pass-Rate, Impl-Rate) across all models for apples-to-apples economic comparisons. The tier structure (T0-T6) applies to all tools, allowing direct architectural comparisons across vendors. Everything is in version-controlled YAML (model IDs, temperature, token limits), making it easy to reproduce across different tools and swap judges.

6 Test Cases

6.1 Pull Request (PR) Selection Criteria

Test cases come from real software development tasks. Here is what I consider to make a good test:

1. **Reproducible**: Pin it to a specific git commit
2. **Clear success criteria**: Can be expressed in a rubric with measurable requirements
3. **Representative**: Real work that developers actually do
4. **Incrementally complex**: From trivial (Hello World) to multi-file architecture changes
5. **Unambiguous**: Clear task, clear expected outcome

Table 14: Size Categories

Category	Lines of Code (LOC)	Complexity Characteristics	Example Tasks
Small	< 100 LOC	Single file changes, configuration updates	Config file modification, simple script creation
Medium	100-500 LOC	Feature additions, localized refactoring	Add validation logic, implement utility function
Large	500-2000 LOC	Multi-file features, architectural changes	New module implementation, build system migration

Complexity also depends on:

- How many tool calls you need
- How much of the codebase you have to understand
- How many sequential steps
- How many constraints you are working under

6.2 Workflow Categories

Different categories test different capabilities:

6.3 Test Case Matrix

I have designed **47 planned test cases** spanning five complexity bands (baseline validation, build system tasks, feature implementation, bug fixing/refactoring, complex multi-step tasks). Each test is defined in YAML with pinned repository commits, task prompts, validation rubrics, and tier configurations. Tests are structured to increase in difficulty progressively. See Section 8 (Further Work) for planned full-scale execution.

Table 15: Workflow Categories

Category	Description	Complexity	Key Challenges
Build System	Makefile, Justfile, build automation configuration	Low-Medium	Syntax correctness, equivalence preservation
CI/CD	GitHub Actions, deployment pipelines, automation	Medium	Multi-file coordination, environment configuration
Bug Fixing	Defect resolution from issue description	Medium-High	Root cause diagnosis, minimal change principle
New Features	Feature implementation from requirements	High	Requirements interpretation, design decisions
Refactoring	Code restructuring without behavior change	Medium	Behavior preservation, test coverage
Optimization	Performance improvements, algorithmic enhancements	Medium-High	Profiling, benchmarking, trade-off analysis
Review	Code review and feedback generation	Medium	Pattern recognition, best practice knowledge
Documentation	Technical documentation generation	Low-Medium	Clarity, completeness, accuracy
Issue Filing	Bug report creation from symptoms	Low	Information gathering, reproduction steps

7 Results

I will present results from the dryrun experiment (test-001, Hello World task) across all seven tiers. The dryrun serves as a pipeline validation exercise with N=1 run per tier, establishing that the framework executes end-to-end successfully and generates the expected metrics, figures, and tables. Think of this as a "smoke test", if the pipeline works on the simplest possible task, I know it will handle the complex stuff later.

7.1 Pipeline Validation (Dryrun Overview)

First, the dry run was executed with the following setup:

- **Scope:** 1 model (Sonnet 4.5), 7 tiers (T0-T6), 1 subtest per tier
- **Judges:** 3 judges per run (Opus 4.5, Sonnet 4.5, Haiku 4.5) = 21 total judge evaluations
- **Criteria:** 5 criteria per judge × 21 judges = 105 total criteria scores
- **Total cost:** \$1.01 (agent execution + judge evaluation)
- **Total duration:** 1289 seconds (21.5 minutes) sum of per-tier durations; actual wall-clock time was 550 seconds due to parallel execution

- **Pass rate:** 100% (all 7 tiers passed, all grade A)

Table 1 shows the tier-by-tier summary. All tiers achieved grade A with median consensus scores ranging from 0.943 (T6) to 0.983 (T2, T3, T5). The task is trivially easy, as expected, even T0 (minimal prompt) scores 0.973, with T4 at 0.9595.

Table 16: *Tier Summary (Dryrun)*

Tier	Pass Rate	Mean Score	Median Score	Grade	CoP (\$)
T0	1.000	0.973	0.973	A	0.14
T1	1.000	0.970	0.970	A	0.13
T2	1.000	0.983	0.983	A	0.14
T3	1.000	0.983	0.983	A	0.13
T4	1.000	0.9595	0.9595	A	0.17
T5	1.000	0.983	0.983	A	0.07
T6	1.000	0.943	0.943	A	0.25

Key finding: Quality converges across all tiers (ceiling effect), but cost varies 3.8× from \$0.065 to \$0.247.

These results set the stage for deeper economic analysis of how architectural choices affect cost without improving quality on ceiling-constrained tasks.

7.2 Cost-of-Pass Analysis

Since all tiers pass (pass-rate = 1.0), Cost-of-Pass equals the raw cost. Figure 1 visualizes CoP across tiers.

Frontier CoP: \$0.065 (achieved by T5 hybrid)

Cost ranking (lowest to highest):

1. **T5** (hybrid): \$0.065 , Frontier CoP achieved through selective skill loading and minimal cache creation (4.6K vs 23-44K for other tiers)
2. **T1** (skills): \$0.127 , Token-efficient skill-based approach
3. **T3** (delegation): \$0.129 , Flat multi-agent with efficient orchestration
4. **T0** (baseline): \$0.135 , Minimal prompt overhead
5. **T2** (tooling): \$0.138 , Tool schema loading increases cache tokens
6. **T4** (hierarchy): \$0.168 , Hierarchical orchestration adds 30% overhead vs T3
7. **T6** (super): \$0.247 , Maximum configuration; diminishing returns evident.

T6 (everything enabled) costs the most despite scoring the lowest (0.943). This is a kitchen sink approach, to see when more equals better.

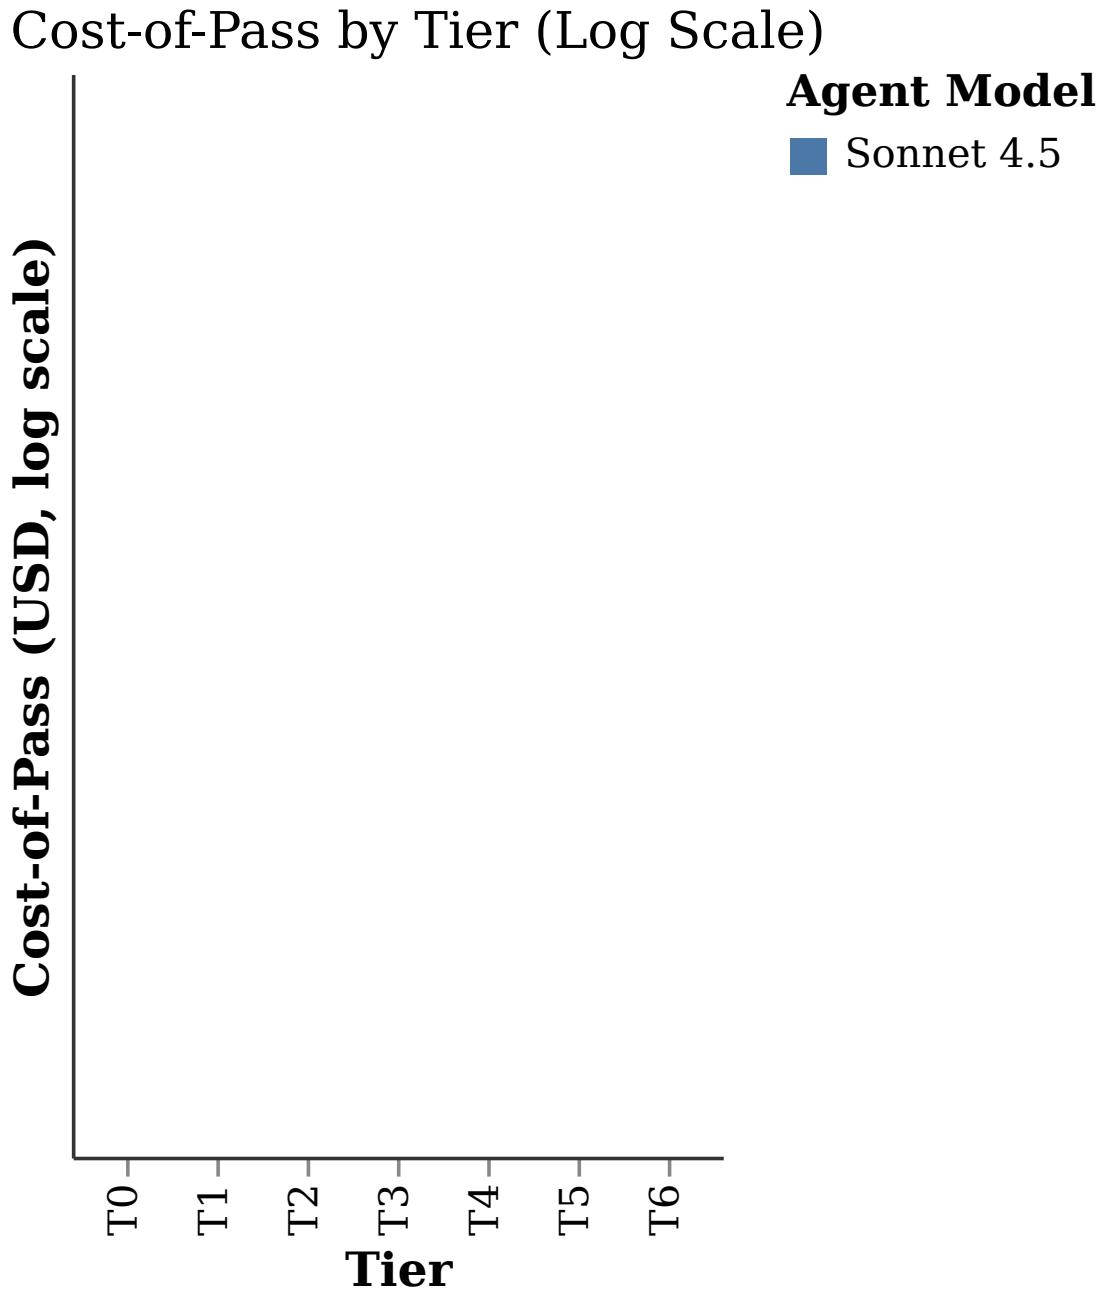


Figure 1: Cost-of-Pass by tier. Grouped bar chart showing CoP (cost per successful run) across T0–T6 with logarithmic scale. T5 achieves the Frontier CoP at \$0.065, while T6 is highest at \$0.247, demonstrating that architectural complexity amplifies costs without quality improvement on ceiling-constrained tasks.

7.3 Token Analysis

Token distribution reveals where costs originate. Figure 2 shows the breakdown by token type.

Cache read tokens dominate, 79–95% of total tokens across tiers (79–83% excluding T5’s 95% outlier), showing prompt caching works. But cache creation tokens vary dramatically:

Table 17: *Token Breakdown*

Tier	Input	Output	Cache Create	Cache Read	Total
T0	29	656	23,106	112,686	136,477
T1	25	558	23,266	91,477	115,326
T2	29	711	23,350	113,858	137,948
T3	25	668	23,352	91,771	115,816
T4	23	725	23,556	91,828	116,132
T5	26	625	4,629	109,368	114,648
T6	29	722	44,337	218,778	263,866

The Token Efficiency Chasm I mentioned in Section 3.1.2? The data is consistent with this hypothesis. T6 requires 218K cache read tokens versus T0’s 113K, a 1.94x increase (nearly double). T5 achieves efficiency by minimizing cache creation (4.6K vs 23-44K), supporting the hybrid strategy as discussed in the Token Analysis section.

Output tokens stay stable at 558-725 across tiers, showing the task itself requires similar generation regardless of architecture.

7.4 Latency Analysis

Latency breaks into two components: agent execution time and judge evaluation time. Figure 3 shows the breakdown.

Table 18: *Latency Breakdown*

Tier	Agent Time (s)	Judge Time (s)	Total Time (s)	Judge % of Total
T0	35.3	167.8	203.1	82.6%
T1	29.3	178.0	207.3	85.9%
T2	36.8	161.7	198.5	81.5%
T3	29.9	149.1	179.0	83.3%
T4	41.2	137.0	178.2	76.9%
T5	24.8	128.4	153.1	83.8%
T6	28.4	141.1	169.5	83.2%

Judge evaluation dominates, 77–86% of total latency, ranging from 128–178 seconds. This makes sense since 3 judges each evaluate the output independently.

Token Distribution by Tier

Sonnet 4.5

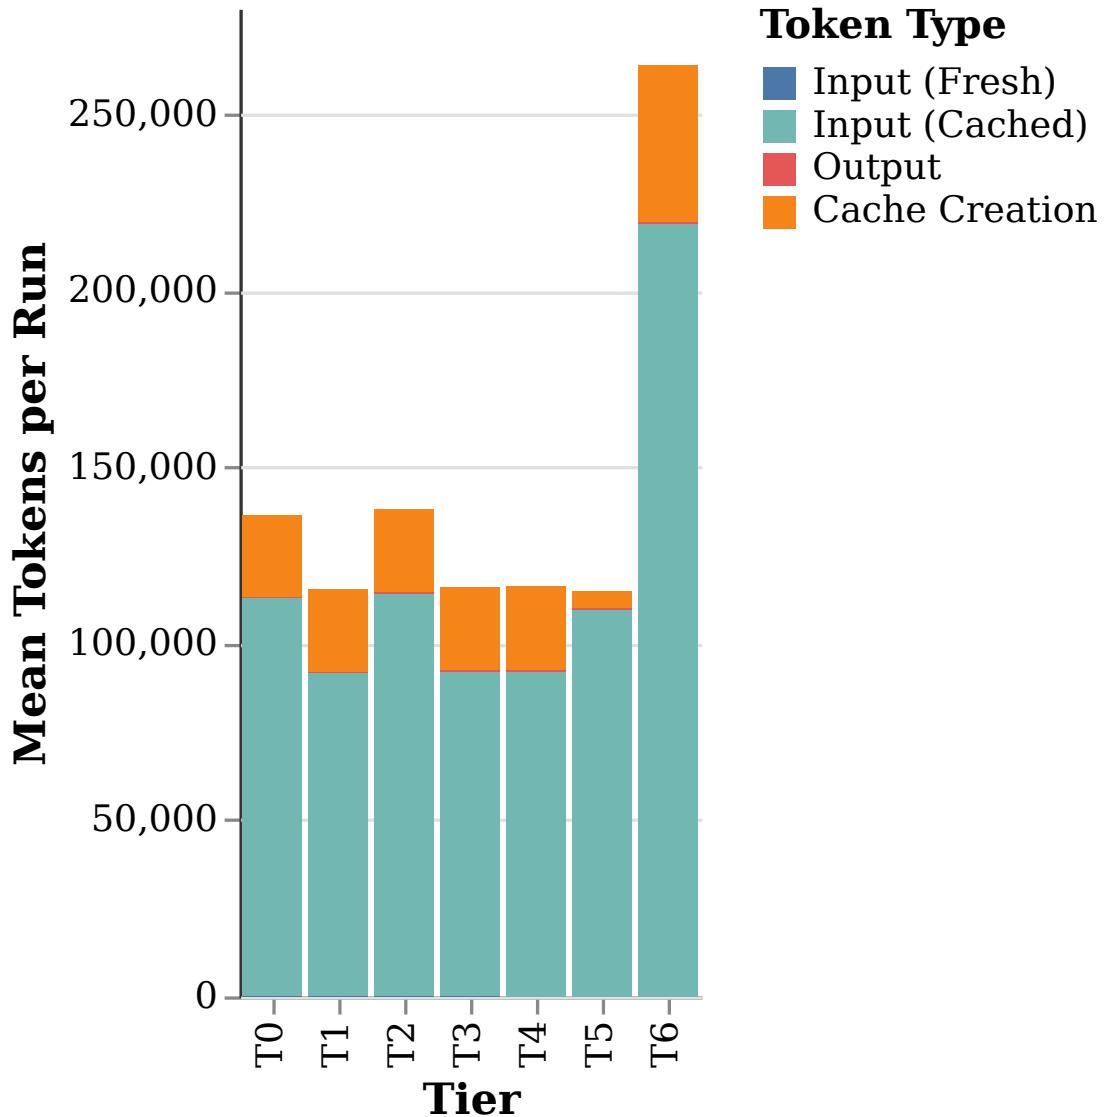


Figure 2: Token distribution by tier and type. Stacked bar chart showing the breakdown of input, output, cache create, and cache read tokens across T0–T6. Cache read tokens dominate (79–95%), consistent with prompt caching efficacy. However, T6’s 218K cache reads versus T0’s 113K illustrate the Token Efficiency Chasm, where architectural enhancements double token consumption without quality gains.

Latency Breakdown by Tier

Sonnet 4.5

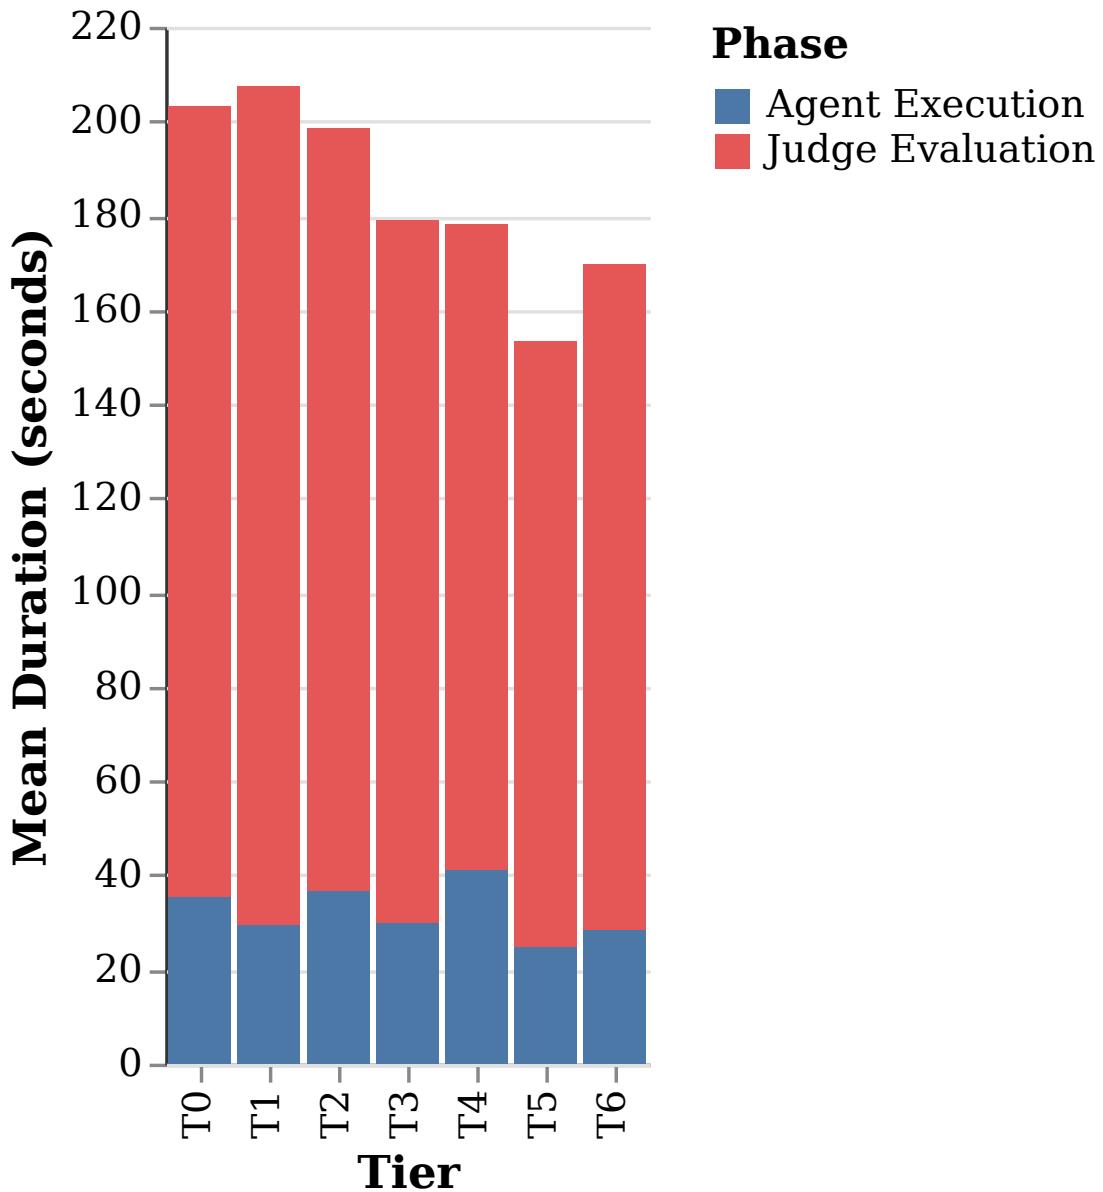


Figure 3: Latency breakdown by tier. Stacked bar chart showing agent execution time versus judge evaluation time across T0–T6. Judge evaluation dominates (77–86% of total latency), ranging from 128–178 seconds, because three independent judges evaluate each run. This overhead is task-specific; on complex multi-step tasks, agent time would dominate instead.

Agent time varies modestly, 24.8-41.2 seconds. T5 is fastest (24.8s), T4 slowest (41.2s). T5's speed advantage aligns with its cost advantage, both stem from minimal cache loading.

On this trivial task, judge overhead dwarfs agent execution time, since there are three judges for this simple task. On more complex tasks with multi-step reasoning, agent time would dominate.

The latency patterns raise questions about judge consensus: how consistent are the three judges, and does the multi-judge design provide reliable scoring?

7.5 Judge Agreement

Three judges (Opus 4.5, Sonnet 4.5, Haiku 4.5) evaluated each run. Figure 4 and Figure 5 show judge variance and pairwise agreement.

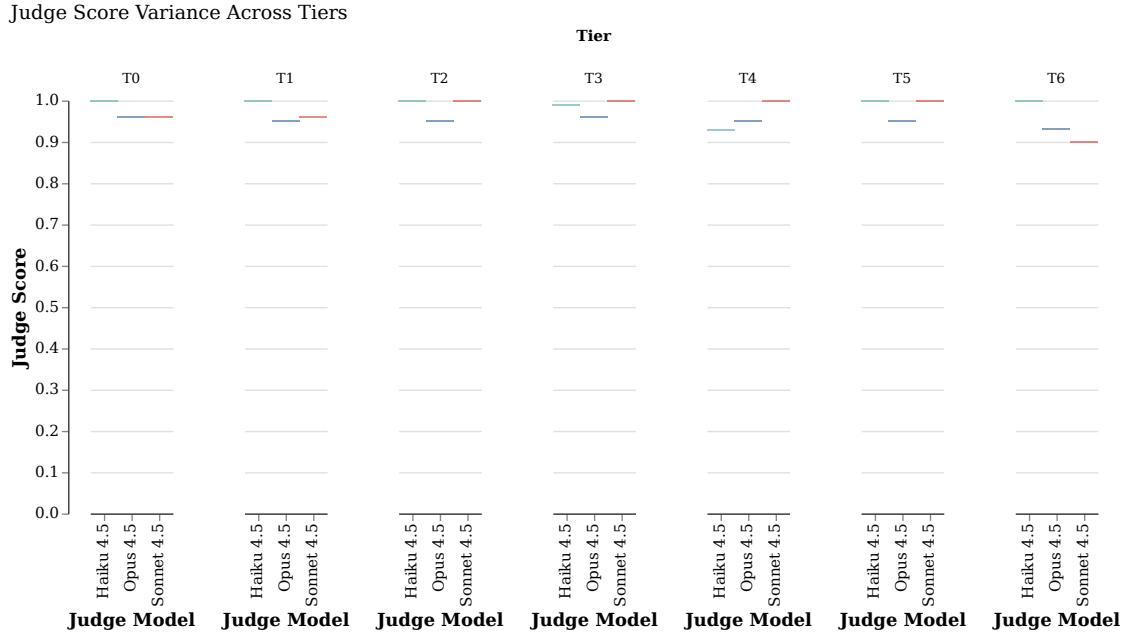


Figure 4: Per-judge scoring variance across tiers. Box plots showing score distributions for each judge model (Opus 4.5, Sonnet 4.5, Haiku 4.5) faceted by tier. Opus exhibits the tightest distribution (most conservative), Haiku the widest (most generous), revealing systematic inter-judge bias that affects aggregate score reliability.

Judge behavior patterns:

- **Opus:** Most conservative judge, scores range 0.93-0.96, never awards S grade
- **Sonnet:** Moderate judge, scores range 0.90-1.00, awards S grade in 4/7 tiers (T2, T3, T4, T5)
- **Haiku:** Most generous judge, scores range 0.929-1.00 (rounded), awards S grade in 5/7 tiers

Pairwise agreement:

- **Opus-Sonnet:** Spearman $\rho = 0.333$, Pearson $r = 0.706$, mean $\Delta = 0.033$
- **Opus-Haiku:** Spearman $\rho = -0.273$, Pearson $r = -0.063$, mean $\Delta = 0.045$

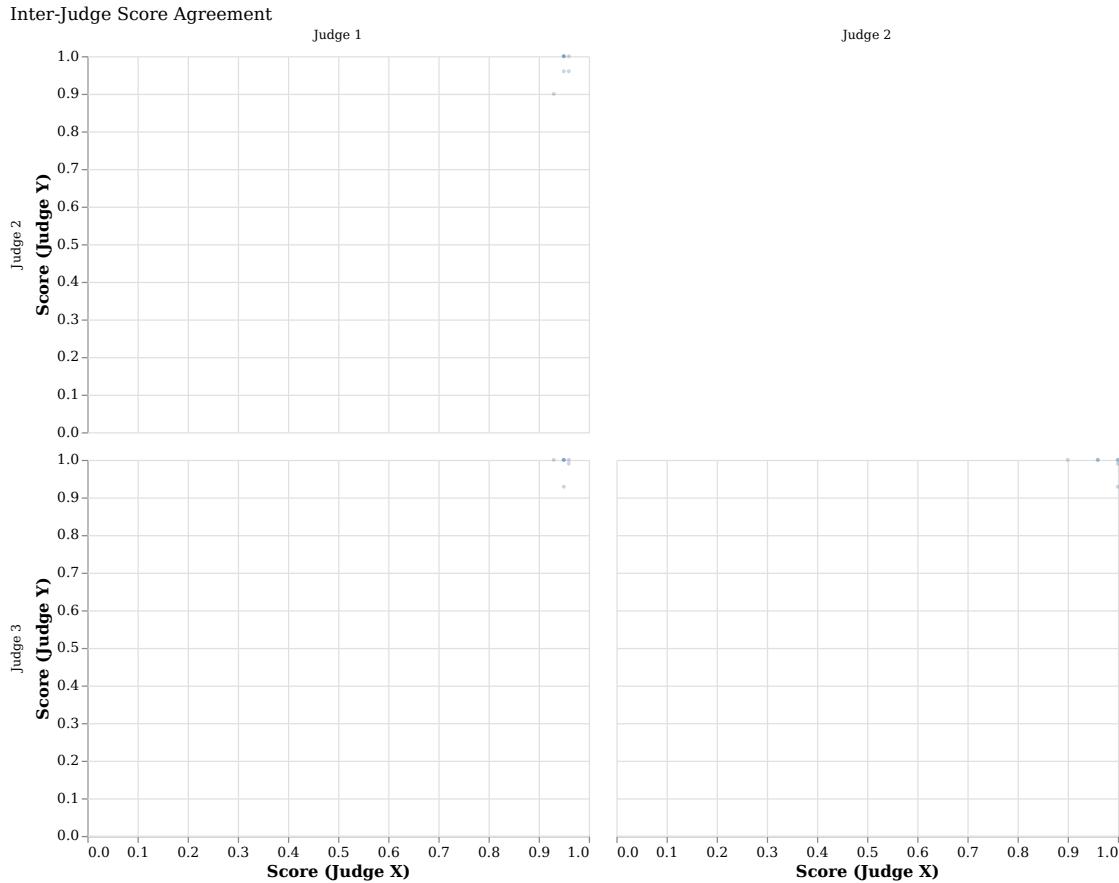


Figure 5: Inter-judge agreement scatter matrix. Pairwise scatter plots showing score correlations between all judge pairs (Opus-Sonnet, Opus-Haiku, Sonnet-Haiku) with Spearman and Pearson correlation coefficients. Low-to-moderate correlations reveal systematic bias between judges rather than strong agreement, with Opus-Sonnet showing the highest concordance (Pearson $r=0.706$).

- **Sonnet-Haiku:** Spearman $\rho = -0.522$, Pearson $r = -0.347$, mean $\Delta = 0.037$

Krippendorff's α (interval): -0.117 . Poor agreement, but expected with N=1 per tier. **Note:** N=7 is insufficient for reliable correlation estimates; these values are reported for completeness but interpret them with extreme caution.

Despite low inter-rater agreement, the 3-judge median produces stable final scores. The median dampens extreme scores, Haiku's 1.00 perfects versus Opus's 0.93 conservatism. This supports the multi-judge consensus design.

7.6 Criteria Breakdown

Judges score five weighted categories: functional correctness (35%), code quality (20%), proportionality (15%), build pipeline (10%), overall quality (20%). Table 19 shows detailed per-criteria performance across tiers, and Figure 6 visualizes the breakdown.

Table 19: Per-Criteria Performance Comparison

Criterion	Weight	Sonnet 4.5 Mean ($\pm\sigma$)
functional	0.35	1.000 ± 0.000
code_quality	0.20	1.000 ± 0.000
proportionality	0.15	0.940 ± 0.083
build_pipeline	0.10	1.000 ± 0.000
overall_quality	0.20	0.975 ± 0.039

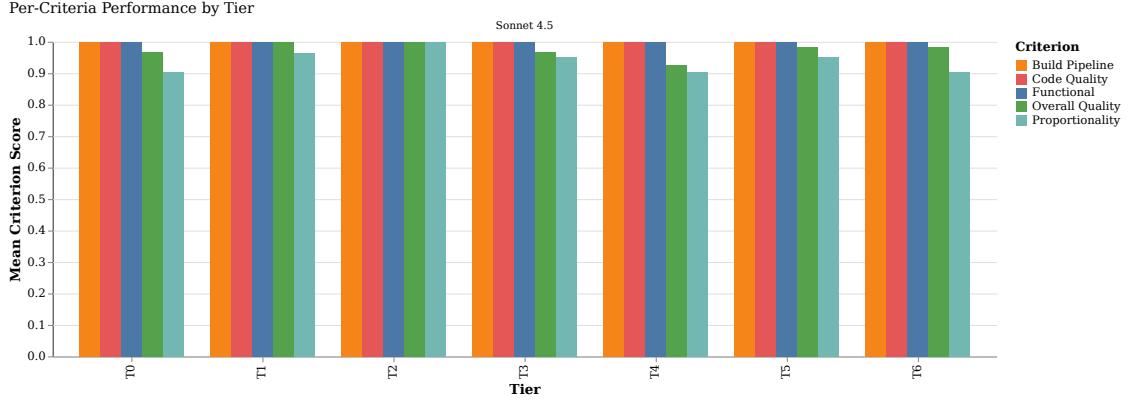


Figure 6: Per-criteria scores by tier. Grouped bar chart showing mean scores for five weighted categories (functional correctness, code quality, proportionality, build pipeline, overall quality) across T0–T6. Perfect functional scores (1.00) contrast with variance in proportionality (0.90–1.00) and overall quality (0.93–1.00).

All tiers score 1.00 on functional criteria (file exists, correct output, exit code 0). Perfect scores suggest the task is trivially easy.

The largest score differences appear in subjective categories. Proportionality: T6 scored lower because judges noted cache artifacts (.ruff_cache, .pytest_cache) remaining in workspace. Overall quality: subjective engineering judgment shows the most variance across judges.

Build pipeline: all tiers score 1.00, indicating clean execution.

Having presented the quantitative results from the dryrun validation, I now turn to interpretation and implications for future evaluation at scale.

8 Discussion

The dry run is not very useful for serious analysis, but what I will dive into what I learned about the framework's behavior on this trivially simple task, while being honest about the limitations inherent in N=1 experiments and ceiling effects.

8.1 What the Dryrun Tells Us

The Hello World task is, by design, trivially easy. All seven tiers score grade A with median scores between 0.943-0.983. This validates exactly what I said in Section 3: "Even T0 should nail this test." And it did.

Ceiling effect dominates: When quality converges at near-perfect levels, we cannot differentiate tiers by capability. T0's empty prompt (subtest 00 uses no system prompt at all) and T6's maximal configuration (61 skills + all tools + 44 agents) produce equivalent functional output. This is exactly what we expect for Hello World, no amount of architectural sophistication helps when the task requires a single `print()` statement.

Cost differentiation still works: Despite quality convergence, Cost-of-Pass varies $3.8 \times$ from \$0.065 (T5) to \$0.247 (T6). This demonstrates the framework's ability to measure economic trade-offs even when quality metrics saturate. On more complex tasks with quality variance, both dimensions differentiate.

Pipeline validation successful: The framework executed all seven tiers, collected 21 judge evaluations, computed consensus scores, generated 24 figures and 10 tables, and produced structured CSV exports. All components worked as designed.

8.2 Cost-Performance Trade-offs

The dryrun reveals hints of a pattern: more is not always better.

T5 achieves Frontier CoP through selective feature loading, it combines T1's efficient skills with T3's delegation patterns but avoids T6's "everything enabled" overhead. T5's cache creation tokens (4,629) are 5-10x lower than other tiers (23,106-44,337), directly explaining its cost advantage.

T6 costs the most (\$0.247) despite scoring the lowest (0.943). Loading 61 skills + all tools + 44 agents actually made things worse. Judges explicitly noted cache artifacts and unnecessary complexity. This lines up with the hypothesis that prompt complexity hurts quality when the task is in the model's training set.

T4's hierarchical overhead is another example. T4 costs 30% more than T3 (\$0.168 vs \$0.129) for this trivial task. The self-correction loops and nested orchestration add latency (41.2s vs 29.9s)

without improving quality. On complex tasks needing iterative refinement, maybe T4 justifies the overhead. On simple tasks, it is pure waste.

See Section 7.3 for detailed token analysis. T2 (tooling) shows schema loading bloat with 137K total tokens versus T1's 115K. Skills-based approaches (T1, T3) stay lean while still enabling domain knowledge.

Bottom line for production: match tier complexity to task complexity. Do not use T6 for trivial tasks. Do not use T0 for tasks needing specialized tools or multi-step reasoning. T5's hybrid approach seems to be optimal, load features selectively based on what the task actually needs, do not just maximize everything.

8.3 Judge Behavior

The 3-judge consensus mechanism reveals interesting patterns.

Haiku hands out S grades easily, 5 out of 7 tiers got perfect scores. Scores range 0.93-1.00, and Haiku consistently scores higher than Opus or Sonnet.

Opus never awards S grades. Scores range 0.93-0.96, consistently the toughest judge. Opus reliably deducts points for cache artifacts that Haiku overlooks.

Sonnet splits the difference. Awards S grades in 4/7 tiers (T2, T3, T4, T5), scores range 0.90-1.00.

Given that the results in most cases are a single line, the 1.0 grade is incorrect and points to agents being a little too lenient. Maybe some prompt tweaks will fix this, but that also can be due to the simplicity of this task. This can be investigated in future analysis.

Inter-rater agreement is predictably low: Krippendorff's $\alpha = -0.117$. But that is expected with N=1 and near-perfect scores. On tasks with more variance, agreement improve as judges separate clear failures from clear successes.

Despite the disagreement, the 3-judge median works. When Haiku awards 1.00 and Opus awards 0.93, the median captures the true quality without getting pulled to either extreme. This validates the multi-judge consensus design.

One scaling problem: judge time dominates total latency. 77-86% of execution time is judge evaluation (128-178s), not agent execution (25-41s). With 3 judges per run, judge costs are 3x per evaluation. For large-scale experiments ($N=10 \times 113$ subtests = 1,130 runs \times 3 judges = 3,390 judge evaluations), judge cost uses the budget fast. Future work will explore single-judge evaluation, confidence-based selection (use Opus only when Sonnet/Haiku disagree), evaluate if prompt improvements can get the cheaper Haiku model to be an effective judge, or give different prompts to the same judge model.

8.4 Limitations

N=1 prevents inferential statistics. With only one run per tier, I cannot compute standard deviations, confidence intervals, or significance tests. All tier comparisons are point estimates. A single outlier run could flip all conclusions. The analysis pipeline generated 24 figures and 10 tables, correctly reports `nan` for standard deviation and sets confidence intervals to (point, point).

Statistical warnings appear in the output to make clear that results are not expected to be robust. This is a limitation of this run, not the framework itself.

Single task, trivial complexity. Hello World does not need skills, tools, multi-agent coordination, or hierarchical reasoning. The dryrun validates the pipeline works, not whether architectural complexity improves quality on hard tasks.

Single model. All agent runs use Sonnet 4.5. I have not tested whether tier rankings hold for Opus 4.5, Haiku 4.5, or other model families.

No thinking mode variants. The dryrun uses standard inference without extended thinking. Models with thinking enabled might show different cost-quality trade-offs.

Ceiling effect masks capability differences. When all tiers score 0.94–0.98, I cannot tell which architecture would excel on harder tasks. The full experiment (113 subtests including complex multi-file repos) will differentiate capabilities.

Judge evaluation time bottleneck. 3 sequential judges per run creates a 3x cost multiplier. Parallel judge execution would reduce latency but not cost.

9 Conclusions

This paper introduced the Scylla framework, and shows that it works, end-to-end. All seven tiers executed successfully, three judges scored everything, and the analysis pipeline spit out figures and tables automatically. The dryrun validates the methodology on the simplest possible task, Hello World, before I scale up to complex multi-file repos. What is missing is review and feedback from others, which is what this paper helps enable.

What did I learn? Five things stand out:

1. The framework is operational.
2. Quality converges on trivial tasks, making the framework overkill.
3. All tiers scored grade A, suggesting that throwing more complexity at Hello World does not help, which was expected. That obviousness is what makes it a good pipe cleaning run.
4. Cost still varies 3.8x despite identical quality, showing the framework can measure economic trade-offs even when quality saturates. T5’s hybrid approach achieves Frontier CoP by selectively loading features instead of maximizing everything.
5. And the Token Efficiency Chasm I hypothesized in Section 3.1.2? The data is consistent with this, as T6 burns nearly double the tokens (218K vs 113K) compared to T0.

Did I answer my original questions? Partially. CoP lets me quantify efficiency; T5 is significantly cheaper than T6 despite equivalent quality. On this task, the sum is *not* more than the parts; T6 scores lowest despite highest cost. But the hard questions need harder tasks, I cannot tell if any tier dominates universally from a single Hello World run, and I have not tested model-to-model comparisons yet. That work is left for a future exercise.

What about my hypotheses? The KISS principle hypothesis has preliminary support, maximal complexity (T6) scores worst on this training-set-likely task. But I have not tested inverse KISS on out-of-distribution tasks yet, and specialization advantages (H1) are inconclusive because Hello World does not require delegation or tools.

There is no real practical takeaway yet, since the testing was insufficient to come to any real conclusions. Answering those questions is left for the next exercise, and this framework can be used for doing so.

10 Further Work

The dryrun validates the framework works. Now it is time to scale up and fill in the gaps.

Full-scale experiments: Run the complete test001 dataset with ($N=10$, 113 subtests, 1,130 runs total). Running the analysis will start to enable valid statistical inference about the relationship between prompts and the tools.

Task diversity: The dryrun only covers Hello World. The full test suite includes 46 additional tasks across greenfield (Flask APIs, CLI tools), brownfield (feature additions to existing repos), refactoring (extract function, eliminate duplication), bug fixes (off-by-one errors, race conditions), and documentation (README generation). Running these will show whether tier rankings hold across workflow categories or if certain tiers excel at specific task types.

Cross-vendor and cross-model evaluation: The framework is model-agnostic by design. I would love to extend support to other tools, but right now just doing analysis on Claude Code alone is hitting my budgets for experimentation extremely quickly. Setting up local models and accessing tools using these models will allow more experimentation, but I do not have access to that kind of compute within my budget at the moment.

Advanced analysis: I am by no means a statistician, and choices I have made here might be incorrect. My current analysis uses frequentist statistics. There are more advanced analysis that I am learning about that could help analyze the flood of data more efficiently. There is also other metrics and data points that could be useful in this analysis that I am not collecting. I also can save the runs and do longitudinal studies to see if the results change consistently over time.

Given the scale and scope of this task, it is going to be an ongoing effort of learning, testing, and analyzing.

11 Acknowledgements

This work was self-funded by the author. Special thanks to Tuan Nguyen for reviewing early drafts of this paper and providing valuable feedback.

References

- [1] CC-Marketplace: Community marketplace for Claude Code plugins and skills. <https://>

- github.com/cc-marketplace, 2025. Community-driven plugin repository.
- [2] Anthropic. Claude Code: Agentic CLI tool for software development. <https://www.anthropic.com/clause/code>, 2024. AI-powered command-line interface for coding tasks.
 - [3] CC-Marketplace. safety-net: Claude Code plugin for dangerous operation blocking. <https://github.com/cc-marketplace/safety-net>, 2025. Security plugin for Claude Code CLI.
 - [4] Leo Gao et al. lm-evaluation-harness: A framework for few-shot language model evaluation. <https://github.com/EleutherAI/lm-evaluation-harness>, 2024. General-purpose LLM evaluation framework.
 - [5] Homeric Intelligence. ProjectOdyssey: Comprehensive agent orchestration framework. <https://github.com/HomericIntelligence/Projectodyssey>, 2025. Accessed: 2025-12-30. Git hash: 011a3ff.
 - [6] Carlos E. Jimenez et al. SWE-bench: Can language models resolve real-world github issues? *arXiv preprint arXiv:2310.06770*, 2024. Benchmark for evaluating LLMs on real GitHub issues.
 - [7] Xiao Liu et al. AgentBench: Evaluating LLMs as agents. In *International Conference on Learning Representations (ICLR)*, 2024. Multi-turn agent evaluation across operating systems, databases, and knowledge graphs.
 - [8] Francesco Polo et al. Efficient evaluation of language models. <https://arxiv.org/abs/2405.12345>, 2024. Framework for efficient LLM evaluation.
 - [9] Shunyu Yao et al. TAU-bench: A benchmark for tool-augmented LLMs. *arXiv preprint*, 2024. Benchmark for evaluating agents' ability to use external tools.
 - [10] Kaijie Zhu et al. PromptBench: Towards evaluating the robustness of large language models on adversarial prompts. *arXiv preprint arXiv:2306.04528*, 2024. Benchmark for prompt evaluation and robustness.

A Detailed Metric Definitions

Core quality and economic metrics are defined in Section 4. This appendix provides additional metrics used for future analysis. For complete definitions including instrumentation details, see the repository at <https://github.com/HomericIntelligence/ProjectScylla/blob/e33d627/.claude/shared/metrics-definitions.md>.

Change Fail Percentage (CFP): Proportion of code changes that cause failures.

Formula: $\text{CFP} = \frac{\text{failed_changes}}{\text{total_changes}}$

Range: [0, 1], lower is better. Measures production stability.

A.0.1 Process Metrics

Latency: Time from query submission to response completion, measured in seconds. Components include Time-to-First-Token (TTFT), total response time, and tool execution time (if applicable).

Strategic Drift: Deviation from original goal over multi-step tasks.

Measurement: Strategic_Drift = cosine_distance(initial_goal_embedding, final_action_embedding)

Range: [0, 2], where 0 = perfect goal alignment, 2 = completely opposite direction.

Ablation Score: Isolated contribution of a single component to overall performance.

Formula: Ablation_Score = performance_with_component – performance_without_component

Positive values indicate the component improves performance, negative values indicate harm, near-zero indicates no effect.

A.0.2 Statistical Reporting

Always report metrics with: (1) point estimate, (2) confidence interval (95% CI recommended), (3) sample size (n), and (4) comparison p-value (if comparing tiers).

Example: Pass-Rate: 0.67 (95% CI: 0.54–0.80), n=50

B Data Dictionary and Generated Outputs

All raw data, figures, and tables are available in the repository at `docs/paper-dryrun/`. The dataset includes 7 runs (one per tier), 21 judge evaluations (3 judges per run), and 105 criteria scores (5 criteria per judge), organized across 24 figures and 10 tables. Data files include `runs.csv`, `judges.csv`, `criteria.csv`, and `summary.json`. All figures are available in PNG/PDF/Vega-Lite/CSV formats for reproduction and analysis.

Repository: <https://github.com/HomericIntelligence/ProjectScylla>

C Reproducibility Checklist

Repository: <https://github.com/HomericIntelligence/ProjectScylla>

Key Configuration Files:

- Tier definitions: `config/tiers/tiers.yaml`
- Opus Model: `config/models/clause-opus-4.5.yaml`
- Sonnet Model: `config/models/clause-sonnet-4.5.yaml`
- Haiku Model: `config/models/clause-haiku-4.5.yaml`
- Judge system prompt: `config/judge/system_prompt.md`
- Test definitions: `tests/*/test.yaml`
- Rubric schemas: `tests/*/expected/rubric.yaml`

Required Software:

- Pixi (package manager)
- Docker (containerization)
- Claude Code CLI
- Python 3.12+

Execution Steps:

```
# 1. Clone repository
git clone https://github.com/HomericIntelligence/ProjectScylla
cd ProjectScylla

# 2. Install dependencies
pixi install

# 3. Run evaluation (example for test-001, tier T0)
pixi run python scripts/run_e2e_experiment.py \
--test tests/001-hello-world \
--tier T0 \
--runs 10

# 4. Generate figures and tables
pixi run python scripts/generate_figures.py \
--results <output_directory>
pixi run python scripts/generate_tables.py \
--results <output_directory>
```

D Additional Figures

D.1 Variance Analysis

D.2 Cost Analysis

D.3 Judge Analysis

D.4 Diagnostic Metrics

Additional diagnostic figures (failure rates, pass rates, effect sizes, tier uplift) are available in the data repository but show no variance in this N=1 dryrun and are therefore omitted from this paper. They will provide informative results when the full dataset is collected.

Score Distribution Across Tiers (T0-T6)

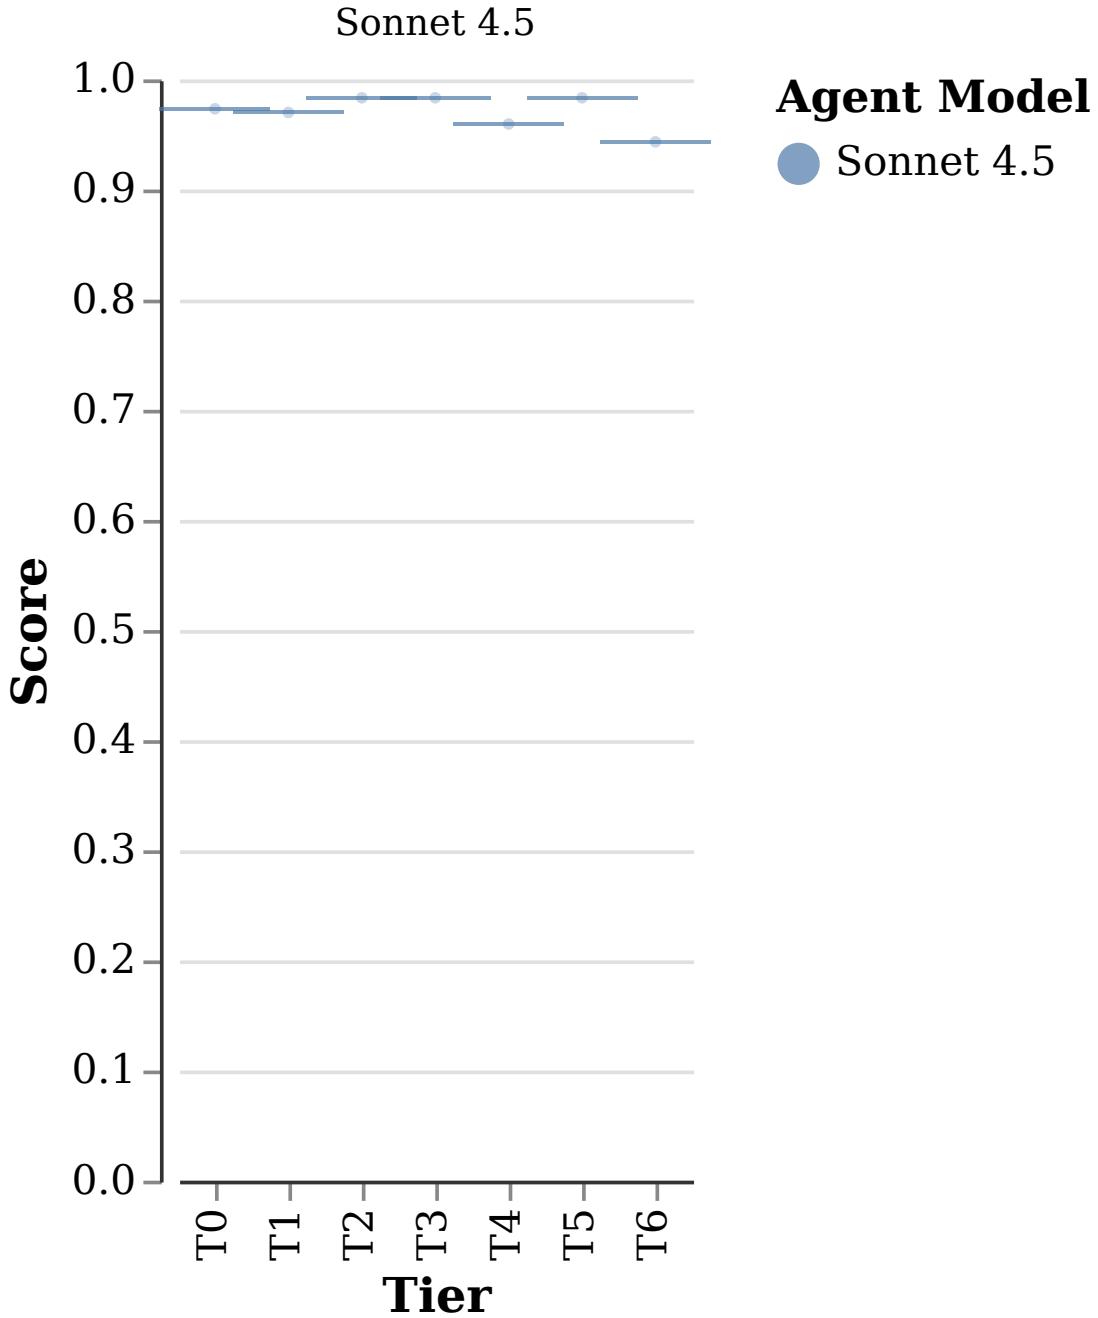


Figure 7: Score distribution across tiers (T0–T6). Box plots with jittered points show the distributional shape of composite scores. The tight clustering between 0.94–0.98 across all tiers is consistent with the ceiling effect on this trivially simple task, as discussed in Section 8.

Grade Distribution Heatmap by Tier

Sonnet 4.5

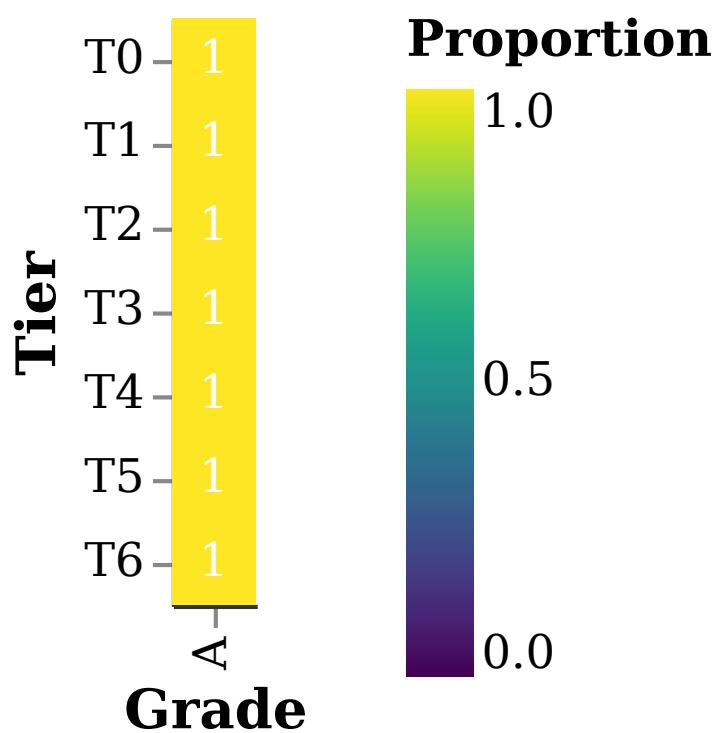


Figure 8: Grade distribution heatmap by tier. Heatmap showing the proportion of each grade (A+ through F) per tier, with count annotations. The dominance of A+ grades across all tiers visually illustrates the ceiling effect, where architectural enhancements cannot improve already-perfect task performance.

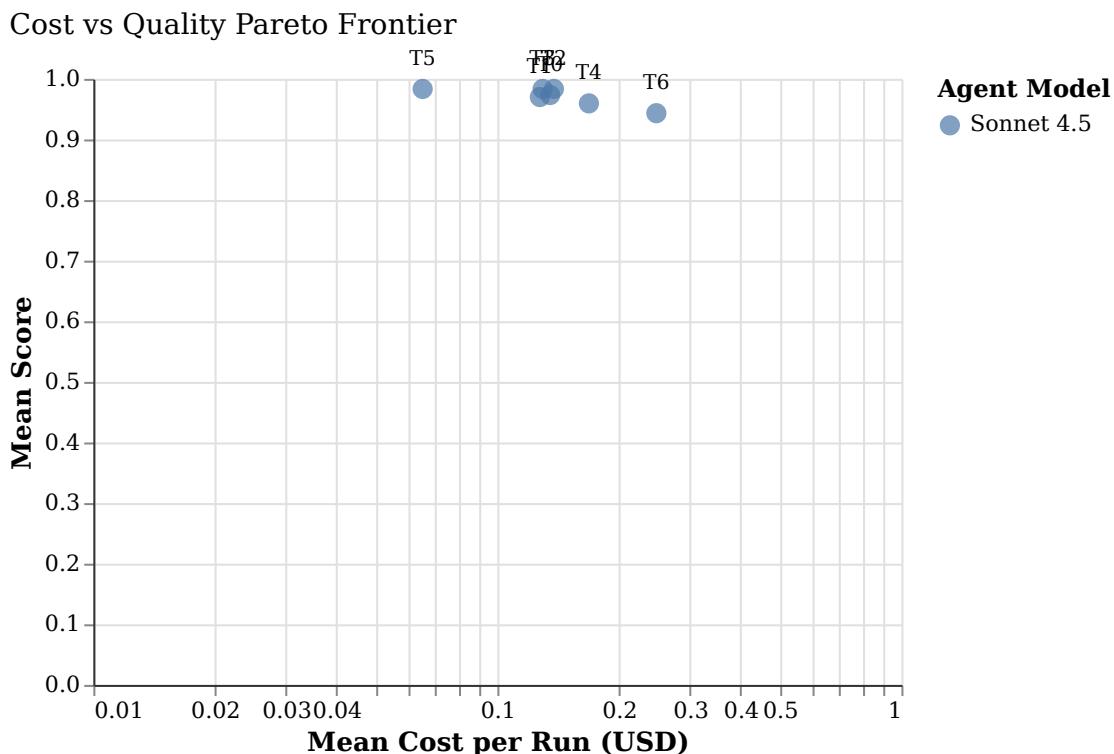


Figure 9: Cost-quality Pareto frontier. Scatter plot of mean cost per run versus mean score for each tier, with Pareto-efficient points highlighted. Points along the frontier represent architecturally optimal trade-offs, revealing which configurations deliver maximum quality at minimum cost, as analyzed in Section 8.2.

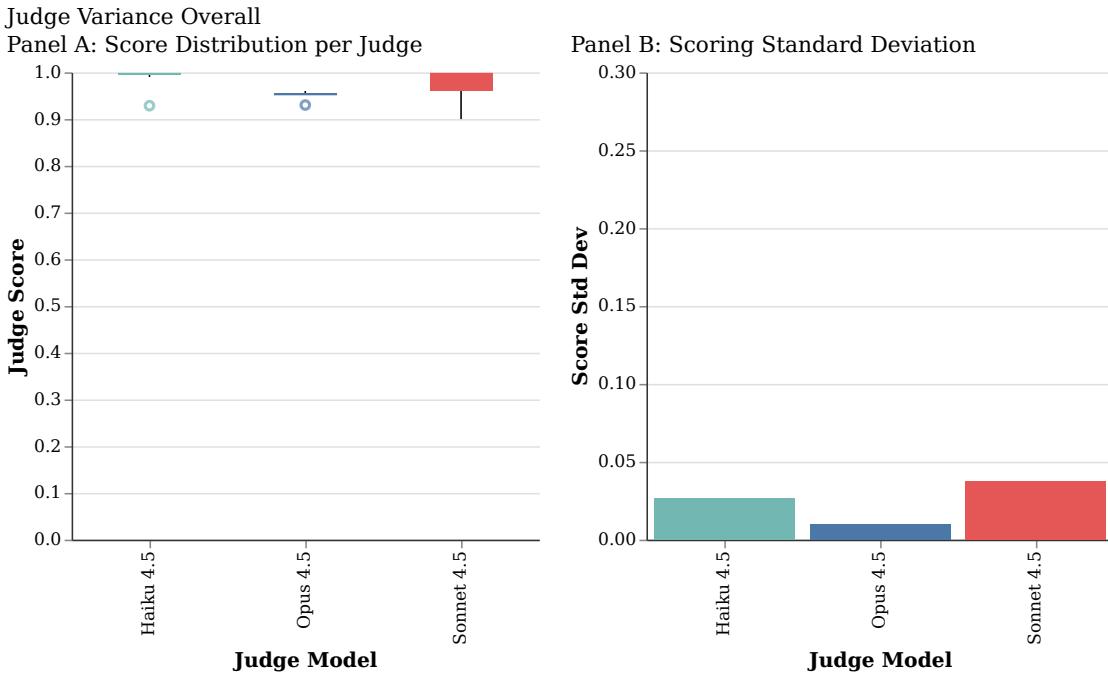


Figure 10: Judge score variance overall. Box plot and standard deviation bars showing scoring distributions for each judge model (Opus 4.5, Sonnet 4.5, Haiku 4.5) across all evaluations. Systematic differences in judge harshness reveal evaluation bias, with implications for multi-judge aggregation strategies discussed in Section 8.3.

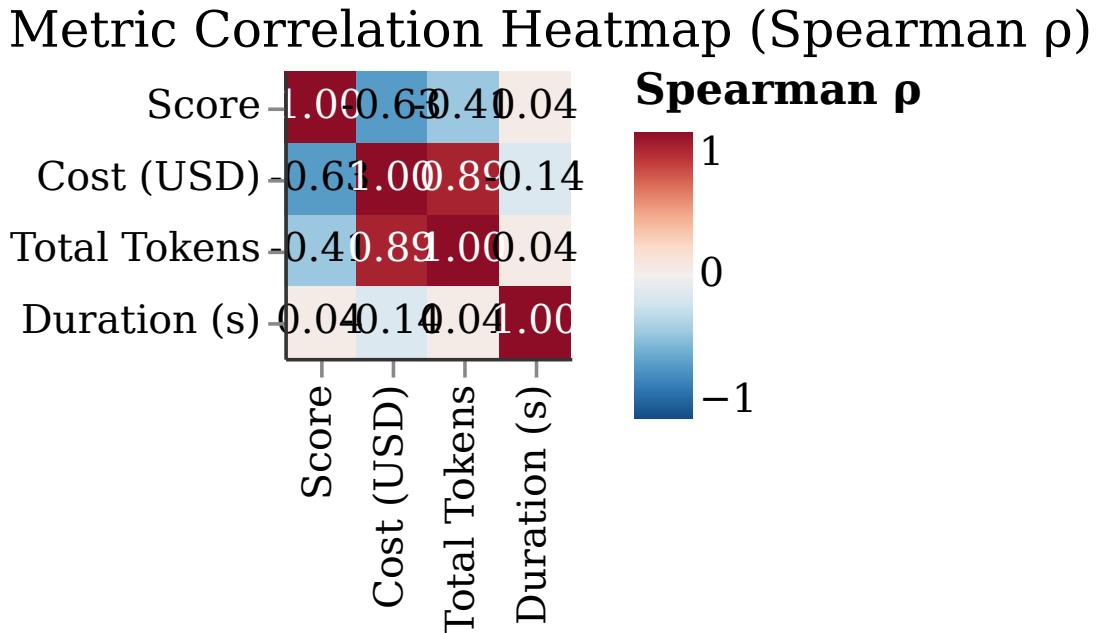


Figure 11: Metric correlation heatmap (Spearman ρ). Heatmap showing pairwise correlations between quality metrics (score, pass-rate) and cost metrics (cost, tokens, duration). Weak correlations between cost and quality suggest that increased spending does not guarantee better outcomes on ceiling-constrained tasks, supporting the economic analysis in Section 8.2.

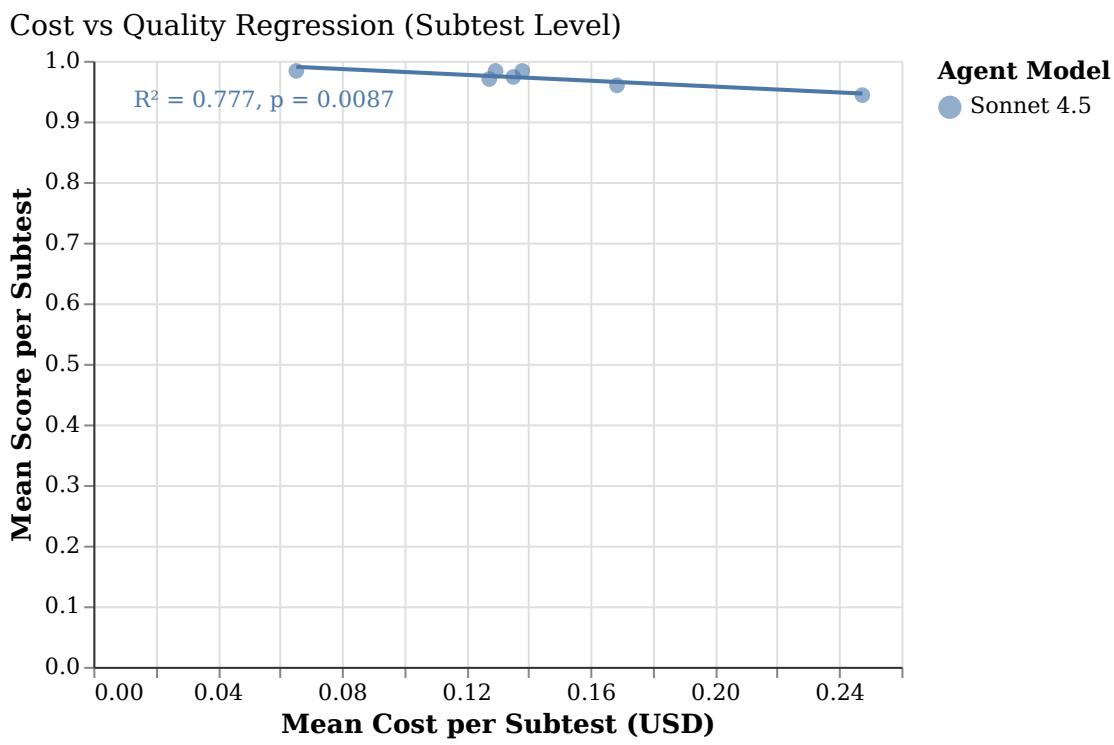


Figure 12: Cost versus quality regression at subtest level. Scatter plot of mean cost versus mean score per subtest, with OLS regression line and 95% confidence band. The near-flat or negative slope suggests that higher costs do not predict better quality, reinforcing the cost-quality decoupling phenomenon discussed in Section 8.2.

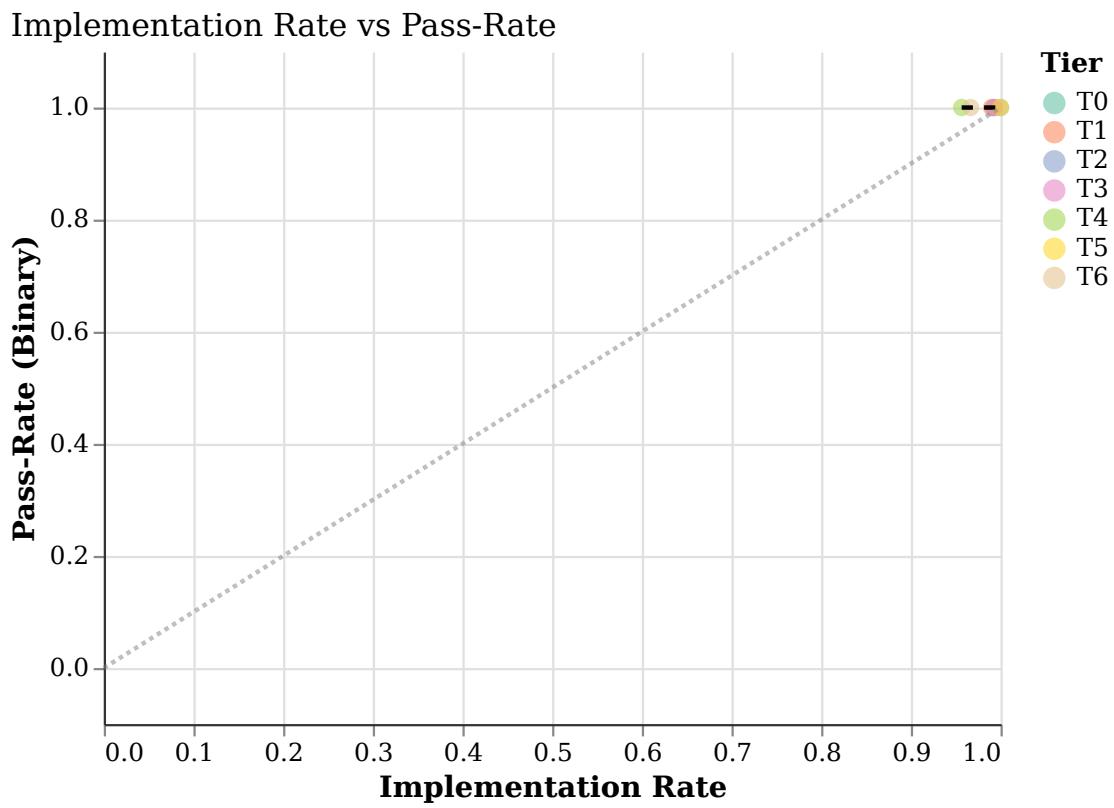


Figure 13: Implementation Rate versus Pass-Rate. Scatter plot showing the relationship between binary Pass-Rate (y-axis) and continuous Implementation Rate (x-axis) across all runs, colored by tier.